



ORACLE

Move to Oracle Database 23ai

Everything you need to know about Oracle Multitenant - part 1



Safe harbor statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, timing, and pricing of any features or functionality described for Oracle's products may change and remains at the sole discretion of Oracle Corporation.

The materials in this presentation pertain to Oracle Health, Oracle, Oracle Cerner, and Cerner Enviza which are all wholly owned subsidiaries of Oracle Corporation. Nothing in this presentation should be taken as indicating that any decisions regarding the integration of any EMEA Cerner and/or Enviza entities have been made where an integration has not already occurred.

Oracle
DBAs
run the world





ROY SWONGER

Vice President

Database Upgrade, Utilities & Patching

 royfswonger





MIKE DIETRICH

Senior Director Product Management
Database Upgrade, Migrations & Patching

mikedietrich

@mikedietrichde

<https://mikedietrichde.com>



DANIEL OVERBY HANSEN

Senior Principal Product Manager
Database Upgrade, Migrations & Patching

dohdatabase

@dohdatabase

<https://dohdatabase.com>



RODRIGO JORGE

Senior Principal Product Manager
Database Upgrade, Migrations & Patching

[rodrigoaraujorge](#)

[@rodrigojorgedb](#)

<https://dbarj.com.br/en>



ALEX ZABALLA

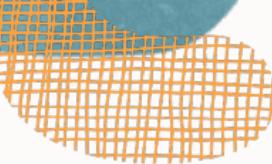
Distinguished Product Manager
Database Upgrade, Migrations & Patching

alexzaballa

@alexzaballa

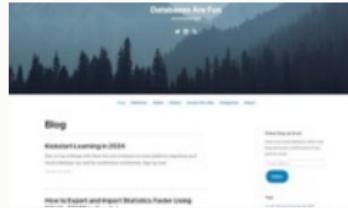
<https://alexzaballa.com>

Find Slides and Much More on Our Blogs



MikeDietrichDE.com

Mike.Dietrich@oracle.com



dohdatabase.com

Daniel.Overby.Hansen@oracle.com



DBArj.com.br

Rodrigo.R.Jorge@oracle.com



AlexZaballa.com

Alex.Zaballa@oracle.com



Get the Slides



<https://dohdatabase.com/slides>

Episode 1

Release and Patching Strategy

105 minutes – Feb 4, 2021



Episode 2

AutoUpgrade to Oracle Database 19c

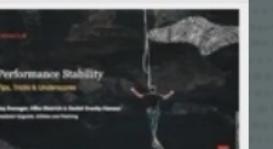
115 minutes – Feb 20, 2021



Episode 3

Performance Stability, Tips and Tricks and Underscores

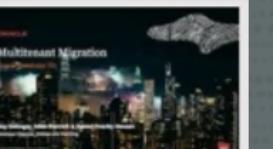
120 minutes – Mar 4, 2021



Episode 4

Migration to Oracle Multitenant

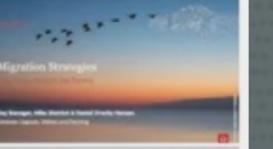
120 minutes – Mar 16, 2021



Episode 5

Migration Strategies – Insights, Tips and Secrets

120 minutes – Mar 25, 2021



Episode 6

Move to the Cloud – Not only for techies

115 minutes – Apr 8, 2021



Recorded Web Seminars

<https://MikeDietrichDE.com/videos>

More than 35 hours of technical content,
on-demand, anytime, anywhere

Part 2

Move to Oracle Database 23ai

- Everything you need to know about Oracle Multitenant

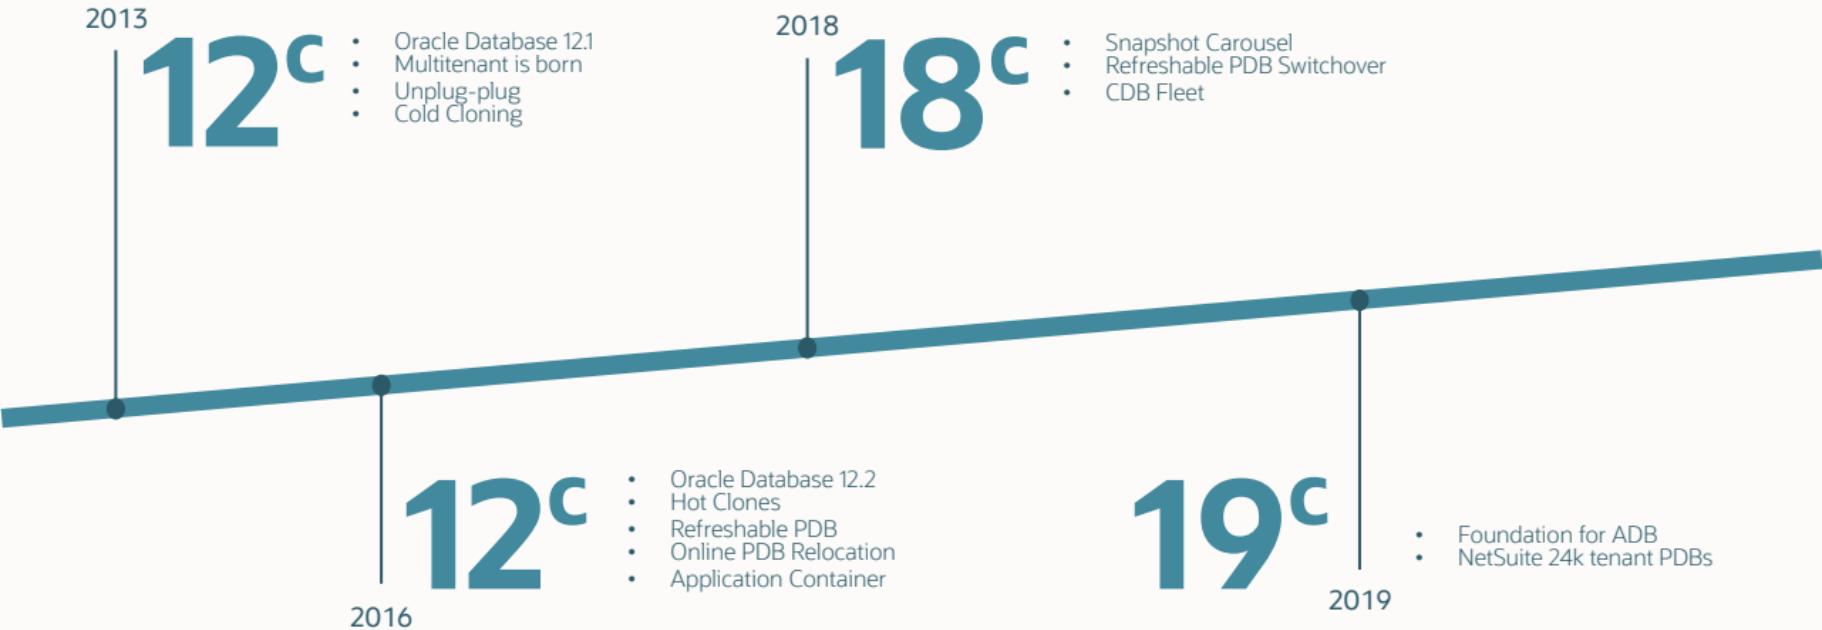
Live on June 27, 14:00 CEST

[Sign up](#)



Architecture

Multitenant Evolution



The background of the slide features a dark grey gradient with abstract white patterns. On the right side, there are several wavy, grid-like lines. In the center, there is a cluster of binary digits (0s and 1s) arranged in a roughly triangular shape.

*Starting with Oracle Database 21c,
installation of non-CDB Oracle Database architecture
is no longer supported*

[Upgrade Guide, 23ai](#)

*Once you upgrade **beyond** Oracle Database 19c,
you must convert to the multitenant architecture*

*Oracle Database 19c is the **last** release
to support the non-CDB architecture*

Next Long Term Support release

Oracle Database 23ai

Upgrade possible only from:

- Oracle Database 19c
- Oracle Database 21c



Multitenant enables an Oracle database
to function as a container database



Generally, you don't need to change your application to use a pluggable database

Single vs. Multitenant



Single Tenant

One PDB
No extra license



Multitenant

Multiple PDBs
Extra license if more than 3 PDBs

--Use up to 3 user-created PDBs
--without a license for Multitenant option.
--Applies to Oracle Database 19c and newer, including SE2

```
alter system set max_pdbs=3;
```



Multitenant

1



2



3



You always create a new CDB

- CREATE DATABASE ... ENABLE PLUGGABLE DATABASE
- Using DBCA

Or clone an existing CDB

Multitenant

1



2



3



When you create a new CDB, it contains:

- The root container
- The seed container

Multitenant

1



2



3



You can create PDBs:

- From the seed container
- By cloning other PDBs
- By converting a non-CDB



Be cautious making changes to *PDB\$SEED*

- Your own customizations do not belong *PDB\$SEED*

Containers

CDB\$ROOT

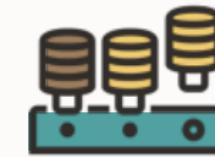
is always 1

PDB\$SEED
is always 2

User-created PDBs
have ID 3 or above

```
SQL> select con_id, name  
      from v$containers;
```

CON_ID	NAME
1	CDB\$ROOT
2	PDB\$SEED
3	PDB1
4	PDB2



```
alter session set container=CDB$ROOT;
show con_id
```

CON_ID

1

```
alter session set container=PDB1;
select sys_context('USERENV', 'CON_ID') as con_id from dual;
```

CON_ID

3



You can switch between containers,
but a transaction belongs to one PDB only

```
alter session set container=PDB1;  
insert into table1 values (...);
```

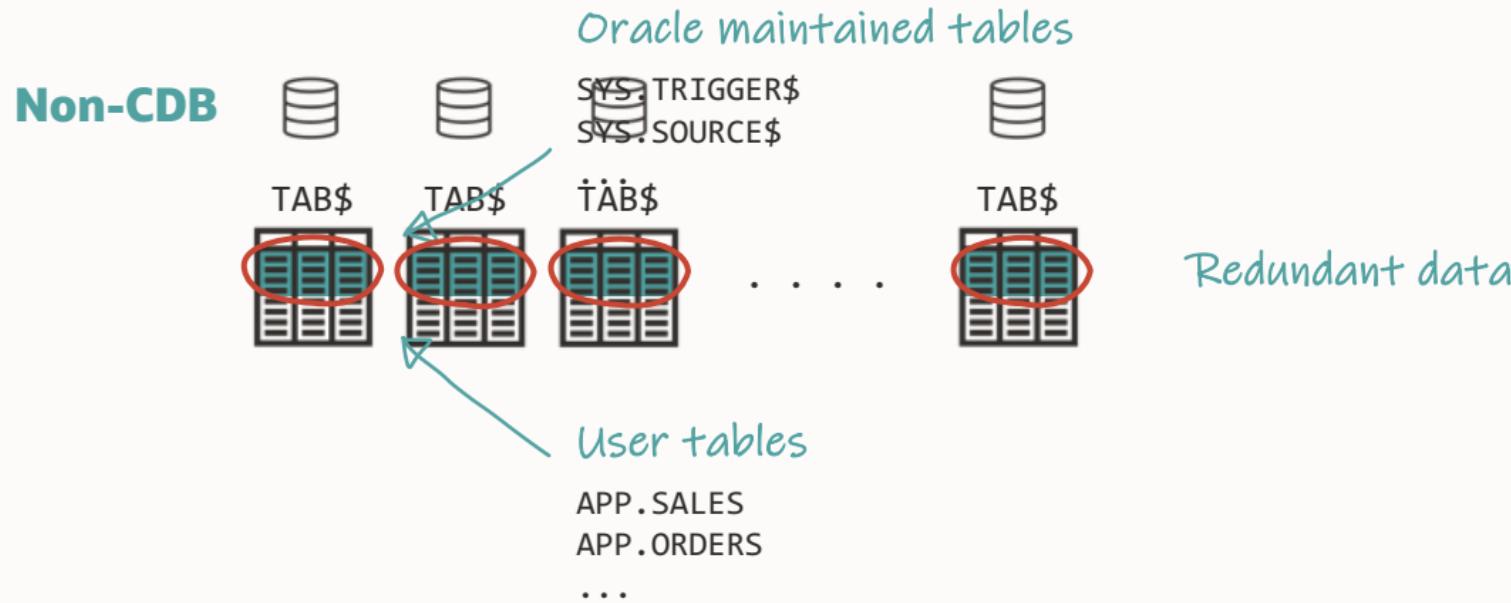
```
alter session set container=PDB2;  
insert into table2 values (...);
```

ORA-65023: active transaction exists in container PDB1

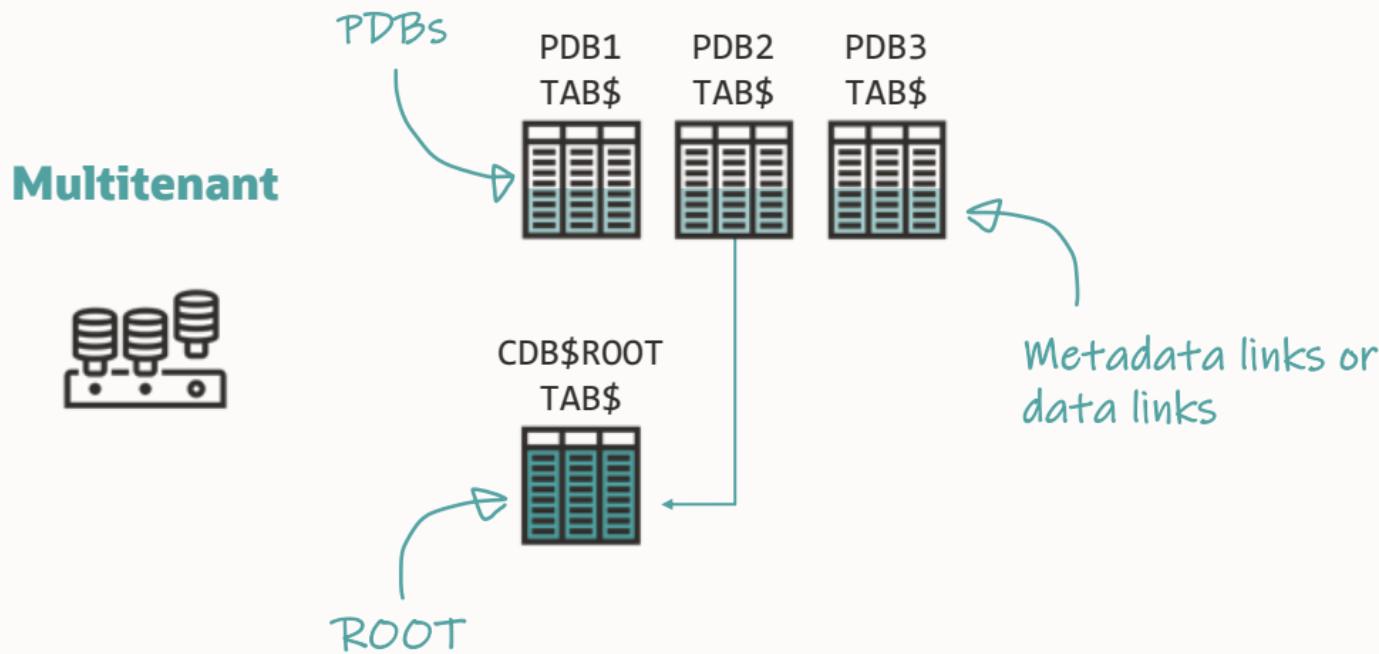


Multitenant implements
data dictionary separation

Data Dictionary Architecture



Data Dictionary Architecture



Data Dictionary Architecture



Deduplication

By storing data just once, you can save space in the data dictionary.



Faster

Smaller dictionaries take less time to patch or upgrade.



Easier

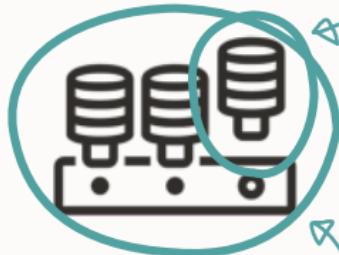
With much metadata stored in root, there is less work for a patch or upgrade.



CDB views describe the entire CDB including all PDBs

- Column **CON_ID** indicates the originating container

Data Dictionary Architecture



```
SELECT * FROM USER_OBJECTS  
ALL_OBJECTS  
DBA_OBJECTS
```

```
SELECT * FROM CDB_OBJECTS
```

Applies to any
data dictionary view

CDB_ALL_TABLES
CDB_ANALYTIC_VIEW_ATTR_CLASS
.
.
.
CDB_XTERNAL_TAB_SUBPARTITIONS

```
alter session set container=CDB$ROOT;
```

```
select con_id, tablespace_name  
from cdb_tablespaces;
```

Originating
container

CON_ID	TABLESPACE_NAME
--------	-----------------

1	SYSTEM
1	SYSAUX
1	UNDOTBS1
1	TEMP
1	USERS
2	SYSTEM
2	SYSAUX
2	UNDOTBS1
2	TEMP
4	SYSTEM
4	SYSAUX
4	UNDOTBS1
4	TEMP

PDB\$SEED information,
hidden by default

```
alter session set container=PDB1;
```

```
select tablespace_name from dba tablespaces;
```

TABLESPACE_NAME

SYSTEM
SYSAUX
UNDOTBS1
TEMP



A PDB never sees information
from other PDBs

```
alter session set container=CDB$ROOT;  
select count(*) from cdb_objects;
```

COUNT(*)

114658

```
alter session set container=PDB1;  
select count(*) from cdb_objects;
```

COUNT(*)

23980



You define common objects in root,
and they are available in all PDBs

- Most options in application containers

```
alter session set container=CDB$ROOT;
```

```
create user C##OPS identified by oracle;  
grant create session to C##OPS container=all;
```

```
create pluggable database pdb1 ... ;  
alter pluggable database pdb1 open;
```

```
conn C##OPS/oracle@pdb1  
Connected.
```

Common Objects

You can define common:

- Profiles
- Roles
- User
- Audit configuration
- Other objects available in application containers

Useful for:

- Maintenance and monitoring users
- Self-service functionality
- Separation of duties
- Enforcing security



You can change the common prefix
with **COMMON_USER_PREFIX**

- We do not recommend changing it
- Use extreme care if you choose to do so



Create same set of common objects in all CDBs to avoid issues during clone/relocate



The database creates common directories

```
alter session set container=PDB1;  
  
select directory_path, origin_con_id  
from cdb_directories  
where directory_name='DATA_PUMP_DIR';  
  
DIRECTORY_PATH  
/u01/app/oracle/admin/CDB2/dpdump/13D6BC6605416ECEE065000000000001
```

PDB GUID

ORIGIN_CON_ID
1

Common directory,
created in root

```
create or replace directory DATA_PUMP_DIR AS '/tmp';
```

ERROR at line 1:
ORA-65040: operation not allowed from within a pluggable database

```
drop directory DATA_PUMP_DIR;
```

ERROR at line 1:
ORA-65040: operation not allowed from within a pluggable database



Use your own directories
if you want to decide on the directory path



You make most configuration
in the root container

```
alter session set container=PDB1;
```

```
alter database backup controlfile to trace;
```

ORA-65040: operation not allowed from within a pluggable database



Non-CDB Compatible

- Some ALTER DATABASE and ALTER SYSTEM commands fail in a PDB
- Enable non-CDB compatibility by setting NONCDB_COMPATIBLE=TRUE
 - When you can't change the application
 - When you accept the reduced security

```
SQL> alter system set noncdb_compatible=true;  
SQL> shutdown immediate  
SQL> startup
```

```
SQL> alter system set noncdb_compatible=true;  
SQL> shutdown immediate  
SQL> startup
```

```
SQL> alter session set container=PDB1;  
SQL> alter database backup controlfile to trace;
```

Database altered.





Fine-tune PDBs with instance parameters

- Parameters apply to PDBs as well
- Some parameters are PDB modifiable

```
SQL> select name from v$system_parameter where ispdb_modifiable='TRUE';
```

NAME

NAME
adg_account_info_tracking
allow_rowid_column_type
approx_for_aggregation
approx_for_count_distinct
approx_for_percentile
.
.
.
xml_handling_of_invalid_chars

246 rows selected.





Use ORAdiff to find PDB modifiable parameters

- Free tool
- <https://oradiff.oracle.com>



A cloned or moved PDB
keeps the changed parameters

- Certain exceptions exist

```
--Find specific parameters that has been defined in a specific PDB  
select name, value from v$system_parameter where con_id=<id>;
```



Share resources between PDBs

Resource Consolidation

Non-CDB
database



Memory



Background processes



Files

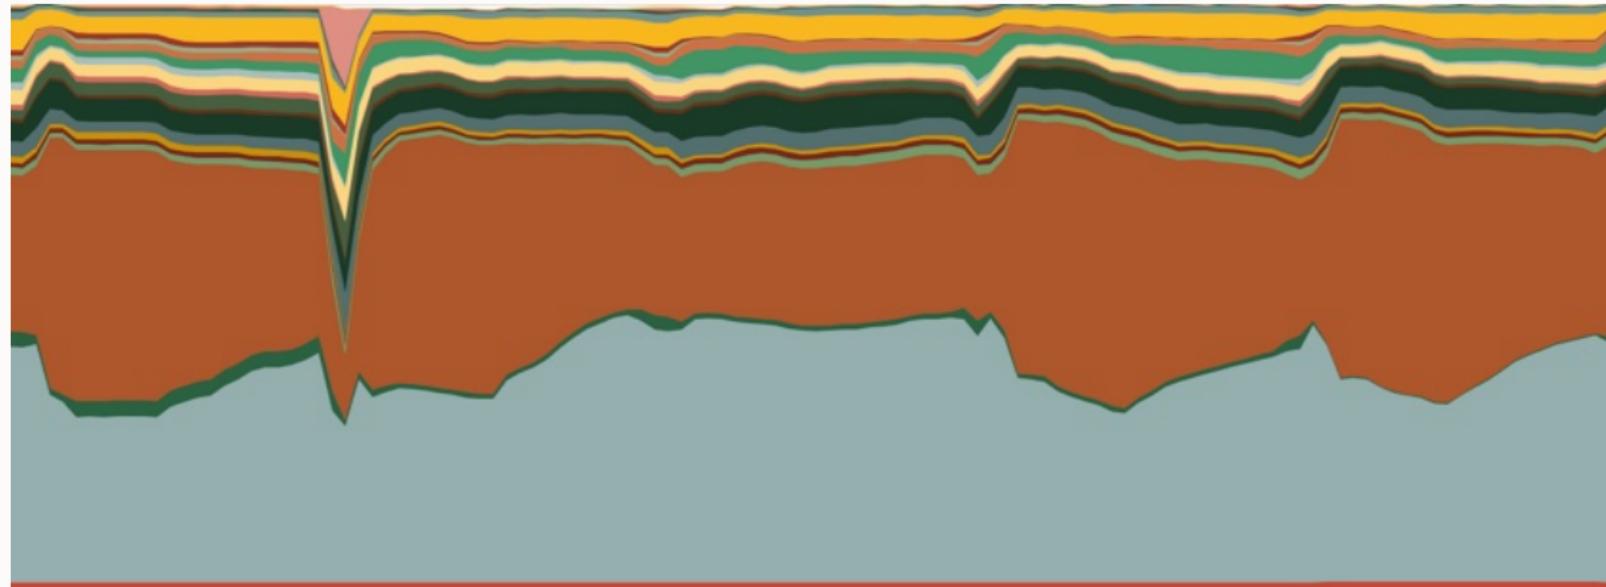


Resource Consolidation



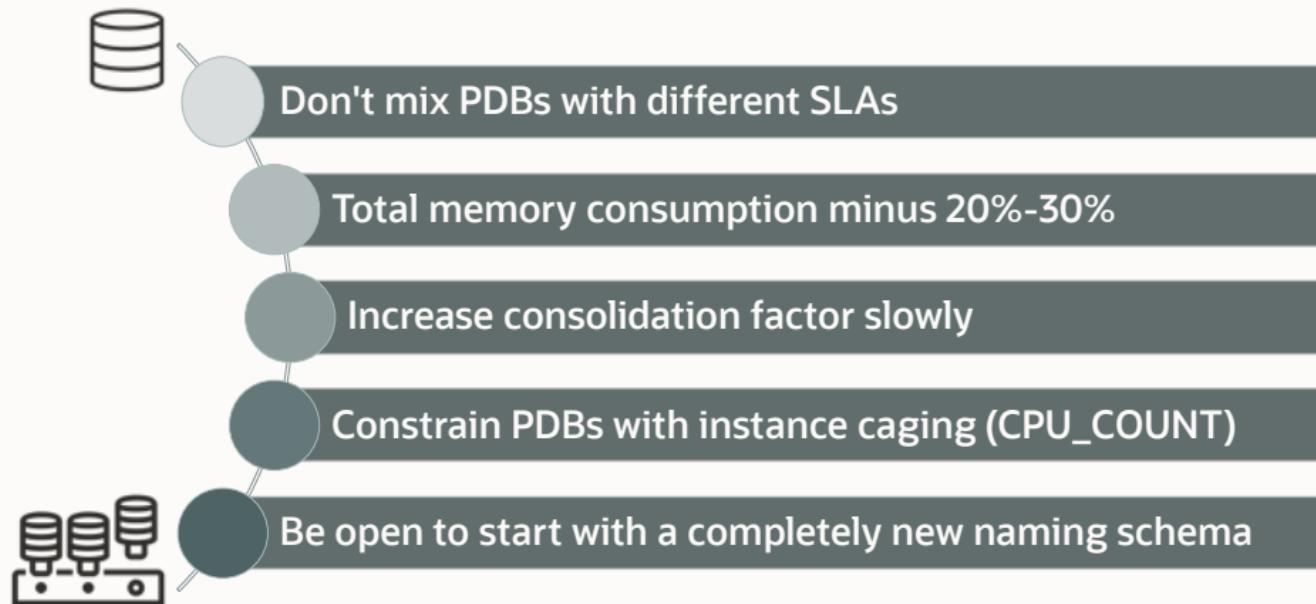


Resource Consolidation



Consolidation Strategies?

There is no "*best*" strategy



Using a Swingbench benchmark,
a single-core machine could host nine non-CDBs
before reaching 75 % CPU utilization

By going multitenant the number of databases reached 123 PDBs

A US Health Care provided managed to

- Reduce the number of database instances by 7x
- Reduce the number of physical servers by 50 %



You can run multiple CDBs on the same host and out of the same Oracle home

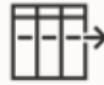
Consolidation



Schema consolidation



PDB consolidation



Virtual Private Database

- Less complexity
- Better isolation
- Operational benefits
- Easier cloning

A global provider of financial services states

The multitenant architecture gives us complete client separation out of the box, without having to maintain a Virtual Private Database setup.

We went away from Virtual Private Database and consolidated our different clients in individual PDBs.

This reduced the complexity of our database implementation and made operations much easier.



The *many-as-one* principle
eases maintenance operations



Many-as-one



Patch databases
one by one

Many-as-one



Patch all containers
in one operation



Many-as-one

Applies to:



- Upgrading
- Patching
- Configuring and performing backups
- Configuring Data Guard
- Configuring RAC
- Monitoring
- ... and many other operations



Benefits

The [multitenant](#) architecture enables



- 1** Self-contained PDBs
- 2** Common and easier operations
- 3** Resource sharing and consolidation



Create

Create Container Database



1 Character set

2 Components

3 COMPATIBLE

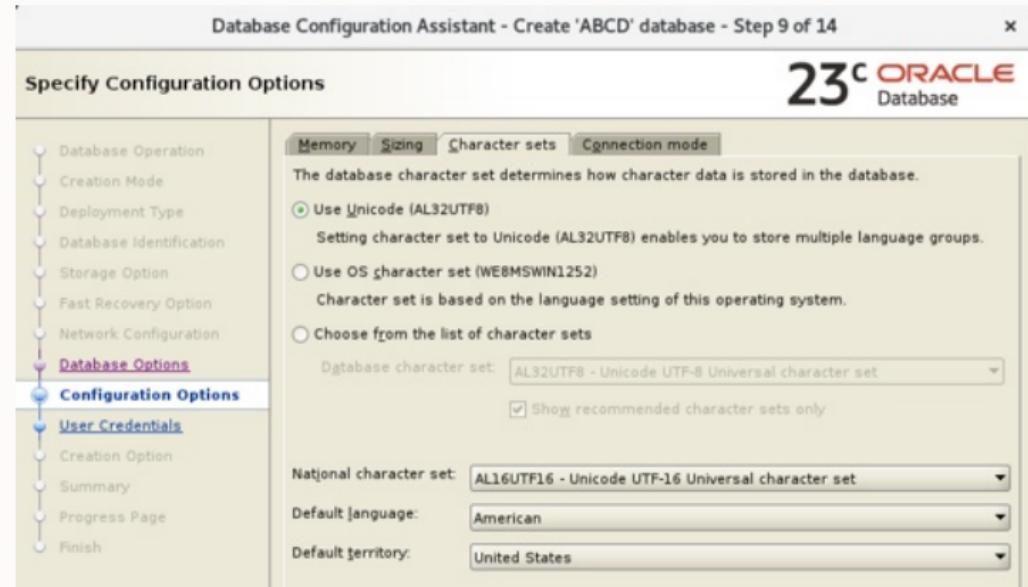
Create Container Database

1 Character set

- Always choose AL32UTF8
- Allows PDBs with any character set

2 Components

3 COMPATIBLE



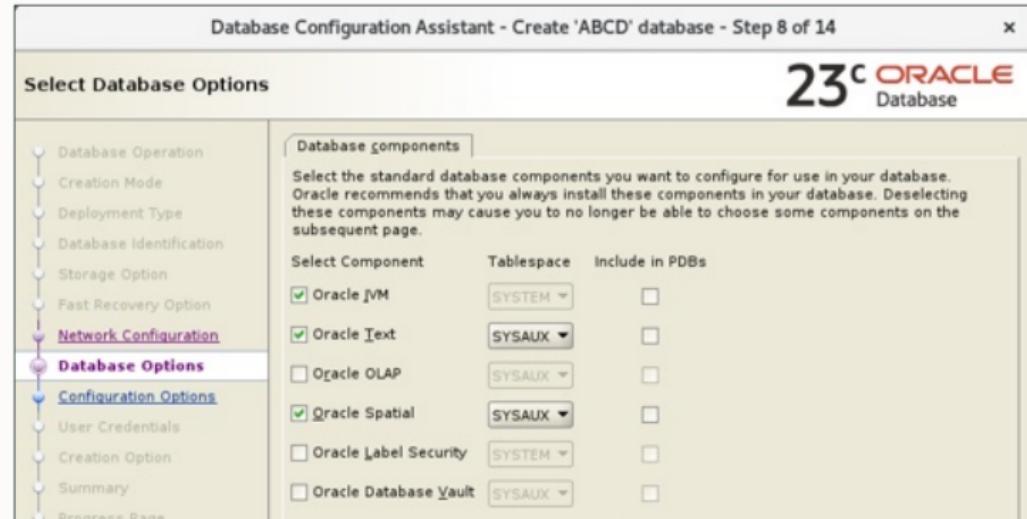
Create Container Database

1 Character set

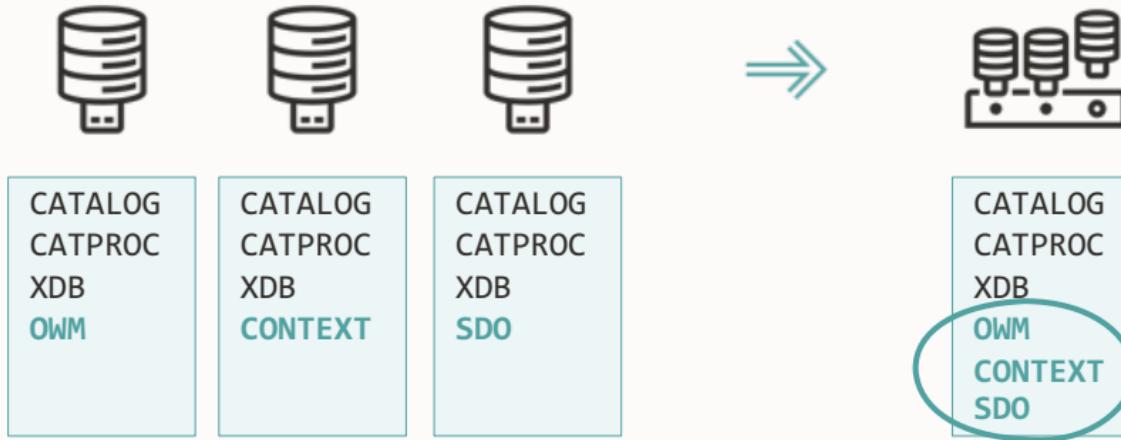
2 Components

- Install as many as you need
- No more than that

3 COMPATIBLE



Components



Create Container Database

1 Character set

2 Components

3 COMPATIBLE

- Keep at the default setting, 23.4.0
- Unless you want the option of downgrade

All initialization parameters				
<input type="checkbox"/> Show advanced parameters				
Name ▲	Value	Include in spfile	Category	
cluster_database	FALSE	<input type="checkbox"/>	Cluster Database	
compatible	19.0.0	<input checked="" type="checkbox"/>	Miscellaneous	
control_files	("/u02/oradata/{DB}_UNI...	<input checked="" type="checkbox"/>	File Configuration	
db_block_size (bytes)	8192	<input checked="" type="checkbox"/>	Cache and I/O	
db_create_file_dest	/u02/oradata/{DB}_UNIQUE ...	<input type="checkbox"/>	File Configuration	

```
-- Always set compatible to the default of a release  
-- Use three digits only
```

```
alter system set compatible='23.4.0' scope=spfile;
```

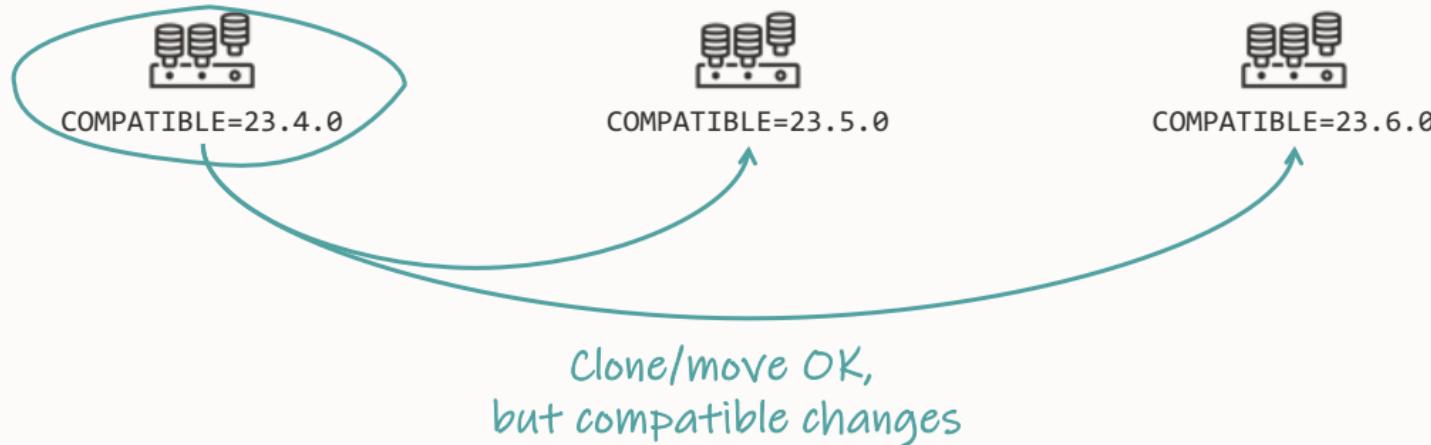
```
-- Should I change compatible when patching?  
-- No, this is a bad idea
```

```
alter system set compatible='23.5.0' scope=spfile;
```

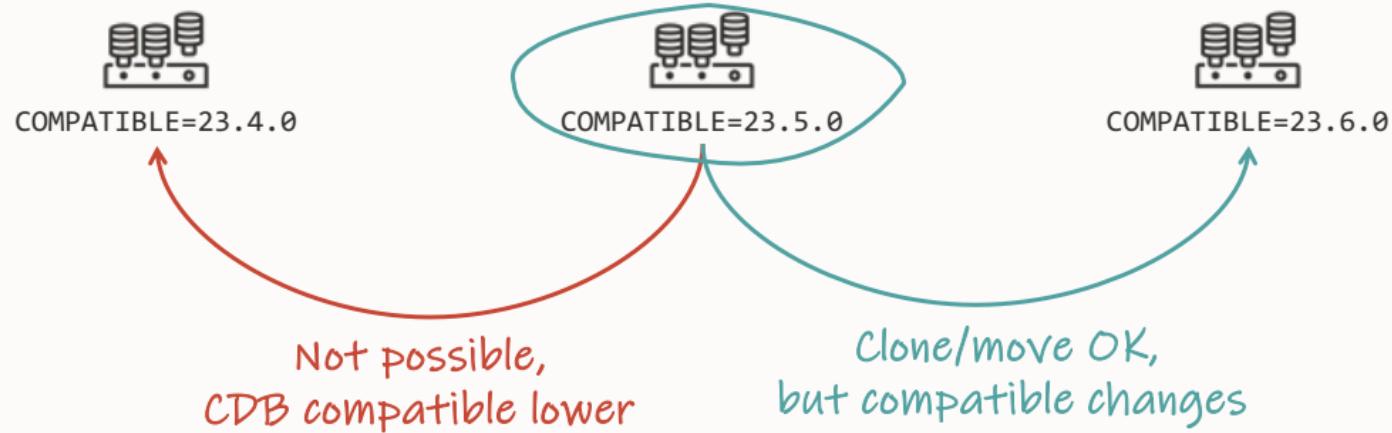
Compatible

- On plug-in, a PDB adopts COMPATIBLE of the CDB **silently** and **without confirmation**
- Changing COMPATIBLE is irreversible
- Might prevent future move or clone of the PDB

Compatible



Compatible





Keep **COMPATIBLE** at the same setting
in all CDBs on the same release

```
--Allows CDB views to include information on PDB$SEED objects.  
--By default, such information is hidden.  
--https://mikedietrichde.com/2017/07/21/why-exclude\_seed\_cdb\_view-is-now-an-underscore-in-oracle-12-2/
```

```
alter system set "_exclude_seed_cdb_view"=false;
```

```
alter session set container=CDB$ROOT;

alter system set "_exclude_seed_cdb_view"=false;

select con_id, tablespace_name
from   cdb_tablespaces;
```

CON_ID	TABLESPACE_NAME
1	SYSTEM
1	SYSAUX
1	UNDOTBS1
1	TEMP
1	USERS
2	SYSTEM
2	SYSAUX
2	UNDOTBS1
2	TEMP
4	SYSTEM
4	SYSAUX
4	UNDOTBS1
4	TEMP

Create Container Database

<https://mikedietrichde.com/2018/08/08/creating-cdbs-non-cdbs-with-less-options/>

<https://mikedietrichde.com/2017/07/11/always-create-custom-database/>

<https://mikedietrichde.com/2017/07/26/remove-clean-components-oracle-11-2-12-2/>



Migration



Migration to multitenant is a one-time operation that requires downtime

- No downtime when using Oracle GoldenGate

Migration



1 Plug in

2 Convert

Migration



1

Plug in

First, check if database is compatible with CDB

1. Generate manifest file in non-CDB
2. Check compatibility in CDB

```
--In source, generate manifest file  
--You can also generate a manifest file of a remote database using a db link  
--You can generate a manifest file on a running database  
  
exec dbms_pdb.describe('/tmp/DB19.xml');
```

Manifest File

What is a **manifest** file

- Data files
- Components
- Parameters
- Services
- Patch level
- Time zone
- ... and more

```
<?xml version="1.0" encoding="UTF-8"?>
<PDB>
  <xmlversion>1</xmlversion>
  <pdbname>DB12</pdbname>
  <cid>0</cid>
  <byteorder>l</byteorder>
  <vsn>203424000</vsn>
  <vsns>
    <vsnnum>12.2.0.1.0</vsnnum>
    <cdbcompt>12.2.0.0.0</cdbcompt>
    <pdbscompt>12.2.0.0.0</pdbscompt>
    <vsnlibnum>0.0.0.24</vsnlibnum>
    <vsnsql>24</vsnsql>
    <vsnbsv>8.0.0.0.0</vsnbsv>
  </vsns>
  <dbid>1852833295</dbid>
  <ncdb2pdb>1</ncdb2pdb>
  <cdbid>1852833295</cdbid>
  <guid>86050C2587337002E0532AB2A8C0A57C</guid>
  <uscnbas>4437941</uscnbas>
  <uscnwrp>0</uscnwrp>
  <undoscn>8</undoscn>
  <rdba>4194824</rdba>
  <tablespace>
    <name>SYSTEM</name>
    <type>0</type>
    <tsn>0</tsn>
    <status>1</status>
    <issft>0</issft>
    <isnft>0</isnft>
    <encts>0</encts>
    <flags>0</flags>
    <bmuunitsize>8</bmuunitsize>
    <file>
      <path>/u02/oradata/DB12/system01.dbf</path>
      <afn>1</afn>
      <rfn>1</rfn>
```

```
--In CDB, check compatibility
--If PDB name changes, add parameter to check_plug_compatibility

set serveroutput on
BEGIN
  IF dbms_pdb.check_plug_compatibility('/tmp/DB19.xml') THEN
    dbms_output.put_line('PDB compatible? ==> Yes');
  ELSE
    dbms_output.put_line('PDB compatible? ==> No');
  END IF;
END;
/
```

--Always check the details

```
select type, message
from   pdb_plug_in_violations
where  name='DB19' and status<>'RESOLVED';
```

TYPE	MESSAGE
ERROR	'19.9.0 Release_Update' is installed in the CDB but no release updates are installed in the PDB
ERROR	DBRU bundle patch 201020: Not installed in the CDB but installed in the PDB
ERROR	PDB's version does not match CDB's version: PDB's version 12.2.0.1.0. CDB's version 19.0.0.0.0.
WARNING	CDB parameter compatible mismatch: Previous '12.2.0' Current '19.0.0'
WARNING	PDB plugged in is a non-CDB, requires noncdb_to_pdb.sql be run.



Always check **PDB_PLUG_IN_VIOLATIONS** after any plug-in operation

- An error prevents the PDB from opening unrestricted



Migration



1

Plug in

Then, perform plug-in

1. Shut down non-CDB
2. Plug into CDB



Migration

1. Restart database in read-only mode

```
SQL> shutdown immediate  
SQL> startup mount  
SQL> alter database open read only;
```

2. Generate manifest file and shut down

```
SQL> exec dbms_pdb.describe('/tmp/DB19.xml');  
SQL> shutdown immediate;
```

3. In CDB, create PDB from manifest file

```
SQL> create pluggable database DB19 using '/tmp/DB19.xml';
```

Migration



1 Plug in

2 Convert

Migration



2

Convert

1. Complete conversion with `noncdb_to_pdb.sql`
2. Requires downtime, but you run it only once
3. Irreversible



Migration

1. Open PDB

```
SQL> alter pluggable database DB19 open;  
SQL> alter session set container=DB19;
```

2. Convert

```
SQL> @?/rdbms/admin/noncdb_to_pdb.sql
```

3. Restart PDB

```
SQL> alter pluggable database DB19 close;  
SQL> alter pluggable database DB19 open;
```

Migration

4. Check plug-in violations

```
SQL> select type, message  
      from  pdb_plug_in_violations  
     where name='DB19' and status<>'RESOLVED';
```

5. Ensure PDB is open READ WRITE and unrestricted

```
SQL> select open_mode, restricted from v$pdbs;
```

6. Configure PDB to auto-start

```
SQL> alter pluggable database DB19 save state;
```

Demo

Multitenant migration
19c non-CDB

[Watch on YouTube](#)



Downtime



*How much downtime
do we need?*

`noncdb_to_pdb.sql`

- Runtime varies, typically 10-20 minutes
- Depends on dictionary complexity
- Not database size

Pre-tasks and post-tasks

- Follow best practices
- Data Guard

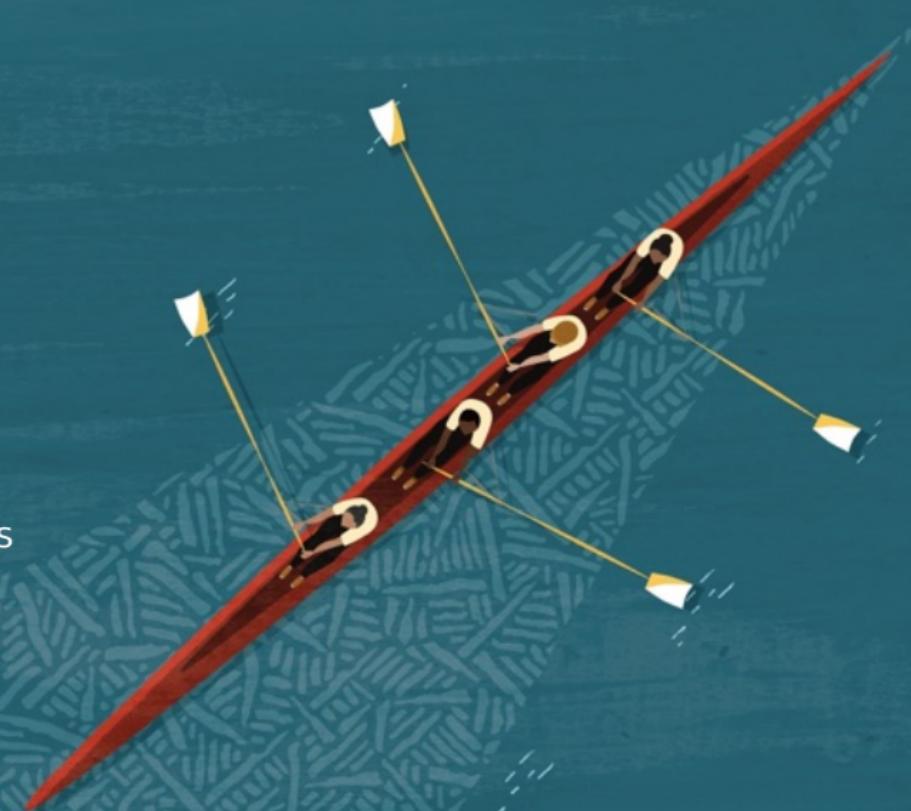
Expect total downtime 1-2 hours

- Simple databases even faster
- It's a migration, don't rush it

Migration

options

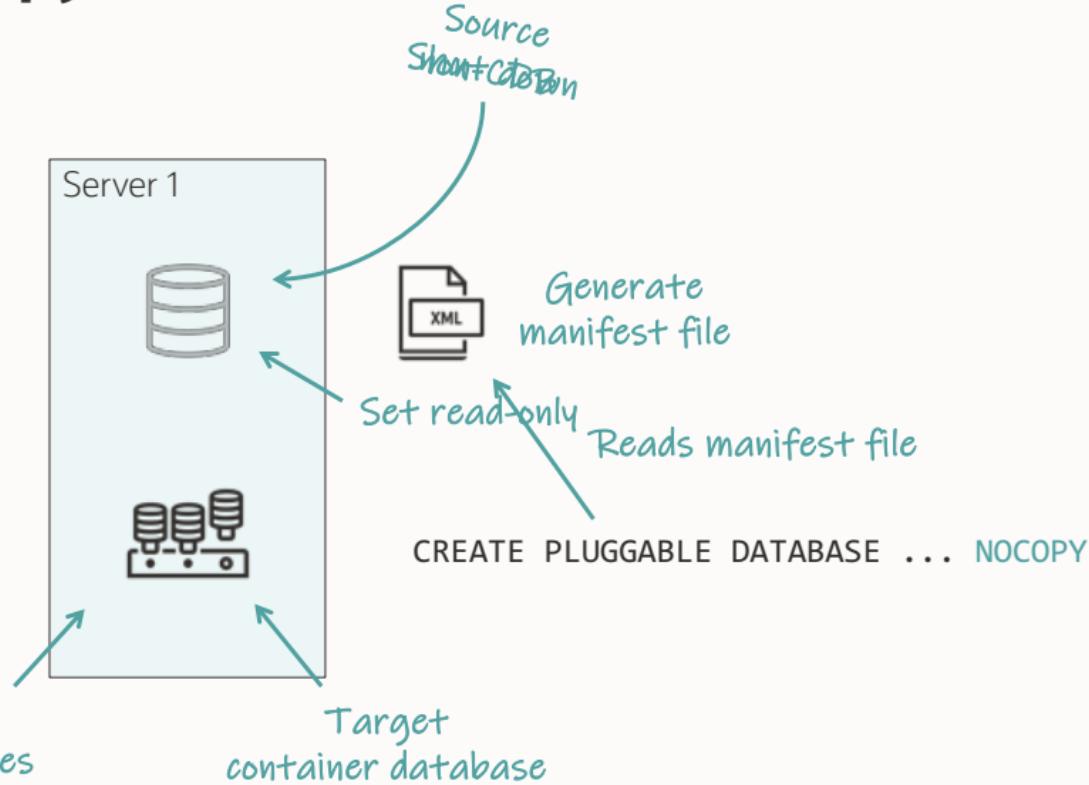
Other Options
Refreshable
Plug-in Copy
Plug-in NoCopy



Plug-in NoCopy

Re-using data files
No rollback!
system01.dbf
sysaux01.dbf
users01.dbf
undo01.dbf
...

Plug-in and
re-use data files



```
#AutoUpgrade config file
#Convert DB12 to PDB in CDB2 - re-use data files
upg1.source_home=/u01/app/oracle/product/19.22.0
upg1.target_home=/u01/app/oracle/product/19.23.0
upg1.sid=DB12
upg1.target_cdb=CDB2
```

```
#Fully automated migration
java -jar autoupgrade.jar -config db12.cfg -mode deploy
```

```
#AutoUpgrade config file
#Convert DB12 to PDB in CDB2 - re-use data files
#Upgrade from Oracle Database 19c to 23ai
upg1.source_home=/u01/app/oracle/product/19.22.0
upg1.target_home=/u01/app/oracle/product/23.4.0
upg1.sid=DB12
upg1.target_cdb=CDB2
```

Demo

Multitenant migration
Including upgrade to Oracle Database 23ai
Using AutoUpgrade

[Watch on YouTube](#)



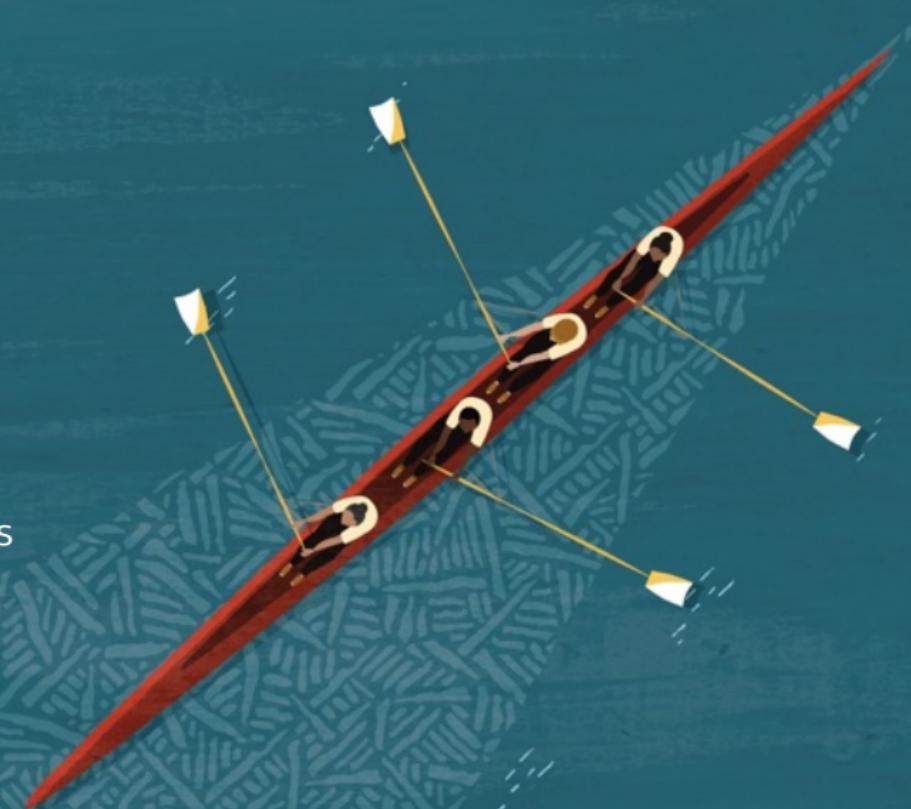
Consider using **MOVE** instead of **NOCOPY** to move files into proper directory

- Use **FILE_NAME_CONVERT** clause

Migration

options

Other Options
Refreshable
Plug-in Copy
Plug-in NoCopy



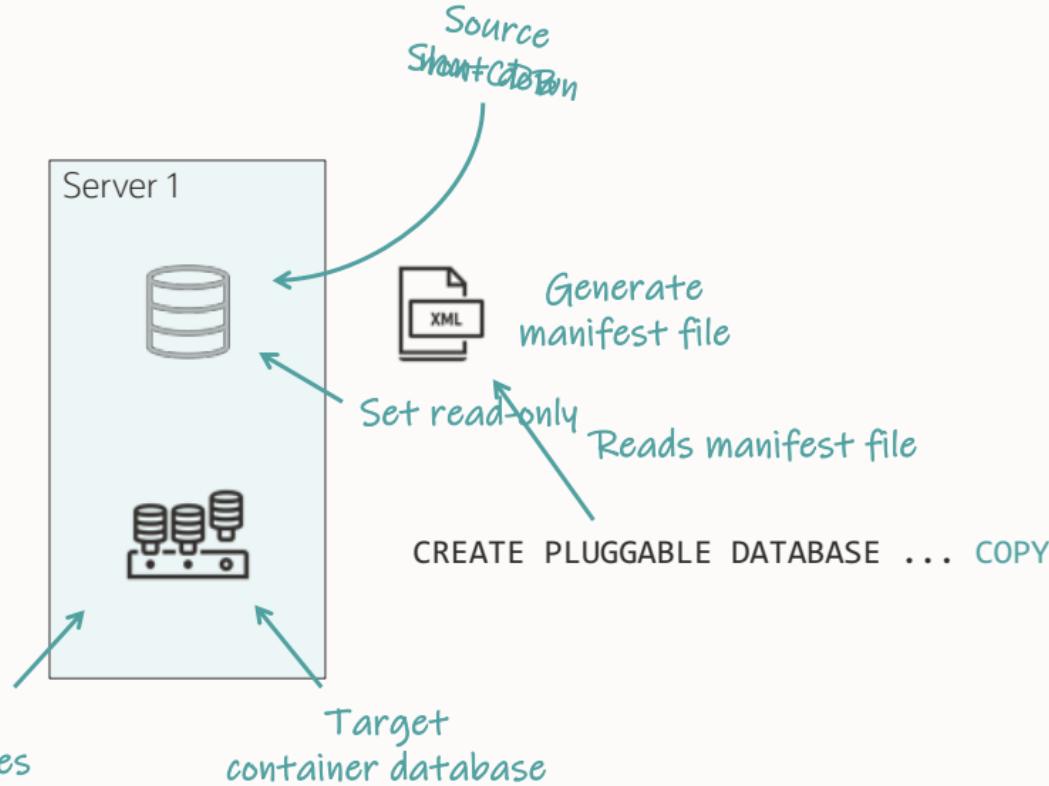
Plug-in Copy

Source preserved
for rollback

system01.dbf
sysaux01.dbf
users01.dbf
undo01.dbf
...

system01.dbf
sysaux01.dbf
users01.dbf
undo01.dbf
...

Plug-in and
copies data files



```
#AutoUpgrade config file
#Convert DB12 to PDB in CDB2 - copy data files
upg1.source_home=/u01/app/oracle/product/19.22.0
upg1.target_home=/u01/app/oracle/product/19.23.0
upg1.sid=DB12
upg1.target_cdb=CDB2
upg1.target_pdb_copy_option db12=file_name_convert('/u01', '/u02')

#Fully automated migration
java -jar autouupgrade.jar -config db12.cfg -mode deploy
```

```
#AutoUpgrade config file
#Convert DB12 to PDB in CDB2 - copy data files
#Generate OMF names - also for ASM
upg1.source_home=/u01/app/oracle/product/19.22.0
upg1.target_home=/u01/app/oracle/product/19.23.0
upg1.sid=DB12
upg1.target_cdb=CDB2
upg1.target_pdb_copy_option.db12=file_name_convert=NONE
```

Generate OMF names
Also for ASM



Be sure to shut down the source database.
Use *prefix.close_source=yes*

Migration

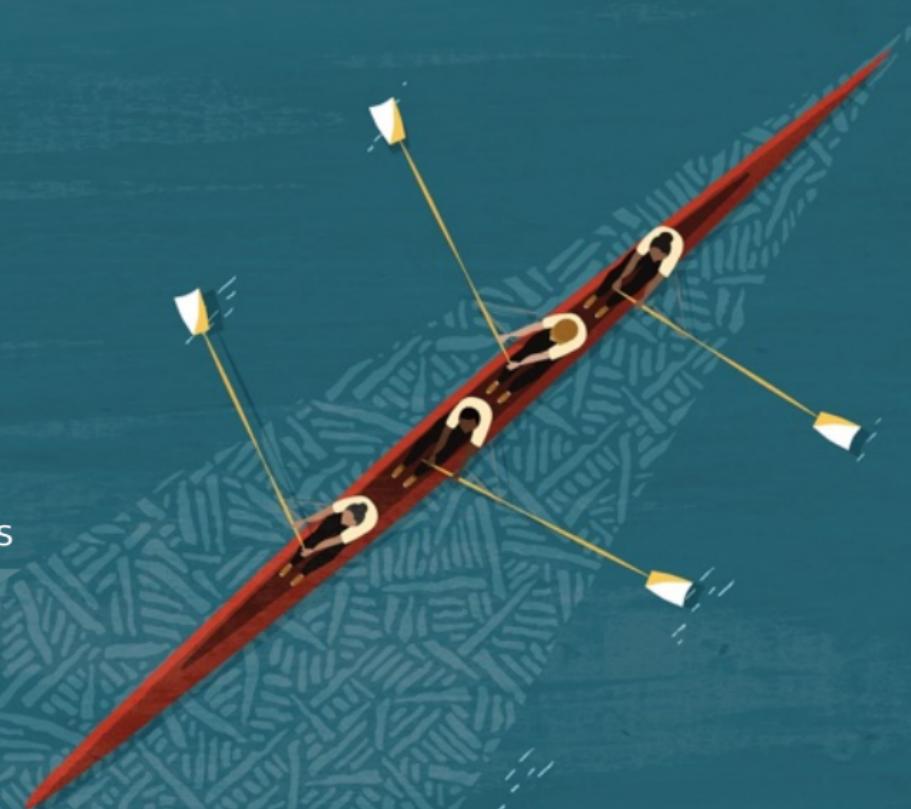
options

Other Options

Refreshable

Plug-in Copy

Plug-in NoCopy



Refreshable Clone



CREATE

Create PDB from non-CDB over a database link



REFRESH

Apply redo from non-CDB to keep PDB up-to-date



OUTAGE

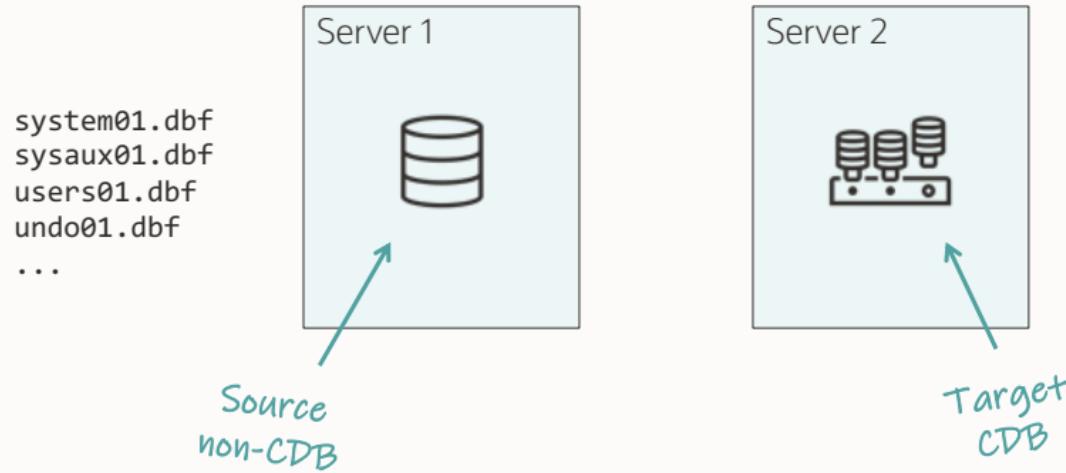
Disconnect users and refresh PDB for the last time



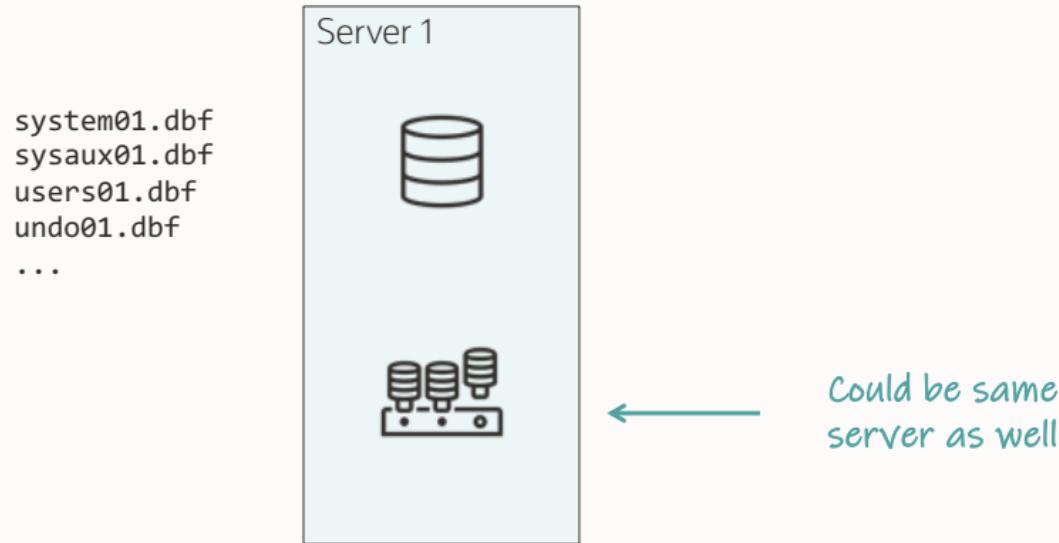
CONVERT

To become a proper PDB, it must be converted

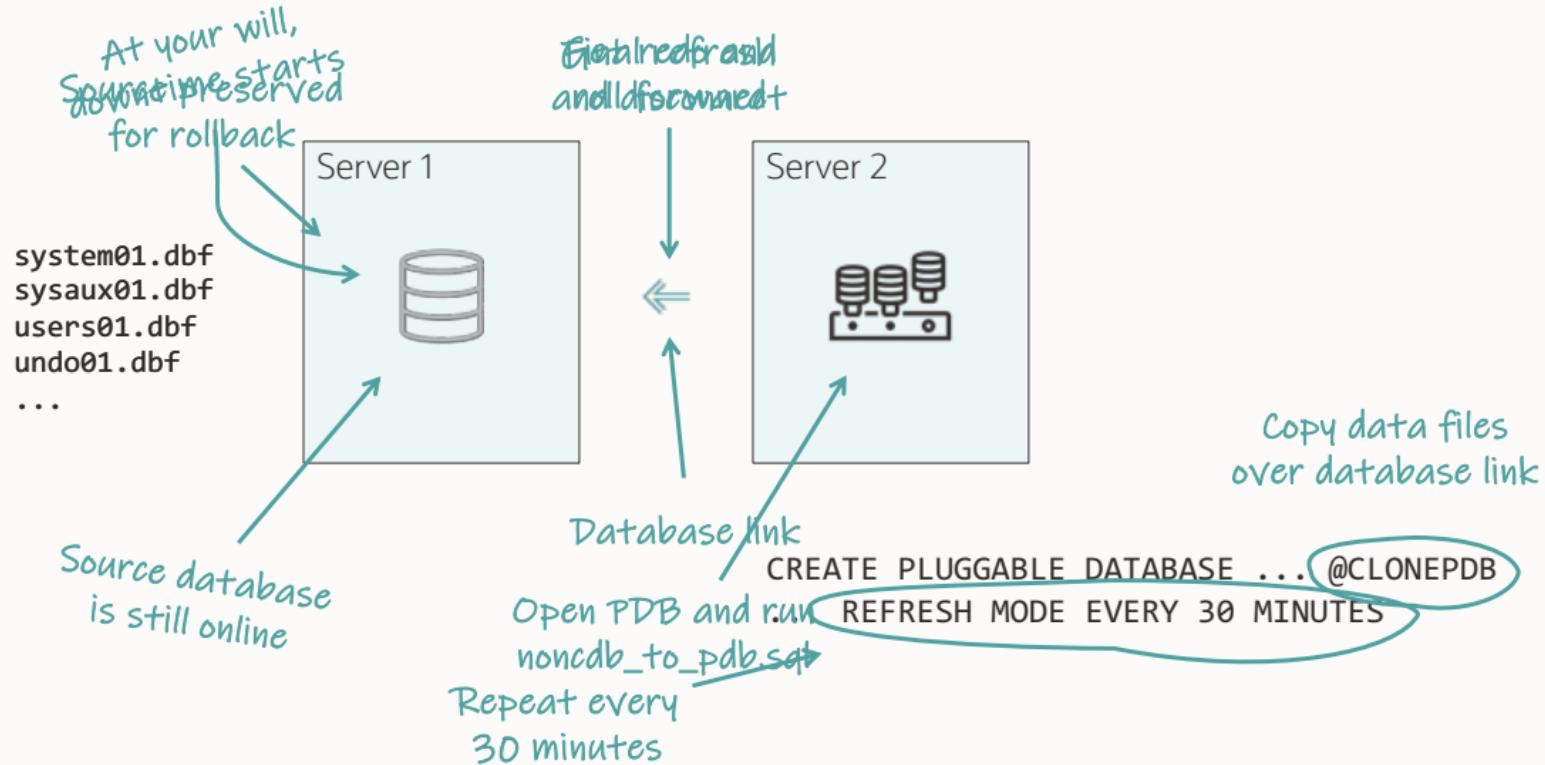
Refreshable Clone PDB



Refreshable Clone PDB



Refreshable Clone PDB



Refreshable Clone

Source non-CDB



Target CDB

```
CREATE USER dblinkuser  
IDENTIFIED BY ... ;
```

```
GRANT CREATE SESSION,  
CREATE PLUGGABLE DATABASE,  
SELECT_CATALOG_ROLE TO dblinkuser;  
  
GRANT READ ON sys.ENC$ TO dblinkuser;
```

```
CREATE DATABASE LINK CLONEPDB  
CONNECT TO dblinkuser  
IDENTIFIED BY ...  
USING 'noncdb-alias';
```

Refreshable Clone

Source non-CDB

Target CDB



```
upg1.source_home=/u01/app/oracle/product/19  
upg1.target_home=/u01/app/oracle/product/23  
upg1.sid=NONCDB1  
upg1.target_cdb=CDB1  
upg1.source_dblink.NONCDB1=CLONEPDB  
upg1.target_pdb_name.NONCDB1=PDB1  
  
--Specify relative start time  
--upg1.start_time=+1h30m
```

Refreshable Clone

Source non-CDB

Target CDB



```
upg1.source_home=/u01/app/oracle/product/19  
upg1.target_home=/u01/app/oracle/product/23  
upg1.sid=NONCDB1  
upg1.target_cdb=CDB1  
upg1.source_dblink.NONCDB1=CLONEPDB 300  
upg1.target_pdb_name.NONCDB1=PDB1  
  
--Specify relative start time  
--upg1.start_time=+1h30m
```

Refreshable Clone

Source non-CDB

Target CDB



```
upg1.source_home=/u01/app/oracle/product/19
upg1.target_home=/u01/app/oracle/product/23
upg1.sid=NONCDB1
upg1.target_cdb=CDB1
upg1.source_dblink.NONCDB1=CLONEPDB 300
upg1.target_pdb_name.NONCDB1=PDB1
upg1.start_time=19/05/2024 02:00:00
--Specify relative start time
--upg1.start_time=+1h30m
```



Refreshable Clone

1

Run on source

```
autoupgrade.jar ... -mode analyze
```

```
autoupgrade.jar ... -mode fixups
```

2

Run on target

```
autoupgrade.jar ... -mode deploy
```



Refreshable Clone



autoupgrade.jar ... -mode deploy

upg1.start_time=19/05/2024 02:00:00



The source non-CDB stays intact
to allow rollback

Demo

Multitenant migration
Upgrade to Oracle Database 19c
Using Refreshable Clone PDBs

[Watch on YouTube](#)



Refreshable clone works only with deferred recovery on standby database

- You must restore the PDB on standby database after disconnect from non-CDB



Ensure archive logs are available
on disk during migration

Cloning



CLONING

AutoUpgrade uses **CREATE PLUGGABLE DATABASE** statement with **PARALLEL** clause which clones the database using multiple parallel processes



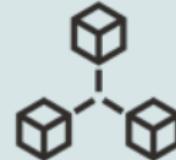
PARALLEL

Based on system resources and current utilization the database automatically determines a proper parallel degree



TRANSFER

A new file transfer protocol that can bypass several layers in the database to achieve very high transfer rates



NETWORK

Watch out for network saturation. Control parallelism in the AutoUpgrade config file



The database applies automatic parallelism based on available system resources

- Control using *prefix.parallel_pdb_creation_clause*

```
SQL> select message, sofar, totalwork, time_remaining as remain, elapsed_seconds as ela
      from v$session_longops
     where opname='kpdbfCopyTaskCbk' and sofar != totalwork;
```

MESSAGE	SOFAR	TOTALWORK	REMAIN	ELA
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 643199 out of 1310720 Blocks done	643199	1310720	134	129
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 443007 out of 1310720 Blocks done	443007	1310720	213	109
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 436351 out of 1310720 Blocks done	436351	1310720	216	108
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 370431 out of 1310720 Blocks done	370431	1310720	256	101



```
SQL> select message, sofar, totalwork,time_remaining as remain, elapsed_seconds as ela
      from v$session_longops
     where opname='kpdbfCopyTaskCbk' and sofar != totalwork;
```

MESSAGE	SOFAR	TOTALWORK	REMAIN	ELA
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 643199 out of 1310720 Blocks done	643199	1310720	134	129
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 443007 out of 1310720 Blocks done	443007	1310720	213	109
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 436351 out of 1310720 Blocks done	436351	1310720	216	108
kpdbfCopyTaskCbk: /u01/app/oracle/oradata/CDB2/EDA 3: 370431 out of 1310720 Blocks done	370431	1310720	256	101

```
SQL> select sql_text
      from v$sql s, v$session_longops l
     where s.sql_id=l.sql_id and l.opname='kpdbfCopyTaskCbk';
```

SQL_TEXT
/* SQL Analyze(256,0) */ SELECT /*+PARALLEL(4) NO_STATEMENT_QUEUEING */ /* FROM X\$KXFTASK /*kpdbfParallelCopyOrMove,PDB_FILE_COPY*/



Zürcher
Kantonalbank

Customer Case | Zürcher Kantonalbank

Customer

- Project
- Constraints
- Preparation
- Migration
- Success?
- Remarks

A reliable partner for over 150 years

- The bank for the people of Zurich since 1870
- With over 5'100 employees one of the largest employers in the canton of Zurich
- Globally networked full-service bank with strong regional and local roots



Customer Case | Zürcher Kantonalbank

Customer	Current situation
Project	<ul style="list-style-type: none">• Oracle databases on old OS and on Oracle Exadata• 2023:<ul style="list-style-type: none">• Migrate everything to Exadata until end of 2023• Consolidation to Multitenant and to the next long-term support release
Constraints	
Preparation	
Migration	
Success?	
Remarks	Planned solution: AutoUpgrade

Customer Case | Zürcher Kantonalbank

Customer

Project

Constraints

Preparation

Migration

Success?

Remarks

Test setup

- 3 non-CDB databases of different size

Source	Size / GB
TEST40 (108)	165
TEST42 (107)	555
TEST41 (106)	18'496

- Exadata X6-2 compute node
- 7 storage cells (2x X6-2L / 3x X7-2L / 2x X8-2L)
- Oracle Database 19.15.0
- No additional options

Customer Case | Zürcher Kantonalbank

Customer

Project

Constraints

Preparation

Migration

Success?

Remarks

Cloning user

```
create user dblinkuser identified by Oracle_4UOracle_4U;
```

Permissions

```
grant CONNECT, RESOURCE, CREATE PLUGGABLE DATABASE,  
      SELECT_CATALOG_ROLE to dblinkuser;  
grant ALL ON SYS.ENC$ to dblinkuser;
```

Database link

```
create database link TEST42.DOMAIN connect to dblinkuser  
identified by oracle_4uoracle_4u using 'test42.domain';
```

Customer Case | Zürcher Kantonalbank

Customer
Project
Constraints
Preparation
Migration

Success?

Remarks

Migration in progress

```
Job#|DB_NAME| STAGE|OPERATION| STATUS|START_TIME| UPDATED| MESSAGE|
-----+-----+-----+-----+-----+-----+-----+
106| [REDACTED]| CLONF0|EXECUTING| RUNNING| 09:50:39|!!!|17801s ago|Creating pluggable database|
107| [REDACTED]| STOPF0|FINISHED| STOPPED| 09:50:40| | |
108| [REDACTED]| STOPF0|FINISHED| 09:50:40| | |
Total job# 3

The command lsnr is running every 60 seconds. PRESS ENTER TO EXIT

Job#|DB_NAME| STAGE|OPERATION| STATUS|START_TIME| UPDATED| MESSAGE|
-----+-----+-----+-----+-----+-----+-----+
106| [REDACTED]| CLONF0|EXECUTING| RUNNING| 09:50:39|!!!|17461s ago|Creating pluggable database|
107| [REDACTED]| STOPF0|FINISHED| STOPPED| 09:50:40| | |
108| [REDACTED]| STOPF0|FINISHED| 09:50:40| | |
Total job# 3

The command lsnr is running every 60 seconds. PRESS ENTER TO EXIT

Job#|DB_NAME| STAGE|OPERATION| STATUS|START_TIME| UPDATED| MESSAGE|
-----+-----+-----+-----+-----+-----+-----+
106| [REDACTED]| CLONF0|EXECUTING| RUNNING| 09:50:39|!!!|17501s ago|Creating pluggable database|
107| [REDACTED]| STOPF0|FINISHED| STOPPED| 09:50:40| | |
108| [REDACTED]| STOPF0|FINISHED| 09:50:40| | |
Total job# 3

The command lsnr is running every 60 seconds. PRESS ENTER TO EXIT
```

Source	Runtime/Min
TEST40 (108)	26
TEST42 (107)	ongoing
TEST41 (106)	ongoing

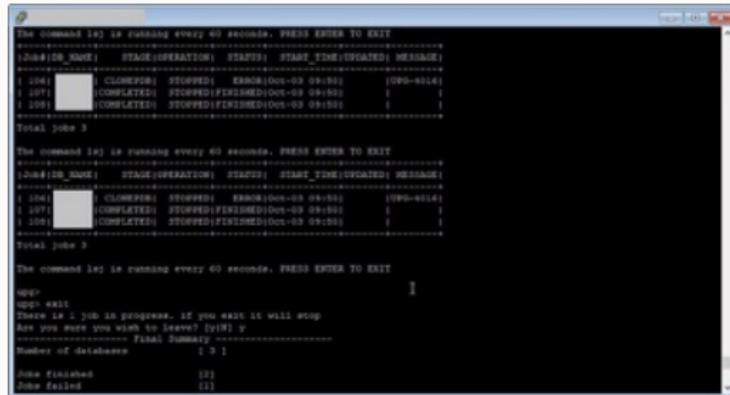
Customer Case | Zürcher Kantonalbank

Customer
Project
Constraints
Preparation
Migration

Success?

Remarks

Migration completed



The command line is running every 60 seconds. PRESS ENTER TO EXIT

ID	DB_NAME	STAGE	OPERATION	STATUS	START_TIME	UPDATED	MESSAGE
1096	[REDACTED]	CLOSED	STOPPED	EXCEPTION	Oct-03 09:50	[REDACTED]	[REDACTED]
1097	[REDACTED]	STOPPED	FINISHED	FINISHED	Oct-03 09:50	[REDACTED]	[REDACTED]
1098	[REDACTED]	STOPPED	FINISHED	FINISHED	Oct-03 09:50	[REDACTED]	[REDACTED]

Total jobs: 3

The command line is running every 60 seconds. PRESS ENTER TO EXIT

ID	DB_NAME	STAGE	OPERATION	STATUS	START_TIME	UPDATED	MESSAGE
1096	[REDACTED]	CLOSED	STOPPED	EXCEPTION	Oct-03 09:50	[REDACTED]	[REDACTED]
1097	[REDACTED]	STOPPED	FINISHED	FINISHED	Oct-03 09:50	[REDACTED]	[REDACTED]
1098	[REDACTED]	STOPPED	FINISHED	FINISHED	Oct-03 09:50	[REDACTED]	[REDACTED]

There is 1 job in progress. If you exit it will stop.
Are you sure you wish to leave? [y/N] y

----- Final Summary -----

Number of databases	1 / 1
Jobs finished	1 / 1
Jobs failed	0 / 1

Source	Runtime/Min
TEST40 (108)	26
TEST42 (107)	226 (~3.5h)
TEST41 (106)	1770 (29h)

Customer Case | Zürcher Kantonalbank

Customer	First non-CDBs migrated successfully
Project	<ul style="list-style-type: none">• Project is ongoing
Constraints	
Preparation	
Migration	
Success?	
Remarks	



Customer Case | Zürcher Kantonalbank

Customer	For large databases, make sure archives aren't cleaned up
Project	<ul style="list-style-type: none">• Solution: restore archivelogs from backup
Constraints	
Preparation	User profile with IDLE_TIME lead to kill of the session
Migration	<ul style="list-style-type: none">• Solution: assign a different profile to the clone user
Success?	
Remarks	

Summary

- Very comfortable to use
 - Everything happens automatically
 - Does not require user interaction
- Simple syntax
- No license costs associated
- Perfect for pre-migration test
- Very Stable



Migration

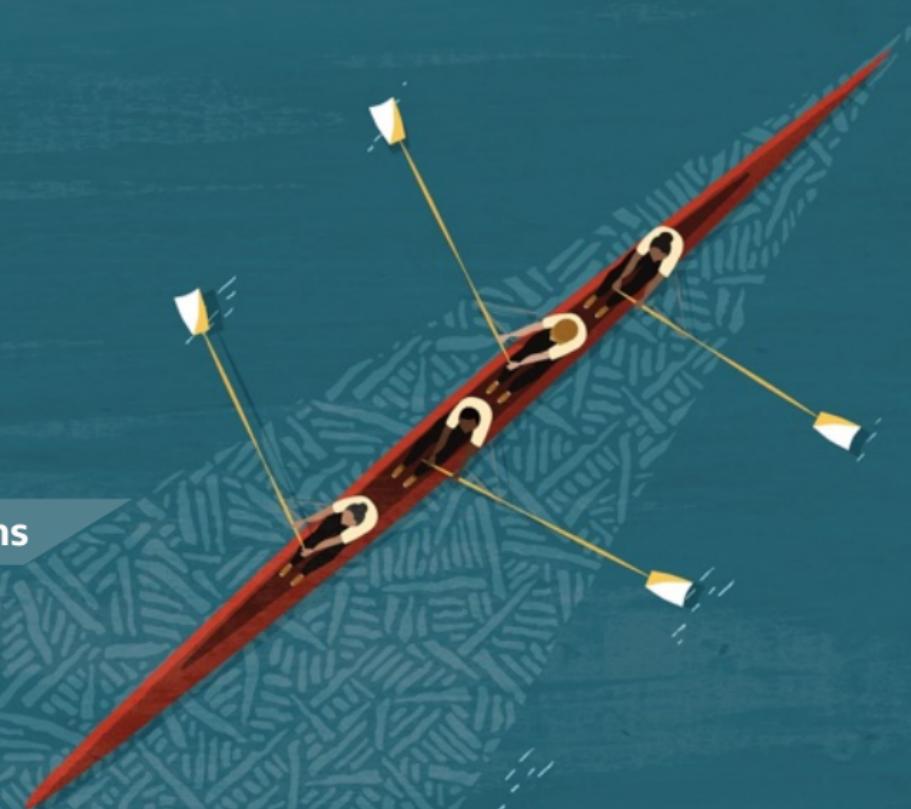
options

Other Options

Refreshable

Plug-in Copy

Plug-in NoCopy





Other Options

It is also possible to migrate using

1 Data Pump

- Well-known and proven method
- Extremely flexible
- Migrate from lower version
- Migrate from cross-Endian
- Preserves source database for fallback

2 Transportable Tablespaces

3 GoldenGate

Other Options

It is also possible to migrate using

1 Data Pump

- Faster for larger databases
- Migrate from lower version
- Migrate from cross-Endian
- Preserves source database for fallback

2 Transportable Tablespaces

3 GoldenGate



Other Options

It is also possible to migrate using

1 Data Pump

- Only zero downtime option
- Migrate from lower version
- Migrate from cross-Endian
- Preserves source database for fallback
- Active-active replication for ultimate solution

2 Transportable Tablespaces

3 GoldenGate



comparing **MIGRATION** options

	Plug-in NoCopy	Plug-in Copy	Refreshable Clone PDB	Data Pump	Transportable	GoldenGate
Downtime	Less	Considerable	Less	Considerable	Minimal	None
Rollback	No	Yes	Yes	Yes	Yes	Yes
Cross-platform (same Endian)	Yes	Yes	Yes	Yes	Yes	Yes
Cross-Endian	No	No	No	Yes	Yes	Yes
Cross-version	No	No	Yes	Yes	Yes	Yes
Complexity	Low	Low	Low	Medium	Medium	High



Best Practices for multitenant migration



Gather dictionary stats before migration

- Preferably also after migration



Perform a dictionary check
before migration

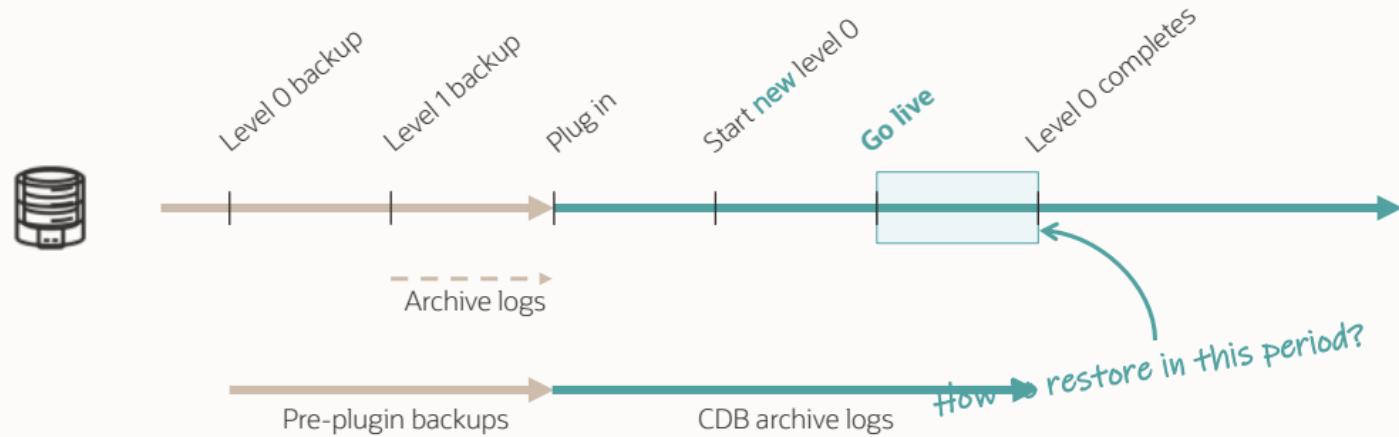
- Use `DBMS_DICTIONARY_CHECK`



Backup your database after migration

- Level 0

Recovery Strategy





Practice restore with pre-plugin backups

- Check `DBMS_PDB.EXPORTRMANBACKUP`



What about Oracle RAC Database

RAC

- Nothing special, it works seamlessly
- Connect to one instance and plug in
- No need to start with `CLUSTER_DATABASE=FALSE`



What happens during plug-in with Oracle Data Guard

Data Guard



1 Enabled recovery

*Plug-in on primary propagates
to standby database via redo*

2 Deferred recovery



Data Guard

1

Enabled recovery

```
create pluggable database ... standbys=all
```

Standby records PDB creation

Standby locates data files

MRP applies redo to PDB

PDB is immediately protected

2

Deferred recovery

Default



Data Guard

1

Enabled recovery

`create pluggable database ... standbys=all`

Standby records PDB creation

Standby locates data files

MRP applies redo to PDB

PDB is immediately protected

2

Deferred recovery

`create pluggable database ... standbys=none`

Standby records PDB creation

Standby ignores data files

MRP skips redo

PDB protected after restore



You can specify recovery mode
for each standby database

- CREATE PLUGGABLE DATABASE ... STANDBYS=STDBY1,STDBY3
- CREATE PLUGGABLE DATABASE ... STANDBYS=ALL EXCEPT STDBY2

--Check the recovery mode of each PDBs
--Possible values: ENABLED, DISABLED

```
select name, recovery_status from v$pdbs;
```



In AutoUpgrade, specify recovery mode using *prefix.manage_standbys_clause*

- Value is inserted directly into **STANDBYS** clause in **CREATE PLUGGABLE DATABASE** statement

Enabled Recovery

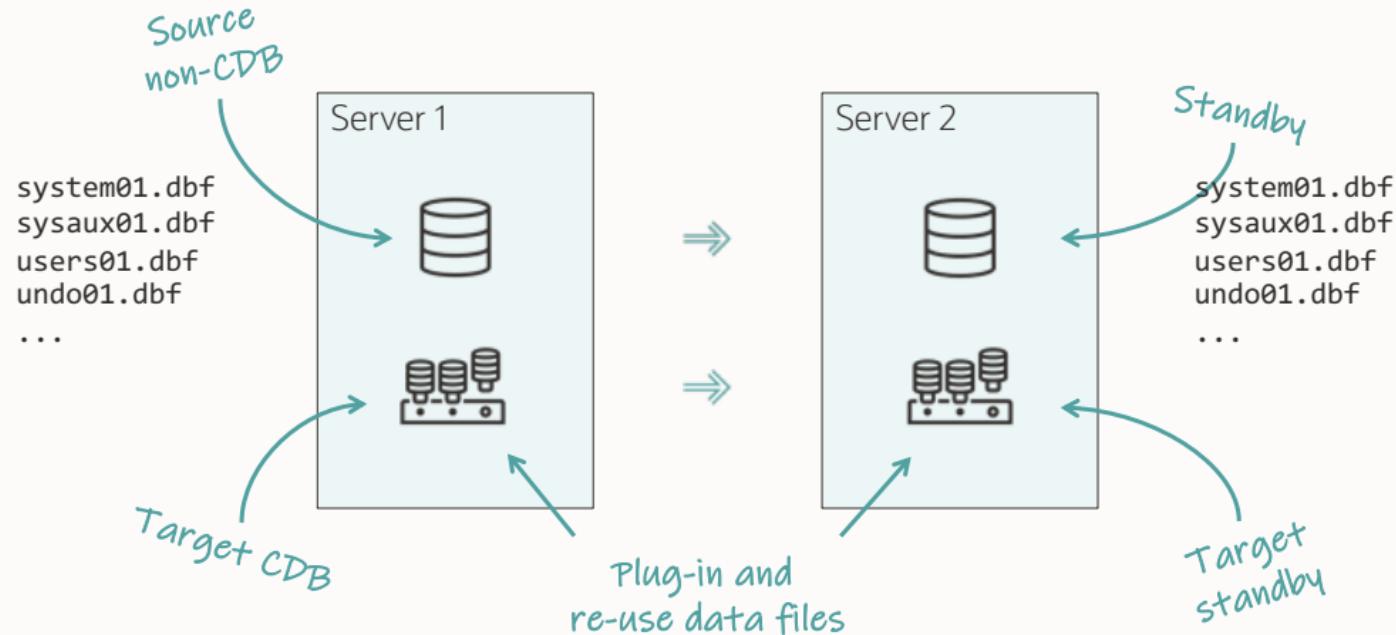
- A more complex approach
- Requires additional work before and during plug-in
- Standby database protects the PDB immediately after plug-in
- When you use STANDBYS=ALL or a list
- Default

Enabled Recovery

Enabled recovery applies to:

- All standby databases when using STANDBYS=ALL
- Or, those standby databases mention in STANDBYS=<list>
- Or, those not mentioned in STANDBYS=ALL EXCEPT <list>

Enabled Recovery





All data files on primary and standby must be at the same SCN

Reusing Data Files

Primary

```
SQL> shutdown immediate
SQL> startup mount
SQL> alter system
  flush redo to stdby no confirm apply;
SQL> alter database open read only;
SQL> select checkpoint_change#
  from v$datafile_header where file#=1;
```

MUST MATCH!

```
SQL> exec dbms_pdb.describe('/home/oracle/prmy.xml');
SQL> shutdown immediate
```

Standby

```
DGMGRL> edit database stdby set state='APPLY-OFF';
SQL> shutdown immediate
SQL> startup
SQL> alter database
  recover managed standby database cancel;
```

Replace with `checkpoint_change#`

```
SQL> alter database recover managed standby database
  until change 2319950;
SQL> select checkpoint_change#
  from v$datafile_header where file#=1;
```

```
SQL> shutdown immediate
```

Enabled Recovery

- The plug-in happens on the primary database
- The plug-in uses the manifest file
- The manifest file contains information on data files from the primary database only

**How does the standby database
know which files to plug in?**

Enabled Recovery

How does the standby database know which files to plug in?

- 1 Regular files
- 2 OMF in regular file system
- 3 ASM

Enabled Recovery

1 Regular files

- Standby search for data files at the same location as the primary
- Override with `DB_FILE_NAME_CONVERT`
- Or, override with `STANDBY_PDB_SOURCE_FILE_DIRECTORY`

Enabled Recovery

2 OMF in regular file system

- Standby search for data files at the OMF location (`DB_CREATE_FILE_DEST`)
- Move data files from non-CDB location into OMF location
- Or, create soft links in OMF location pointing to data file location

Enabled Recovery

3 ASM

- Standby search for data files at the OMF location (`DB_CREATE_FILE_DEST`)
- Now, it becomes a little tricky...

Enabled Recovery | ASM

23ai
Non-CDB
Primary



23ai
Non-CDB
Standby



```
SQL> select name from v$datafile;
```

```
NAME
```

```
-----  
+DATA/DB_BOSTON/DATAFILE/system.269.1103046537  
+DATA/DB_BOSTON/DATAFILE/sysaux.270.1103046537  
+DATA/DB_BOSTON/DATAFILE/users.273.1103046827
```



```
SQL> select name from v$datafile;
```

```
NAME
```

```
-----  
+DATA/DB_CHICAGO/DATAFILE/system.265.1103050007  
+DATA/DB_CHICAGO/DATAFILE/sysaux.266.1103050007  
+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009
```



Same file,
but different name

Enabled Recovery | ASM

23ai
Non-CDB
Primary



23ai
Non-CDB
Standby



The manifest file contains

```
SQL>exec dbms_pdb.set_manifest_database('mymanifest_DB.xml');
```

- Not standby database

⋮

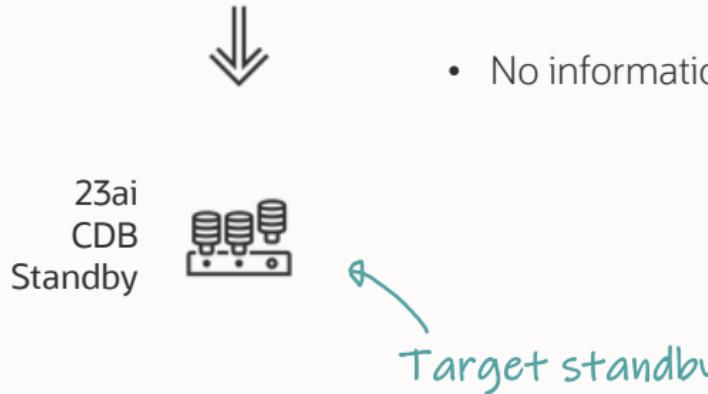
```
<?xml version="1.0" encoding="UTF-8"?>
<PDB>
  <xmlversion>1</xmlversion>
  <pdbname>PDB1</pdbname>
  ...
  <guid>DDB49CFEFD8ED4FCE053E801000A078C</guid>
  ...
  <tablespace>
    <name>USERS</name>
  ...
  <file>
    <path>+DATA/DB_BOSTON/DATAFILE/users.273.1103046827</path>
```

Enabled Recovery | ASM



```
SQL> create pluggable database PDB1 using '/tmp/manifest_DB.xml' ... ;
```

- Manifest file lists the location of data files on primary
- No information about standby databases



Enabled Recovery | ASM

23ai
CDB
Primary



+DATA/DB_BOSTON/DATAFILE/users.273.1103046827

Redo record says:
Plug in this data file



No good, data file
has a different name

23ai
CDB
Standby



+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009



Enabled Recovery | ASM

23ai
CDB
Primary



+DATA/DB_BOSTON/DATAFILE/users.273.1103046827



23ai
CDB
Standby



+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009

+DATA/CDB1_CHICAGO/<PDB_GUID>/DATAFILE

OK, let's check the OMF directory

It's empty

Enabled Recovery | ASM

23ai
CDB
Primary



+DATA/DB_BOSTON/DATAFILE/users.273.1103046827



23ai
CDB
Standby



+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009
+DATA/CDB1_CHICAGO/<PDB_GUID>/DATAFILE

OK, let's check the OMF directory

It's empty



I'll just move the file in ASM



There's no move command in ASM.
How about copying?



```
ASMCMD> cp users.269.1103050009  
+DATA/DB_CHICAGO/.../users.273.1103046827
```

ASMCMD-8016: copy source '+DATA/DB_BOSTON/.../users.269.1103050009' and target '+DATA/DB_CHICAGO/.../users.273.1103046827' failed

ORA-15056: additional error message

ORA-15046: **ASM file name 'users.273.1103046827' is not in single-file creation form**

ORA-06512: at "SYS.X\$DBMS_DISKGROUP", line 617

ORA-06512: at line 3 (DBD ERROR: OCISqlExecute)





Only a database can produce files
with ASM/OMF data file names



There's no **move** command in ASM.
But you can create *aliases*

- Similar to file system soft links



```
SQL> alter diskgroup data add alias  
  ' +DATA/DB_CHICAGO/DATAFILE/users.269.1103050009 '  
for  
  ' +DATA/CDB1_CHICAGO/<PDB_GUID>/DATAFILE/users.269.1103050009 ':
```

Non-CDB standby data file



```
SQL> alter diskgroup data add alias  
  '+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009'  
for  
  '+DATA/CDB1_CHICAGO/<PDB_GUID>/DATAFILE/users.269.1103050009':
```

```
SQL> alter diskgroup data add alias  
  '+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009'  
for  
  '+DATA/CDB1_CHICAGO/<PDB_GUID>/DATAFILE/users.269.1103050009':
```



Standby PDB OMF location

```
SQL> alter diskgroup data add alias  
  '+DATA/DB_CHICAGO/DATAFILE/users.269.1103050009'  
for  
  '+DATA/CDB1_CHICAGO/<PDB_GUID>/DATAFILE/users.269.1103050009':
```



Name does not matter.
Standby scans all files in OMF directory

Data Guard | Re-use Data Files

23ai
CDB
Primary



- Standby database scans its own OMF directory for data files
- Standby ignores file names and look at file headers
- Standby will find aliases and find the real file locations

23ai
CDB
Standby



Data Guard | Re-use Data Files

23ai
CDB
Primary



23ai
CDB
Standby



Looking for file like on primary

```
Recovery scanning directory +DATA/DB_BOSTON/... for any matching files
Deleted Oracle managed file +DATA/DB_BOSTON/...
Successfully added datafile 37 to media recovery
Datafile #37: +DATA/DB_CHICAGO/DATAFILE/users.269.1103050009
```

Follows alias and finds the real file

Demo

Multitenant migration
19c non-CDB to 23ai
With Data Guard and re-using data files

[Watch on YouTube](#)



Move data files into proper OMF location
and get rid of aliases

Data File Location

23ai
CDB
Primary



23ai
CDB
Standby



```
SQL> select file#, name from v$datafile;  
  
FILE# NAME  
-----  
14    +DATA/DB_BOSTON/DATAFILE/system.269.1103046537  
15    +DATA/DB_BOSTON/DATAFILE/sysaux.270.1103046537  
16    +DATA/DB_BOSTON/DATAFILE/users.273.1103046827
```



Non-CDB OMF location



The database does not care.
But humans might care

- A database can use files from a non-OMF location

Online Data File Move

23ai
CDB
Primary



23ai
CDB
Standby



```
SQL> ALTER DATABASE MOVE DATAFILE 14;  
  
SQL> select file#, name from v$datafile;  
  
FILE# NAME  
-----  
14    +DATA/CDB_BOSTON/<PDB GUID>/DATAFILE/system.269.1103046537  
15    +DATA/DB_BOSTON/DATAFILE/sysaux.270.1103046537  
16    +DATA/DB_BOSTON/DATAFILE/users.273.1103046827
```

- The database copies the file, then drops the alias and original file
- Requires additional disk space
- Online operation

Online Data File Move

23ai
CDB
Primary



23ai
CDB
Standby



- On standby, online datafile move works only when CDB and PDB are open
- Stop redo apply before opening, unless you have a license for Active Data Guard
- Requires time and disk space to hold a copy of the data file
- Removes ASM alias and original file upon completion

```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-off';
```



```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-off';
```

```
SQL> alter database open read only;
SQL> alter pluggable database PDB1 open read only;
SQL> alter session set container=PDB1;
SQL> alter database move datafile <file#>;
SQL> alter database move datafile <file#>;
SQL> ...
SQL> alter database move datafile <file#>;
```

```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-off';
```

```
SQL> alter database open read only;
SQL> alter pluggable database PDB1 open read only;
SQL> alter session set container=PDB1;
SQL> alter database move datafile <file#>;
SQL> alter database move datafile <file#>;
SQL> ...
SQL> alter database move datafile <file#>;
```

```
srvctl stop database -d $ORACLE_UNQNAME -o immediate
srvctl start database -d $ORACLE_UNQNAME -o mount
```



```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-off';
```

```
SQL> alter database open read only;
SQL> alter pluggable database PDB1 open read only;
SQL> alter session set container=PDB1;
SQL> alter database move datafile <file#>;
SQL> alter database move datafile <file#>;
SQL> ...
SQL> alter database move datafile <file#>;
```

```
srvctl stop database -d $ORACLE_UNQNAME -o immediate
srvctl start database -d $ORACLE_UNQNAME -o mount
```

```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-on';
```



```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-off';
```

```
SQL> alter database open read only;
```

```
SQL> alter pluggable database PDB1 open read only;
```

```
SQL> alter session set container=PDB1;
```

```
SQL> alter database move datafile <file#>;
```

```
SQL> alter database move datafile <file#>;
```

```
SQL> ...
```

```
SQL> alter database move datafile <file#>;
```

Just move data files,
if Active Data Guard

```
srvctl stop database -d $ORACLE_UNQNAME -o immediate
```

```
srvctl start database -d $ORACLE_UNQNAME -o mount
```

```
DGMGRL> edit database <STANDBY_UNQNAME> set state='apply-on';
```



While relocating data files,
apply lag increases if redo apply is off

- Redo transport is still active
- Switchover/failover time increases



What happens with enabled recovery
if the standby fails to find the data files?

Enabled Recovery | Missing Data Files

What if a standby database fails to find data files?

- If Active Data Guard and PDB Recovery Isolation is turned on
 - New feature in Oracle Database 21c
 - Recovery disabled for PDB
 - Recovery proceeds in the entire CDB, except in specific PDB
 - Standby automatically restores data files from primary and re-enables recovery afterward
 - PDB protected after auto-restore
- If not, recovery halts in the **entire** CDB
 - **This is a critical situation**



What about AutoUpgrade
and enabled recovery?

Enabled Recovery | AutoUpgrade

The current version (24.1) does not support plugging in with enabled recovery

- Enabled recovery requires work on both primary and standby hosts
- You must execute commands at specific times
- It's complicated - but we're working on it



What about **AS CLONE** clause

- CREATE PLUGGABLE DATABASE ... AS CLONE

Enabled Recovery | As Clone

- When you plug in a non-CDB the GUID doesn't change
 - Get GUID from the manifest file or `V$CONTAINERS`
 - GUID only changes when you use `AS CLONE` keyword
- Impossible to re-use standby data files when
 - Using OMF and `CREATE PLUGGABLE DATABASE ... AS CLONE`
 - Regardless of whether you use regular file system or ASM
 - Only works with regular files and non-OMF

Enabled Recovery | As Clone

- Regular file system and non-OMF
 - Put data files at the same location as primary data files
 - Take `FILE_NAME_CONVERT` into account (`CREATE PLUGGABLE DATABASE`)
 - Adjust according to `DB_FILE_NAME_CONVERT`

Data Guard | Enabled Recovery

[Reusing the Source Standby Database Files When Plugging a PDB into the Primary Database of a Data Guard Configuration \(Doc ID 2273829.1\)](#)

★ **Reusing the Source Standby Database Files When Plugging a PDB into the Primary Database of a Data Guard Configuration (Doc ID 2273829.1)**

In this Document

[Goal](#)
[Solution](#)
[Prerequisites](#)
[Steps](#)
[Resolving Errors](#)
[References](#)

APPLIES TO:

Oracle Database Cloud Service - Version N/A and later
Oracle Database Exadata Express Cloud Service - Version N/A and later
Oracle Database - Enterprise Edition - Version 12.1.0.2 and later
Oracle Database Cloud Schema Service - Version N/A and later
Gen 1 Exadata Cloud at Customer (Oracle Exadata Database Cloud Machine) - Version N/A and later
Information in this document applies to any platform.

GOAL

To plug in an existing 12.1.0.2 or later PDB residing in a CDB as part of a Data Guard configuration into another CDB that is part of a different Data Guard configuration where the current Primary CDB and the target CDB both have standby databases and allow you to use the original Standby database's data files to update the destination CDB's Standby.

This note describes a multitenant migration option for maintaining standby databases when the source database is a PDB. If your source database is a non-CDB, please see [Document 2273304.1](#).

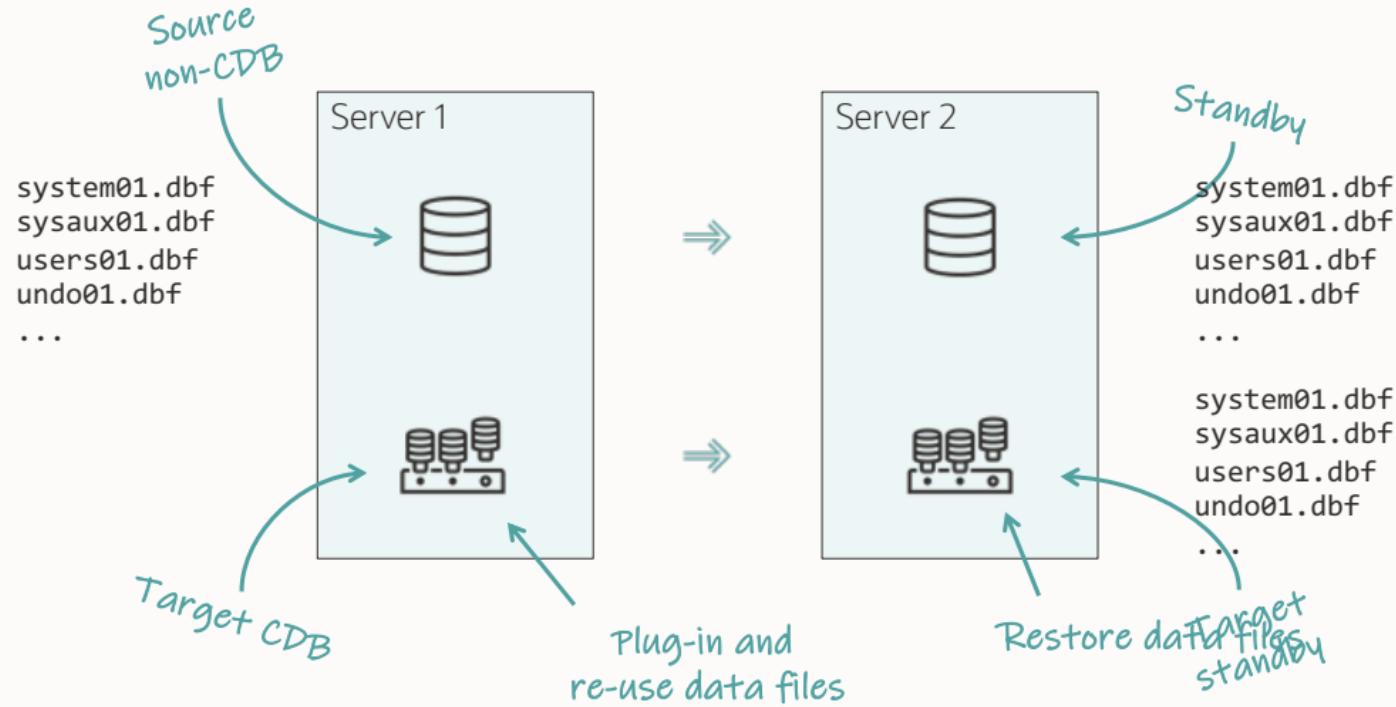
For Oracle RDBMS 19.15 and later, the Data Guard broker MIGRATE command has been enhanced to execute the steps contained in this document. It will manage configurations of the destination CDB containing a single physical standby database and will handle TDE enabled databases. Please see High Availability Overview and Best Practices - PDB Switchover and Failover in a Multitenant Configuration for more information on this feature.

Always test the steps in a dev/test environment prior to using in production. Since the original files are being modified directly by the plugin on the primary and by the consumption

Deferred Recovery

- The simplest approach
- Requires additional work after plug-in
- You must restore the PDB and re-enabled recovery
- Standby database protects the PDB after restore
- When you use STANDBYS=NONE

Deferred Recovery

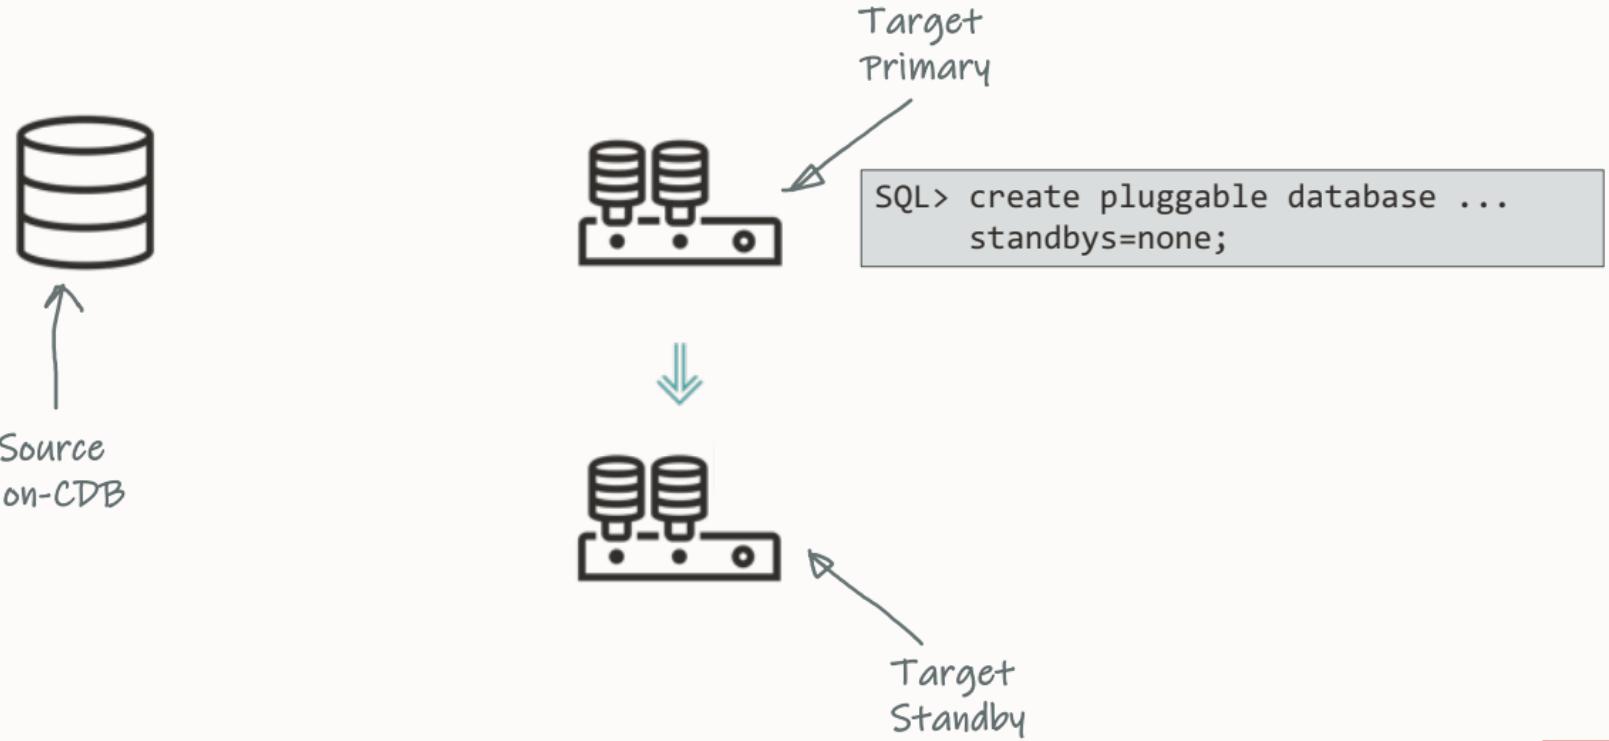


Deferred Recovery

Deferred recovery applies to:

- All standby databases when using STANDBYS=NONE
- Or, those standby databases mention in STANDBYS=ALL EXCEPT <list>
- Or, those not mentioned in STANDBYS=<list>

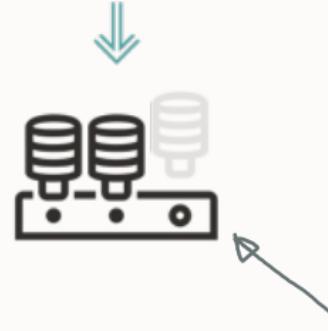
Deferred Recovery



Deferred Recovery



```
SQL> create pluggable database ...  
standbys=none;
```



PDB created
Data files missing

Deferred Recovery



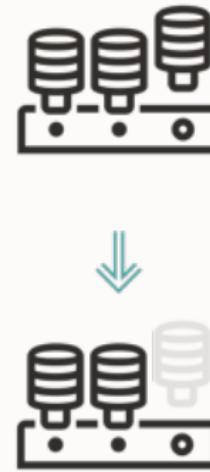
```
SQL> show pdbs
```

CON_NAME	OPEN MODE
PDB1	READ WRITE

```
SQL> show pdbs
```

CON_NAME	OPEN MODE
PDB1	MOUNTED

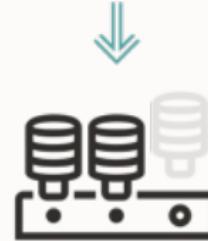
Deferred Recovery



```
SQL> select name, recovery_status  
  from v$pdbs;
```

NAME	RECOVERY_STATUS
PDB1	DISABLED

Deferred Recovery



```
RMAN> restore pluggable database  
... from service ... ;  
  
SQL> alter pluggable database  
enable recovery;  
SQL> alter database datafile  
... online;
```

Deferred Recovery



```
RMAN> restore pluggable database  
... from service ... ;  
  
SQL> alter pluggable database  
enable recovery;  
SQL> alter database datafile  
... online;
```

- Automated process in Oracle Database 21c
- PDB Recovery Isolation
- Requires Active Data Guard

Data Guard | Deferred Recovery

[Making Use Deferred PDB Recovery and the STANDBYS=NONE Feature with Oracle Multitenant \(Doc ID 1916648.1\)](#)

★ Making Use Deferred PDB Recovery and the STANDBYS=NONE Feature with Oracle Multitenant (Doc ID 1916648.1)

In this Document

[Goal](#)

[Solution](#)

[Creating a PDB with the STANDBYS=NONE clause in a Data Guard configuration with 1 physical standby](#)

[Showing how the cloned PDB will appear in certain tables and views on the physical standby](#)

[Performing a Data Guard Role Transition with a PDB in DISABLED RECOVERY](#)

[The zero downtime instantiation process using RMAN for copying the files from the primary to standby](#)

[Steps required for enabling recovery on the PDB after the files have been copied](#)

[Steps to DISABLE RECOVERY of a Pluggable Database](#)

[Conclusion](#)

[References](#)

APPLIES TO:

Oracle Cloud Infrastructure - Database Service - Version N/A and later

Oracle Database Cloud Service - Version N/A and later

Oracle Database - Enterprise Edition - Version 12.1.0.2 and later

Oracle Database Cloud Schema Service - Version N/A and later

Oracle Database Exadata Express Cloud Service - Version N/A and later

Information in this document applies to any platform.

Data Guard | Additional Information

Data Guard Impact on Oracle Multitenant Environments (Doc ID 2049127.1)

The physical standby database and redo apply will normally expect a new PDB's datafiles to have been pre-copied to the standby site and be in such a state that redo received from the primary database can be immediately applied. The standby database ignores any file name conversion specification on the CREATE PLUGGABLE DATABASE statement and relies solely on the standby database's initialization parameter settings for DB_CREATE_FILE_DEST and DB_FILE_NAME_CONVERT for locations and file naming.

For these cases, Oracle recommends deferring recovery of the PDB using the STANDBYS=NONE clause on the CREATE PLUGGABLE DATABASE statement. Recovery of the PDB can be enabled at some point in the future once the PDB's data files have been copied from the primary database to the standby database in a manner similar to that documented in Document 1916648.1.



Don't jeopardize your Data Guard

- Test the procedure and verify before go-live



How a customer handled the migration

- Customer case



Customer Case



Customer

Project

Result

Learnings

- **Swisscom** - Switzerland's leading telco
- One of the leading IT companies in Switzerland
- One of Switzerland's most sustainable and innovative companies



Customer Case





Customer Case



Customer

Project

Result

Learnings

Oracle Siebel CRM



Non-CDB

- Database size: 18 TB
- Release: 19.17.0
- Average active users: 3000



Single-tenant architecture



Customer Case



Customer

Project

Result

Learnings

Oracle Siebel CRM



Non-CDB

- 6,000 tables
- 29,000 indexes
- Partitioning
- LOBs
- 51 bigfile tablespaces



Single-tenant architecture

Customer Case



Data Guard with **five** standby databases



Each database is a 4-node **RAC** database



Running on **Exadata** Database Machine



Streaming data to microservices using **GoldenGate**



After consulting the business,
a plan was made for the standby databases

Customer Case



Standby 1: DR



Re-use data files, ASM alias



Standby 2-3: Auxiliary DR



Restore data files



Standby 4-5: Reporting



Restore data files



Customer Case



Customer

Project

Result

Learnings

- Total downtime was 3 hours 30 minutes
 - Planned maintenance window was 4 hours
 - Most time spent on application work and GoldenGate configuration
- Database migration
 - `noncdb_to_pdb.sql`: 18 minutes



Customer Case



Customer

Project

Result

Learnings

- 1** Test, test, and test
- 2** Create a detailed runbook
- 3** Remove complexity from the project to the extent possible
- 4** Work together and double-check all actions



Converting on Exadata Database Service and Exadata Cloud@Customer

Exadata Database Service

1. Use tooling to create a new CDB - or use an existing one
2. Plug in and convert using a method of choice
3. Tooling automatically adapts the new PDB after a while



Converting Oracle E-Business Suite

E-Business Suite

- Oracle E-Business Suite Release 12.2 and 12.1.3 support multitenant architecture
- But only in single-tenant architecture
- Useful MOS notes
 - FAQ: Oracle E-Business Suite and the Oracle Multitenant Architecture (Doc ID [2567105.1](#))
 - Interoperability Notes: Oracle E-Business Suite Release 12.2 with Oracle Database 19c (Doc ID [2552181.1](#))
 - Getting Started with Oracle E-Business Suite on Oracle Cloud Infrastructure (Doc ID [2517025.1](#))
 - Using Oracle 19c Oracle RAC Multitenant (Single PDB) with Oracle E-Business Suite Release 12.2 (Doc ID [2530665.1](#))



Plugging in copies of the same database

As Clone

Each PDB has a unique GUID

- Check V\$CONTAINERS

If you plug in the same database multiple times,
there are conflicting GUIDs

Use CREATE PLUGGABLE DATABASE ... AS CLONE to generate new GUIDs on
plug-in



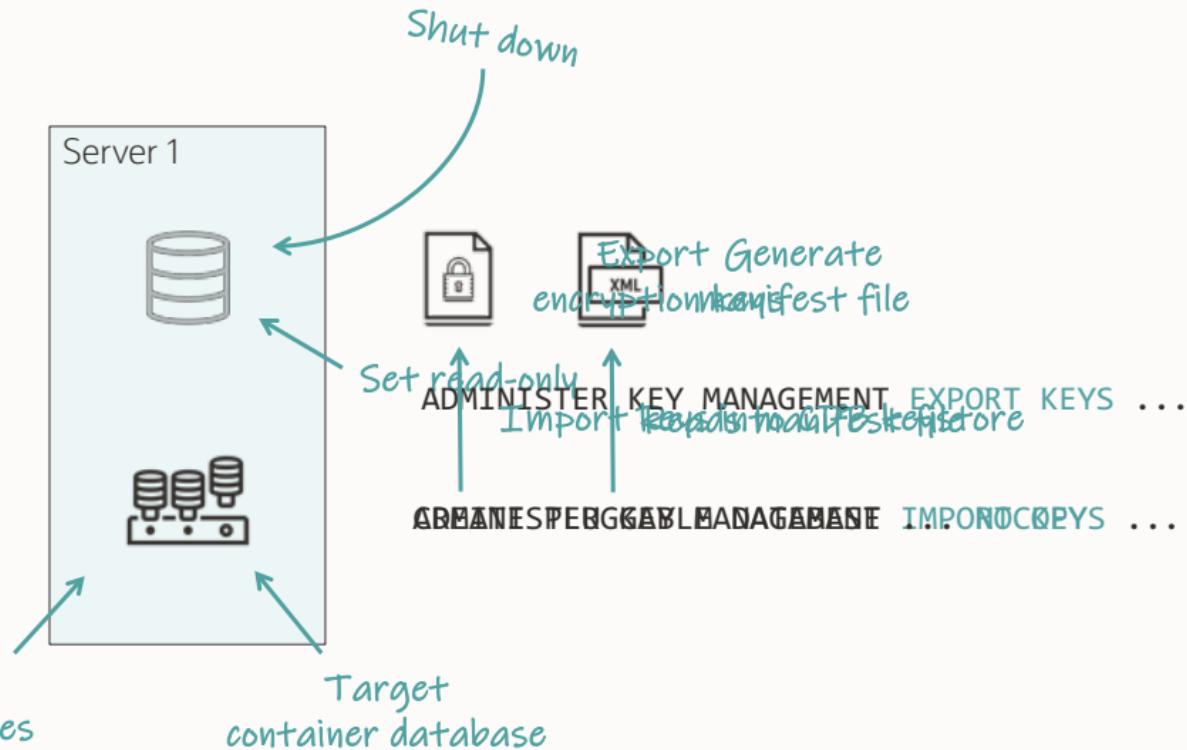
Migrating encrypted databases

- TDE Tablespace Encryption

Encrypted Database

Encrypted database

 system01.dbf
sysaux01.dbf
users01.dbf
undo01.dbf
...





AutoUpgrade fully supports
encrypted databases

Encrypted Database

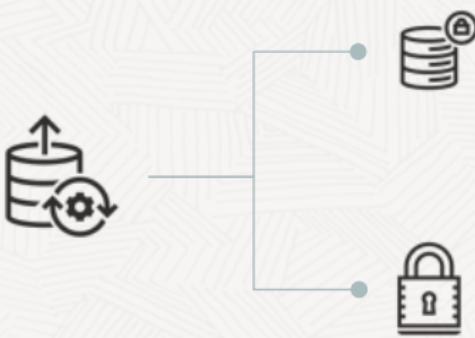
Certain database operations require passwords or secrets

```
CREATE PLUGGABLE DATABASE ... KEYSTORE IDENTIFIED BY <password>
```

```
ALTER PLUGGABLE DATABASE ... UNPLUG INTO ... ENCRYPT USING <secret>
```

```
CREATE PLUGGABLE DATABASE ... DECRYPT USING <secret>
```

```
ADMINISTER KEY MANAGEMENT ... KEYSTORE IDENTIFIED BY <password>
```



Secure External Password Store

Operator stores database keystore password
in a Secure External Password Store

AutoUpgrade Keystore

Operator loads database keystore password
into AutoUpgrade keystore ahead of upgrade

```
# Configure AutoUpgrade to work on encrypted databases
# Specify path for AutoUpgrade keystore

global.keystore=/etc/oracle keystores/autouupgrade/DB12
global.autoapg_log_dir=/u01/app/oracle/cfgtoollogs/autouupgrade
upg1.source_home=/u01/app/oracle/product/12.2.0.1
upg1.target_home=/u01/app/oracle/product/19
upg1.sid=DB12
```

Encrypted Database

Analyze the database for upgrade readiness

```
$ java -jar autoupgrade.jar -config PDB1.cfg -mode analyze
```

Summary report will show which keystore passwords are needed

REQUIRED ACTIONS

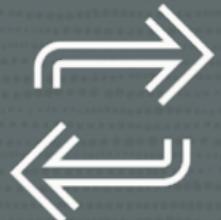
1. Perform the specified action ...

ORACLE_SID	Action Required
CDB1	Add TDE password
CDB2	Add TDE password

Demo

Multitenant migration
Including upgrade to Oracle Database 19c
Using AutoUpgrade

[Watch on YouTube](#)



In the unlikely event of ...

- Rollback and fallback options



PDB conversion is irreversible

- Not even Flashback Database can help

Rollback Options | Before Go-Live

1 Leave a copy



2 RMAN Restore

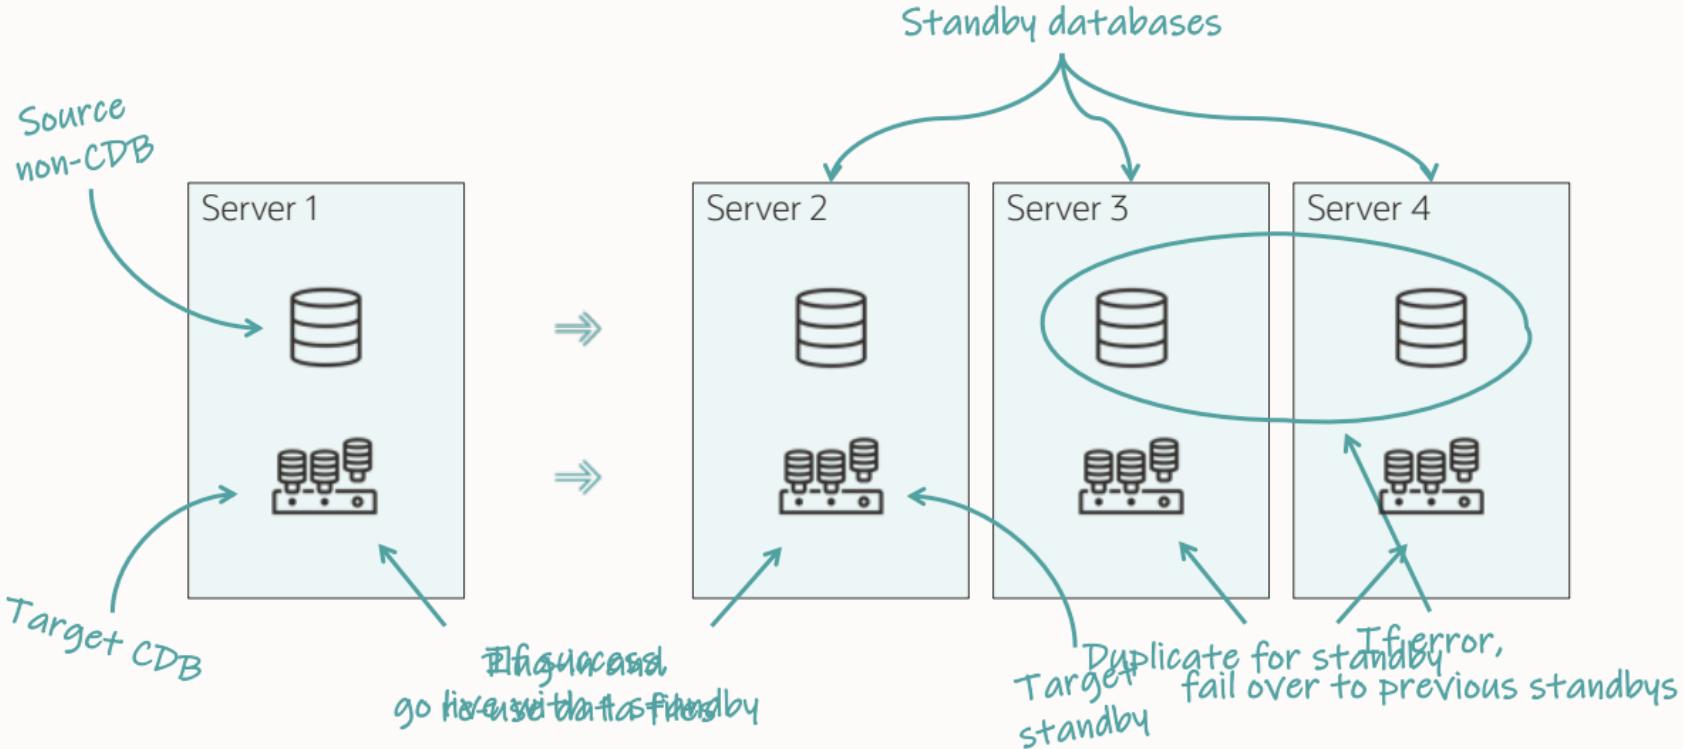
Rollback Options | Before Go-Live



1 Leave a copy

- CREATE PLUGGABLE DATABASE ... COPY
- Refreshable clone PDBs
- Incrementally roll forward data files copies
- Leave a standby database behind

Rollback Options | Before Go-Live



Rollback Options | Before Go-Live



2 RMAN Restore

- Time-consuming
- May not satisfy business requirements

Fallback Options | After Go-Live



1

Back to non-CDB

- Data Pump
- Transportable Tablespaces
- GoldenGate



An alternative option to fall back
from upgrade and PDB conversion



Fallback Options | After Go-Live

If you upgraded and converted

- From Oracle Database 19c to 23ai

1 Back to 19c, stay multitenant

- Downgrade
- COMPATIBLE must be 19.0.0 in 23ai CDB

2 Back to 19c, back to non-CDB

- Follow 1
- Transportable tablespace back into non-CDB
- Alternatively, Data Pump from 23ai directly to 19c non-CDB



Wrapping Up



Words of Advice

-  1 Start with simple databases
- 2 Leave time to learn and adjust
- 3 Proceed with bigger databases

Further Reading

Oracle Support:

- Oracle Multitenant: Frequently Asked Questions (Doc ID [1511619.1](#))
- How to migrate a non pluggable database that uses TDE to pluggable database ? (Doc ID [1678525.1](#))

Blog posts:

- [Database Migration from non-CDB to PDB – Typical Plugin Issues and Workarounds](#)
- [Upgrade & Plug In: With ASM, Data Guard, TDE and no Keystore Password](#)



There are many **benefits** to explore

- Application containers
- Faster cloning
- Faster provisioning
- Faster redeployment
- Sparse clones
- Resource consolidation
- Save resources
- Improved functionality
- Better management
- More secure
- Separation of duties
- Easier operations
- Many-as-one
- Faster updates
- Faster patching
- Enables self-service



*It's better to fail in our lab,
than in production*

Try multitenant migration in our [Hands-On Lab](#)

For free using Oracle LiveLabs

Part 2

Move to Oracle Database 23ai

- Everything you need to know about Oracle Multitenant

Live on June 27, 14:00 CEST

[Sign up](#)

Episode 1

Release and Patching Strategy

105 minutes – Feb 4, 2021



Episode 2

AutoUpgrade to Oracle Database 19c

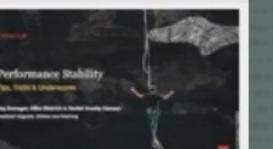
115 minutes – Feb 20, 2021



Episode 3

Performance Stability, Tips and Tricks and Underscores

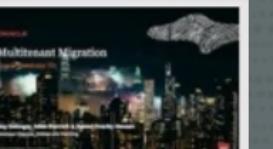
120 minutes – Mar 4, 2021



Episode 4

Migration to Oracle Multitenant

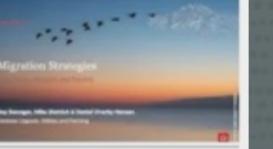
120 minutes – Mar 16, 2021



Episode 5

Migration Strategies – Insights, Tips and Secrets

120 minutes – Mar 25, 2021



Episode 6

Move to the Cloud – Not only for techies

115 minutes – Apr 8, 2021

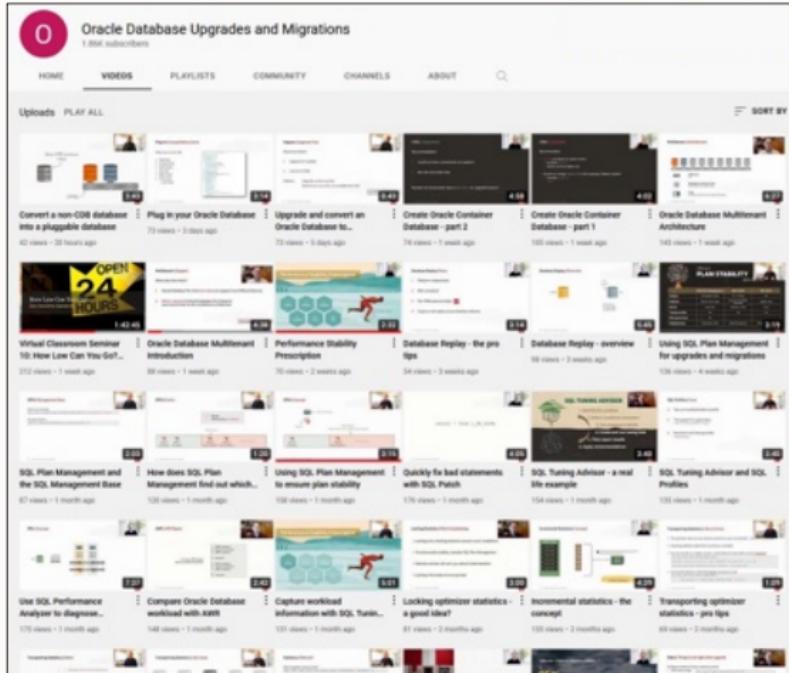


Recorded Web Seminars

<https://MikeDietrichDE.com/videos>

More than 35 hours of technical content,
on-demand, anytime, anywhere

YouTube | @UpgradeNow



[Link](#)

- 300+ videos
- New videos every week
- No marketing
- No buzzwords
- All tech



Thank You

