

**ORACLE**

# CloudWorld

## Hints and Tips for Fast and Predictable Query Performance with Star Schemas

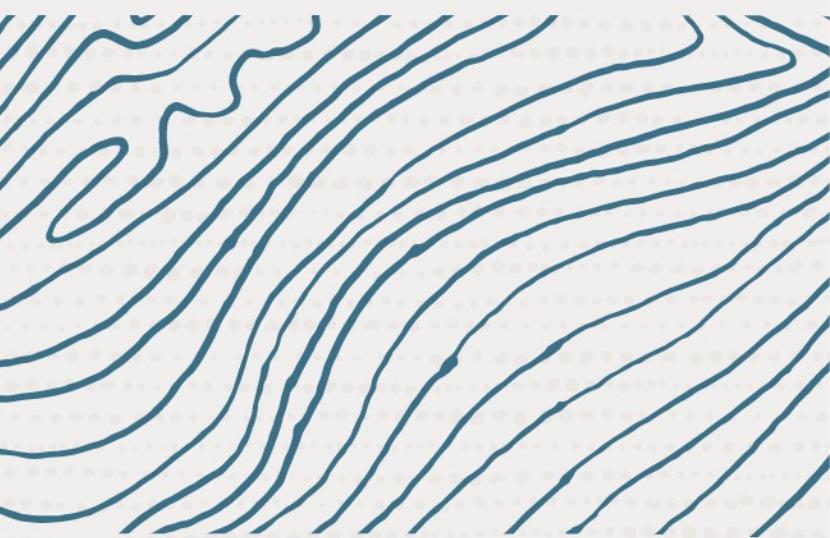
Session LRN3511

Juergen Mueller

Mihajlo Tekic

October 17–20, 2022

Caesars Forum and The Venetian  
Las Vegas, NV



# Real-World Performance Sessions

<b>Tuesday</b>	LRN3511
<b>11:00-11:45AM</b>	Hints and Tips for Fast and Predictable Query Performance with Star Schemas
<b>Wednesday</b>	LRN1483
<b>02:30-03:15PM</b>	Solving Performance Problems Using Automatic Workload Repository
<b>Online</b>	<a href="http://oracle.com/goto/oll/rwp">oracle.com/goto/oll/rwp</a>



# What is Real-World Performance?

- Getting the most out of Software and Hardware
- Achieving Performance Excellence



# Real-World Performance Team

## Who We Are

- Part of Oracle Database Development
- Team members in USA, Europe and Asia
- Over one hundred years of experience combined

## How We Work

- Use the product as designed
- Take a holistic view
- Aim for best performance
- Apply data-driven analysis
- Share what we learn

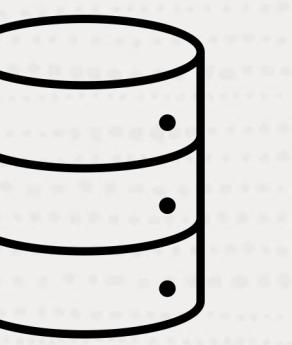


# What We Do



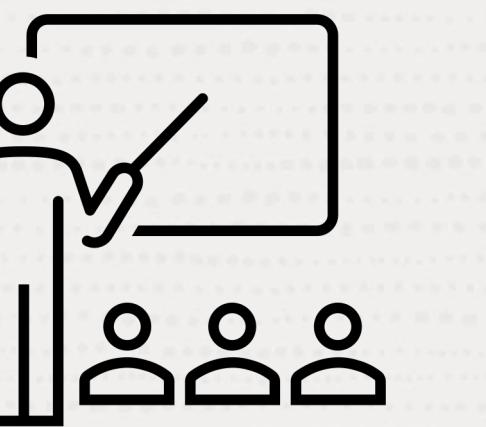
## Customer Engagement

- Design Review
- Performance Project



## Database Development

- Applications
- Tools



## Customer Education

- Online
- Focused



**ORACLE**

# CloudWorld

## Hints and Tips for Fast and Predictable Query Performance with Star Schemas

Session LRN3511

Juergen Mueller

Mihajlo Tekic

October 17–20, 2022

Caesars Forum and The Venetian  
Las Vegas, NV



# Agenda

## Dimensional Model

Overview

## Worked Example

From minutes to seconds

What you must do

## Star Query Execution on ADW

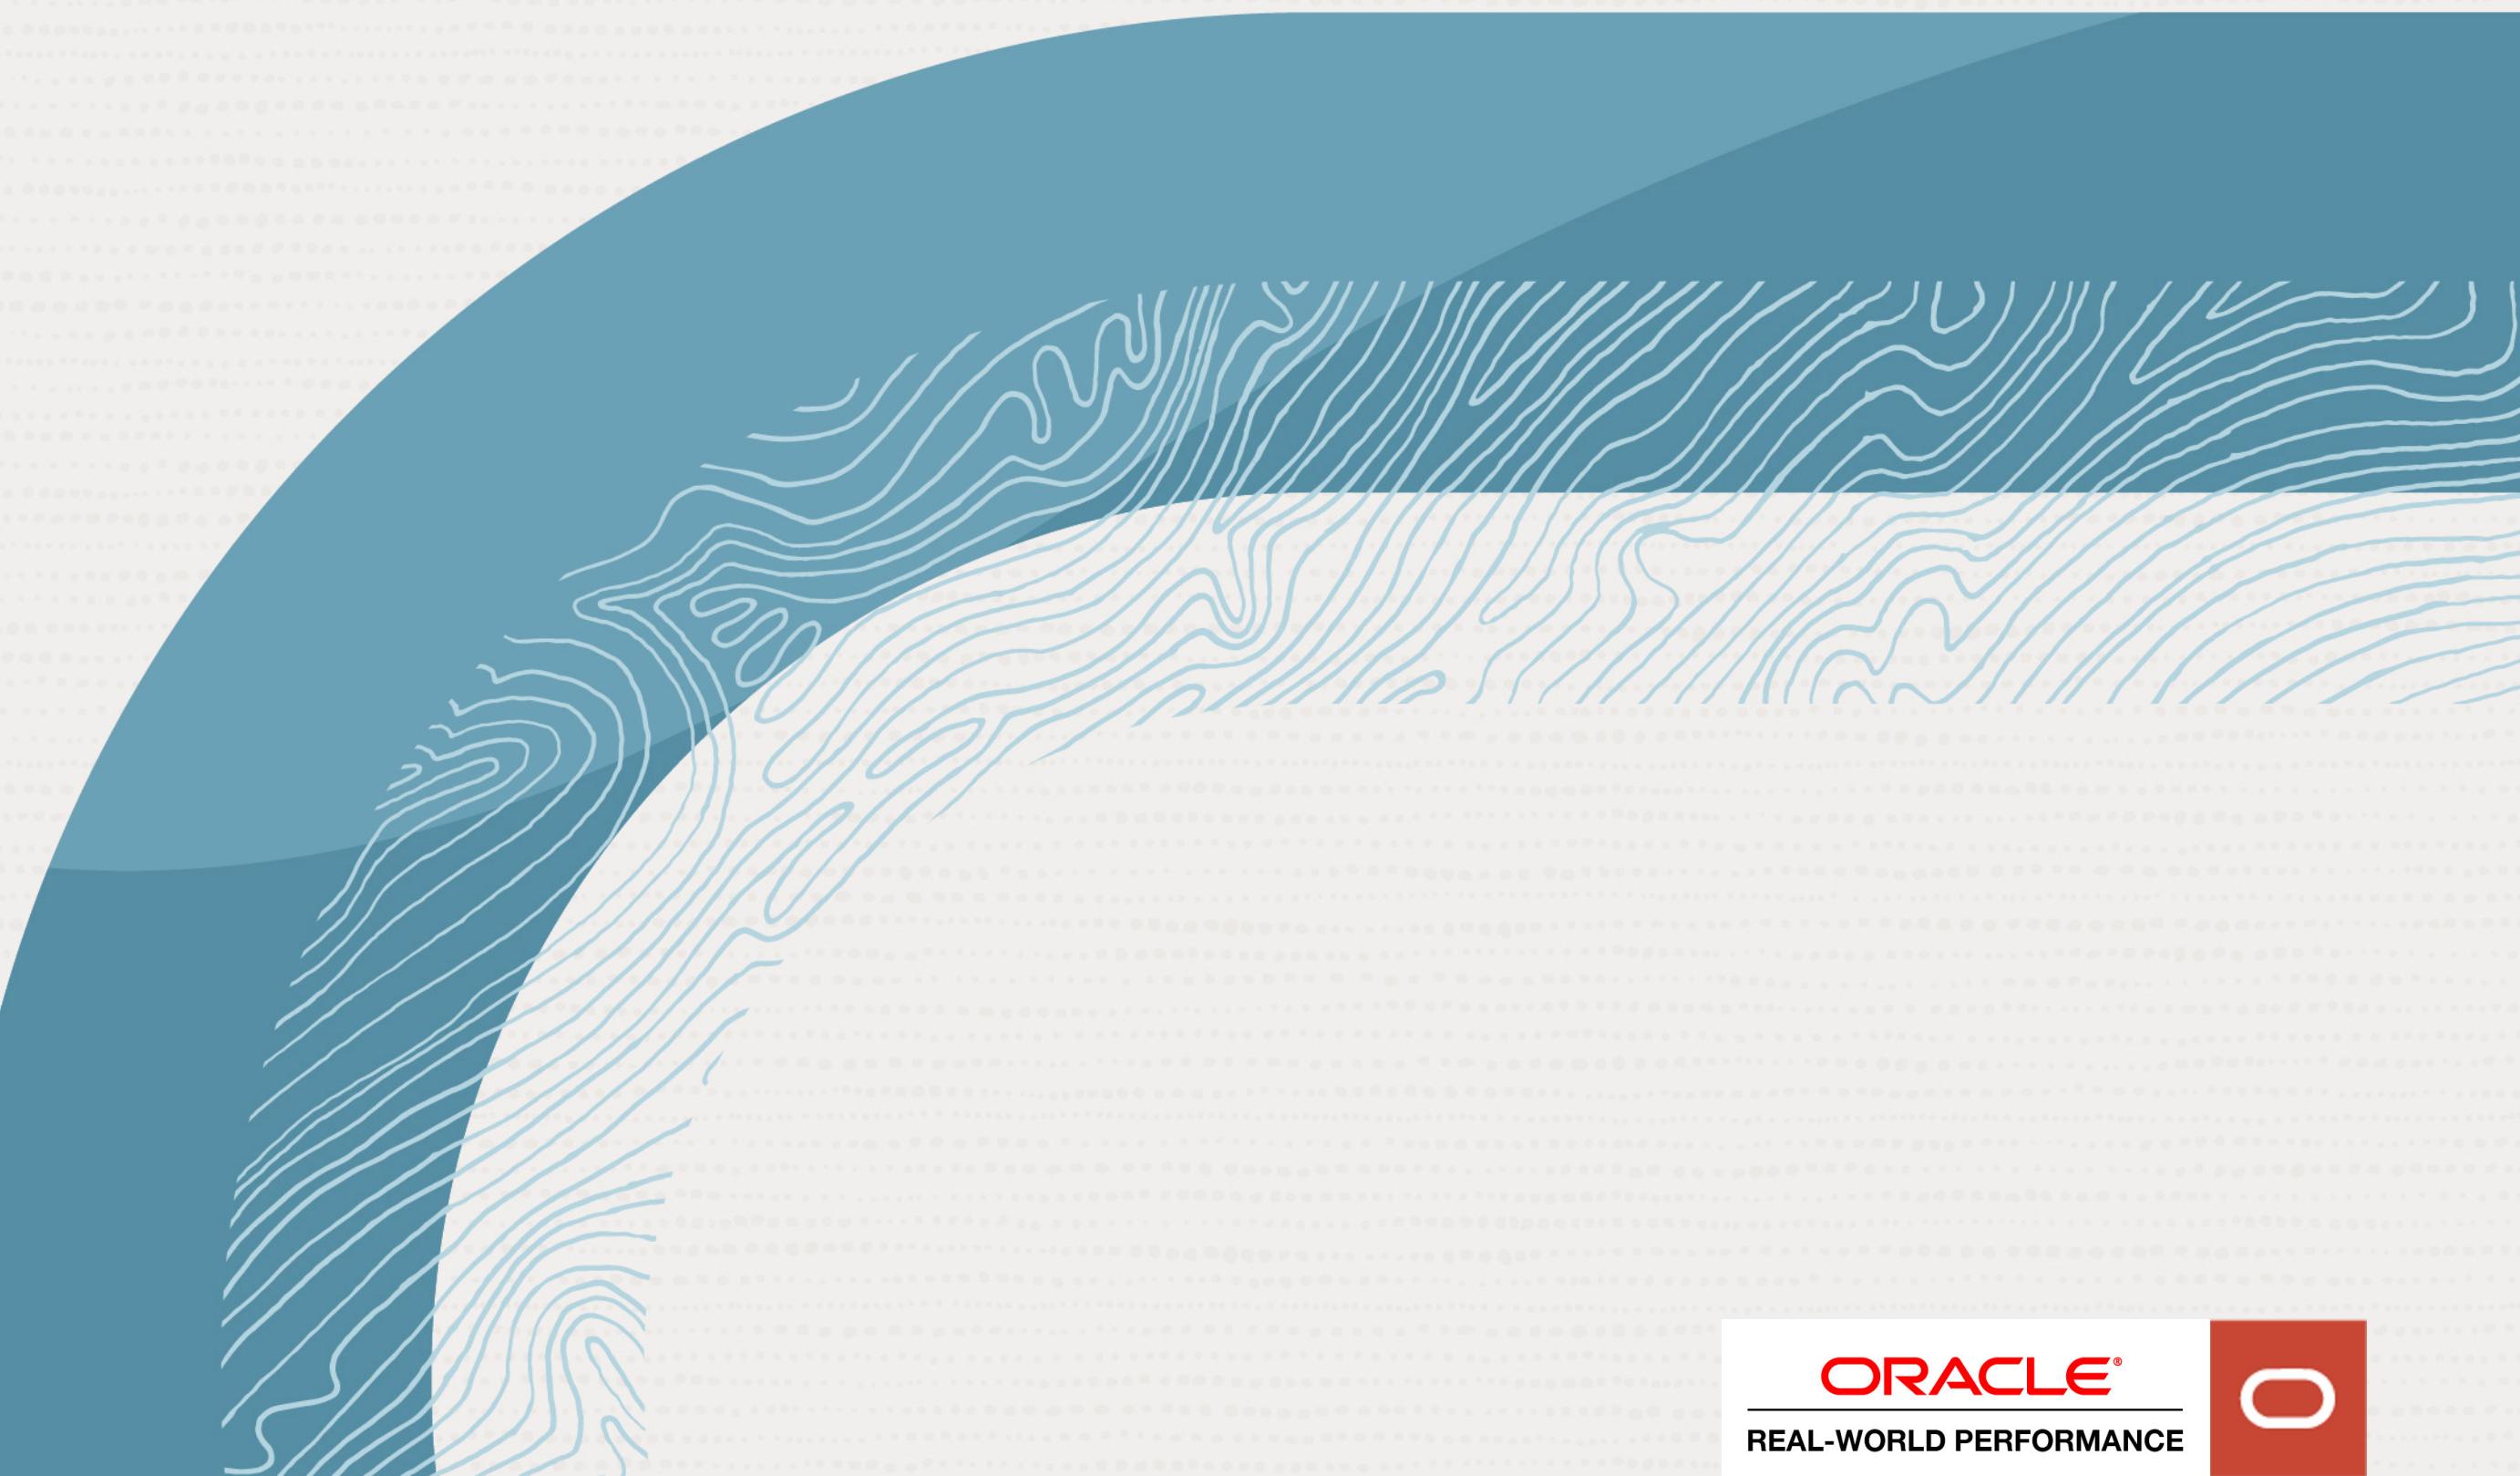
Analysis of our worked example on ADW

## Summary

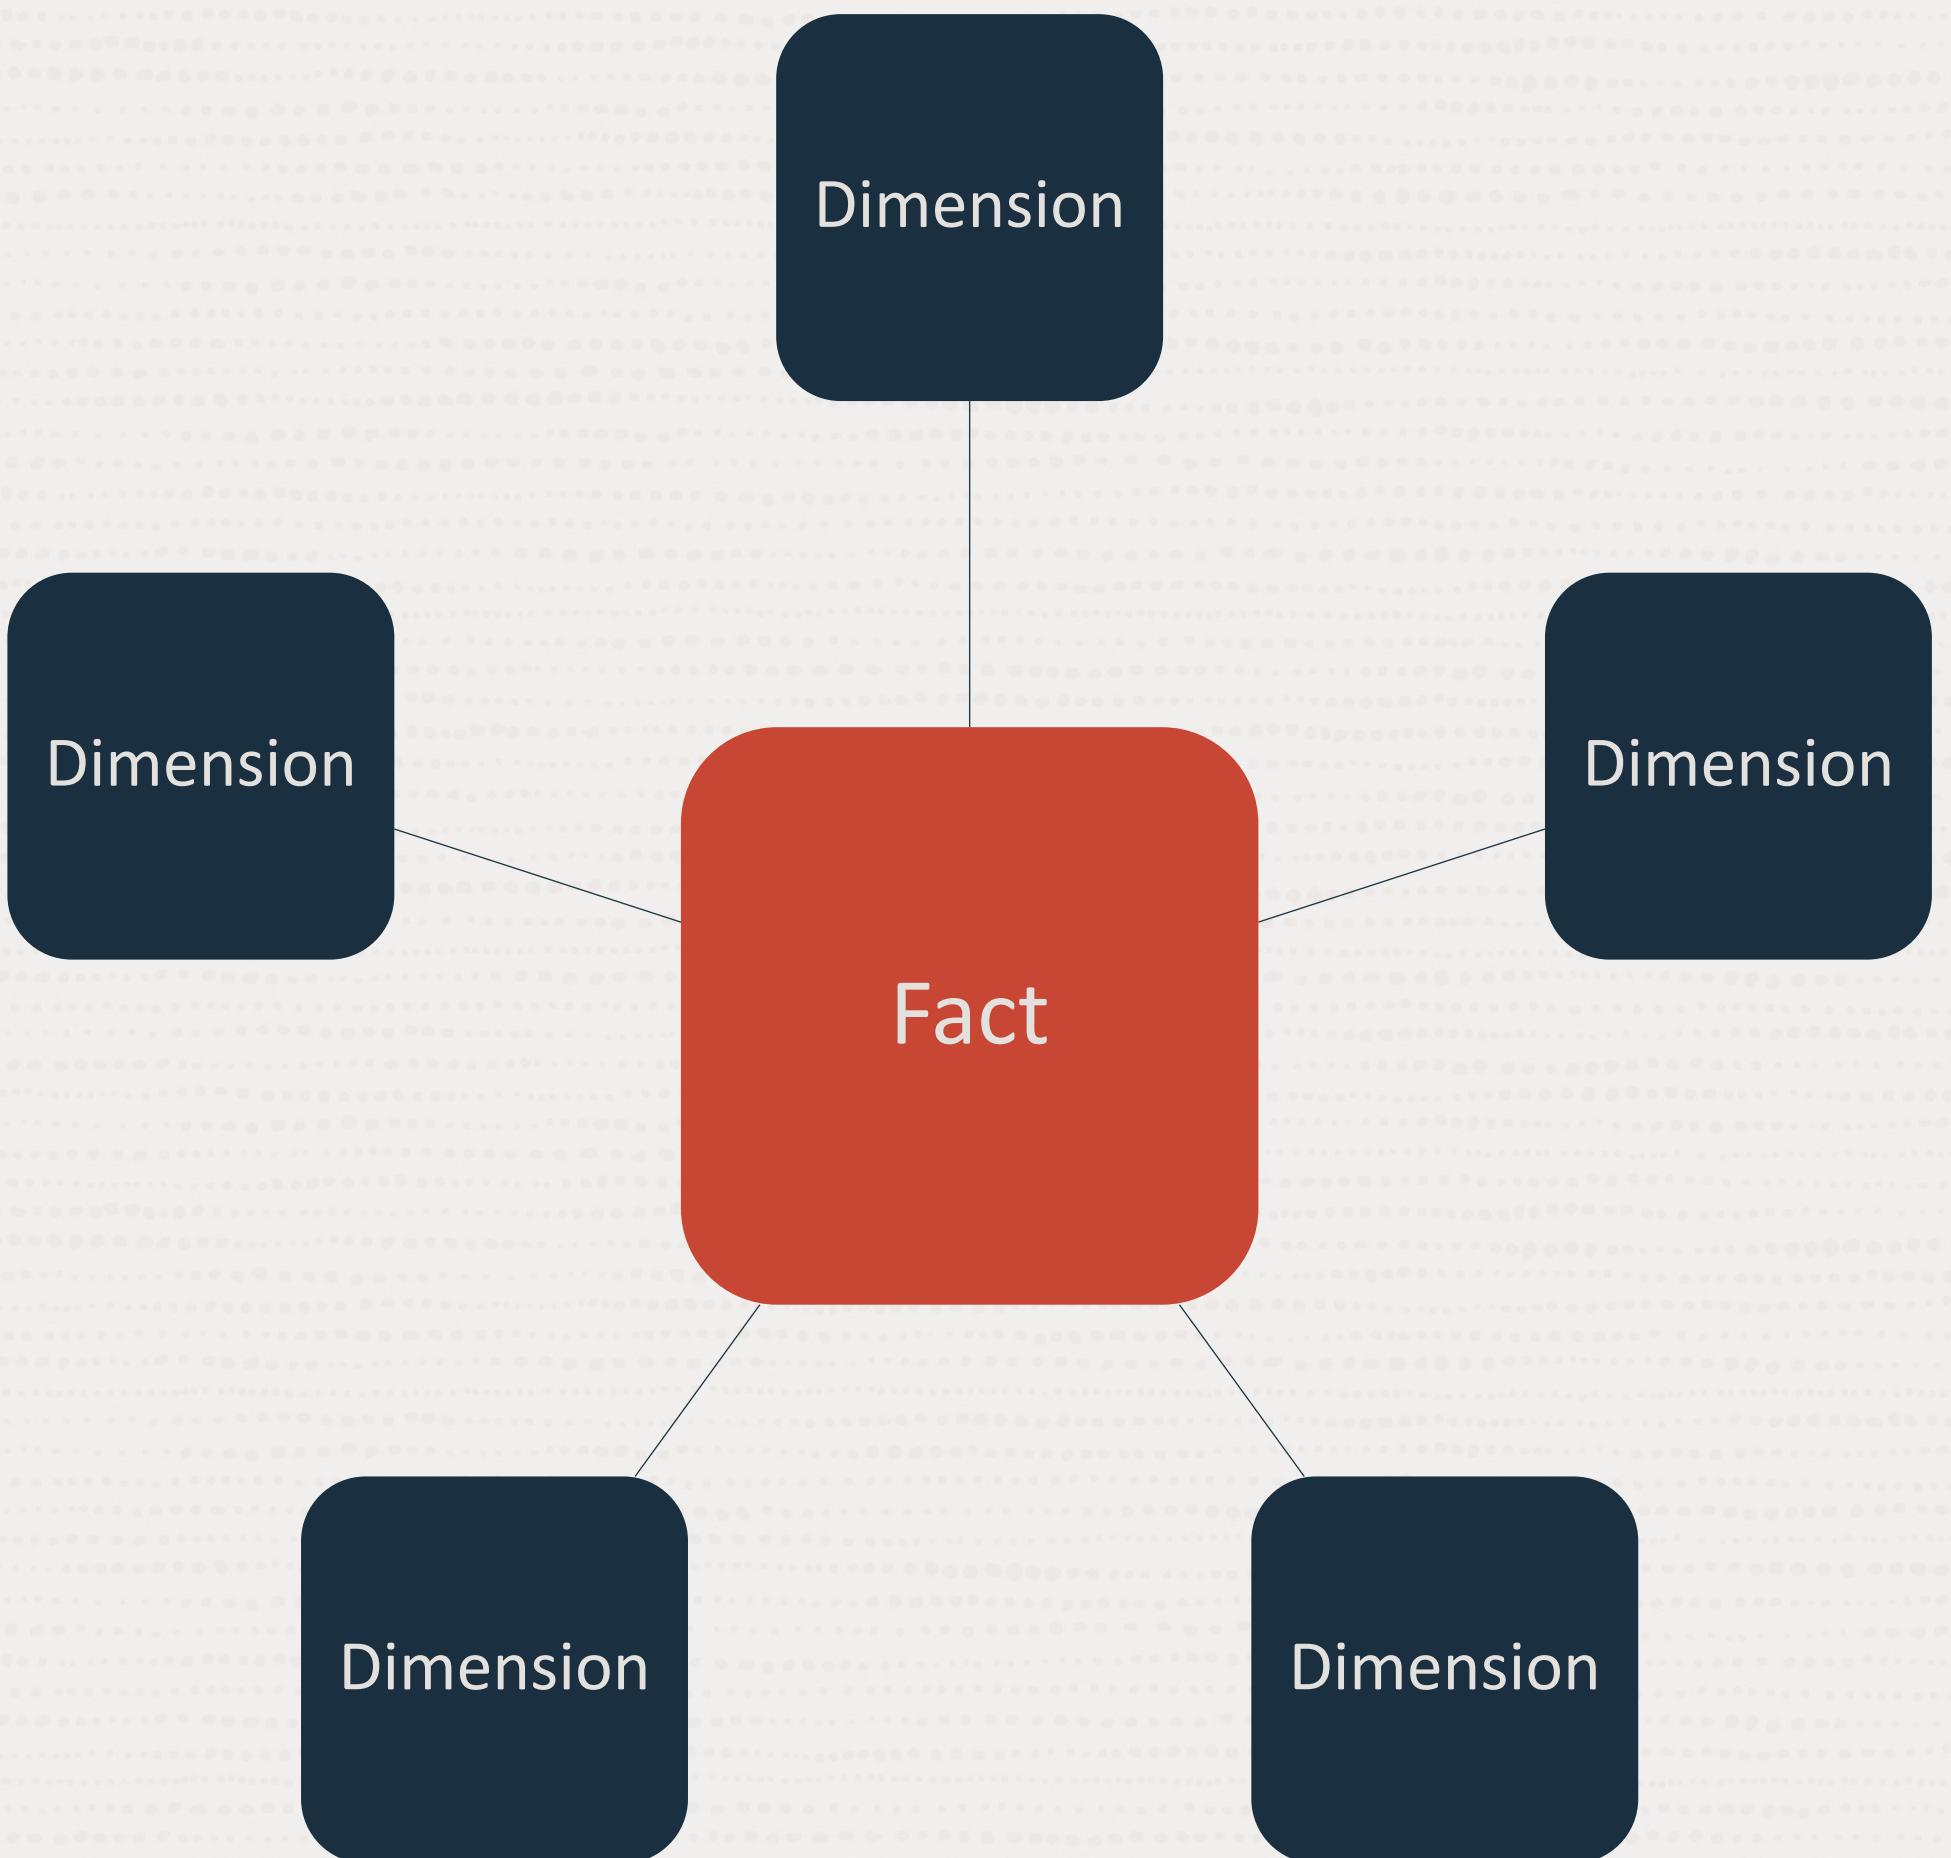


# Dimensional Model

## Overview



# Why a Dimensional Model?



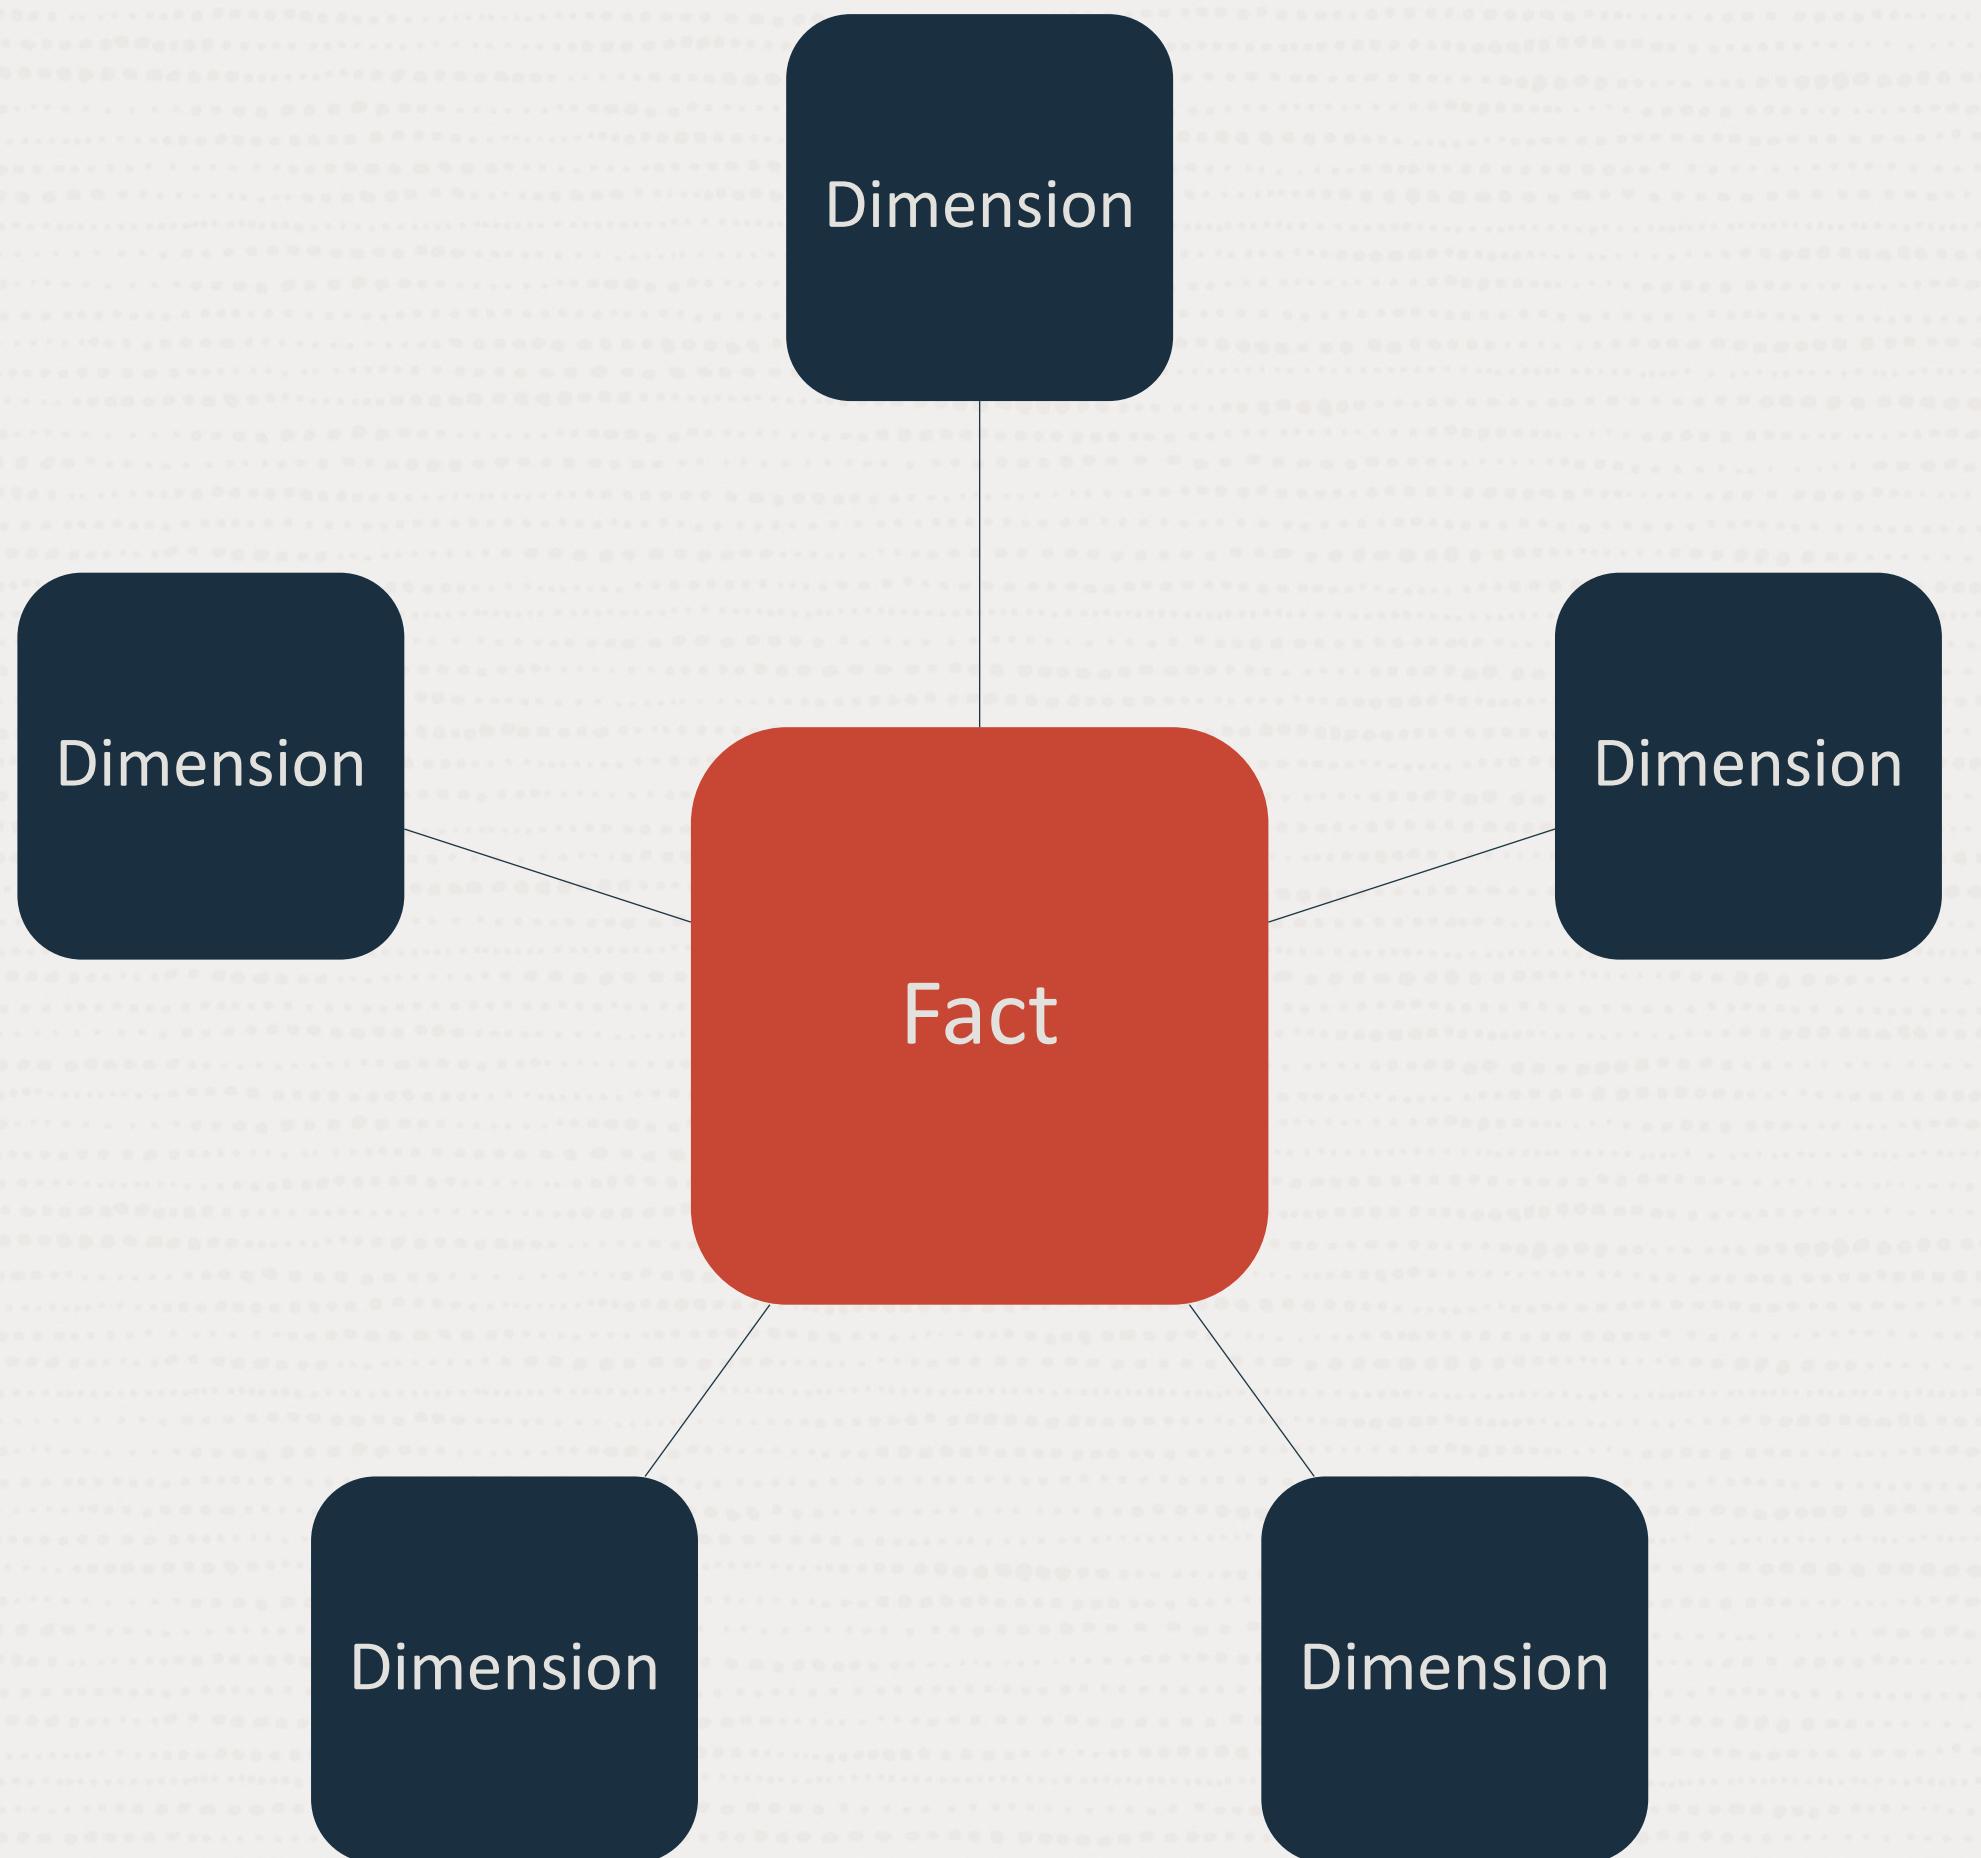
Predictable and standard format

Star join framework withstands unexpected changes in user behavior

## Performance

- Well established performance algorithms for this model

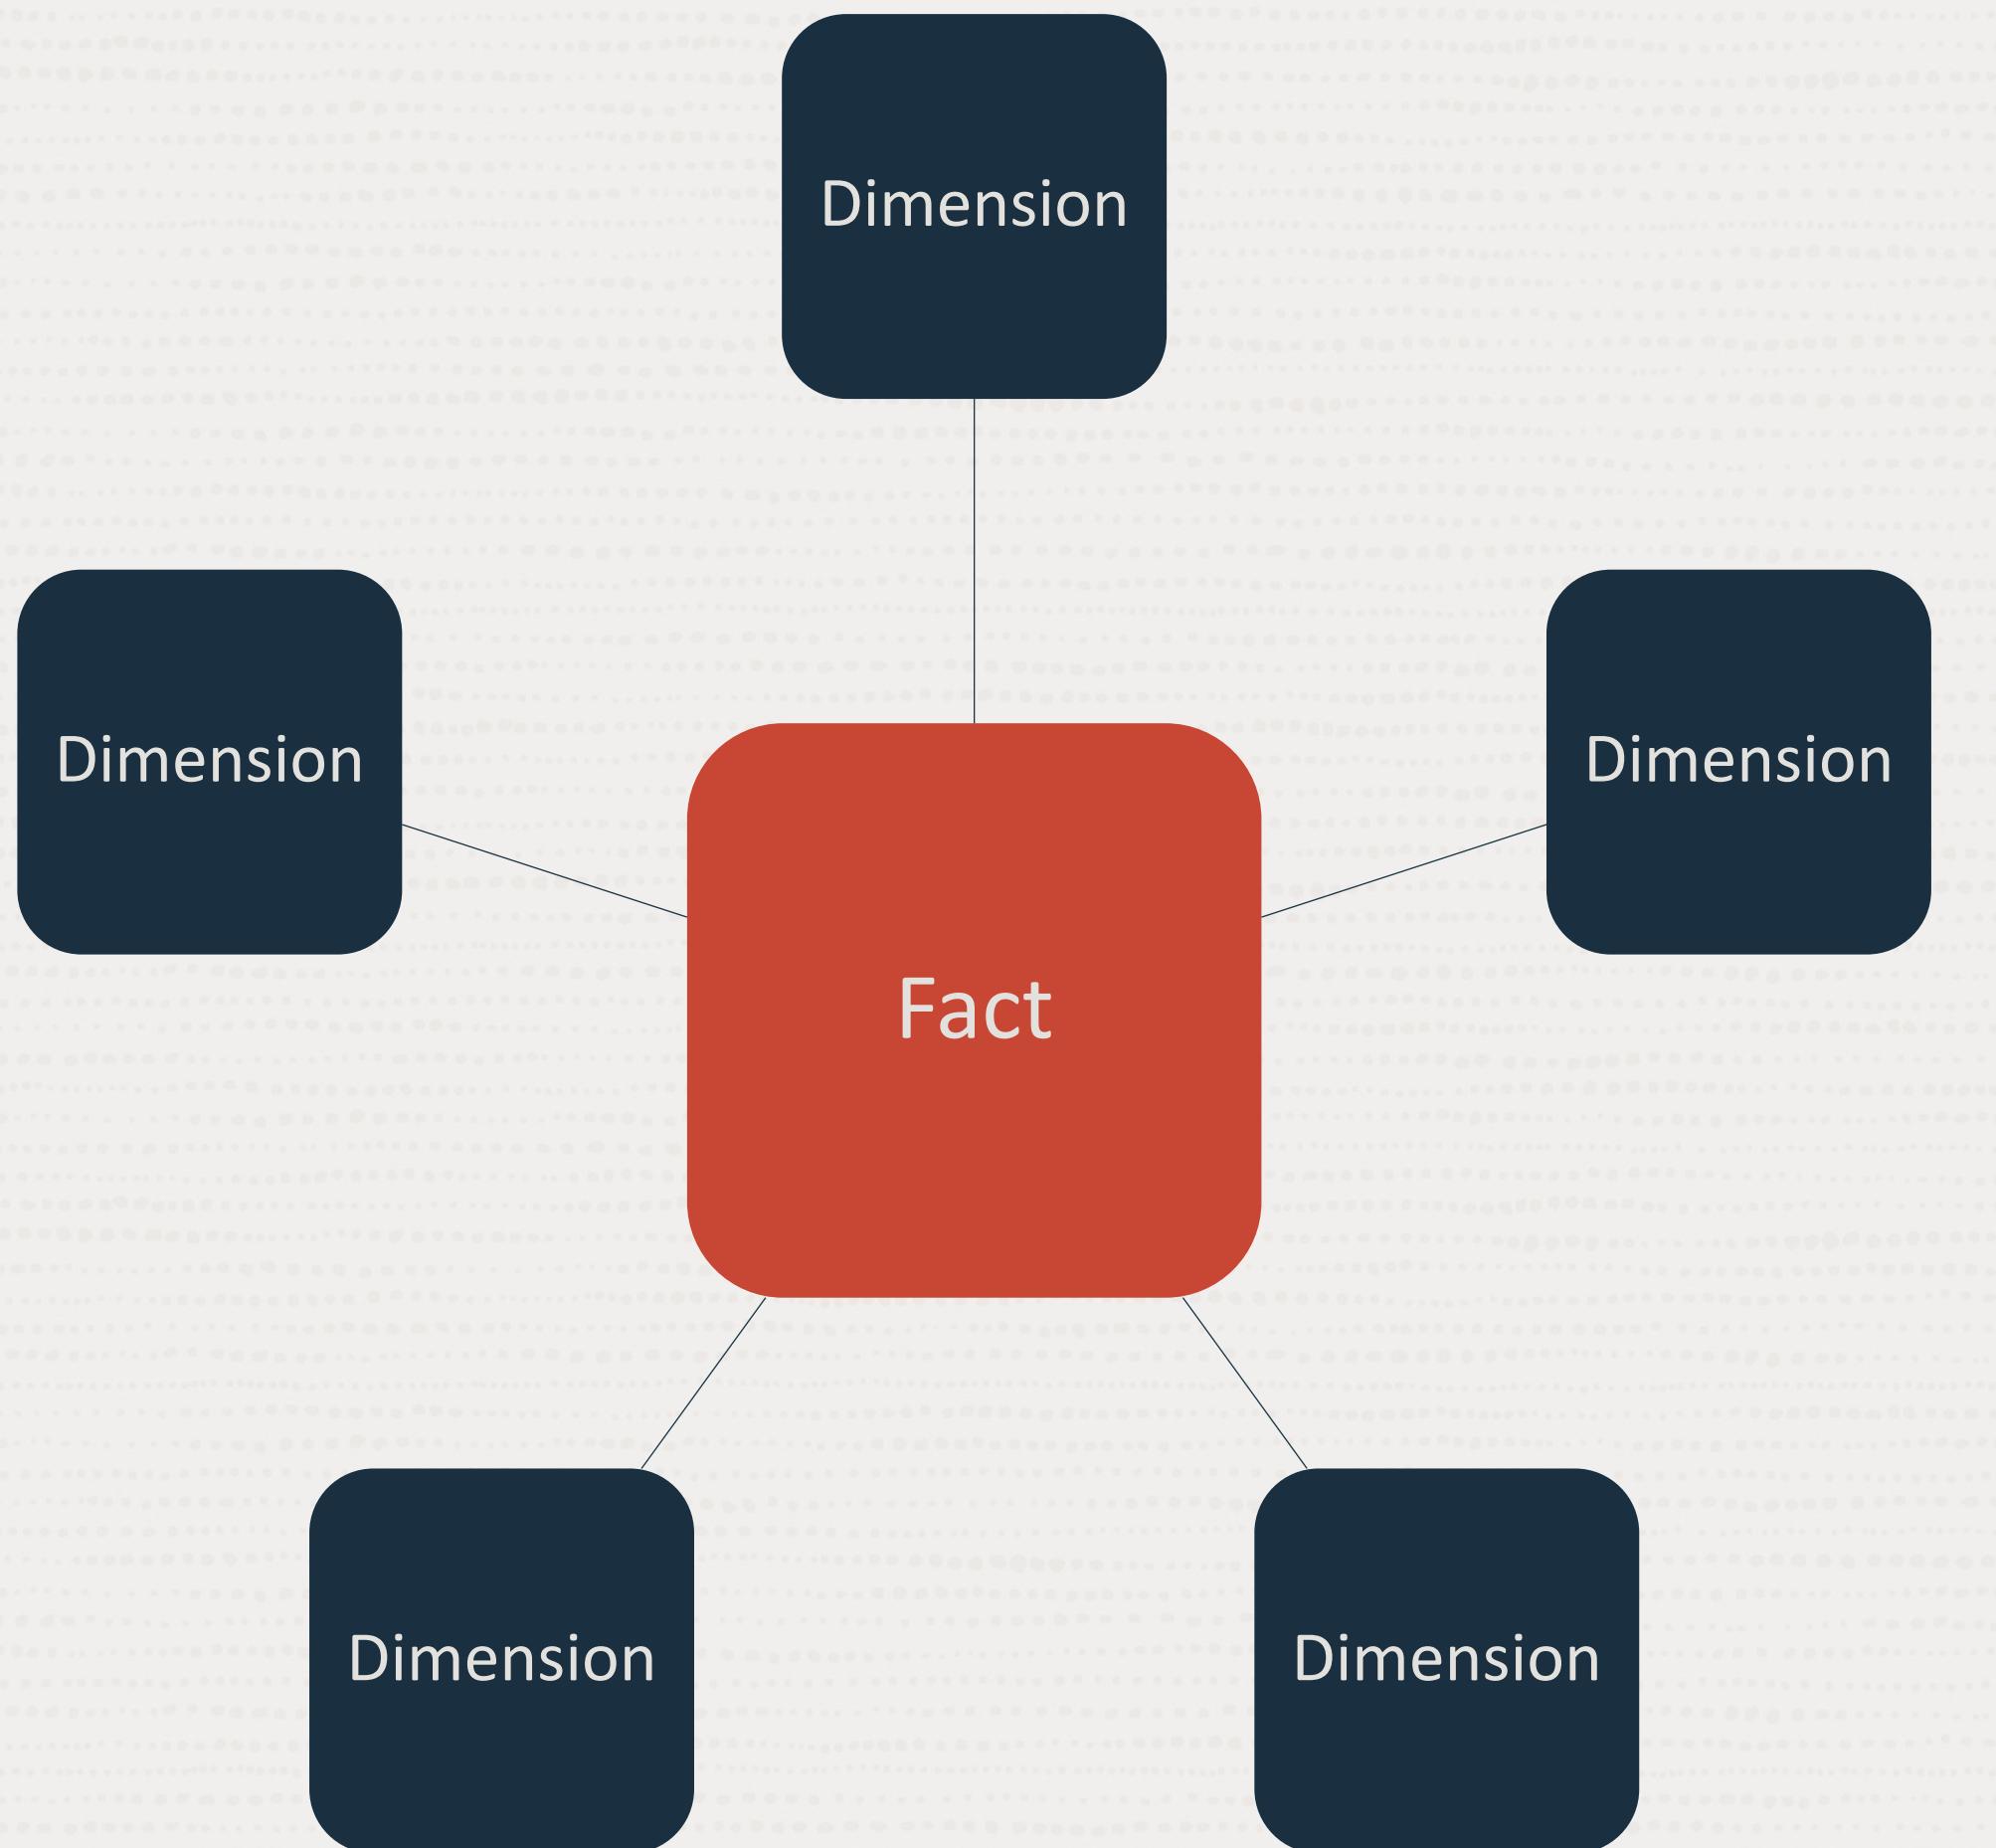
# Star Schema Characteristics



## Fact tables

- Contain business events called measures
  - Generally numerical values
    - Sales
    - Orders
    - Logins
- Large
- Append only
- Historical in nature

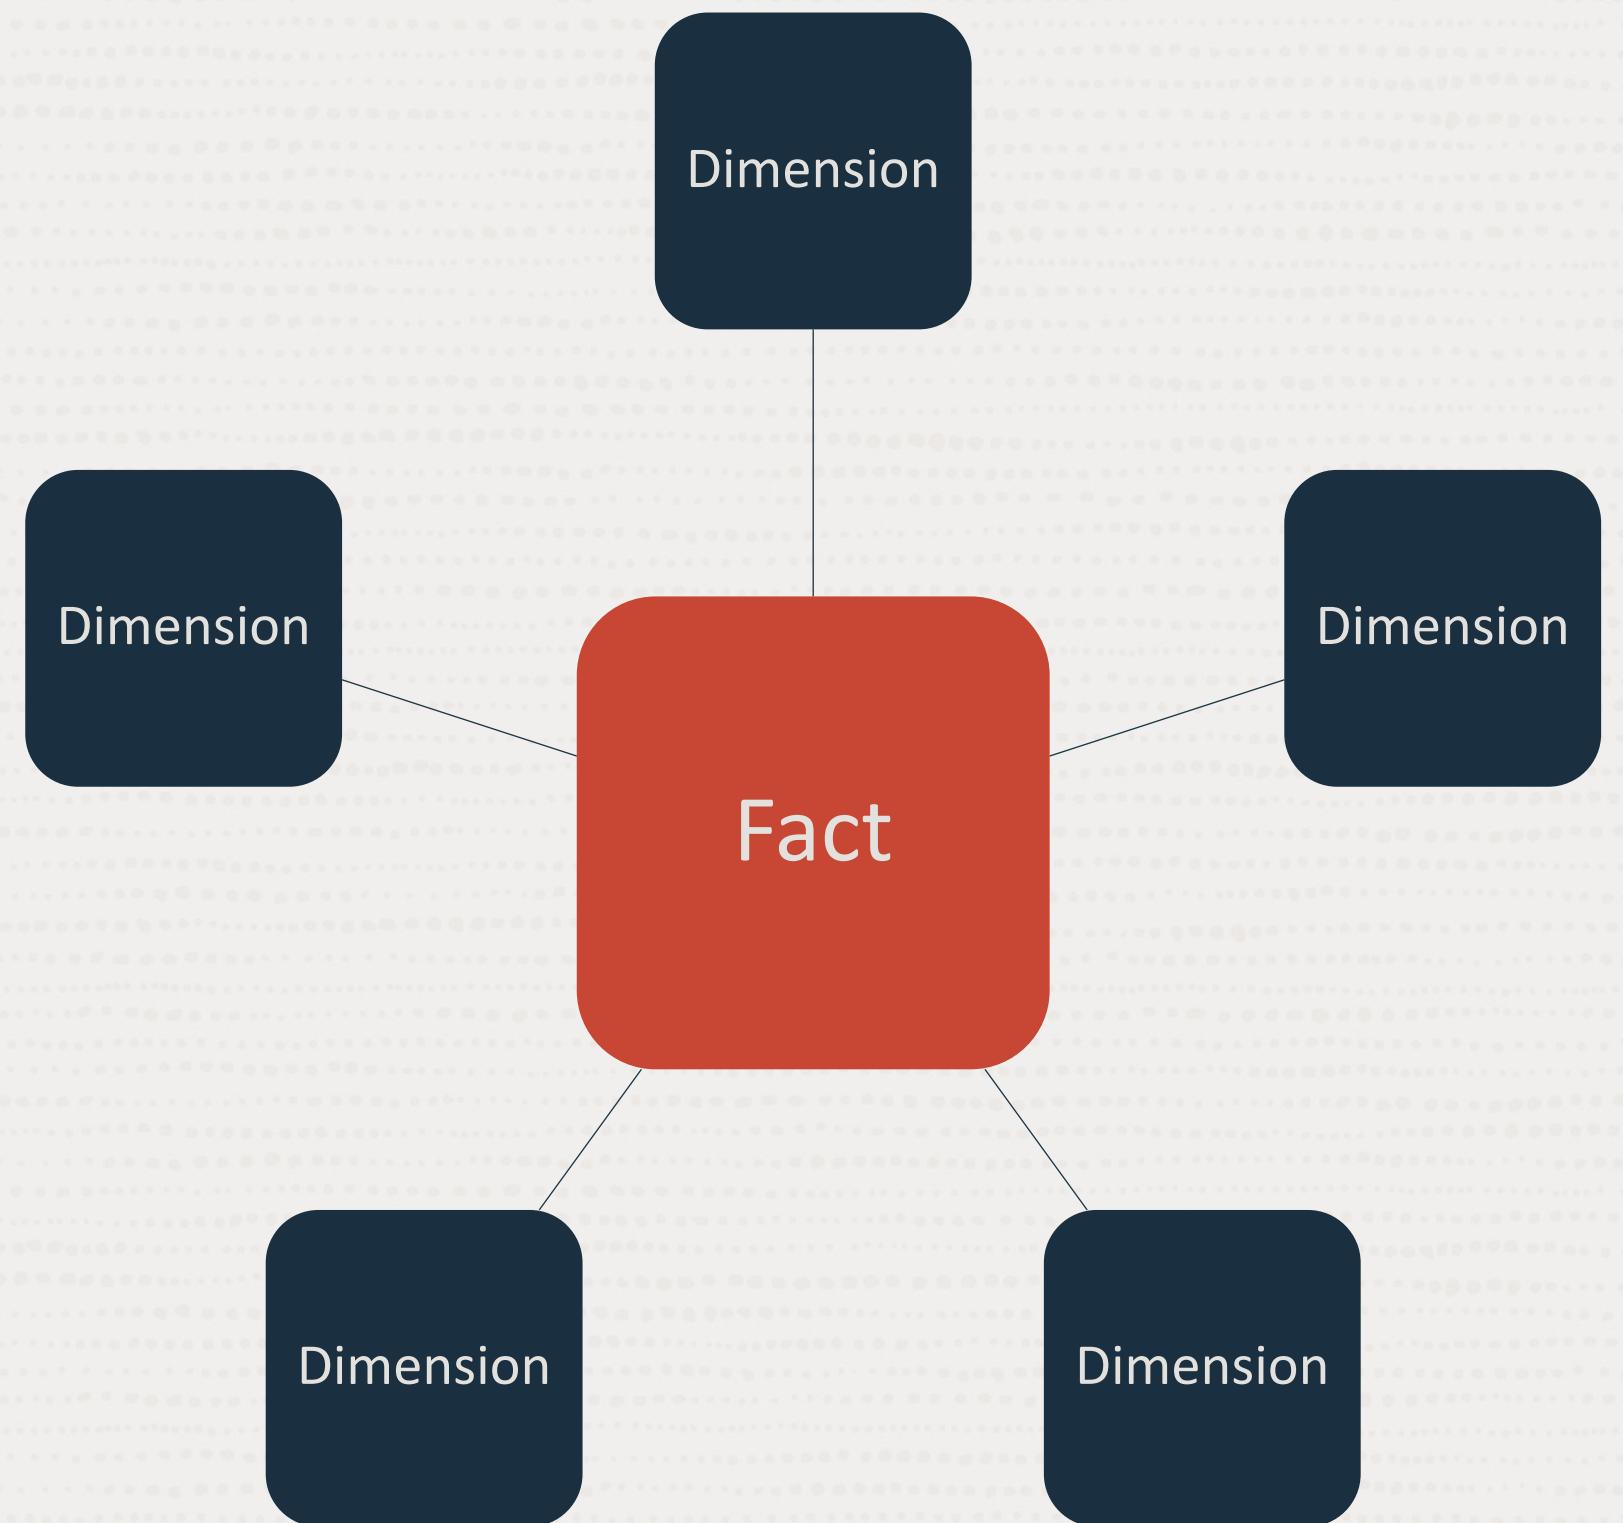
# Star Schema Characteristics



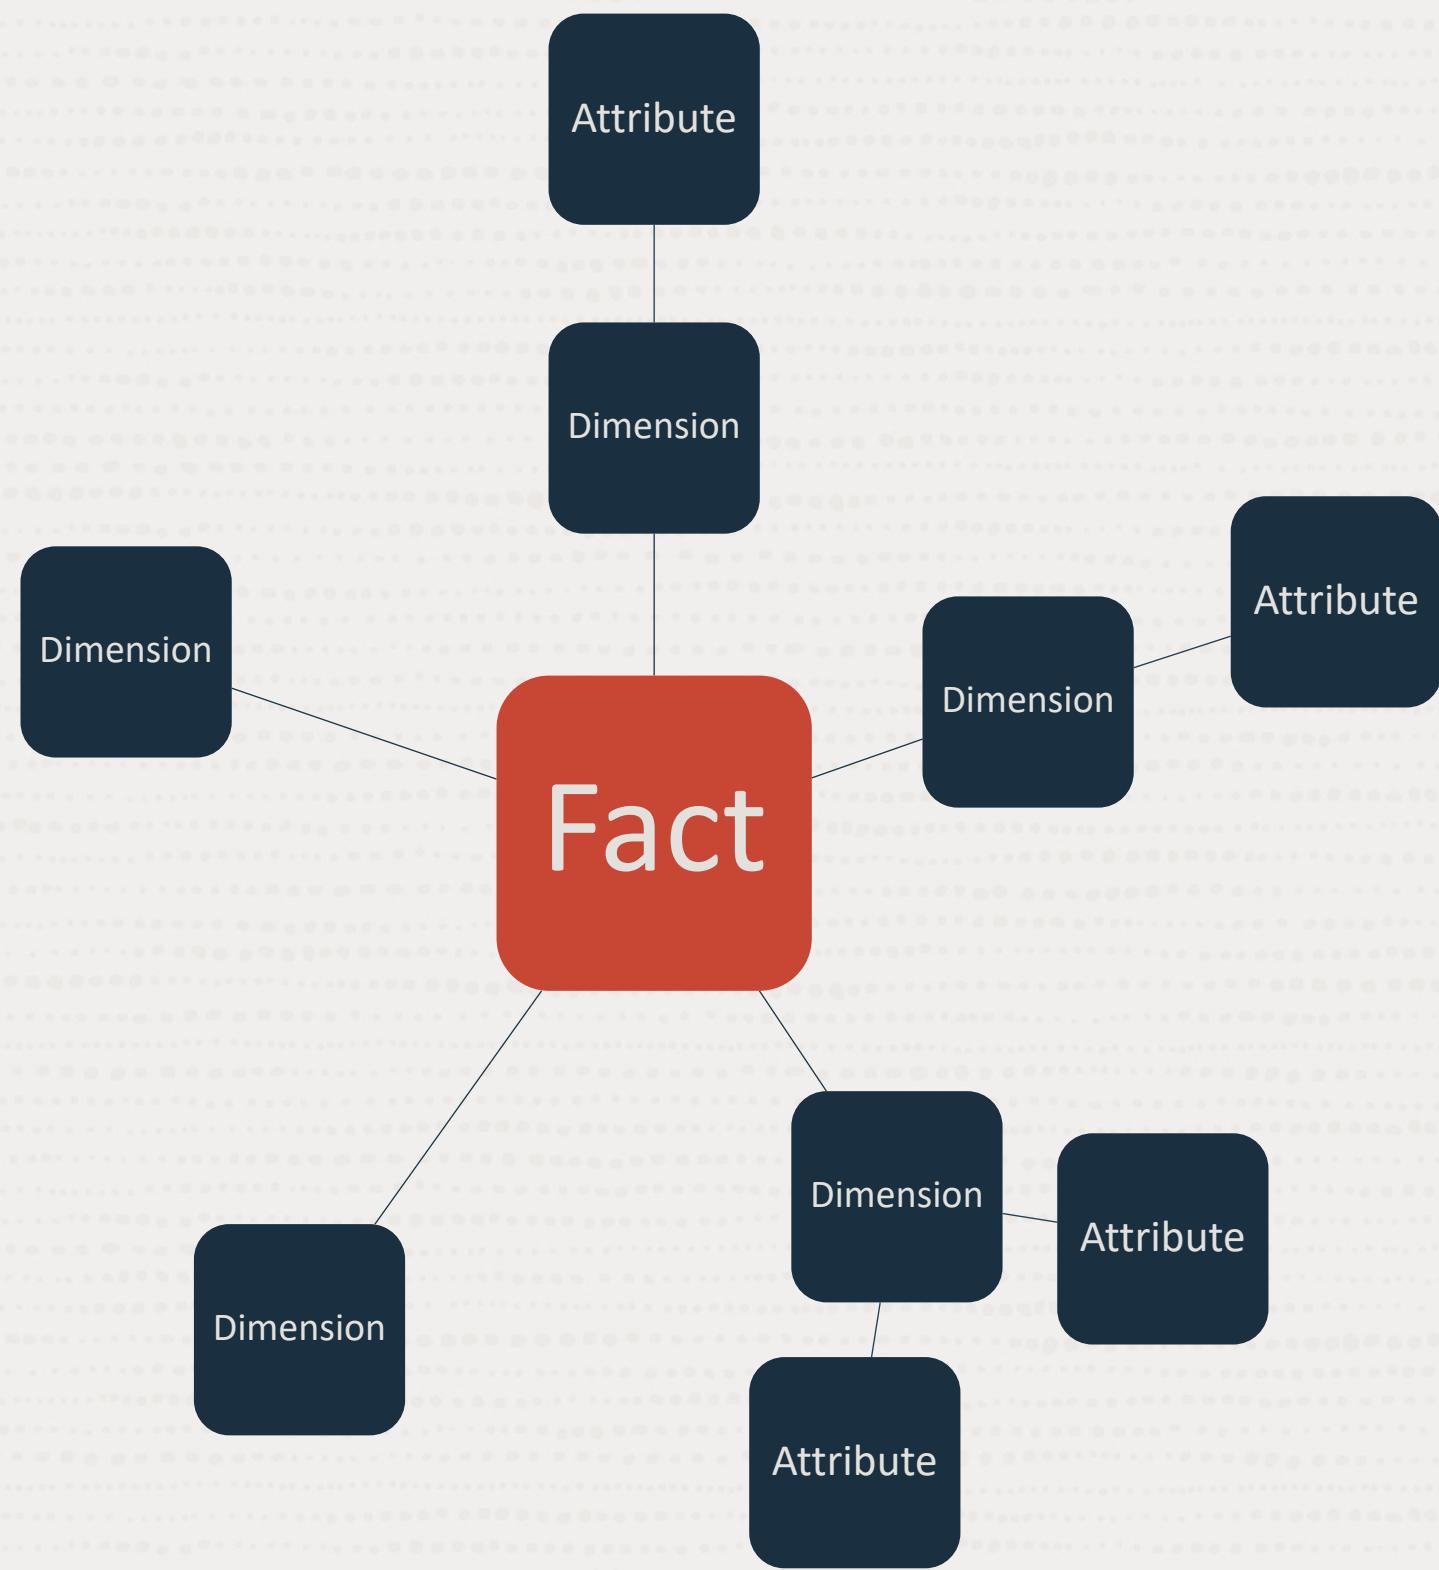
## Dimension tables

- Contain reference information about facts called attributes
  - when
  - what
  - where
  - who
  - how
  - why
- Small compared to Fact
- Static or slowly changing

# Dimensional Schemas



Star Schema



Star Schema derivative

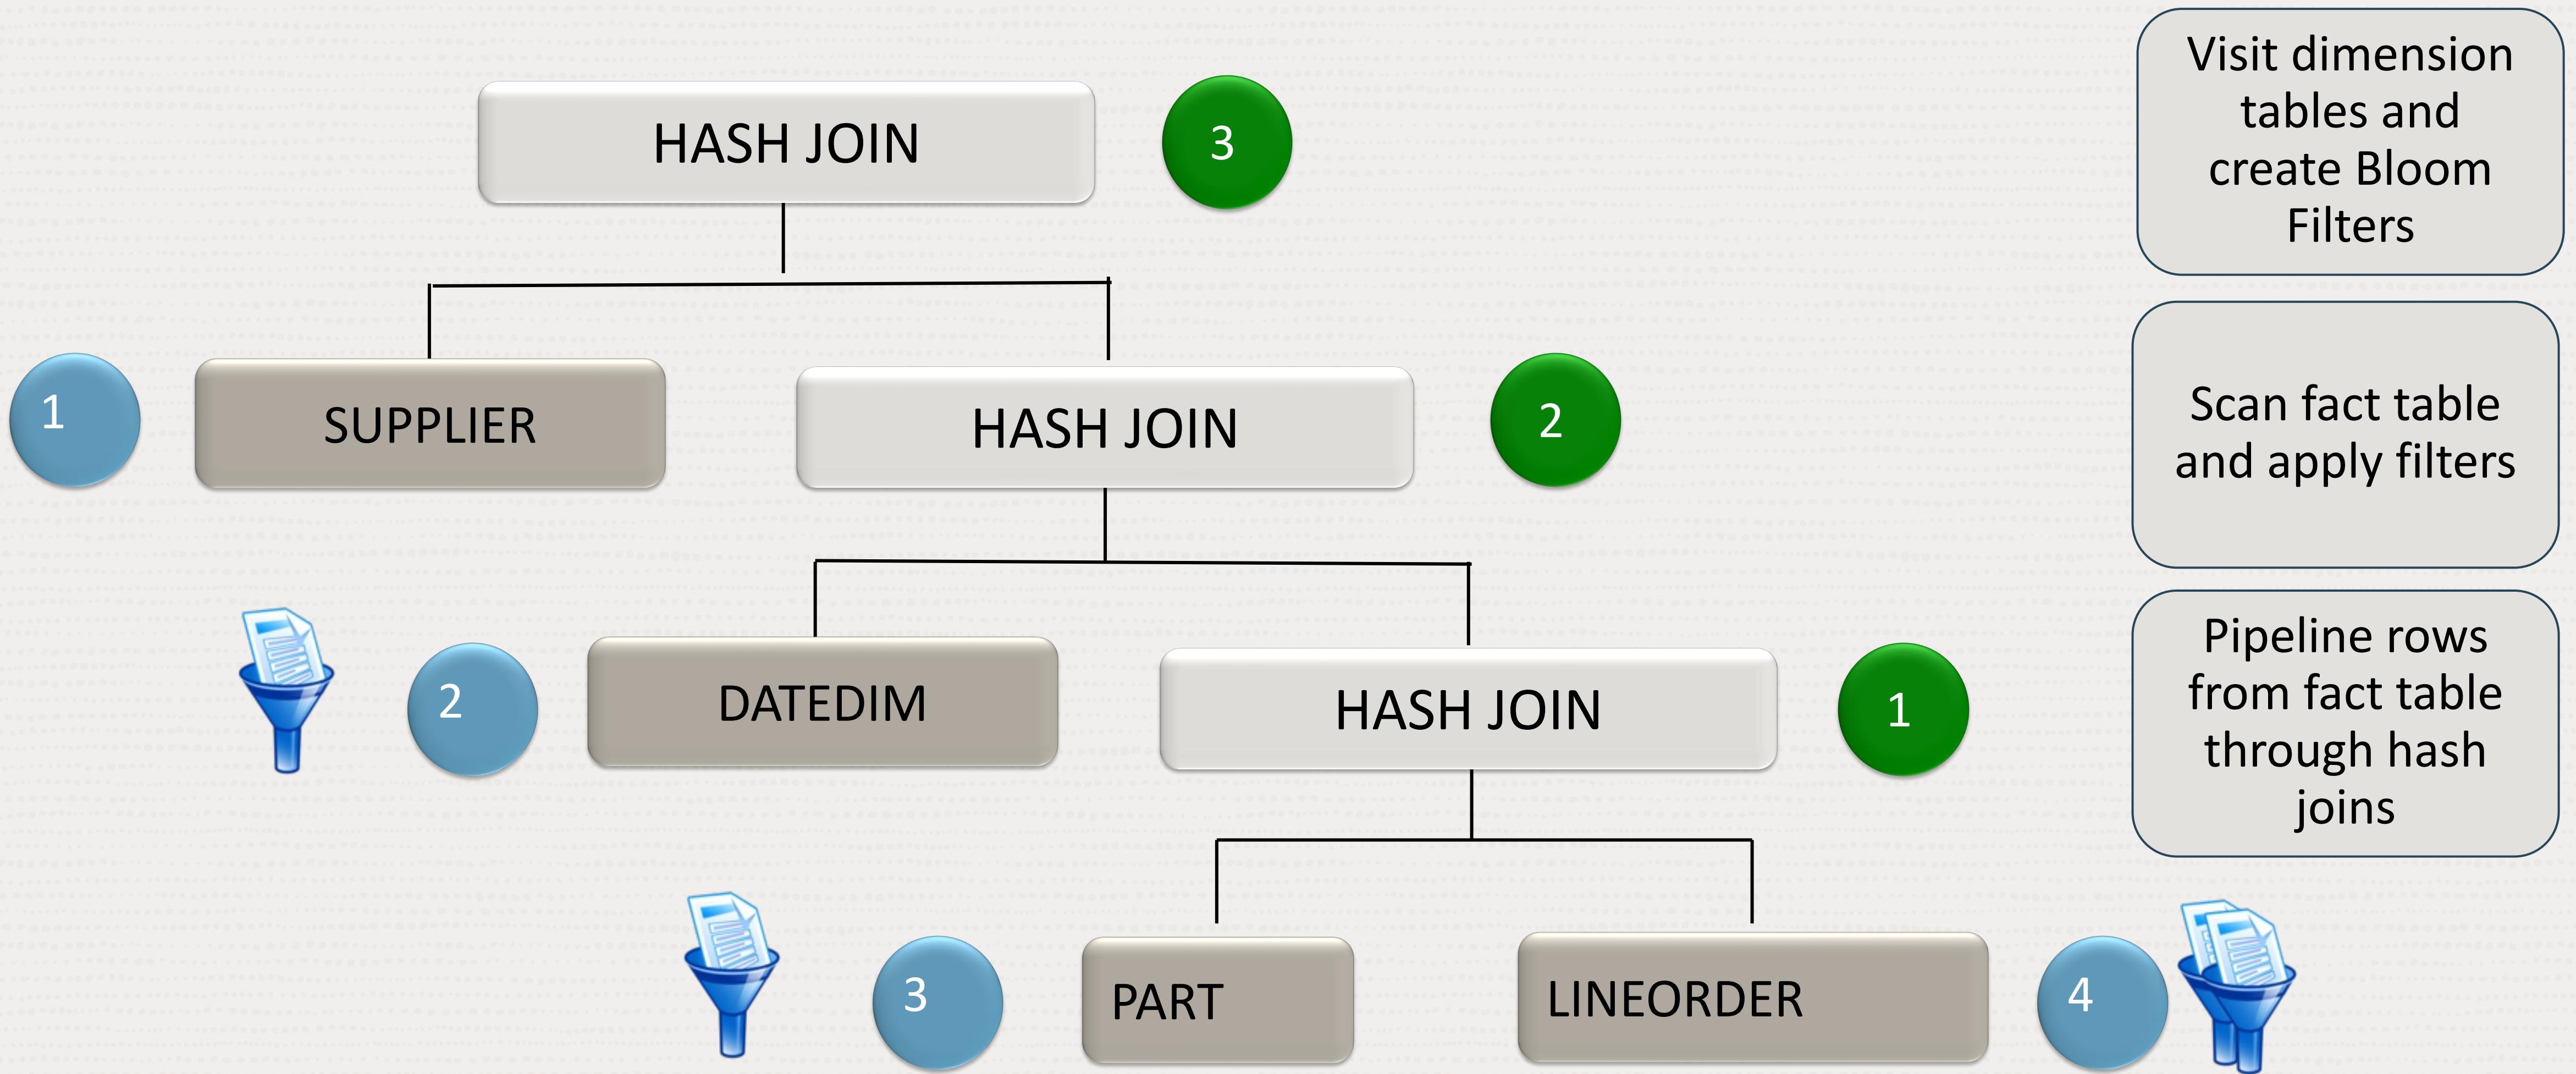
# Shape and Structure of a Typical Dimensional Query

```
SELECT d_sellingseason, p_category, s_region,
       SUM(lo_extendedprice)
    FROM lineorder
        JOIN customer          ON lo_custkey = c_custkey
        JOIN date_dim           ON lo_orderdate = d_datekey
        JOIN part                ON lo_partkey = p_partkey
        JOIN supplier            ON lo_suppkey = s_suppkey
   WHERE d_year IN (1993, 1994, 1995)
     AND p_container in ('JUMBO PACK')
GROUP BY d_sellingseason, p_category, s_region
ORDER BY d_sellingseason, p_category, s_region
```

- Choose your **fact** table
- Complete the star by defining relationships with **joins** to dimension tables
- Choose **filter** criteria based upon dimension attributes
- Choose **measures** for aggregation
- Choose **segmentation/roll up** columns
- Choose **grouping** requirements
- Choose **ordering** requirements



# Shape of the Execution Plan – Right Deep Tree



# Shape of the Execution Plan – Right Deep Tree

Operation	Name	Lin...
└ SELECT STATEMENT		0
└ SORT GROUP BY		1
└ HASH JOIN		2
└ TABLE ACCESS STORAGE FULL	SUPPLIER	3
└ HASH JOIN		4
└ JOIN FILTER CREATE	:BF0001	5
└ PART JOIN FILTER CREATE	:BF0000	6
└ TABLE ACCESS STORAGE FULL	DATE_DIM	7
└ HASH JOIN		8
└ JOIN FILTER CREATE	:BF0002	9
└ TABLE ACCESS STORAGE FULL	PART	10
└ JOIN FILTER USE	:BF0001	11
└ JOIN FILTER USE	:BF0002	12
└ PARTITION RANGE JOIN-FILTER		13
└ TABLE ACCESS STORAGE FULL	LINEORDER	14

Visit dimension  
tables and  
create Bloom  
Filters

Scan fact table  
and apply filters

Pipeline rows  
from fact table  
through hash  
joins

# Things You Must Do to Ensure Optimal Execution Plans

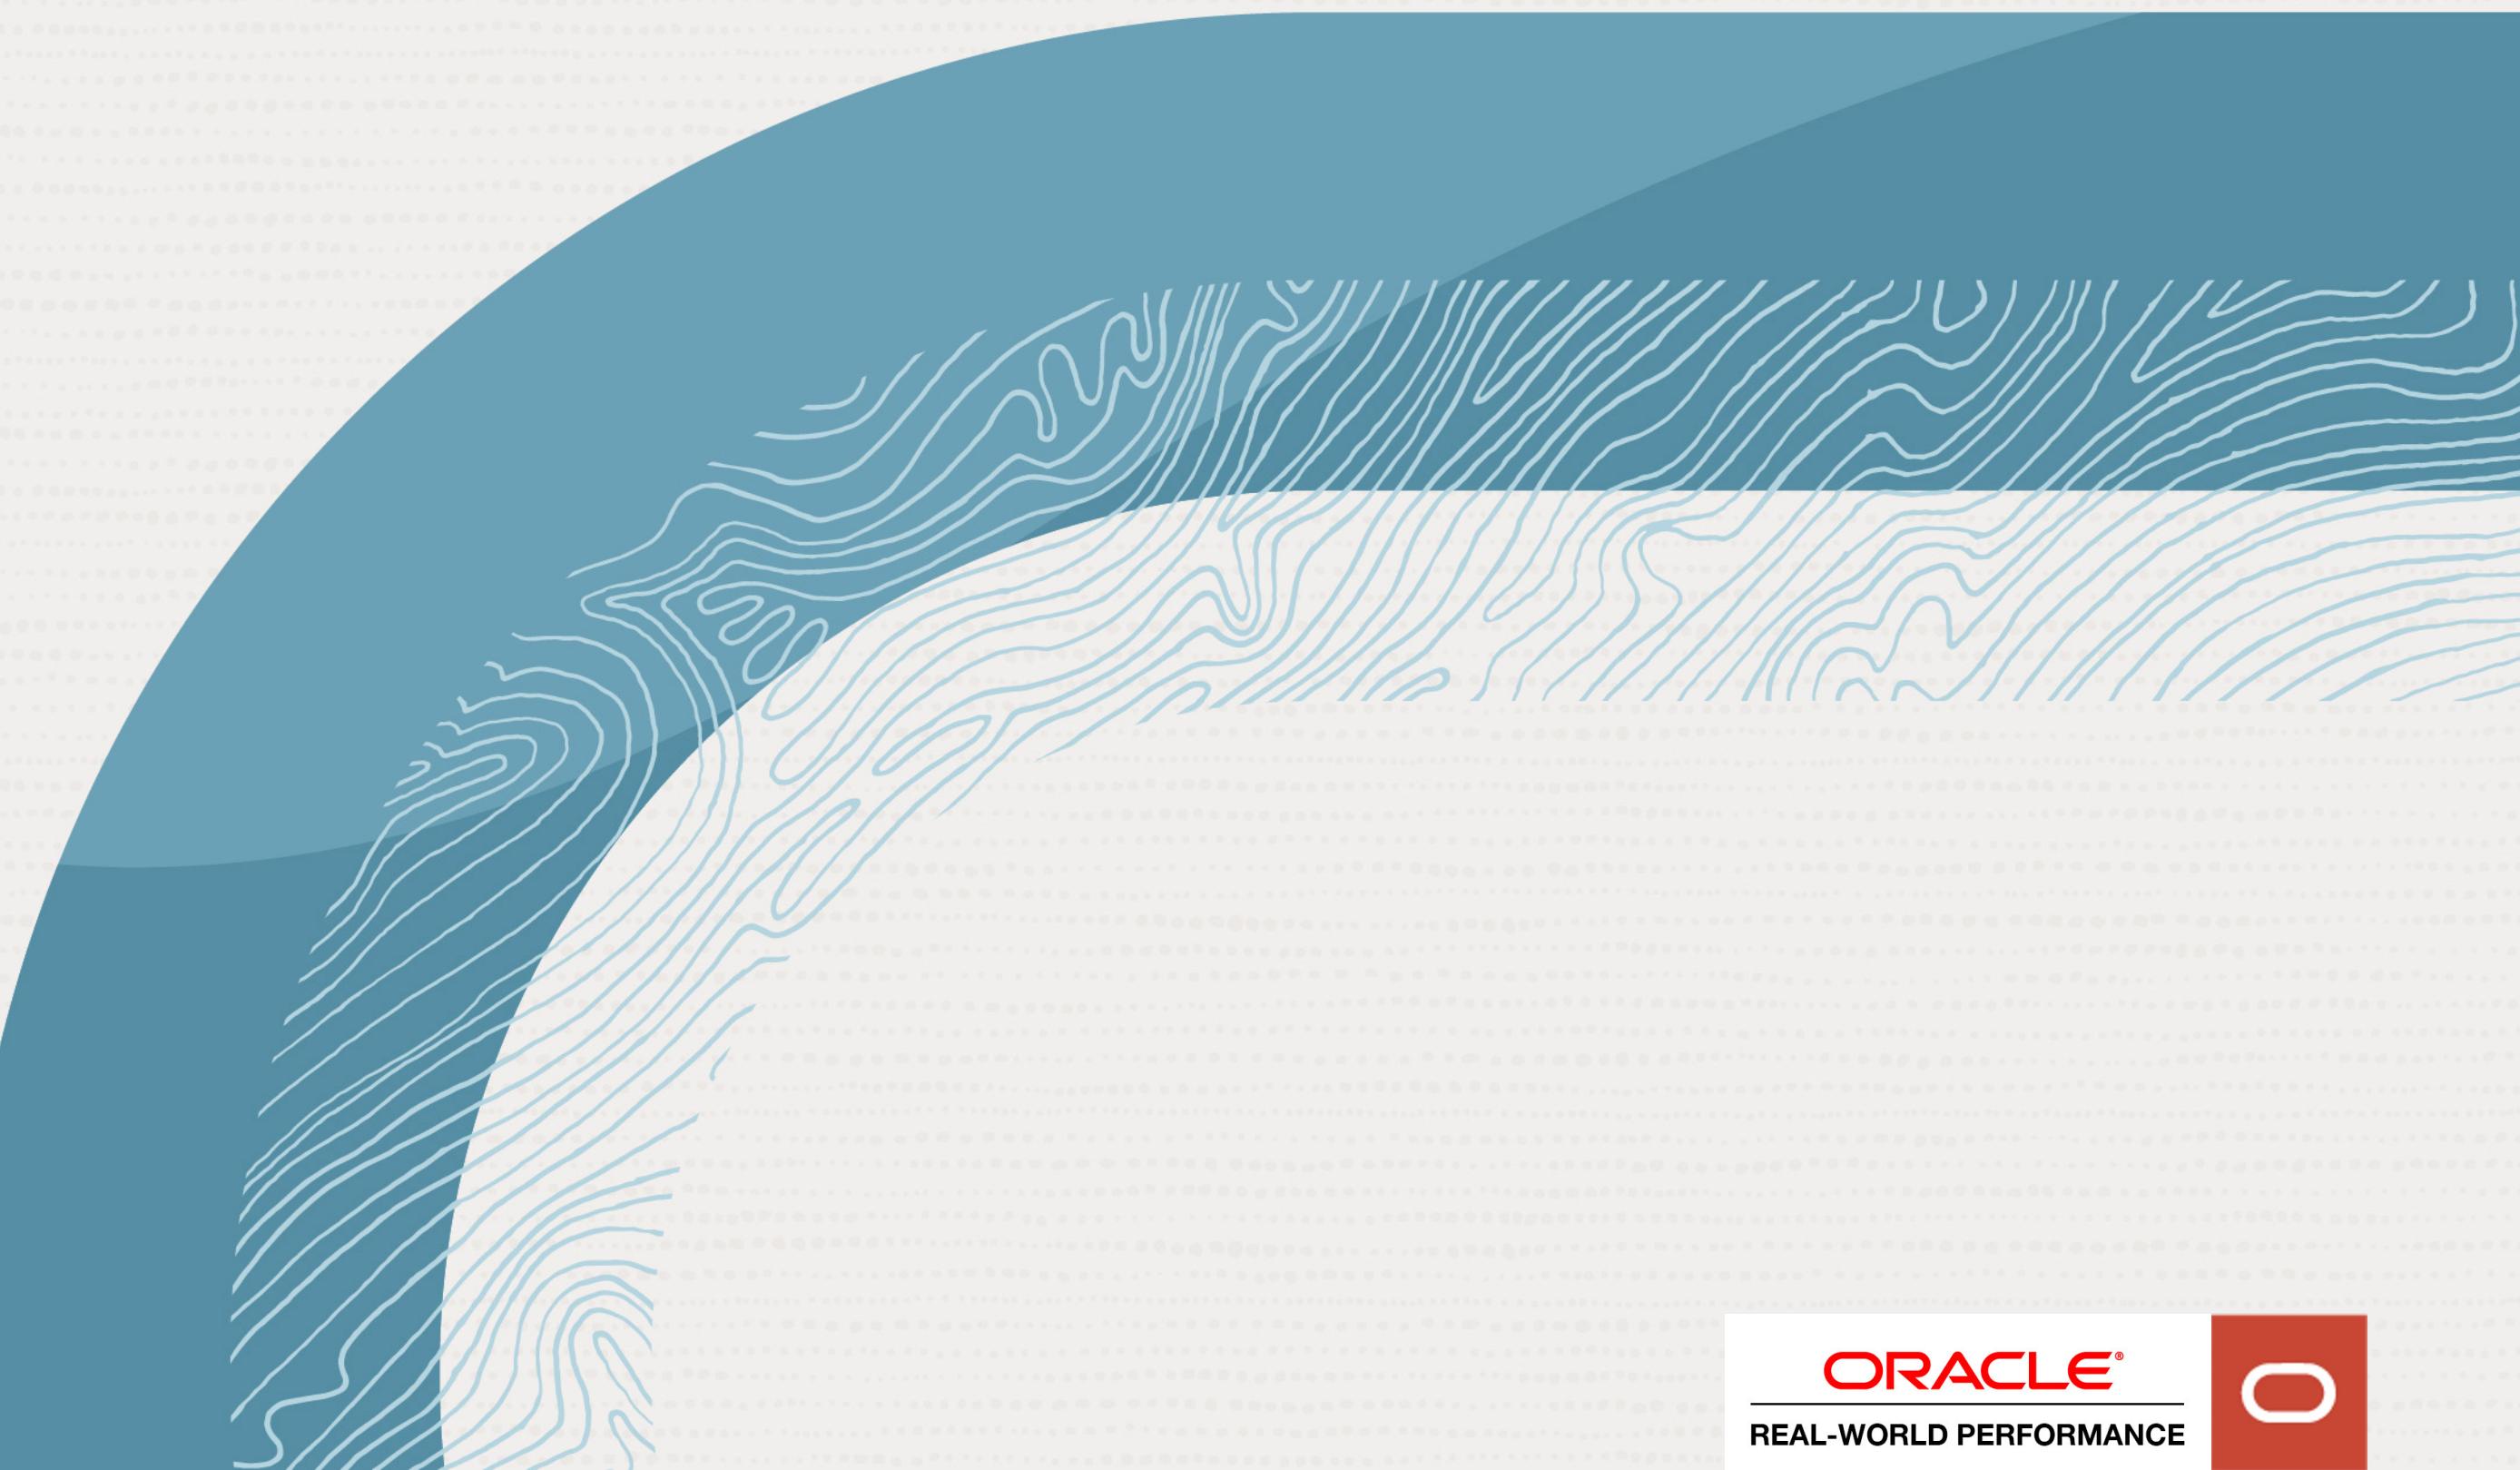


- Statistics
  - Constraints
  - Partitioning
  - Data Types
- ... all of them are equally important

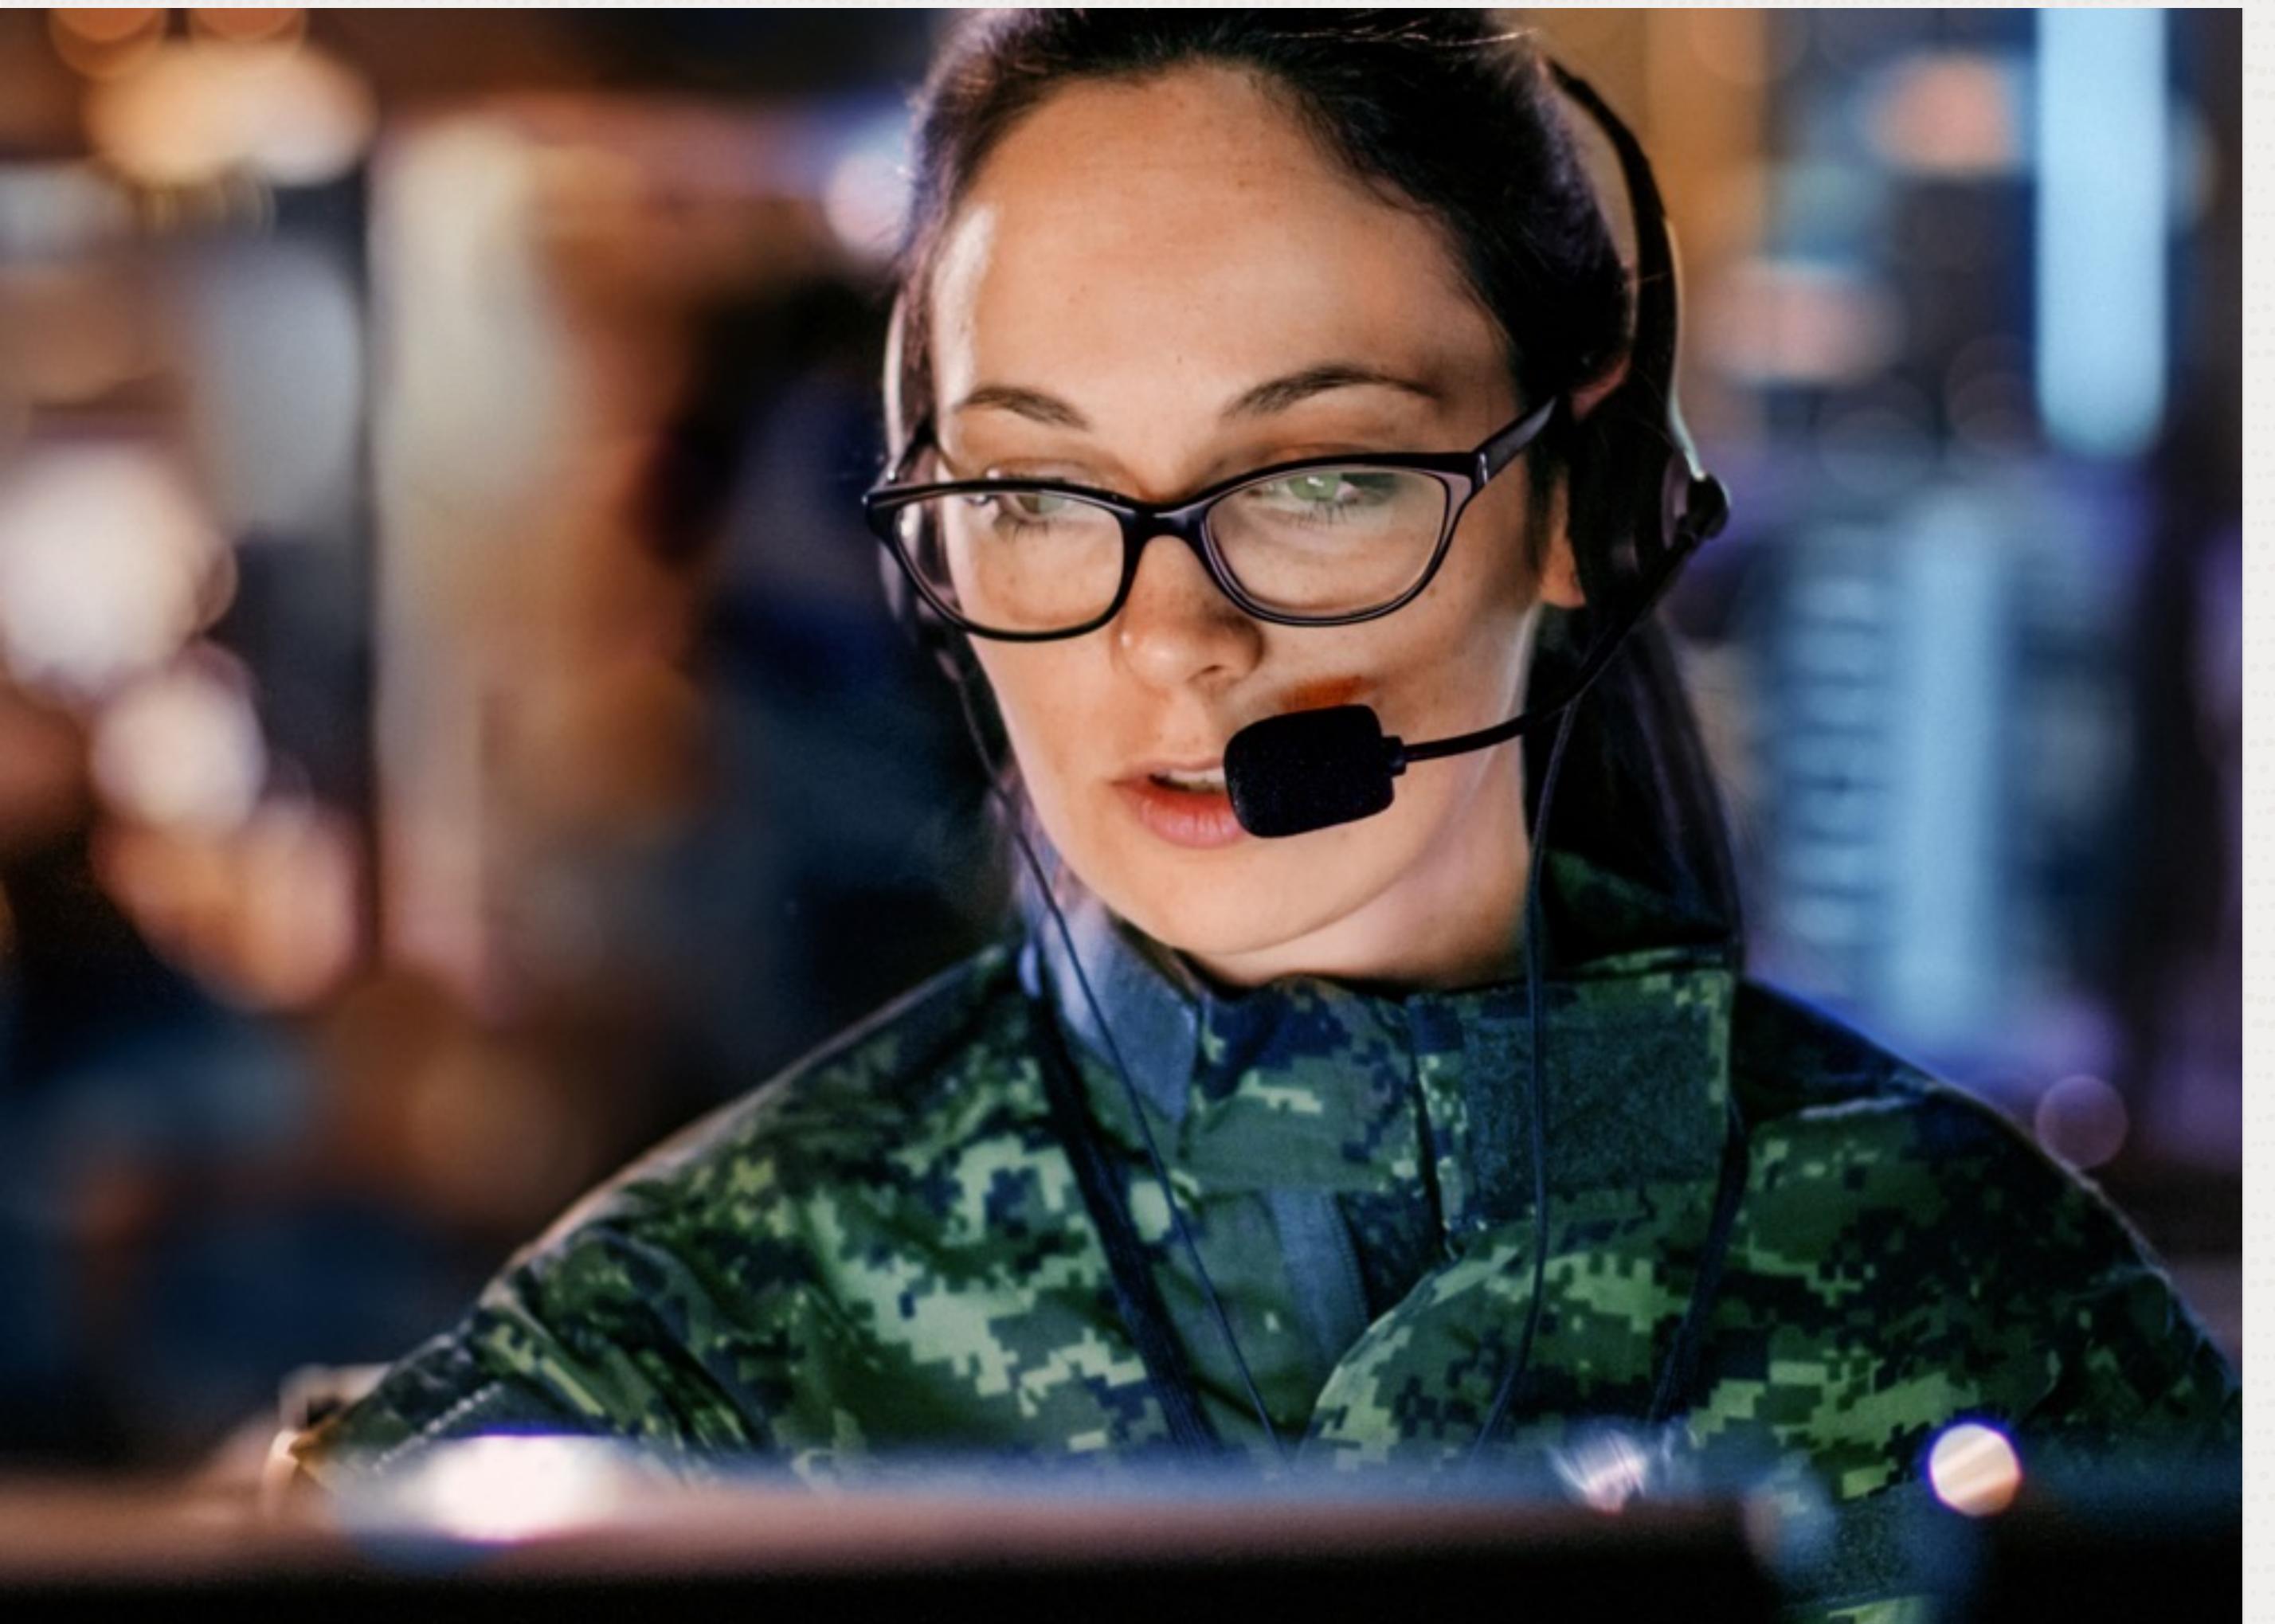
# Star Query Execution

Worked example: From minutes to seconds

What you must do?



# Scenario



A star query completes in about 5 minutes.

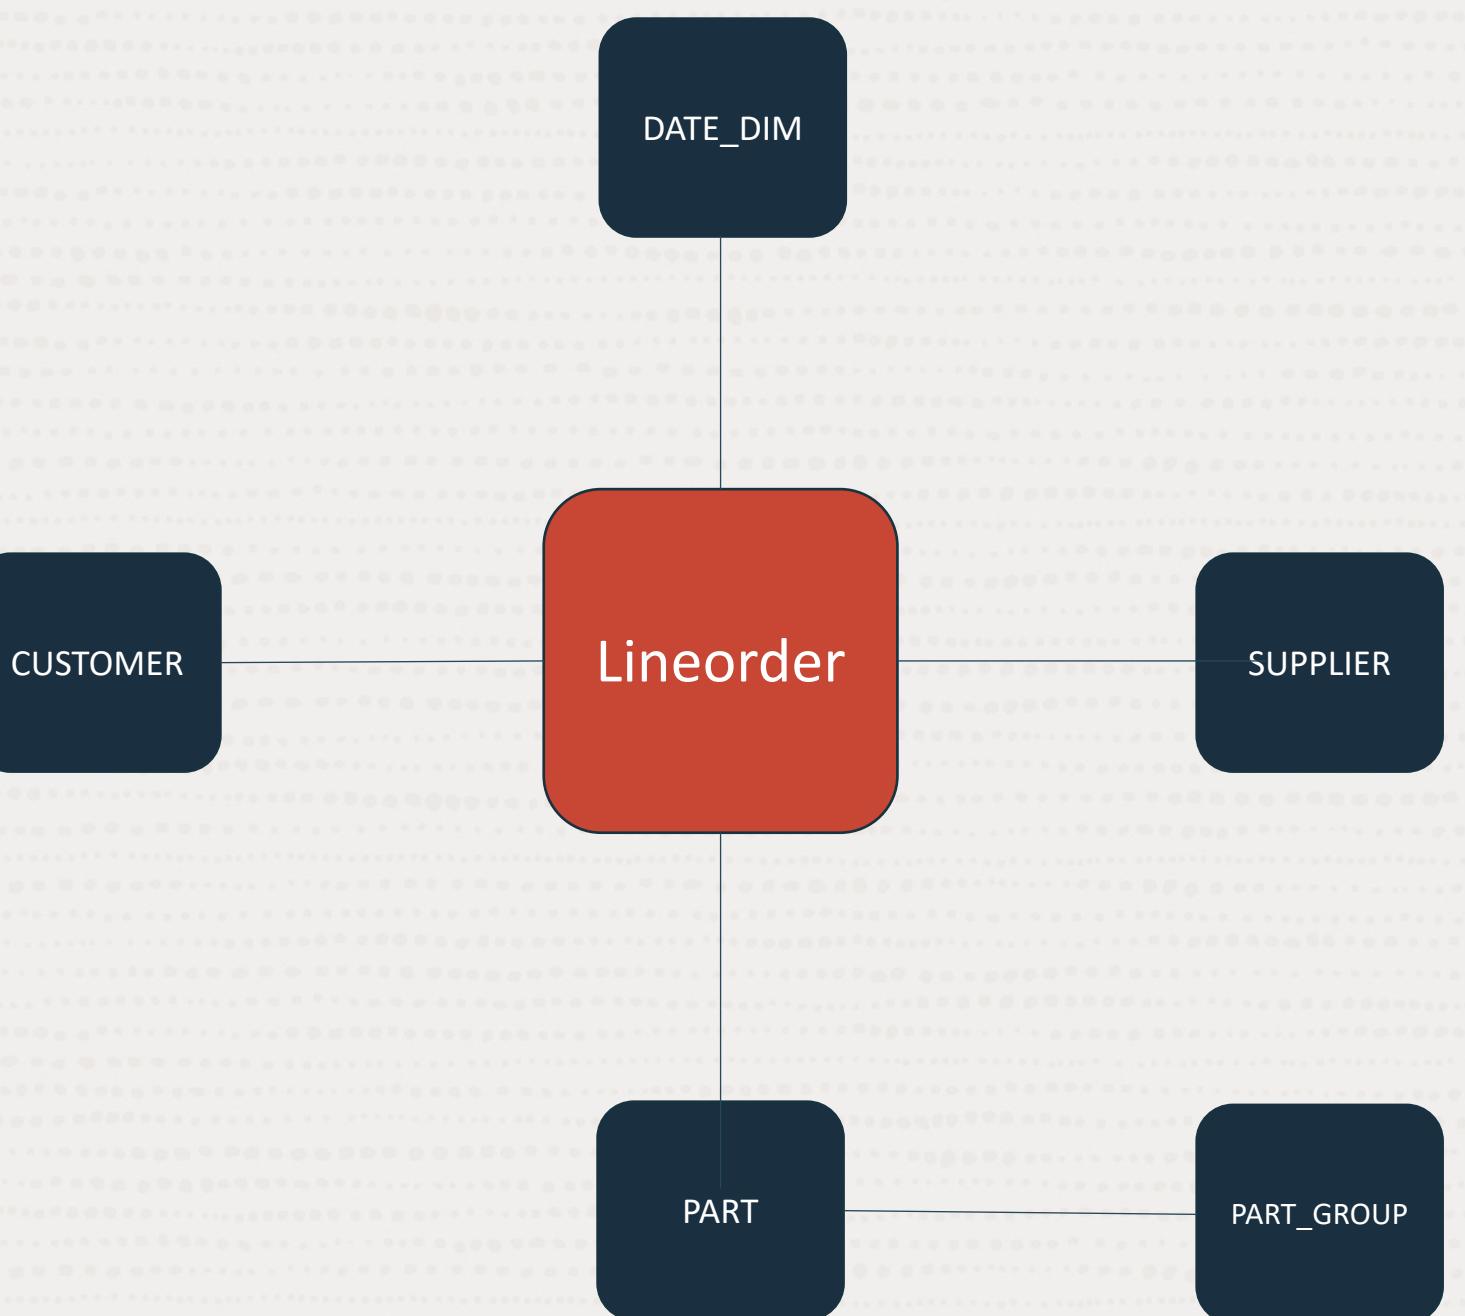
Goal: It needs to complete in 1s or better

Our task:

Optimize the query to complete in 1s or better.  
Query modifications are not allowed.

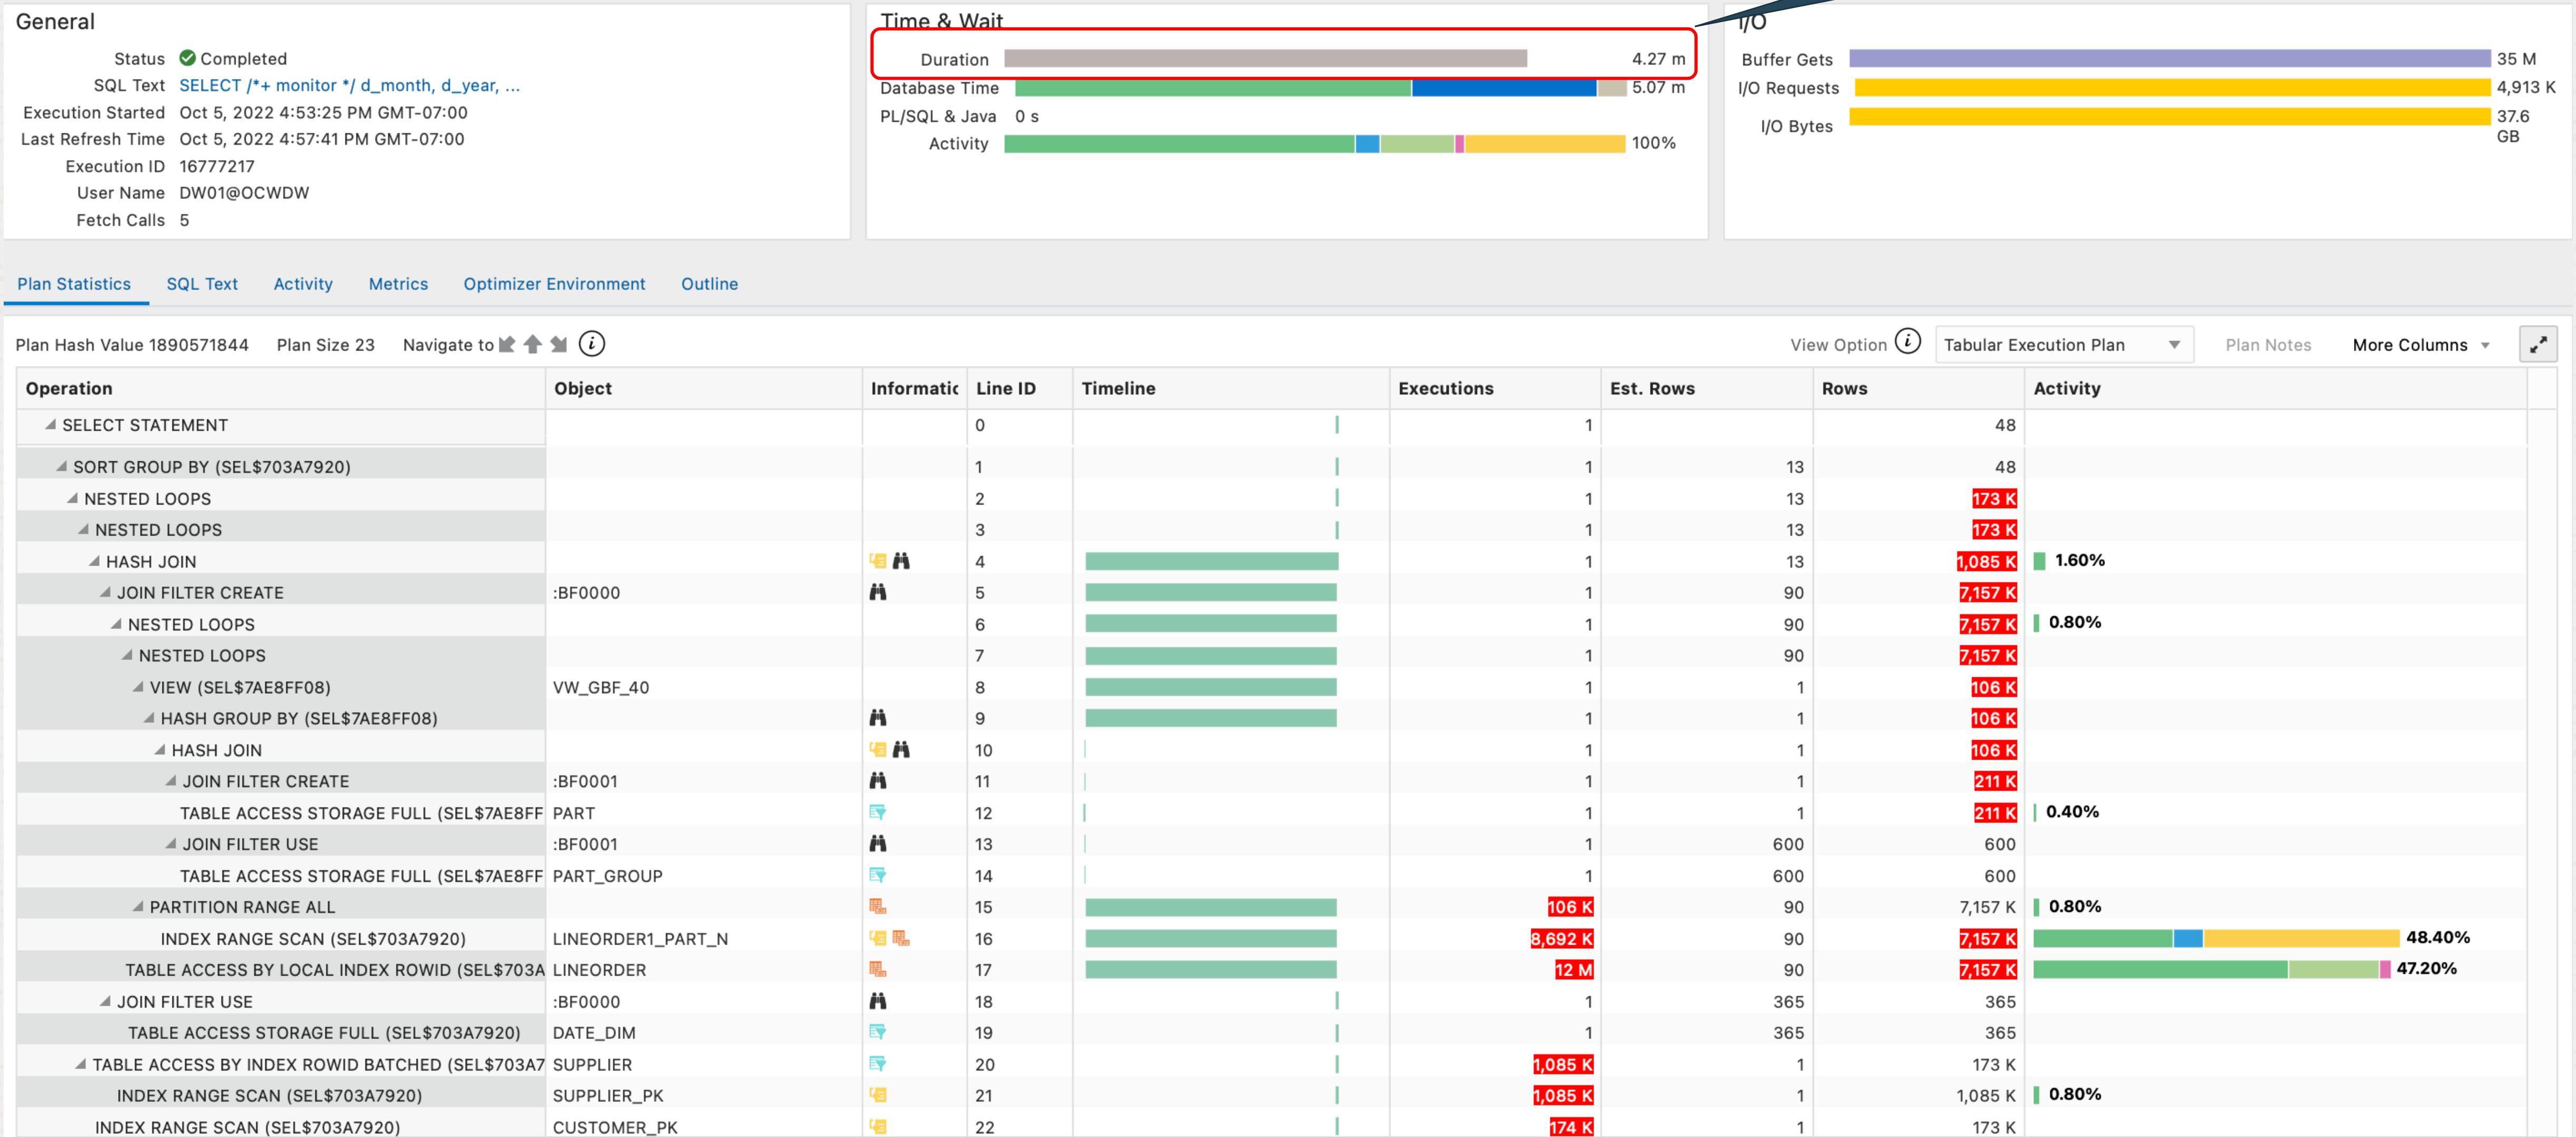
# Our Query

```
SELECT d_month, d_year, s_nation,  
       SUM(lo_quantity) lo_quantity,  
       SUM(lo_revenue) lo_revenue,  
       SUM(lo_supplycost) lo_supplycost  
  FROM lineorder  
    JOIN customer          ON lo_custkey = c_custkey  
    JOIN date_dim           ON lo_orderdate = d_datekey  
    JOIN part                ON lo_partkey = p_partkey  
    JOIN part_group          ON p_mfgr=pg_mfgr  
                             AND p_category=pg_category  
                             AND p_brand1=pg_brand1  
    JOIN supplier            ON lo_suppkey = s_suppkey  
 WHERE d_year = 2007  
   AND s_region = 'ASIA'  
   AND s_nation in ('CHINA','INDIA','JAPAN','VIETNAM')  
   AND pg_type in (1,2,3)  
   AND p_size<12  
 GROUP BY d_month, d_year, s_nation  
 ORDER BY d_month, d_year, s_nation;
```



Initial running time: 256 seconds

# Initial performance

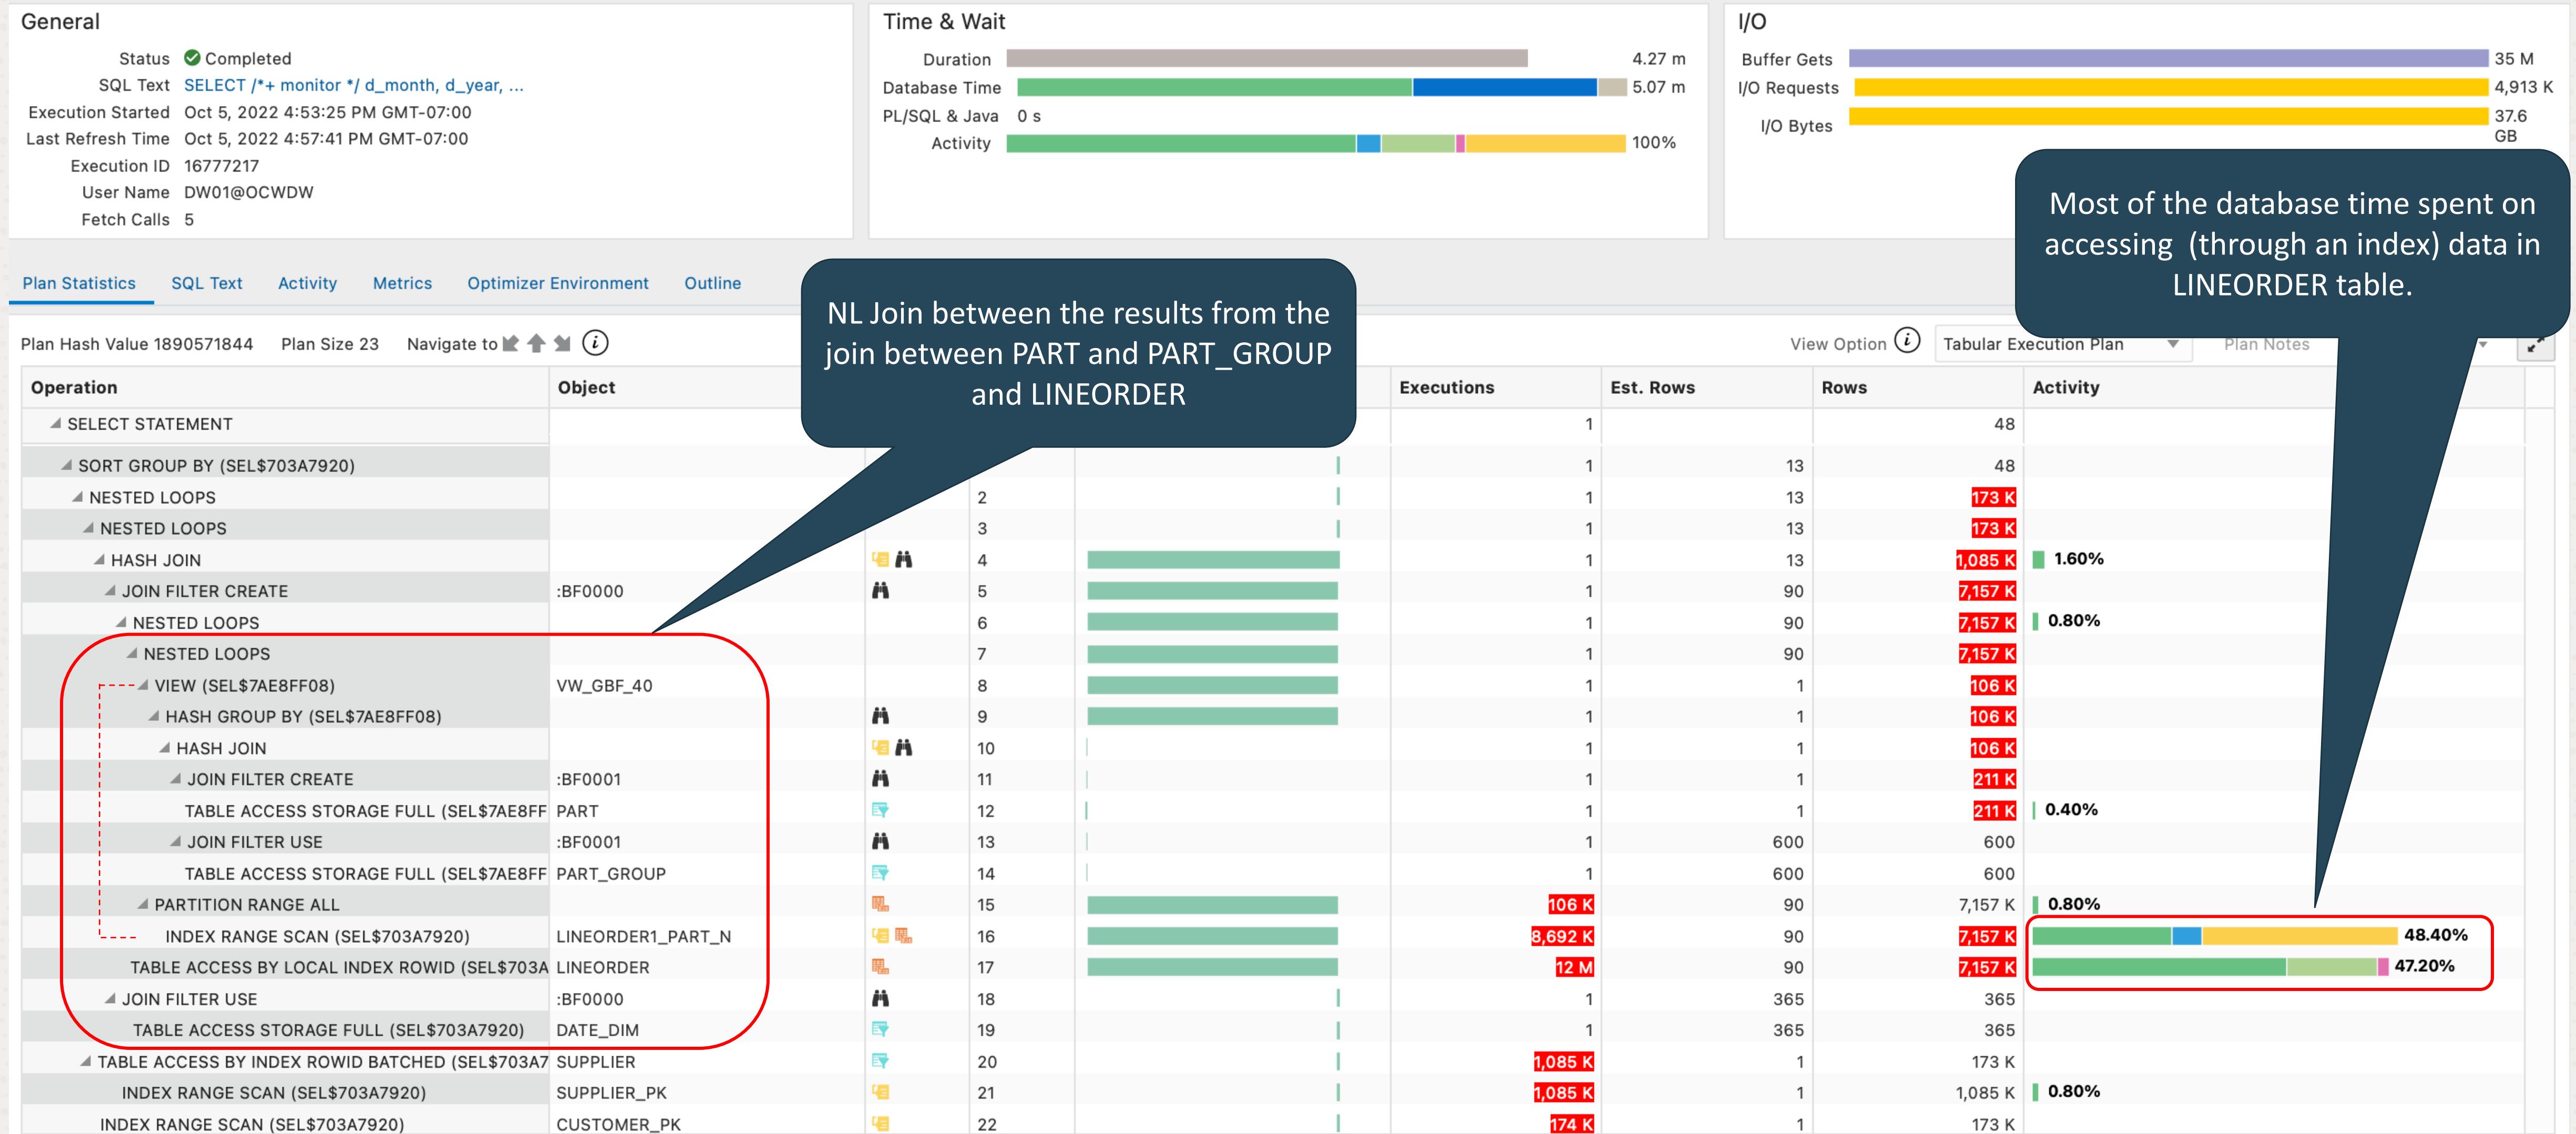


# Worked example: Progress

	<b>Fix Category</b>	<b>Running Time</b>
Initial Performance		256 seconds



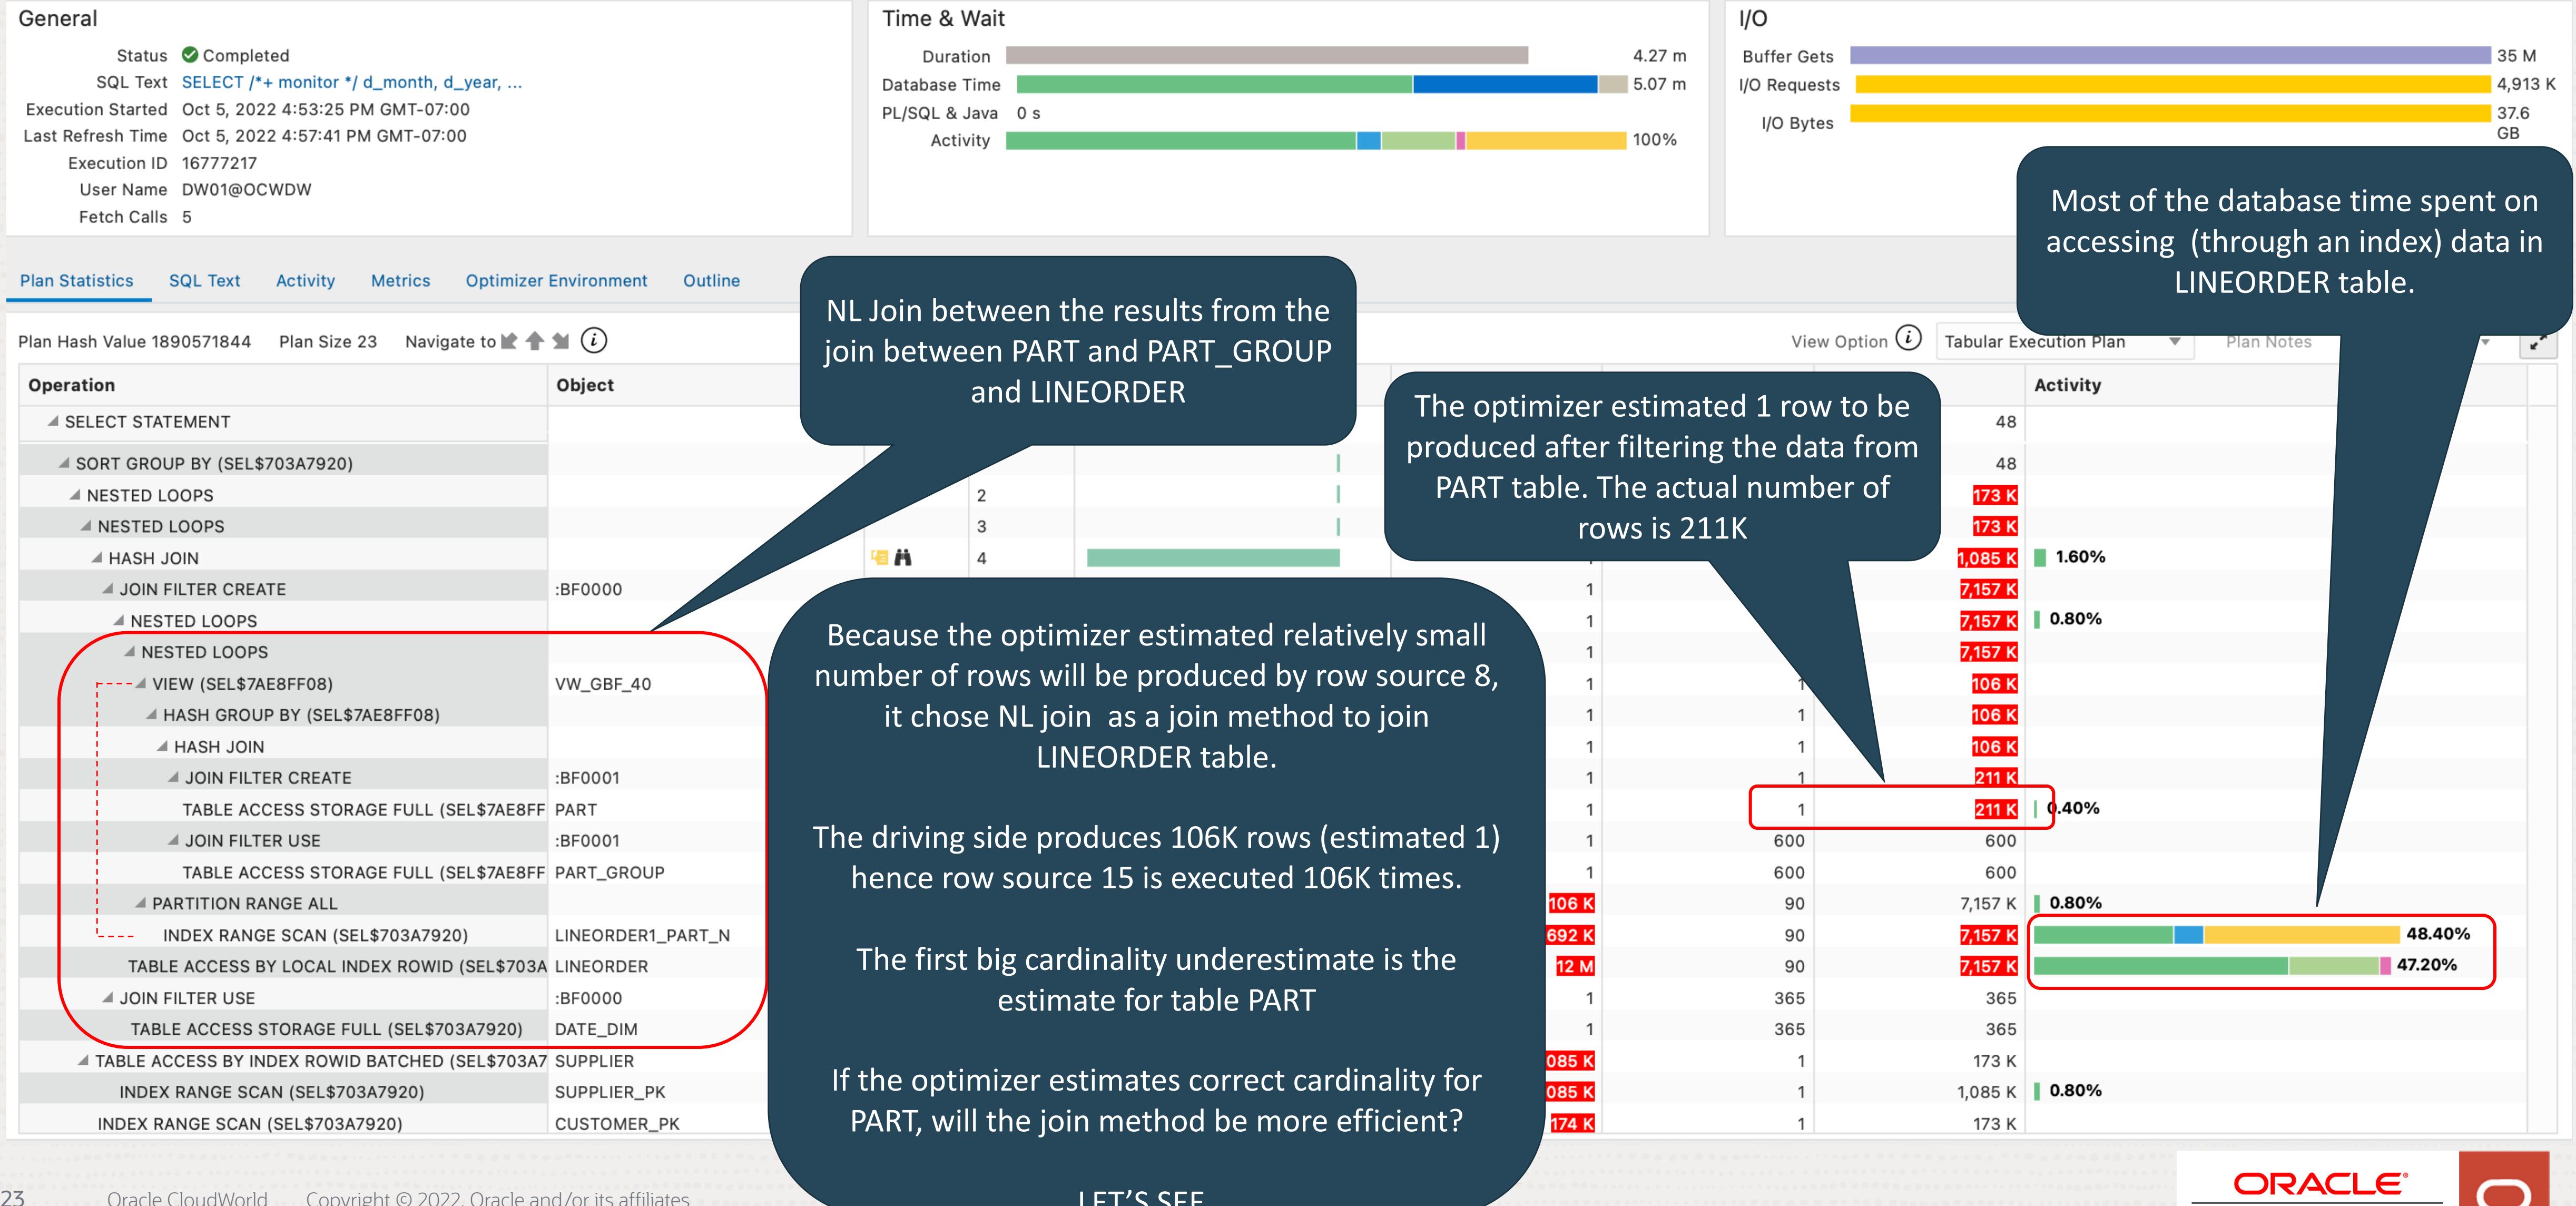
# Initial performance



Most of the database time spent on processing (through an index) data in LINEORDER table.



# Initial performance



# What You Must Do: Statistics

Think about **skew**

Think about **correlation**

Do not rely on **Dynamic Statistics** alone

Think about **how** and **when** to gather statistics

Statistics must be **representative**



# Step 1: Fix the cardinality estimate for PART

Fix the cardinality estimate for table PART (estimate 1, actual 211K)

- When addressing a problem with single table cardinality estimates, the first place to look is the state of the table statistics. Are they representative?

```
SELECT table_name, num_rows FROM user_tables WHERE table_name='PART';
```

TABLE_NAME	NUM_ROWS
PART	1

Statistics for table PART indicate there is 1 row in the table

```
SELECT COUNT(*) FROM part;
```

COUNT(*)
959616

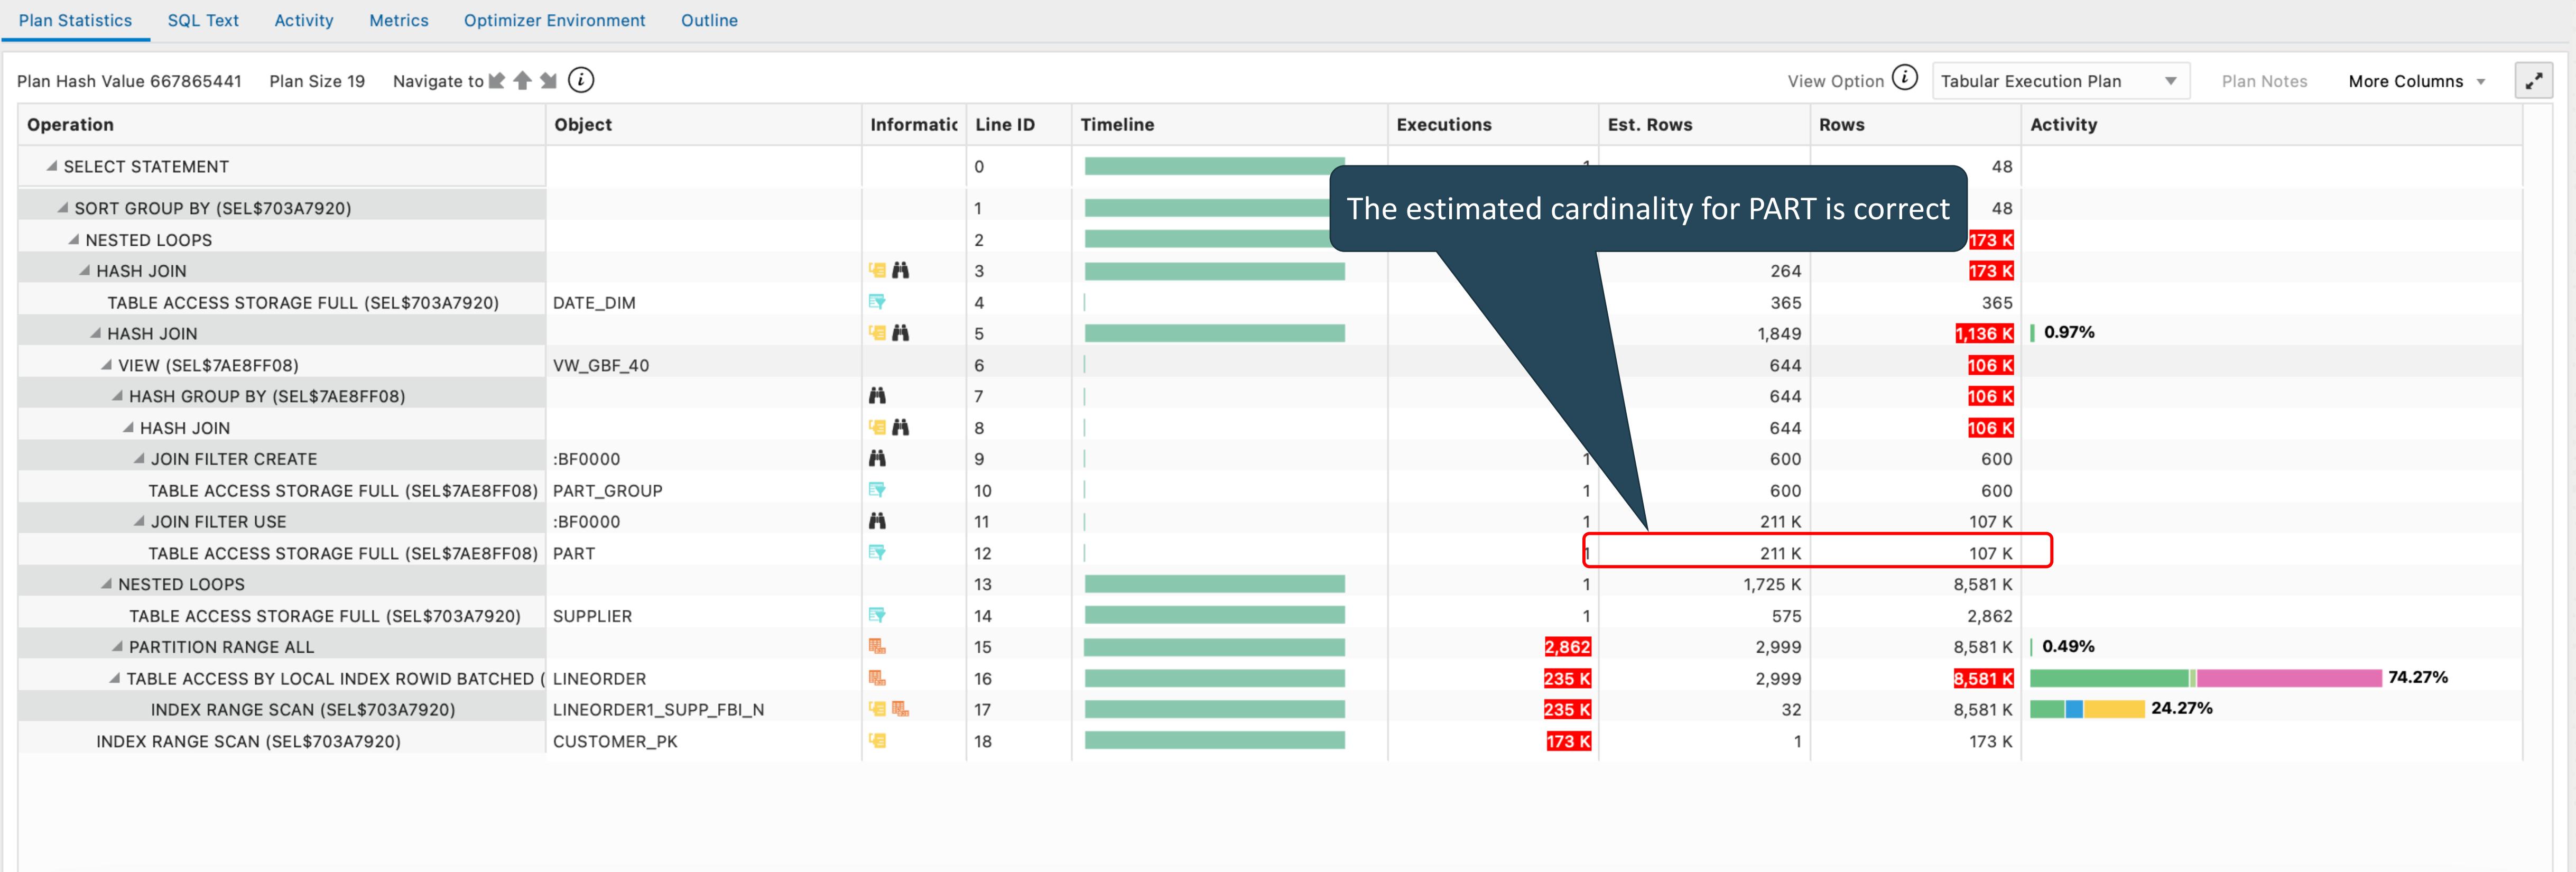
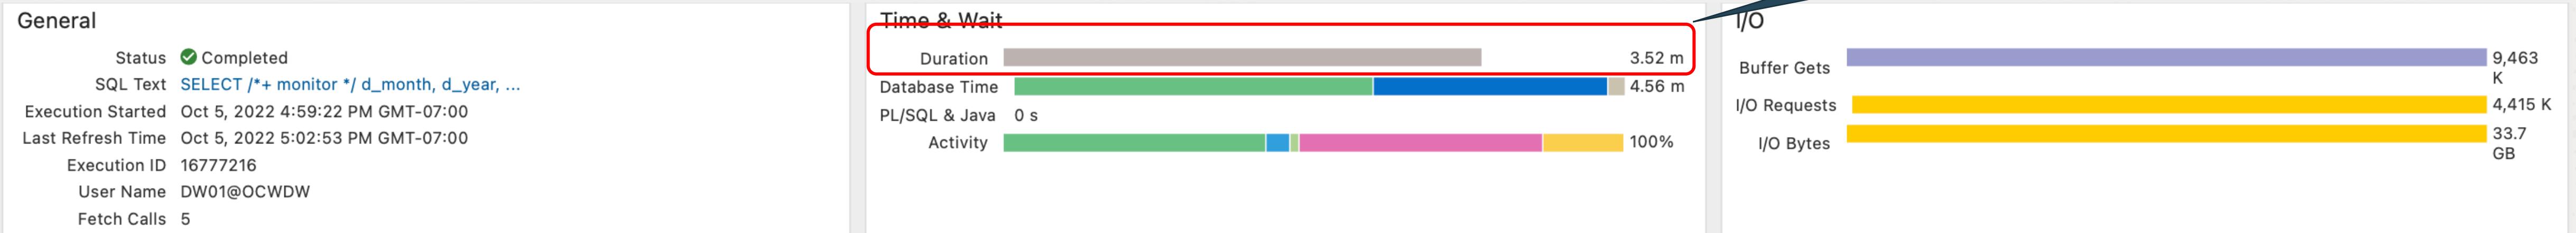
The actual number of rows in table PART is 959,616

```
EXEC dbms_stats.gather_table_stats(ownname => user,tabname => 'PART');
```

Gather statistics for PART table

Step1: running time: 211 seconds

# Step 1: Statistics on PART

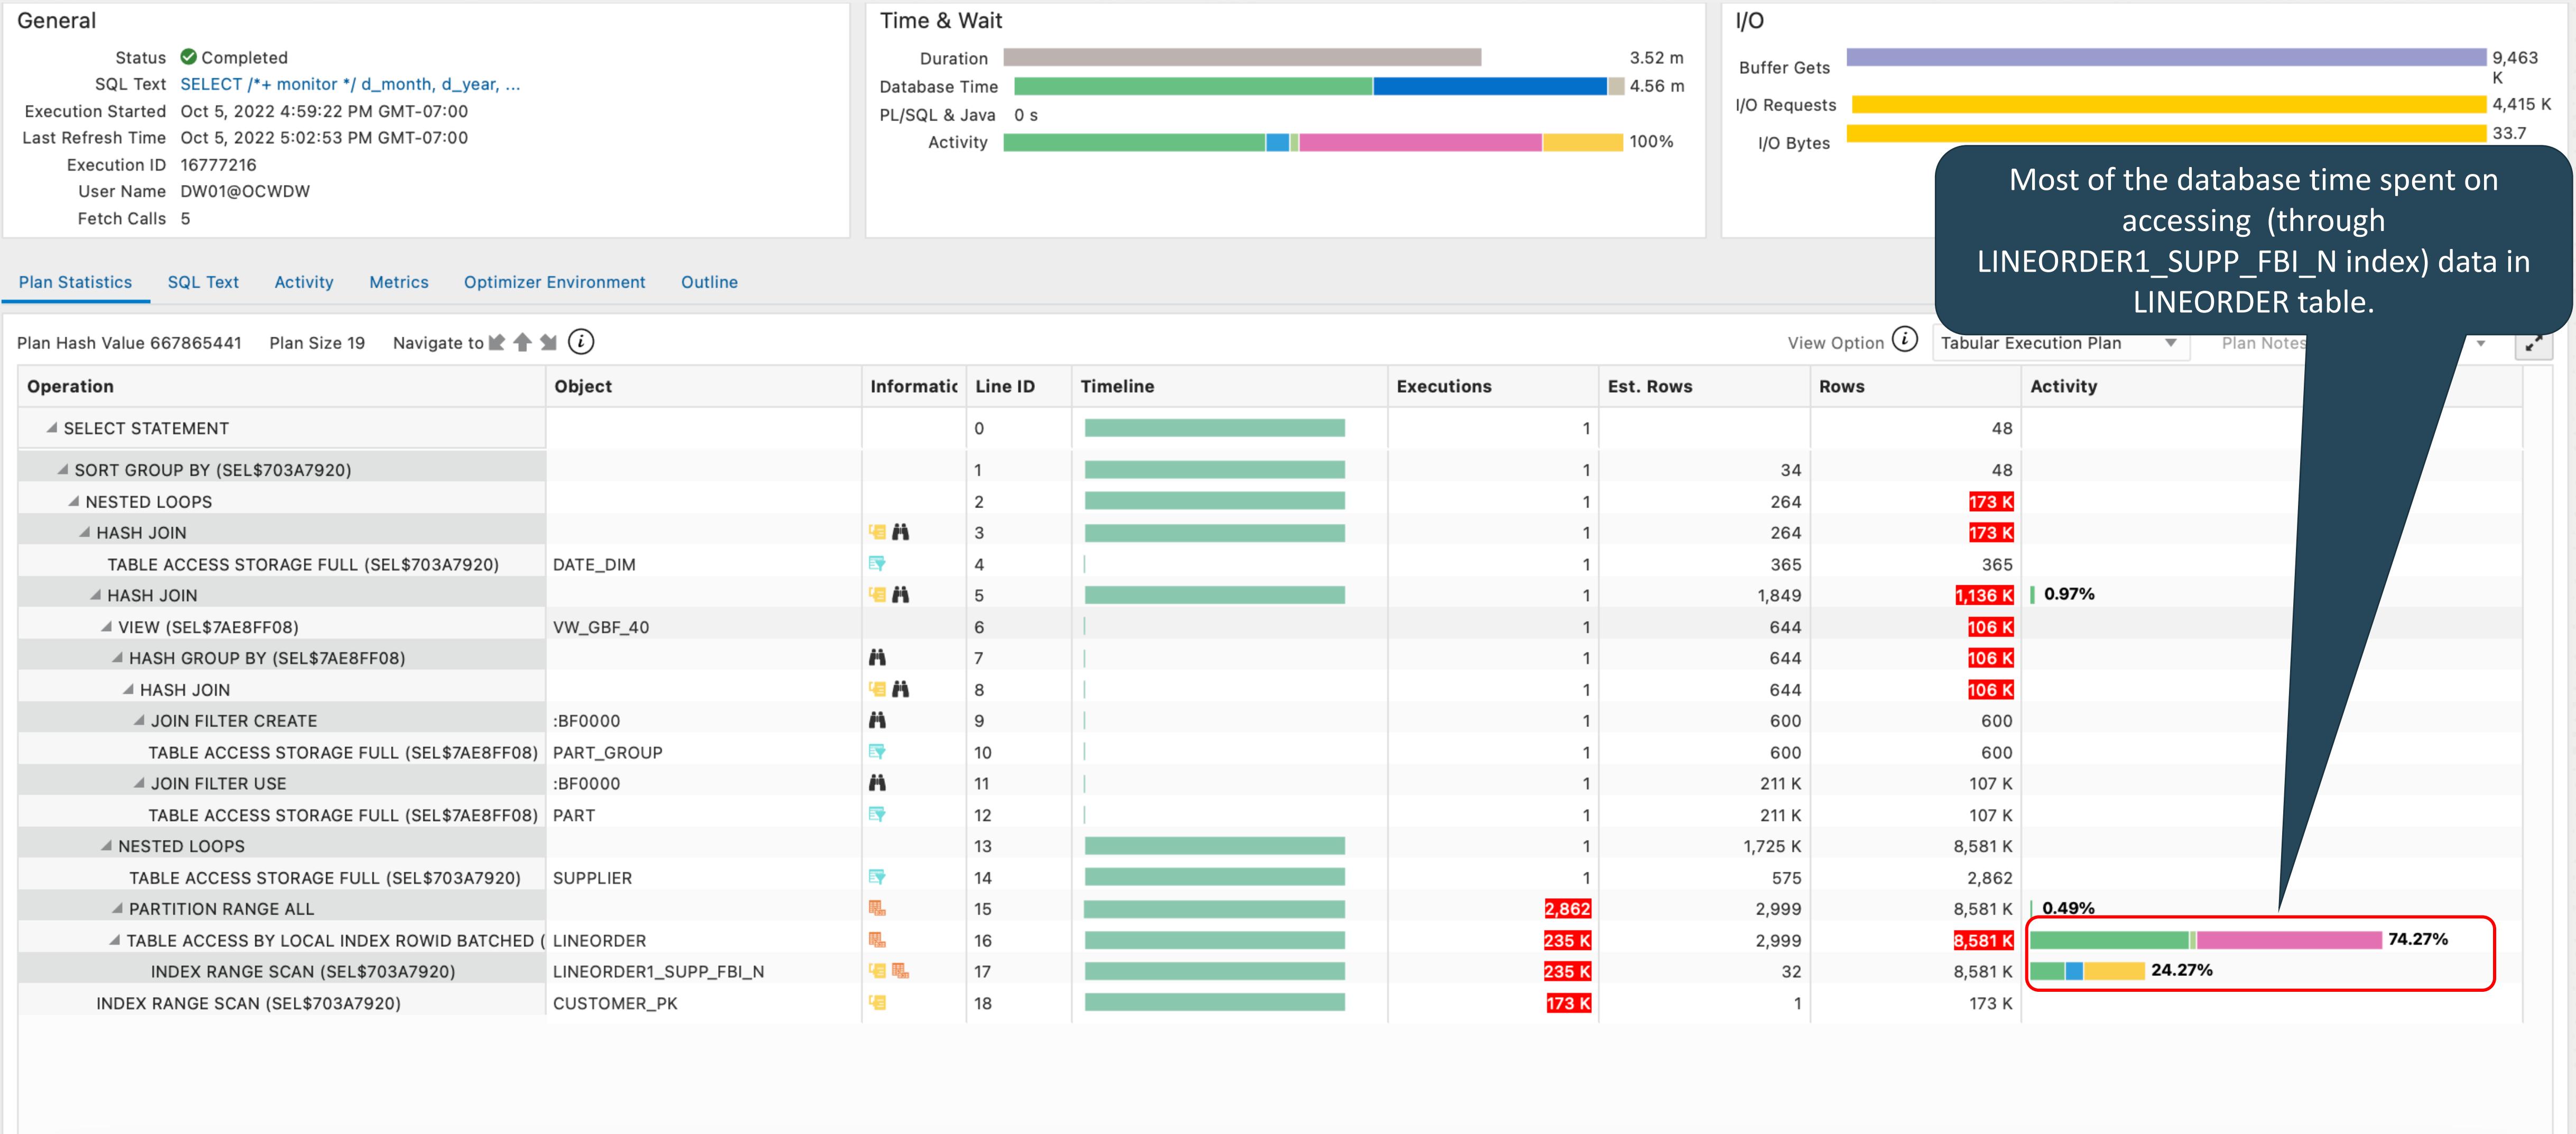


# Worked example: Progress

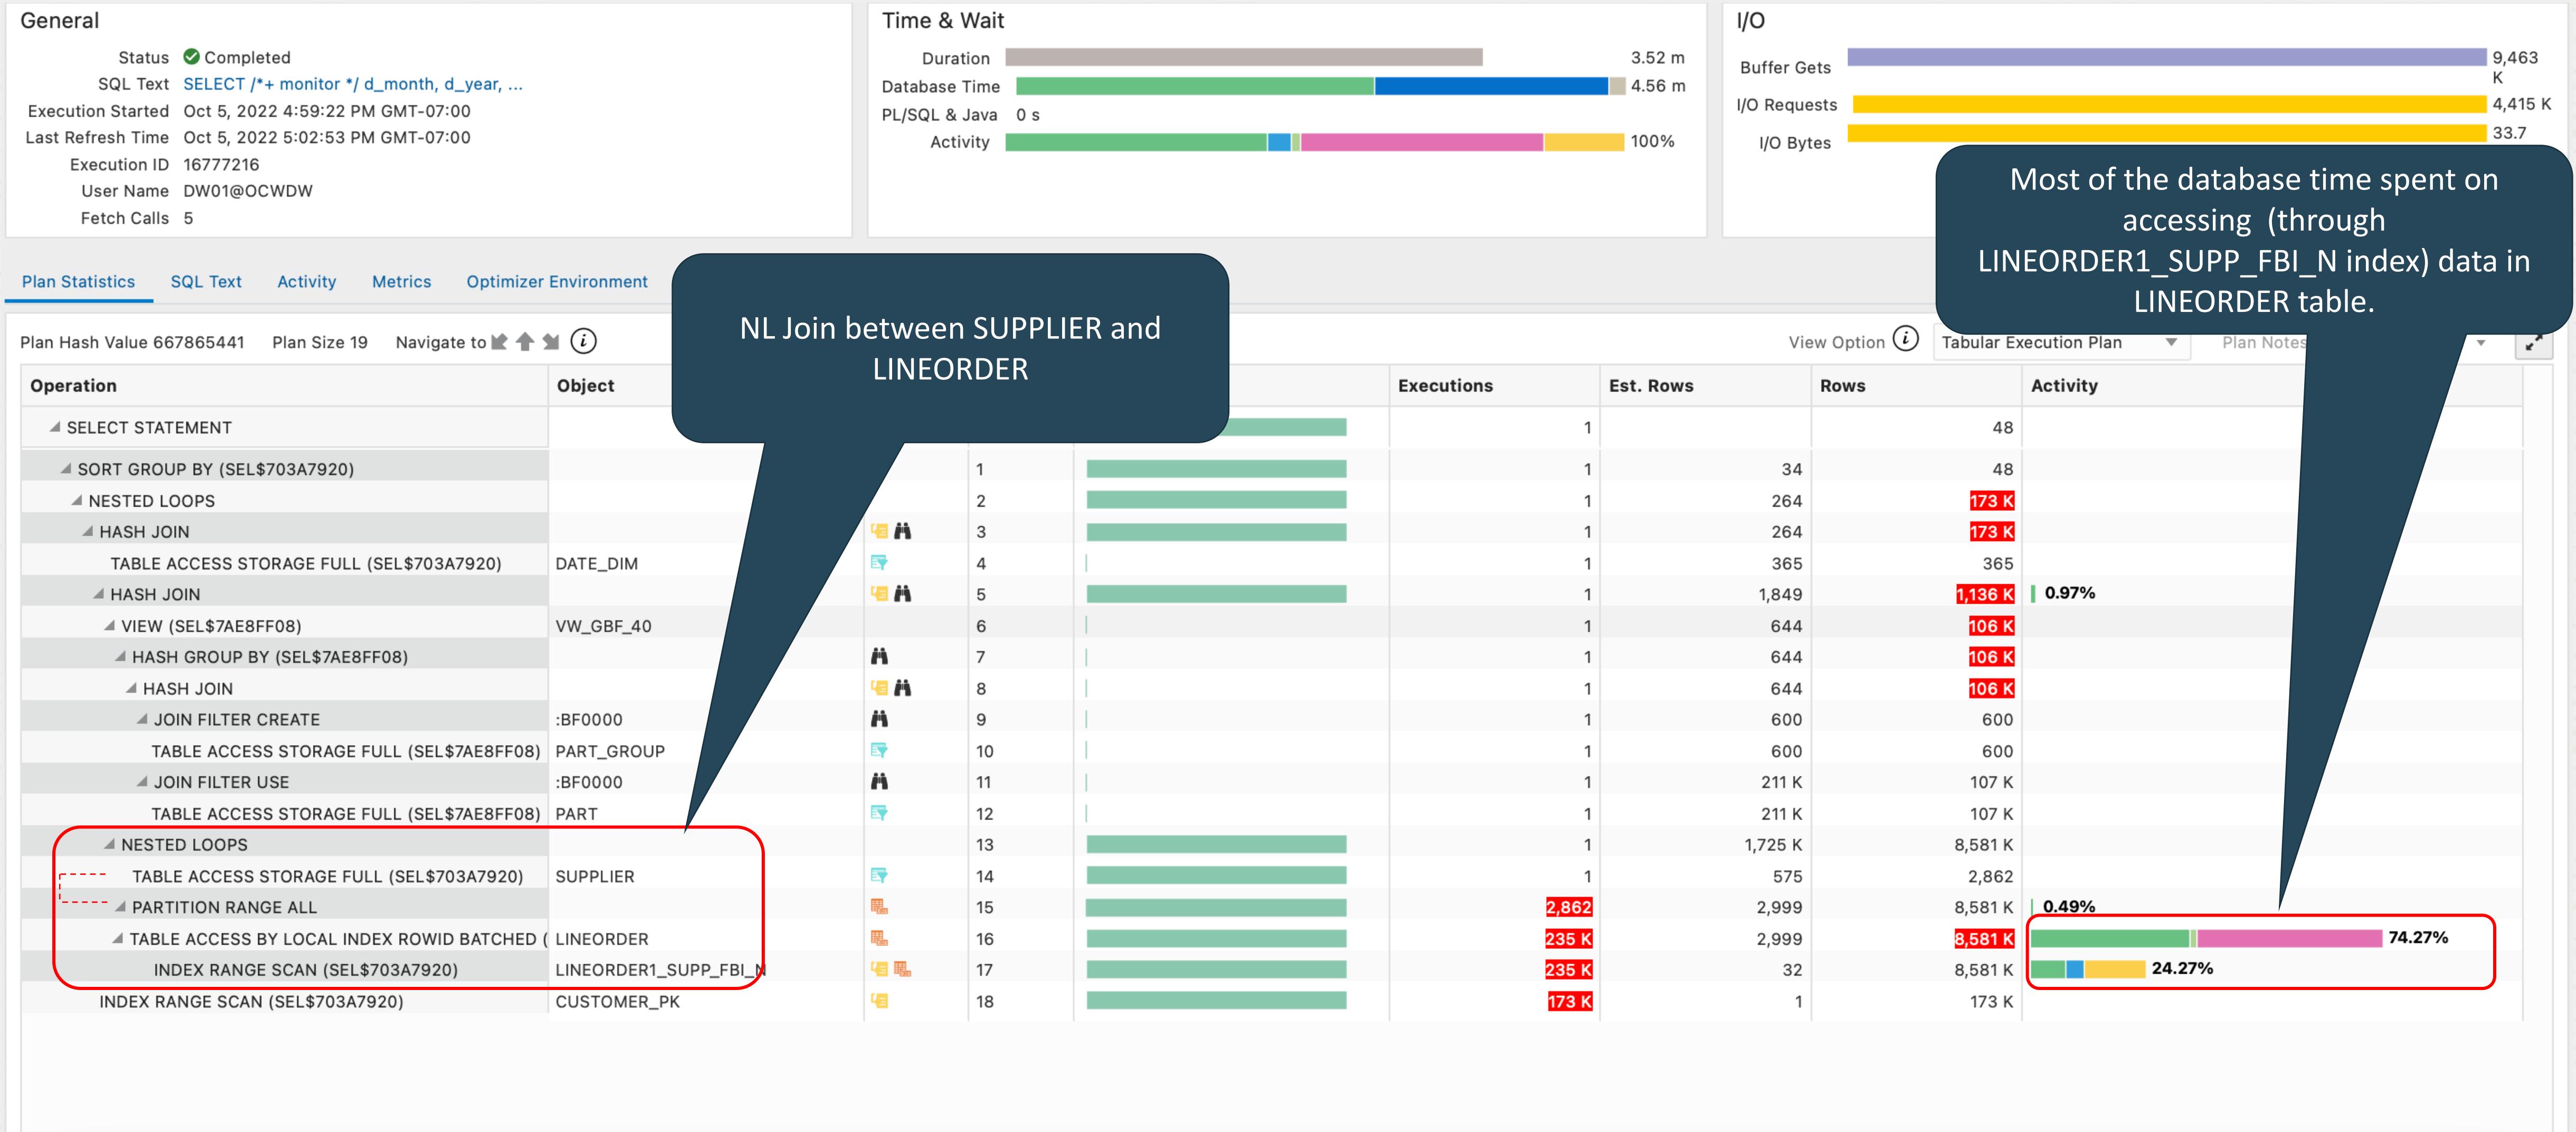
	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds



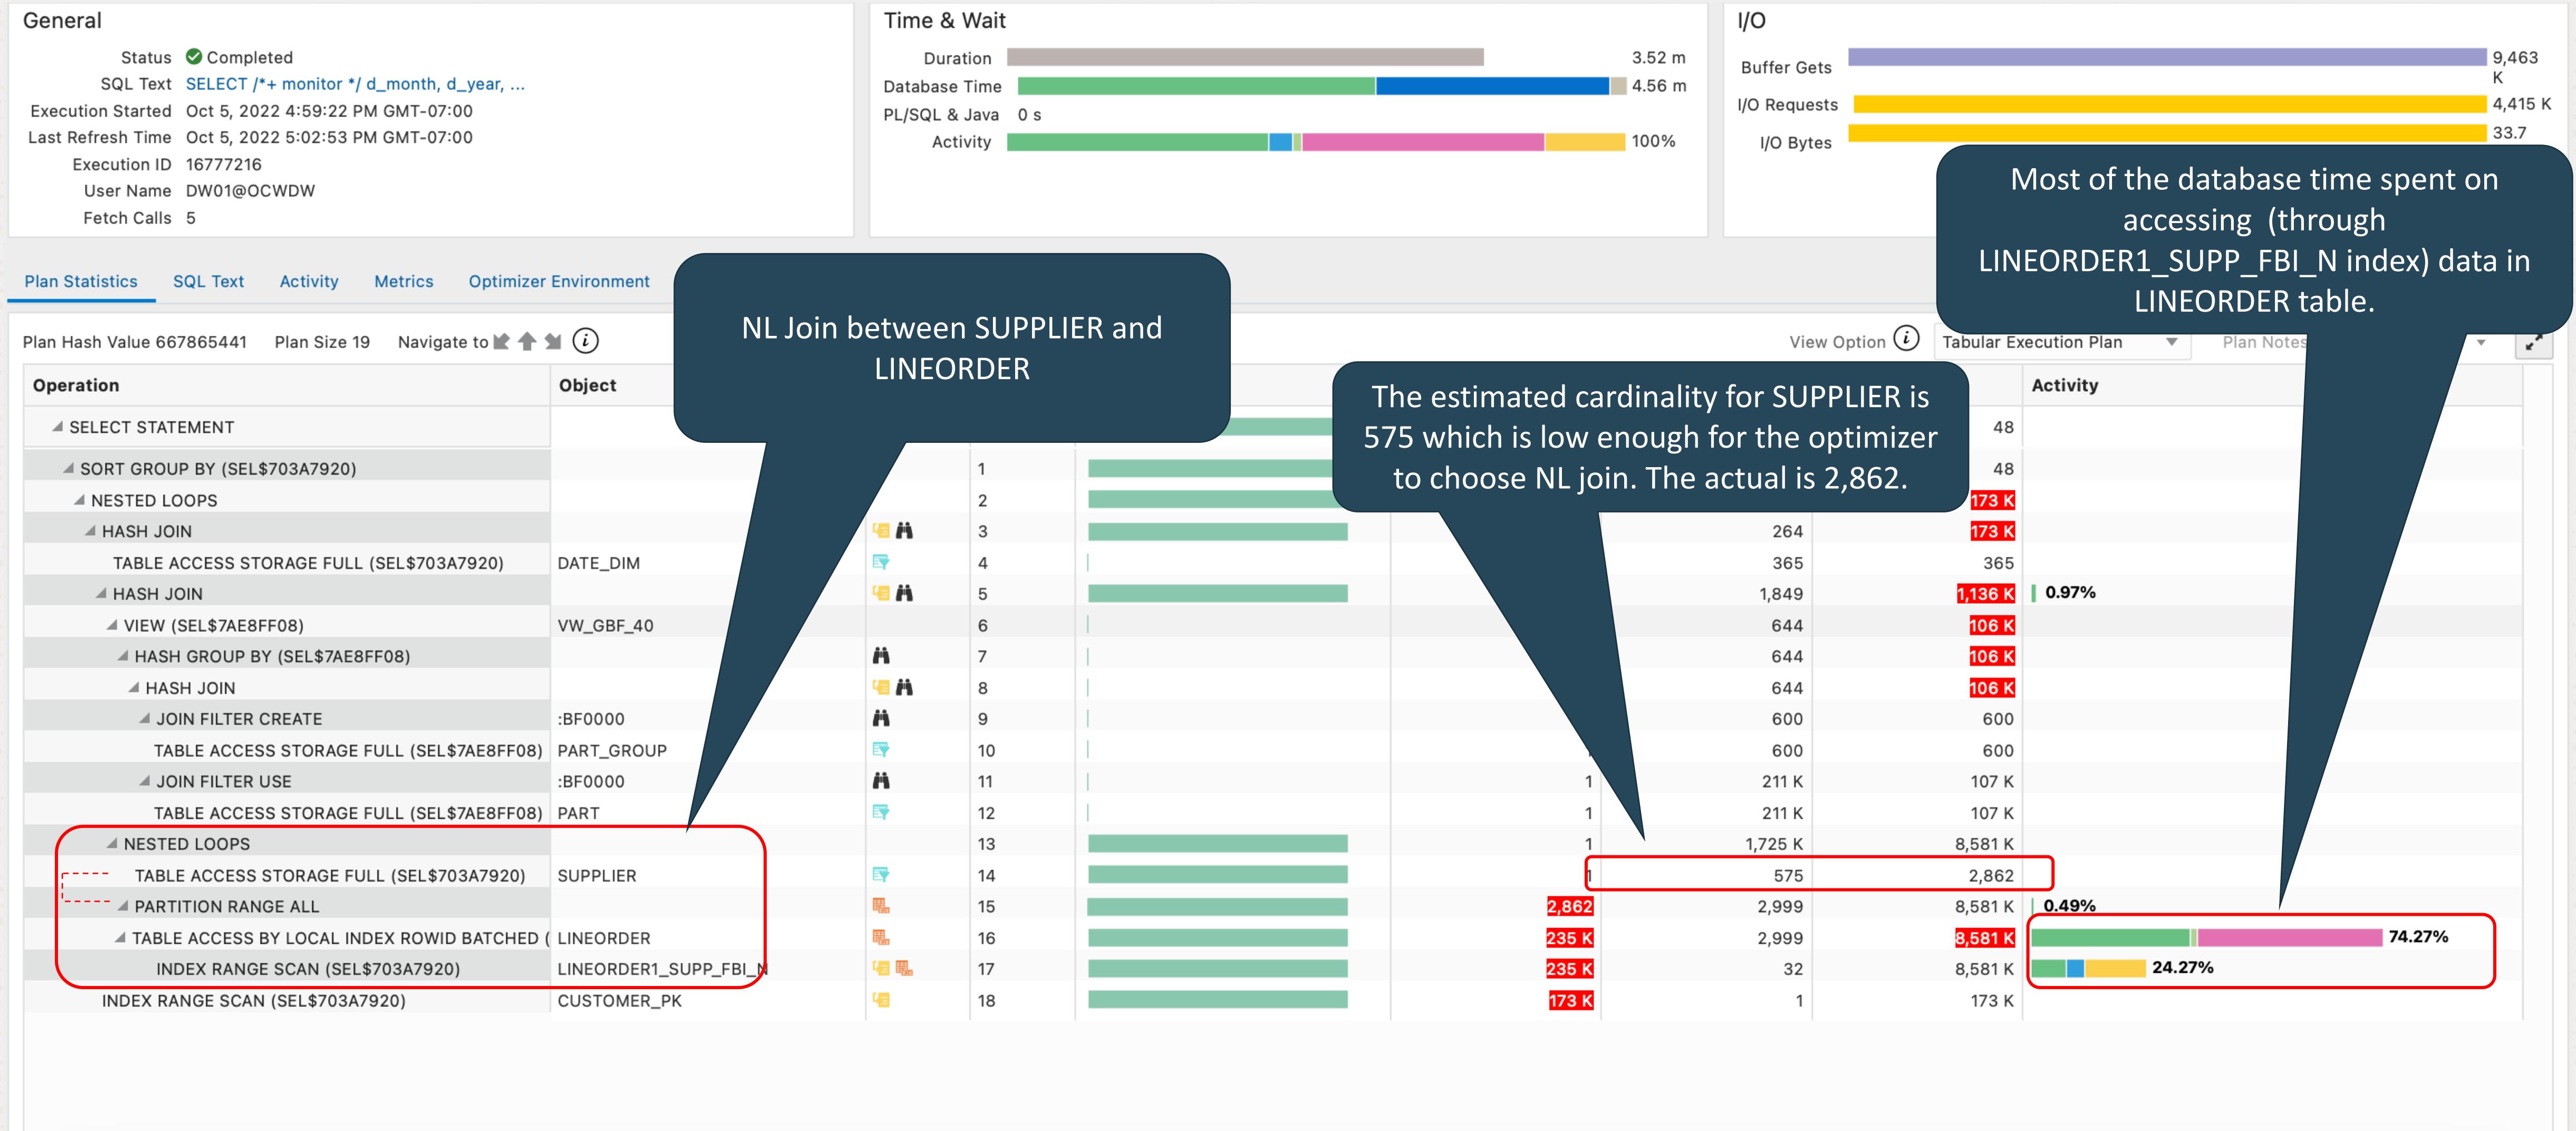
# Step 1: Statistics on PART



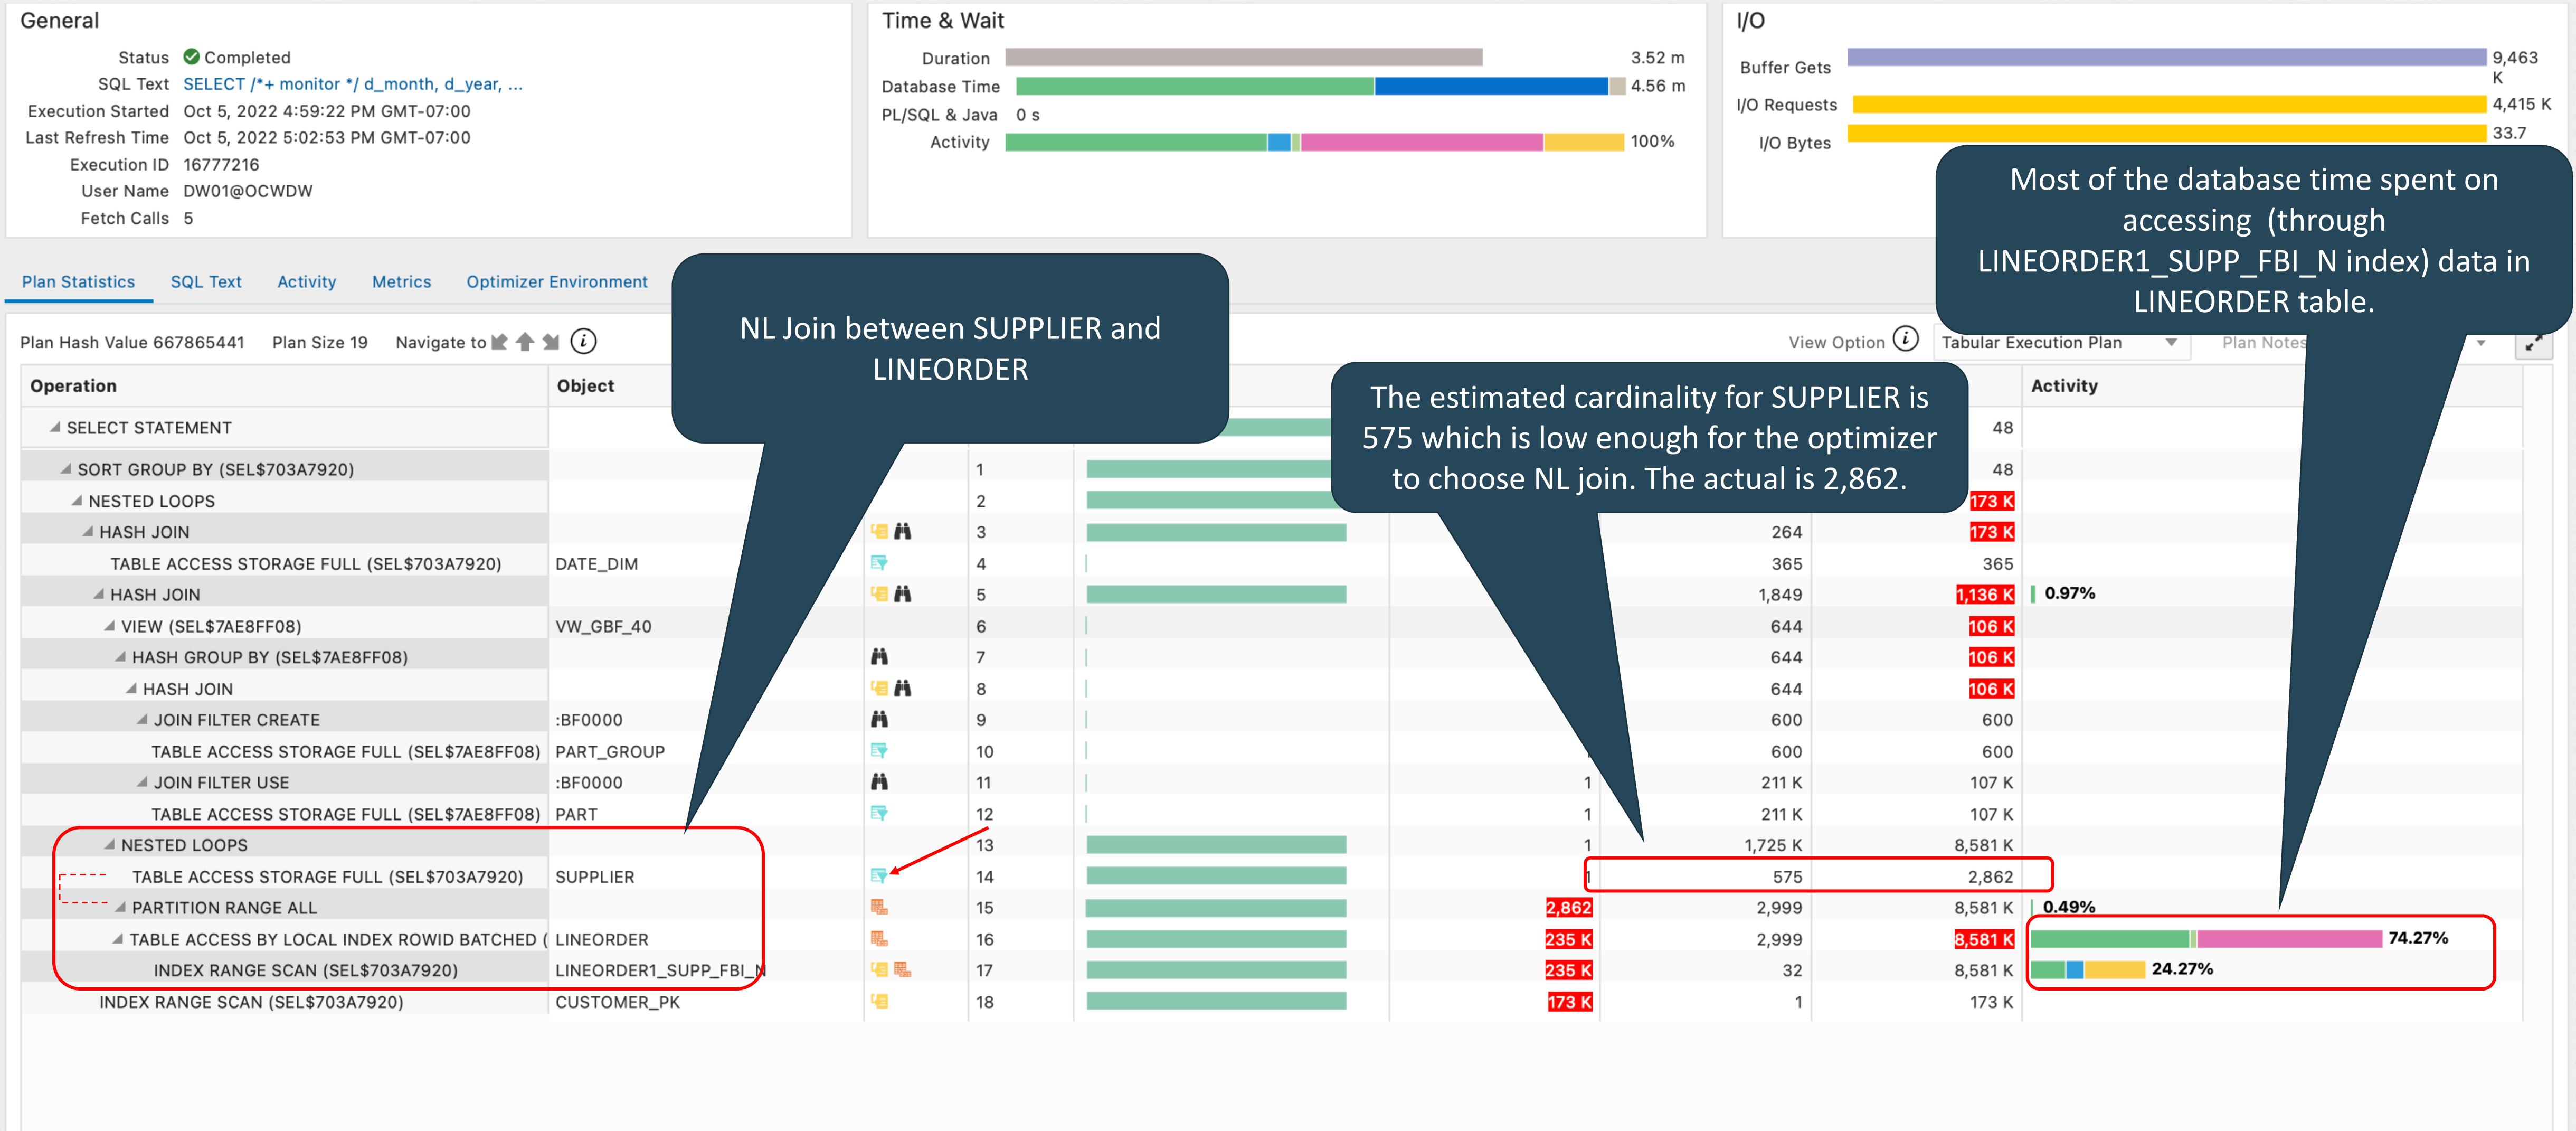
# Step 1: Statistics on PART



# Step 1: Statistics on PART



# Step 1: Statistics on PART



# Step 1: Statistics on PART

**General**

- Status: Completed
- SQL Text: `SELECT /*+ monitor */ d_month, d_year, ...`
- Execution Started: Oct 5, 2022 4:59:22 PM GMT-07:00
- Last Refresh Time: Oct 5, 2022 5:02:53 PM GMT-07:00
- Execution ID: 16777216
- User Name: DW01@OCWDW
- Fetch Calls: 5

**Time & Wait**

Duration	3.52 m
Database Time	4.56 m
PL/SQL & Java	0 s
Activity	100%

**I/O**

Buffer Gets	9,463 K
I/O Requests	4,415 K 33.7 GB

The estimated cardinality represents the number of rows that optimizer expects to be returned from SUPPLIER after the filter predicates are applied:

```
s_region = 'ASIA'  
AND s_nation in ('CHINA', 'INDIA', 'JAPAN', 'VIETNAM')
```

**Plan Statistics** **SQL Text** **Activity** **Metrics** **Optimizer Environment** **Outline**

Plan Hash Value 667865441 Plan Size 19 Navigate to

Operation	Object	Informatic	Line ID	Timeline	Executions
SELECT STATEMENT			0		
SORT GROUP BY (SEL\$703A7920)			1		
NESTED LOOPS			2		
HASH JOIN			3		
TABLE ACCESS STORAGE FULL (SEL\$703A7920)	DATE_DIM		4		
HASH JOIN			5		
VIEW (SEL\$7AE8FF08)	VW_GBF_40		6		
HASH GROUP BY (SEL\$7AE8FF08)			7		
HASH JOIN			8		
JOIN FILTER CREATE	:BF0000		9		
TABLE ACCESS STORAGE FULL (SEL\$7AE8FF08)	PART_GROUP		10		
JOIN FILTER USE	:BF0000		11		
TABLE ACCESS STORAGE FULL (SEL\$7AE8FF08)	PART		12		
NESTED LOOPS			13		
TABLE ACCESS STORAGE FULL (SEL\$703A7920)	SUPPLIER		14		
PARTITION RANGE ALL			15		
TABLE ACCESS BY LOCAL INDEX ROWID BATCHED (LINEORDER)			16		
INDEX RANGE SCAN (SEL\$703A7920)	LINEORDER1_SUPP_FBI_N		17		
INDEX RANGE SCAN (SEL\$703A7920)	CUSTOMER_PK		18		

**Filter Predicates**

(`"SUPPLIER"."S_REGION"='ASIA' AND INTERNAL_FUNCTION("SUPPLIER"."S_NATION")'`)

**Close**

# Step 2: Fix the cardinality estimate for SUPPLIER

Fix the cardinality estimate for table SUPPLIER (estimated 575, actual 2,862)

- When addressing a problem with single table cardinality estimates, the first place to look is the state of the table statistics. Are they representative?

```
SELECT table_name, num_rows FROM user_tables WHERE table_name='SUPPLIER';
```

TABLE_NAME	NUM_ROWS
SUPPLIER	18000

```
SELECT COUNT(*) FROM supplier;
```

COUNT(*)
18000

The statistics for SUPPLIER indicate there 18,000 rows in the table

The actual number is also 18,000.

The table level statistics for SUPPLIER seem to be up to date.

But, are they representative ?



# Step 2: Fix the cardinality estimate for SUPPLIER

Fix the cardinality estimate for table SUPPLIER (estimated 575, actual 2,862)

- When addressing a problem with single table cardinality estimates, the first place to look is the state of the table statistics. Are they representative?
- A **selectivity** of a filter predicate indicates the estimated proportion of rows in the table that will be returned after the filter predicate is applied.
- The **estimated cardinality**, or the number of rows expected to be produced from the table after a filter predicate is applied, is a product of the number of rows in the table and the selectivity of the filter predicate.
- When filters on multiple columns are used, by default, the final selectivity is calculated by multiplying the individual selectivities of all filter predicates.



# Step 2: Fix the cardinality estimate for SUPPLIER

Fix the cardinality estimate for table SUPPLIER (estimated 575, actual 2,862)

```
SELECT column_name, num_distinct
FROM user_tab_col_statistics
WHERE table_name='SUPPLIER'
  AND column_name IN ('S_REGION', 'S_NATION');
```

COLUMN_NAME	NUM_DISTINCT
S_NATION	25
S_REGION	5

*Assuming no histograms are present.*

Selectivity for s\_nation in ('CHINA') is 1/NUM\_DISTINCT, that is 1/25

Selectivity for s\_nation in ('CHINA', 'INDIA', 'JAPAN', 'VIETNAM') is 4/25

*Assuming no histograms are present.*

Selectivity for s\_region in ('ASIA') is 1/NUM\_DISTINCT, that is 1/5

Selectivity when both filter predicates are evaluated is  $4/25 \times 1/5 = 0.032$

The estimated cardinality when both filter predicates are evaluated is 18,000 (rows)  $\times 0.032 = 576$  (approx)

```
EXPLAIN PLAN FOR
SELECT * FROM supplier
WHERE s_region = 'ASIA'
  AND s_nation in ('CHINA', 'INDIA', 'JAPAN', 'VIETNAM')
```

Id   Operation	Name	Rows
0   SELECT STATEMENT		575
1   TABLE ACCESS STORAGE FULL	SUPPLIER	575

With histograms in place the calculations are different



# Step 2: Fix the cardinality estimate for SUPPLIER

Fix the cardinality estimate for table SUPPLIER (estimated 575, actual 2,862)

Think about correlation

- What is the possibility that a row in SUPPLIER will have s\_region='ASIA' and s\_nation in ('CHINA','INDIA','JAPAN','VIETNAM')?
- Is there a country named Japan outside of Asia?

```
SELECT COUNT(DISTINCT s_region) as distinct_region  
      , COUNT(DISTINCT s_nation) as distinct_nation  
      , COUNT(DISTINCT s_region||' '||s_nation) as distinct_pairs  
FROM supplier;
```

DISTINCT_REGION	DISTINCT_NATION	DISTINCT_PAIRS
5	25	25

If the number of distinct pairs of values between s\_region and s\_nation is close to the number of distinct values of either s\_region or s\_nation, we can say that the values in these two columns are correlated.

s\_region and s\_nation are strongly correlated.

Create a correlated column group on s\_region and s\_nation!

[Link: Managing Extended Statistics](#)



# Step 2: Fix the cardinality estimate for SUPPLIER

Fix the cardinality estimate for table SUPPLIER (estimated 575, actual 2,862)

```
SELECT DBMS_STATS.CREATE_EXTENDED_STATS(user, 'SUPPLIER', '(S_REGION,S_NATION)') FROM DUAL;
```

Extended statistic (column group) created, but no stats have been gathered.

```
DBMS_STATS.CREATE_EXTENDED_STATS(USER, 'SUPPLIER', '(S_REGION,S_NATION)')
```

Gather statistics for SUPPLIER table.

```
-----  
SYS_STU8TQAAMA8C62YIA0QMZ_A26F
```

```
EXEC DBMS_STATS.GATHER_TABLE_STATS(user, 'SUPPLIER');
```

```
SELECT column_name, histogram FROM user_tab_col_statistics WHERE table_name='SUPPLIER' AND column_name IN ('S_REGION', 'S_NATION', 'SYS_STU8TQAAMA8C62YIA0QMZ_A26F');
```

No histogram for the column group

COLUMN_NAME	HISTOGRAM
SYS_STU8TQAAMA8C62YIA0QMZ_A26F	NONE
S_NATION	FREQUENCY
S_REGION	FREQUENCY

If any of the columns that are used in the column group has histograms, in order for the column group to be considered by the optimizer, it needs to have histogram too.

```
EXPLAIN PLAN FOR  
SELECT * FROM supplier  
WHERE s_region = 'ASIA'  
AND s_nation in ('CHINA', 'INDIA', 'JAPAN', 'VIETNAM');
```

Run EXPLAIN PLAN of a query that uses filter predicates on the columns used in the column group will populate SYS.COL\_USAGE\$

```
EXEC DBMS_STATS.GATHER_TABLE_STATS(user, 'SUPPLIER');
```

```
SELECT column_name, histogram FROM user_tab_col_statistics WHERE table_name='SUPPLIER'  
('S_REGION', 'S_NATION', 'SYS_STU8TQAAMA8C62YIA0QMZ_A26F');
```

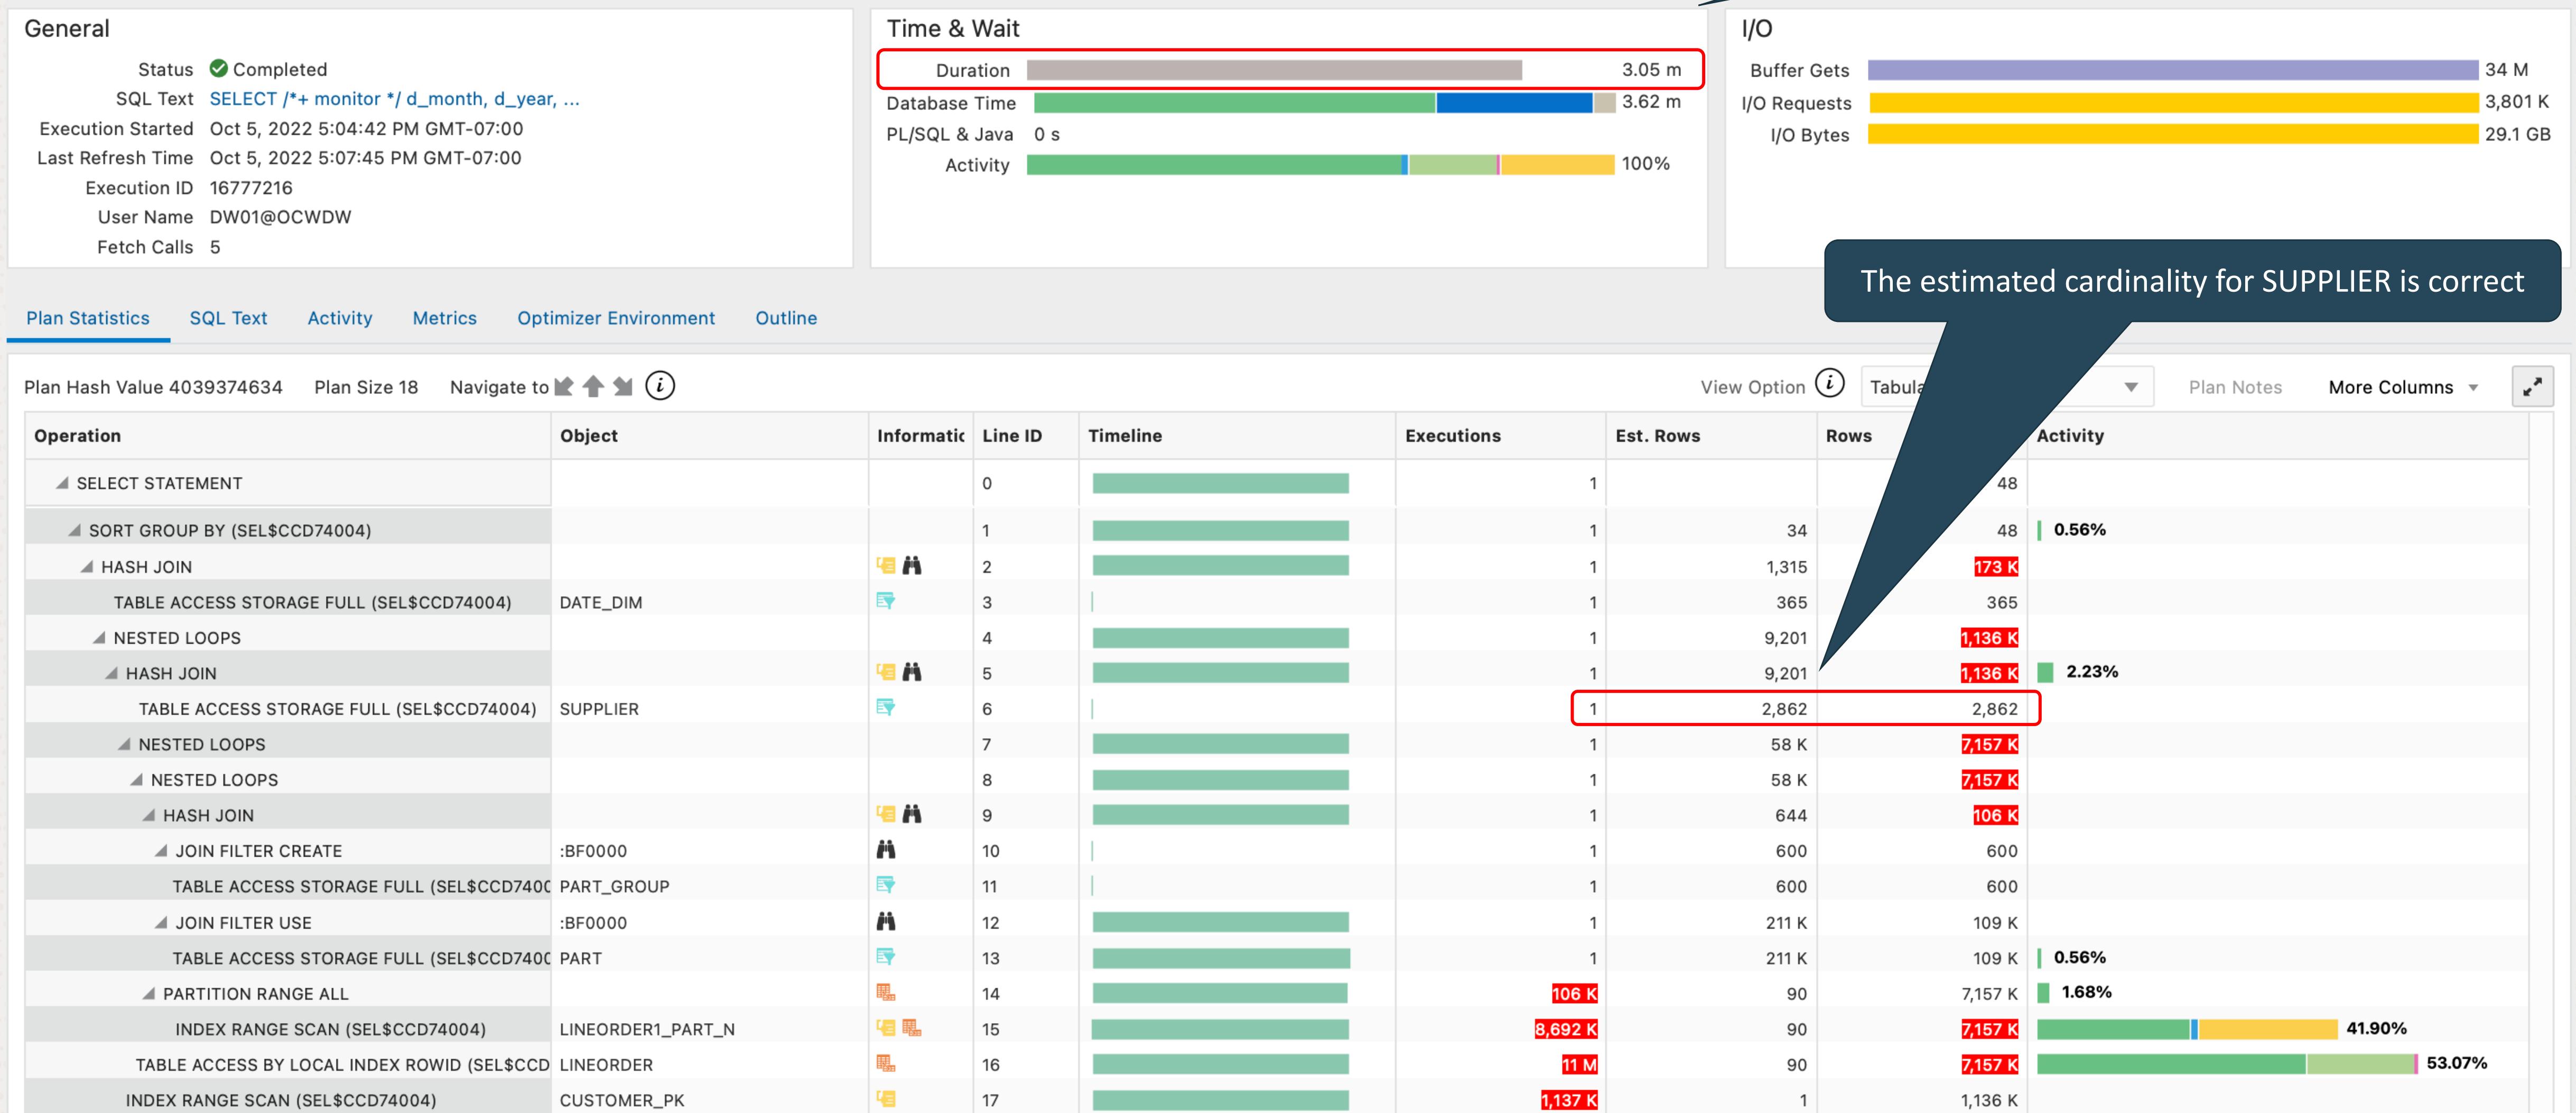
COLUMN_NAME	HISTOGRAM
SYS_STU8TQAAMA8C62YIA0QMZ_A26F	FREQUENCY
S_NATION	FREQUENCY
S_REGION	FREQUENCY

When stats are gathered Oracle uses SYS.COL\_USAGE\$ to find columns that will benefit from histograms. In this case, it will find the column group that we created and will create a histogram on it

The histogram on the column group has been created

Step2: running time: 185 seconds

# Step 2: Column group on SUPPLIER(S\_REGION, NATION)

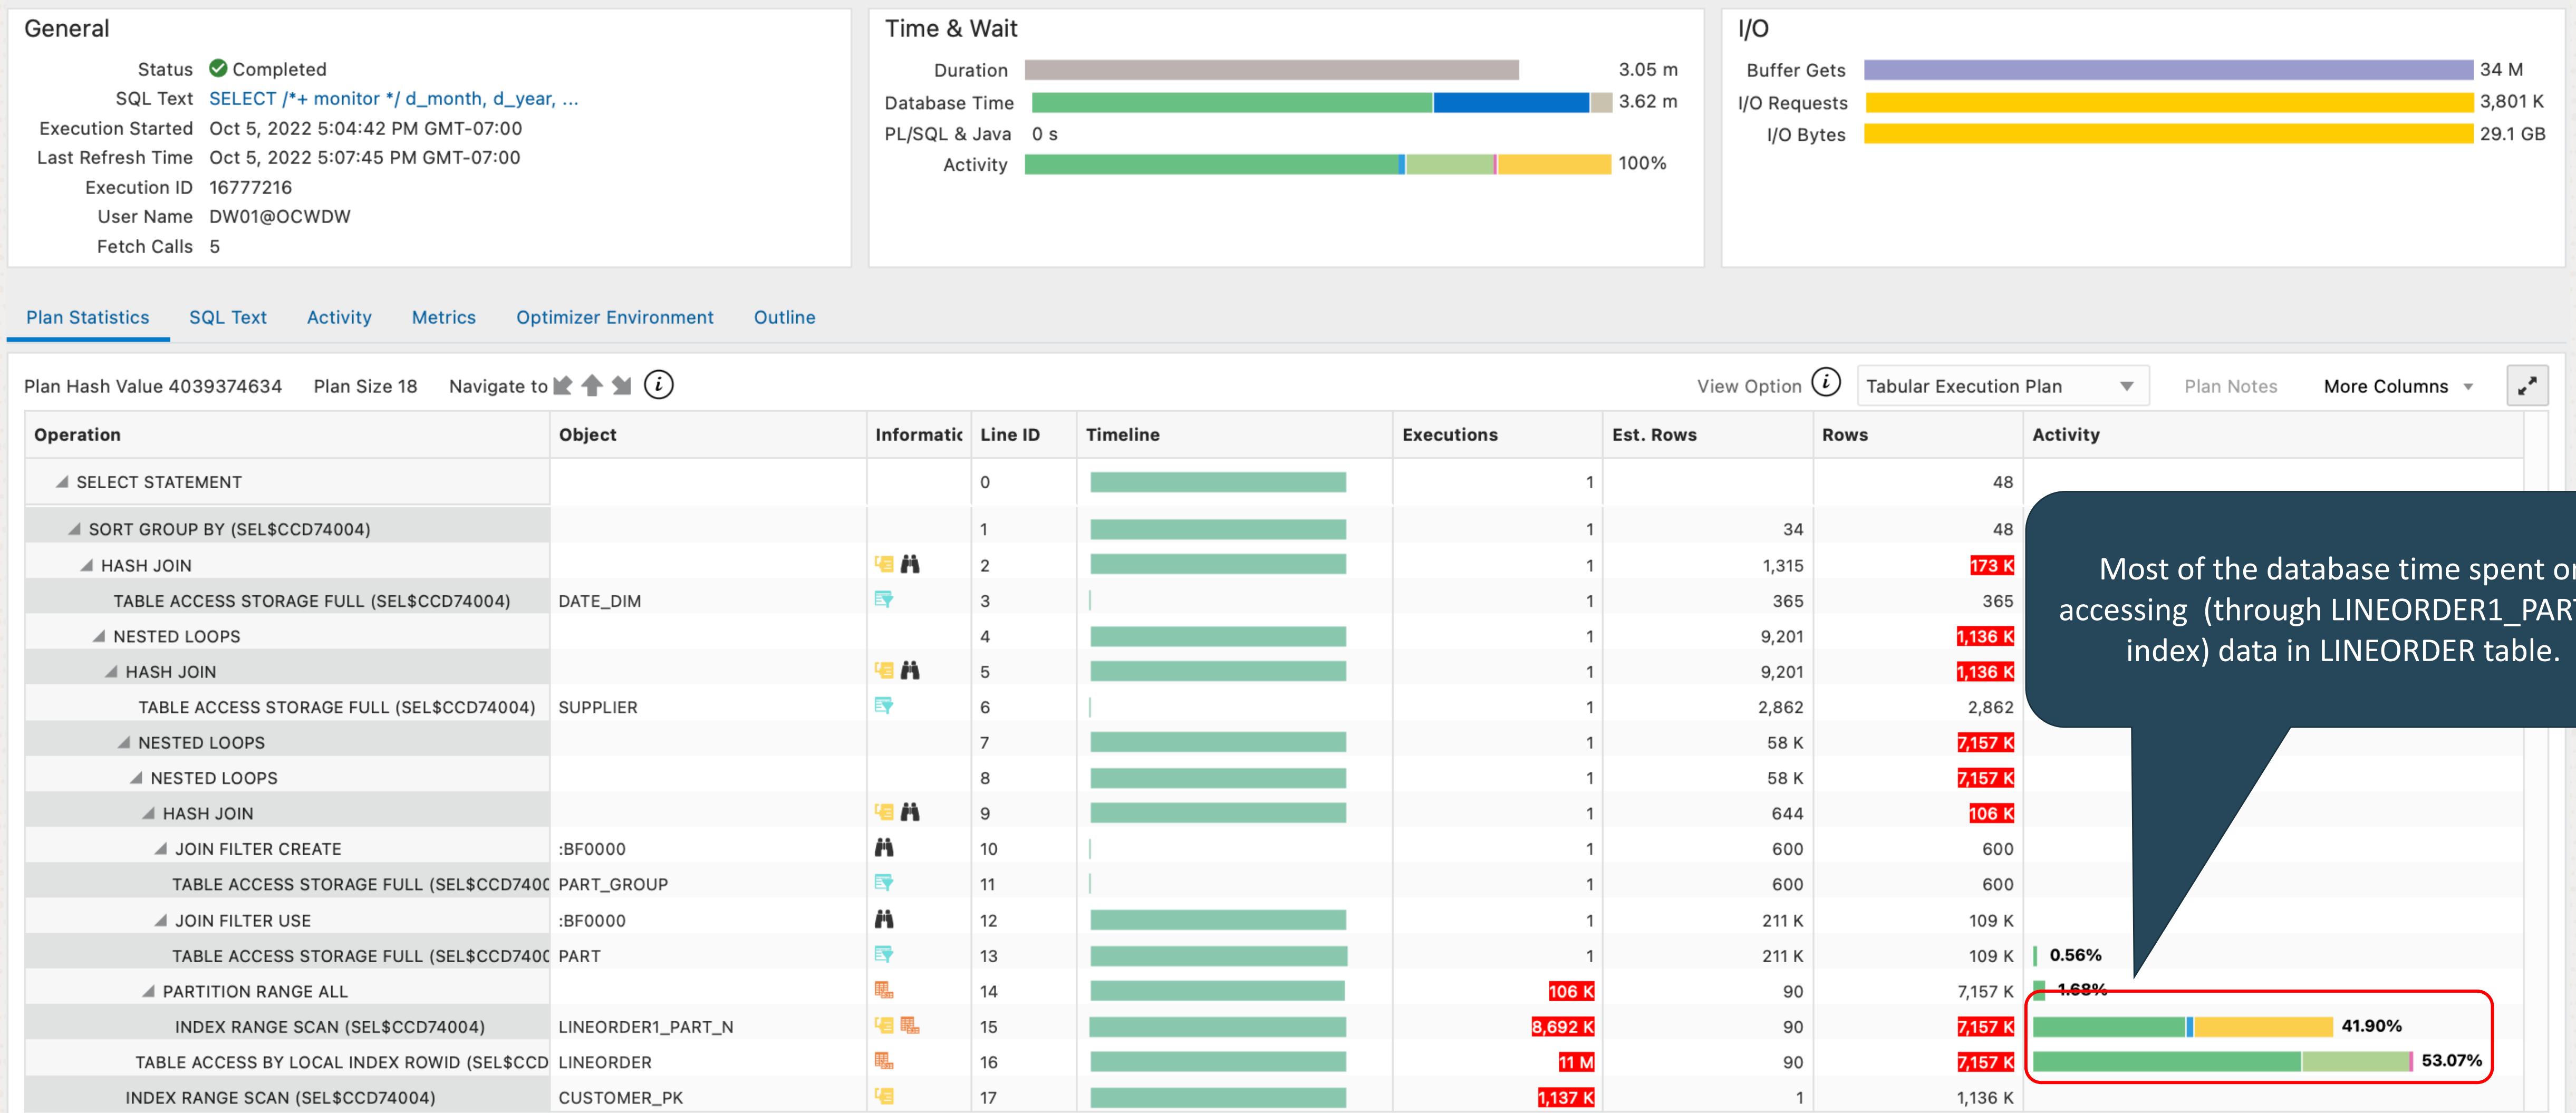


# Worked example: Progress

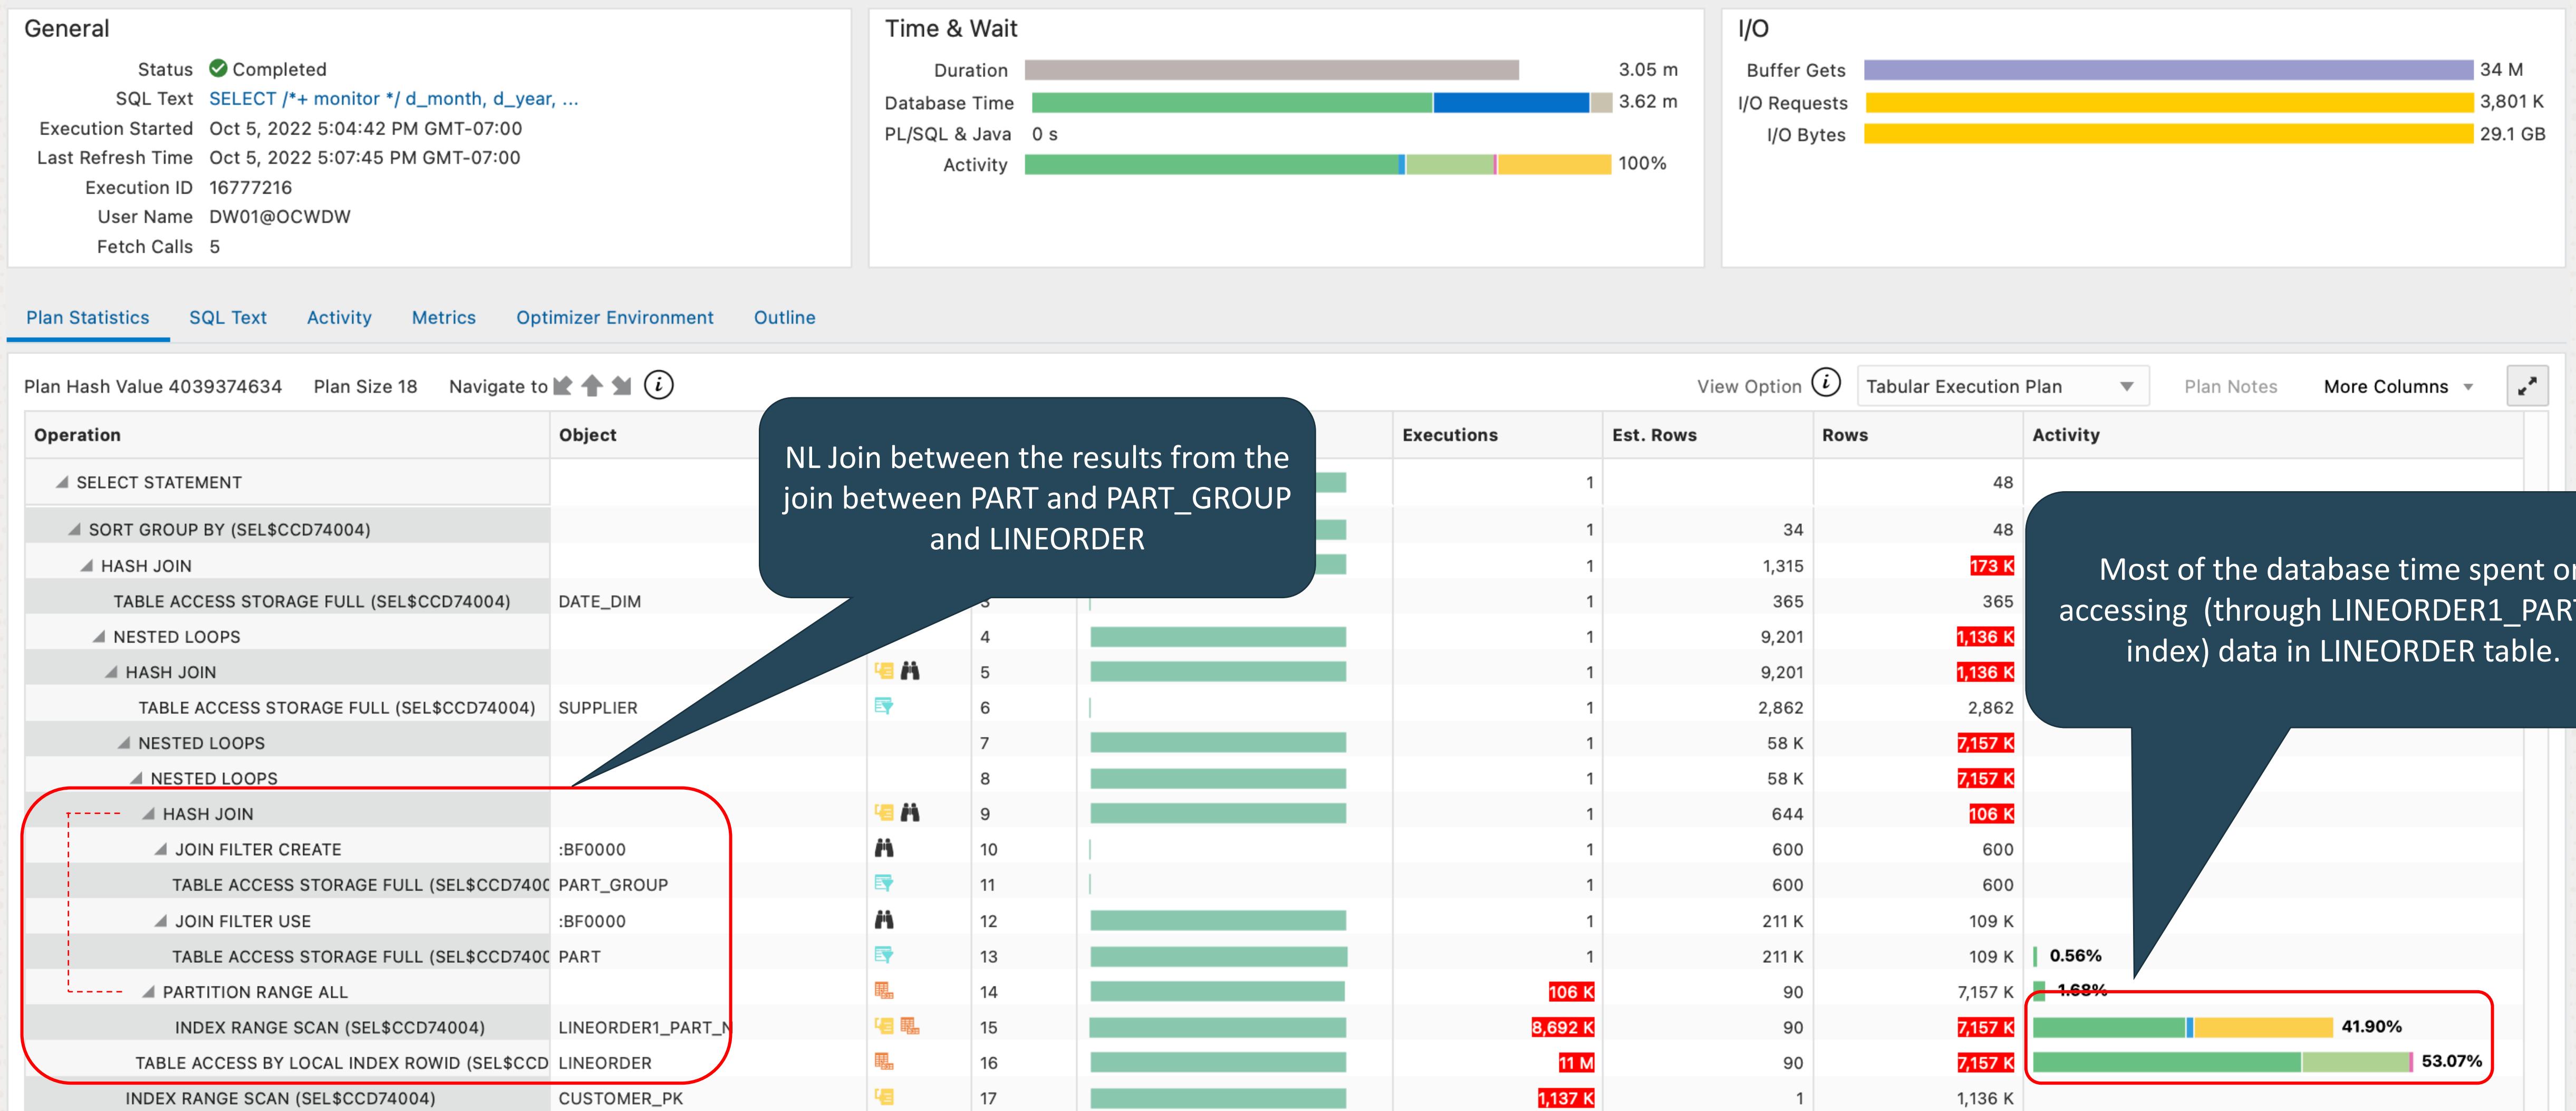
	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds
Step 2: Column group on SUPPLIER	Statistics	185 seconds



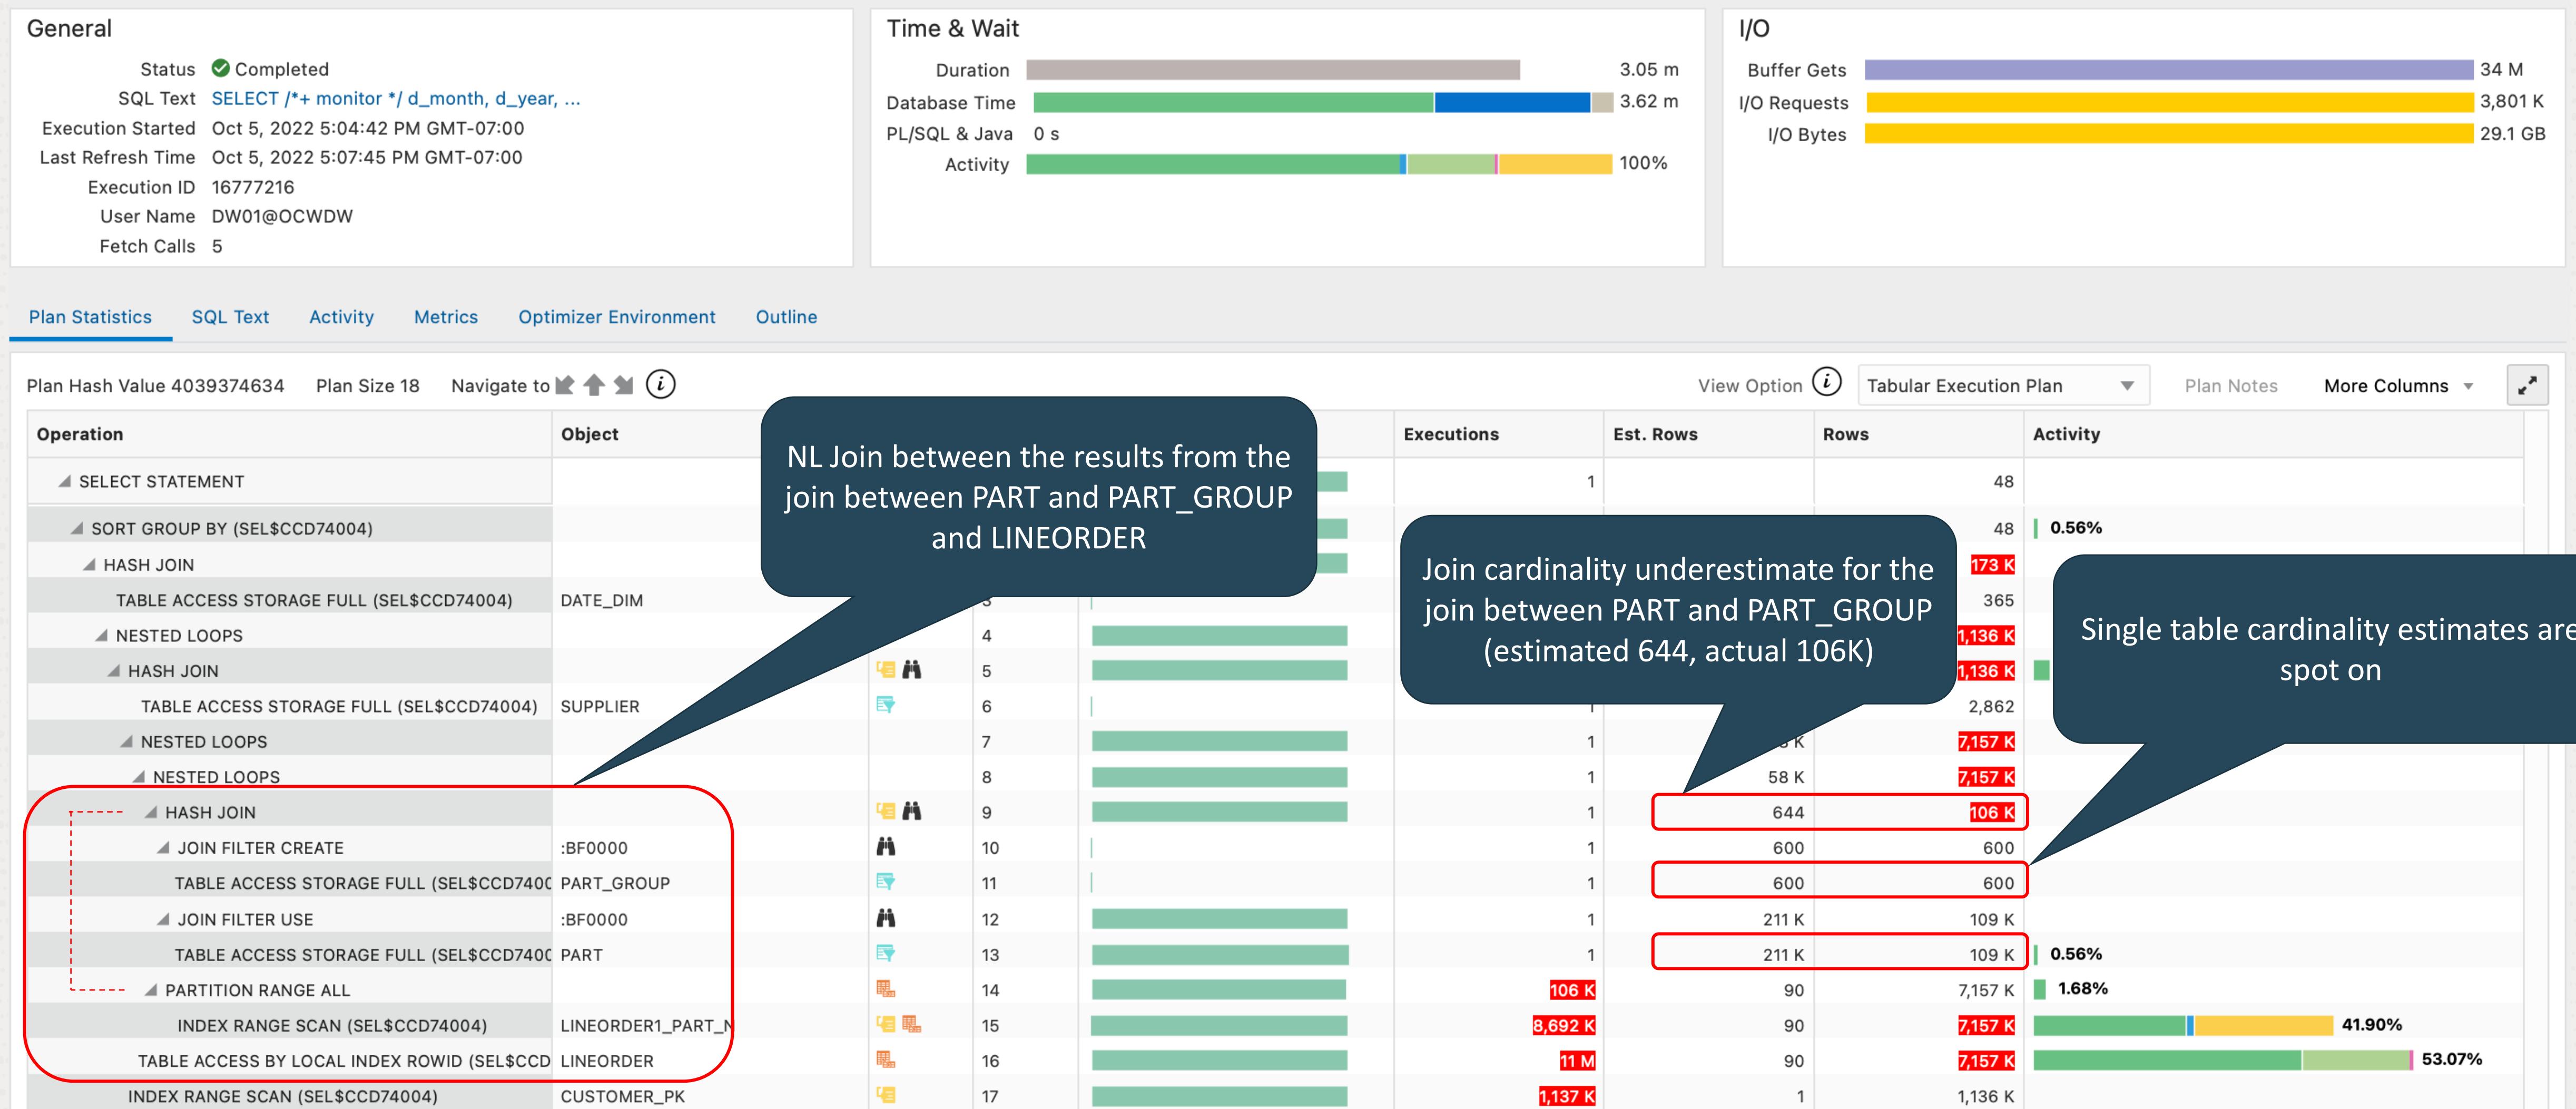
# Step 2: Column group on SUPPLIER(S\_REGION,S\_NATION)



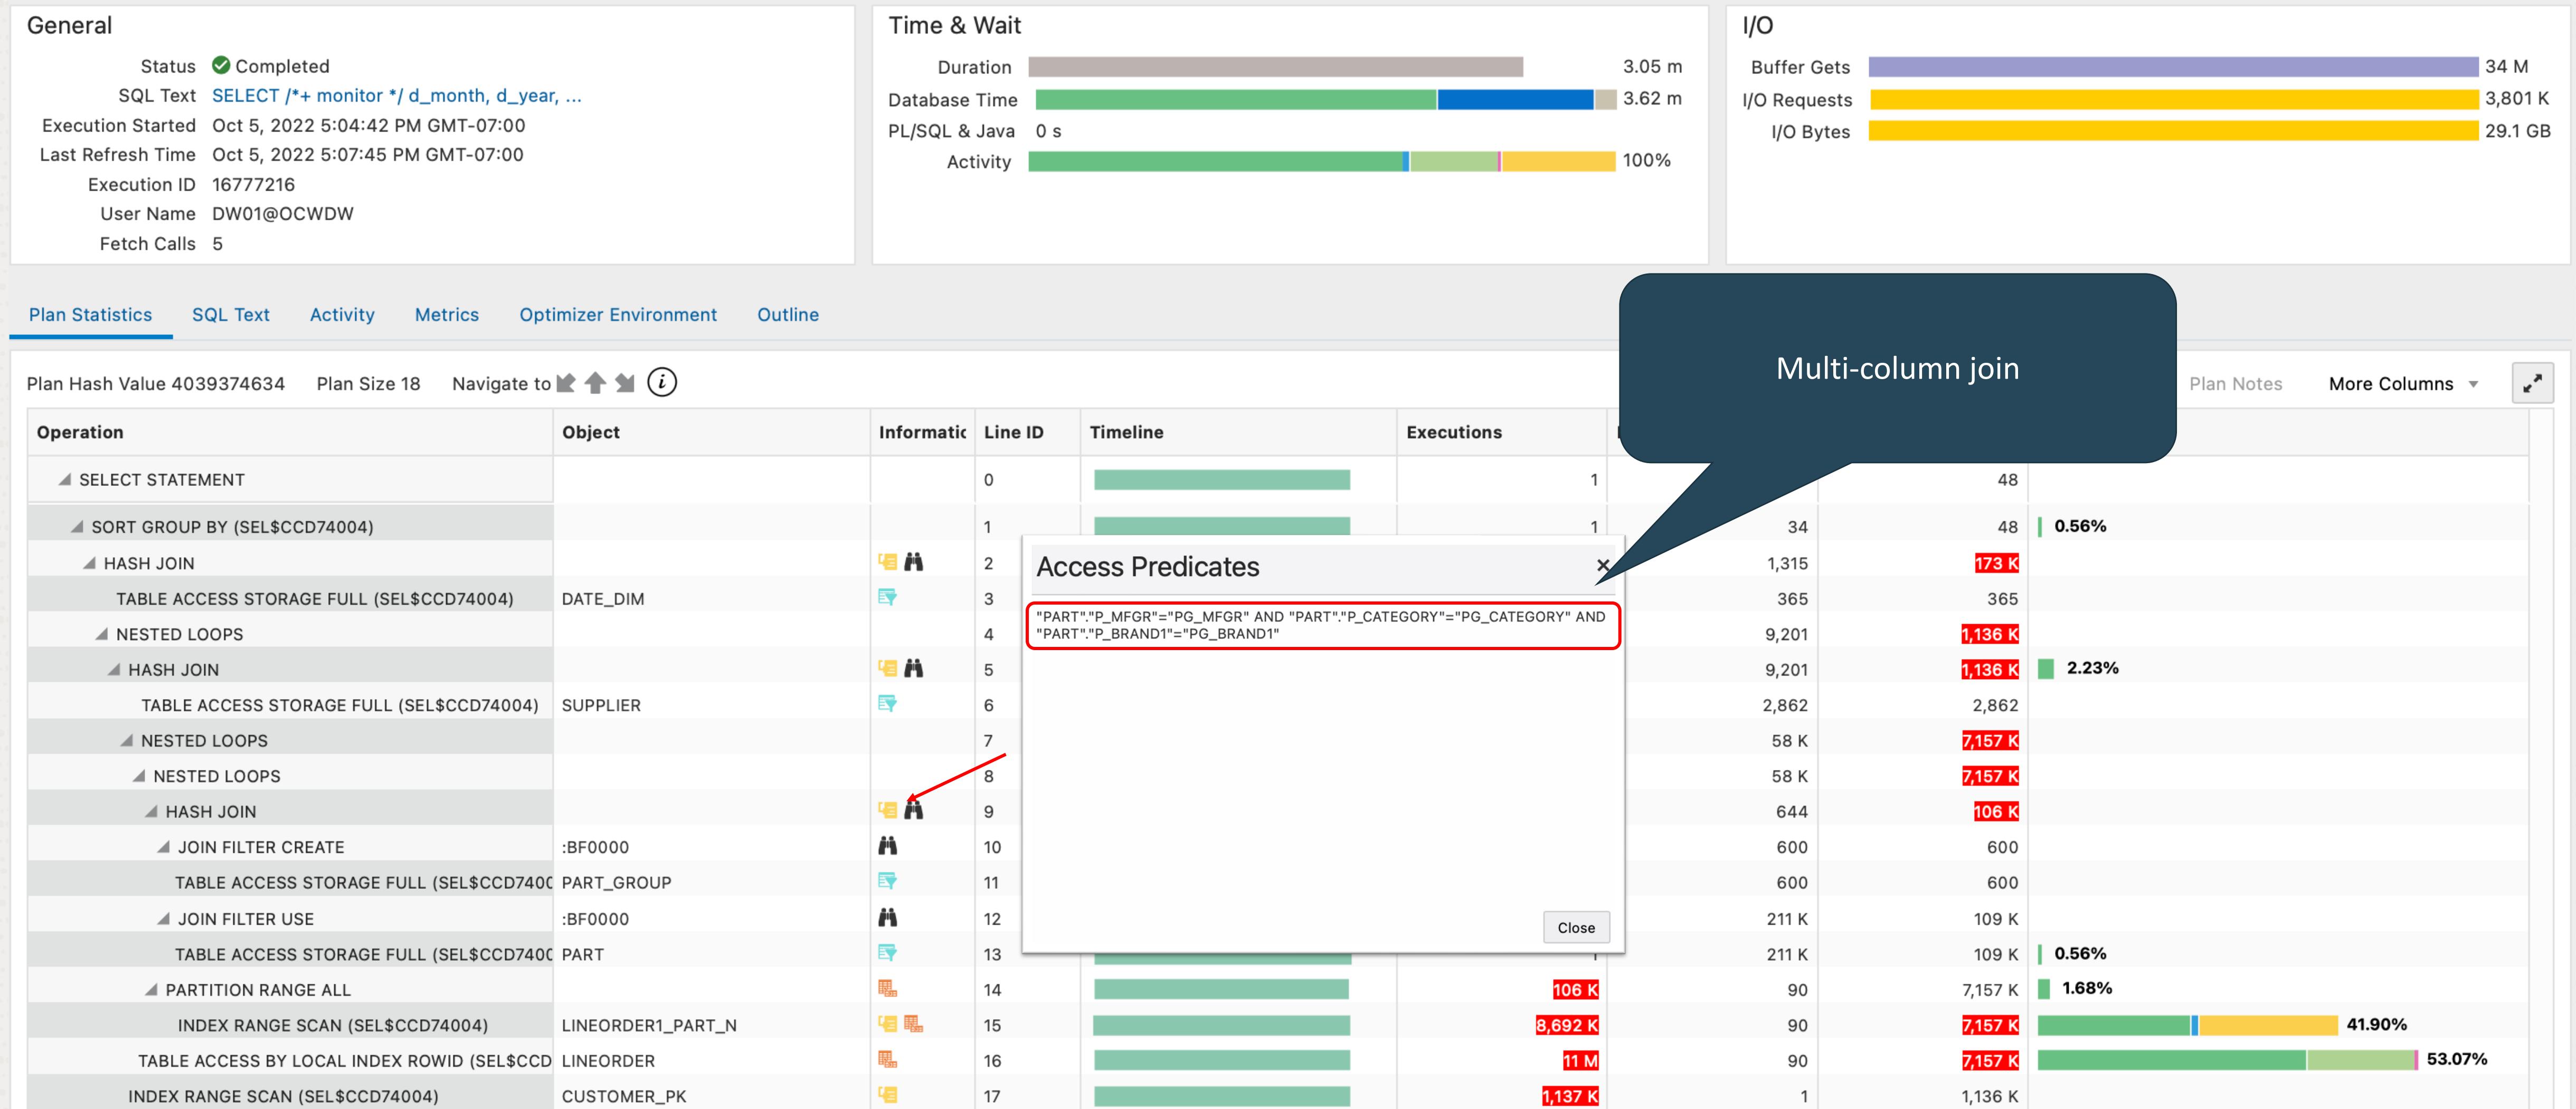
## Step 2: Column group on SUPPLIER(S\_REGION,S\_NATION)



## Step 2: Column group on SUPPLIER(S\_REGION,S\_NATION)



## Step 2: Column group on SUPPLIER(S\_REGION,S\_NATION)



## Multi-column join



# What You Must Do: Constraints



- NOT NULL Constraints on Join Keys
- Primary Key Constraints on Dimension Join Keys
- Foreign Key Constraints on Fact Join Keys

# What You Must Do: Primary Key and Foreign Key Constraints

```
alter table customer
add constraint customer_pk
primary key (c_custkey)
RELY;
```

```
alter table lineorder
add constraint lo_customer_pk
foreign key (lo_custkey)
references
customer (c_custkey)
RELY
DISABLE NOVALIDATE;
```

```
alter system
set query_rewrite_integrity=TRUSTED;
```

- There must be a primary key on the dimension table
- There must be a foreign key on the fact table
- The state of the constraint depends on trust in the ETL process and volume of data
- Constraints must be in RELY state
- It is *not* necessary to enforce constraints on the fact table
- You need to tell the optimizer you can trust constraints in the RELY state

# What You Must Do: Validating ETL/ELT

```
SELECT *  
FROM lineorder  
LEFT OUTER JOIN customer  
    ON lo_custkey = c_custkey  
WHERE c_custkey IS NULL;
```

- How do we validate our data when our constraints are not enforced?
- In other words, when constraints are in RELY mode, how do we ensure we can rely on the quality of data being inserted into our fact table?
- This SQL checks for rows in lineorder for values of lo\_custkey which do not exist in the customer dimension table



# What You Must Do: Validating ETL/ELT

```
SELECT *
FROM lineorder
LEFT OUTER JOIN customer
  ON lo_custkey = c_custkey
LEFT OUTER JOIN date_dim
  ON lo_orderdate = d_datekey
LEFT OUTER JOIN part
  ON lo_partkey = p_partkey
LEFT OUTER JOIN supplier
  ON lo_suppkey = s_suppkey
WHERE c_custkey IS NULL
  OR d_datekey IS NULL
  OR p_partkey IS NULL
  OR s_suppkey IS NULL;
```

- We can also validate the rows in lineorder against multiple dimensions
- Check the lineorder table for rows which contain keys that do not exist in the dimension tables



# What You Must Do: Constraints

How the constraints help the optimizer?

- Improve cardinality estimates
- Opens possibility for additional transformation such as : join elimination and other query transformation techniques



# Step 3: Fix the join cardinality estimate

Fix the cardinality estimate for the join between PART and PART\_GROUP

- Create a primary key on PART\_GROUP
- Create a foreign key on PART referencing PART\_GROUP
- Create a primary key on PART
- Create a foreign key on LINEORDER referencing PART

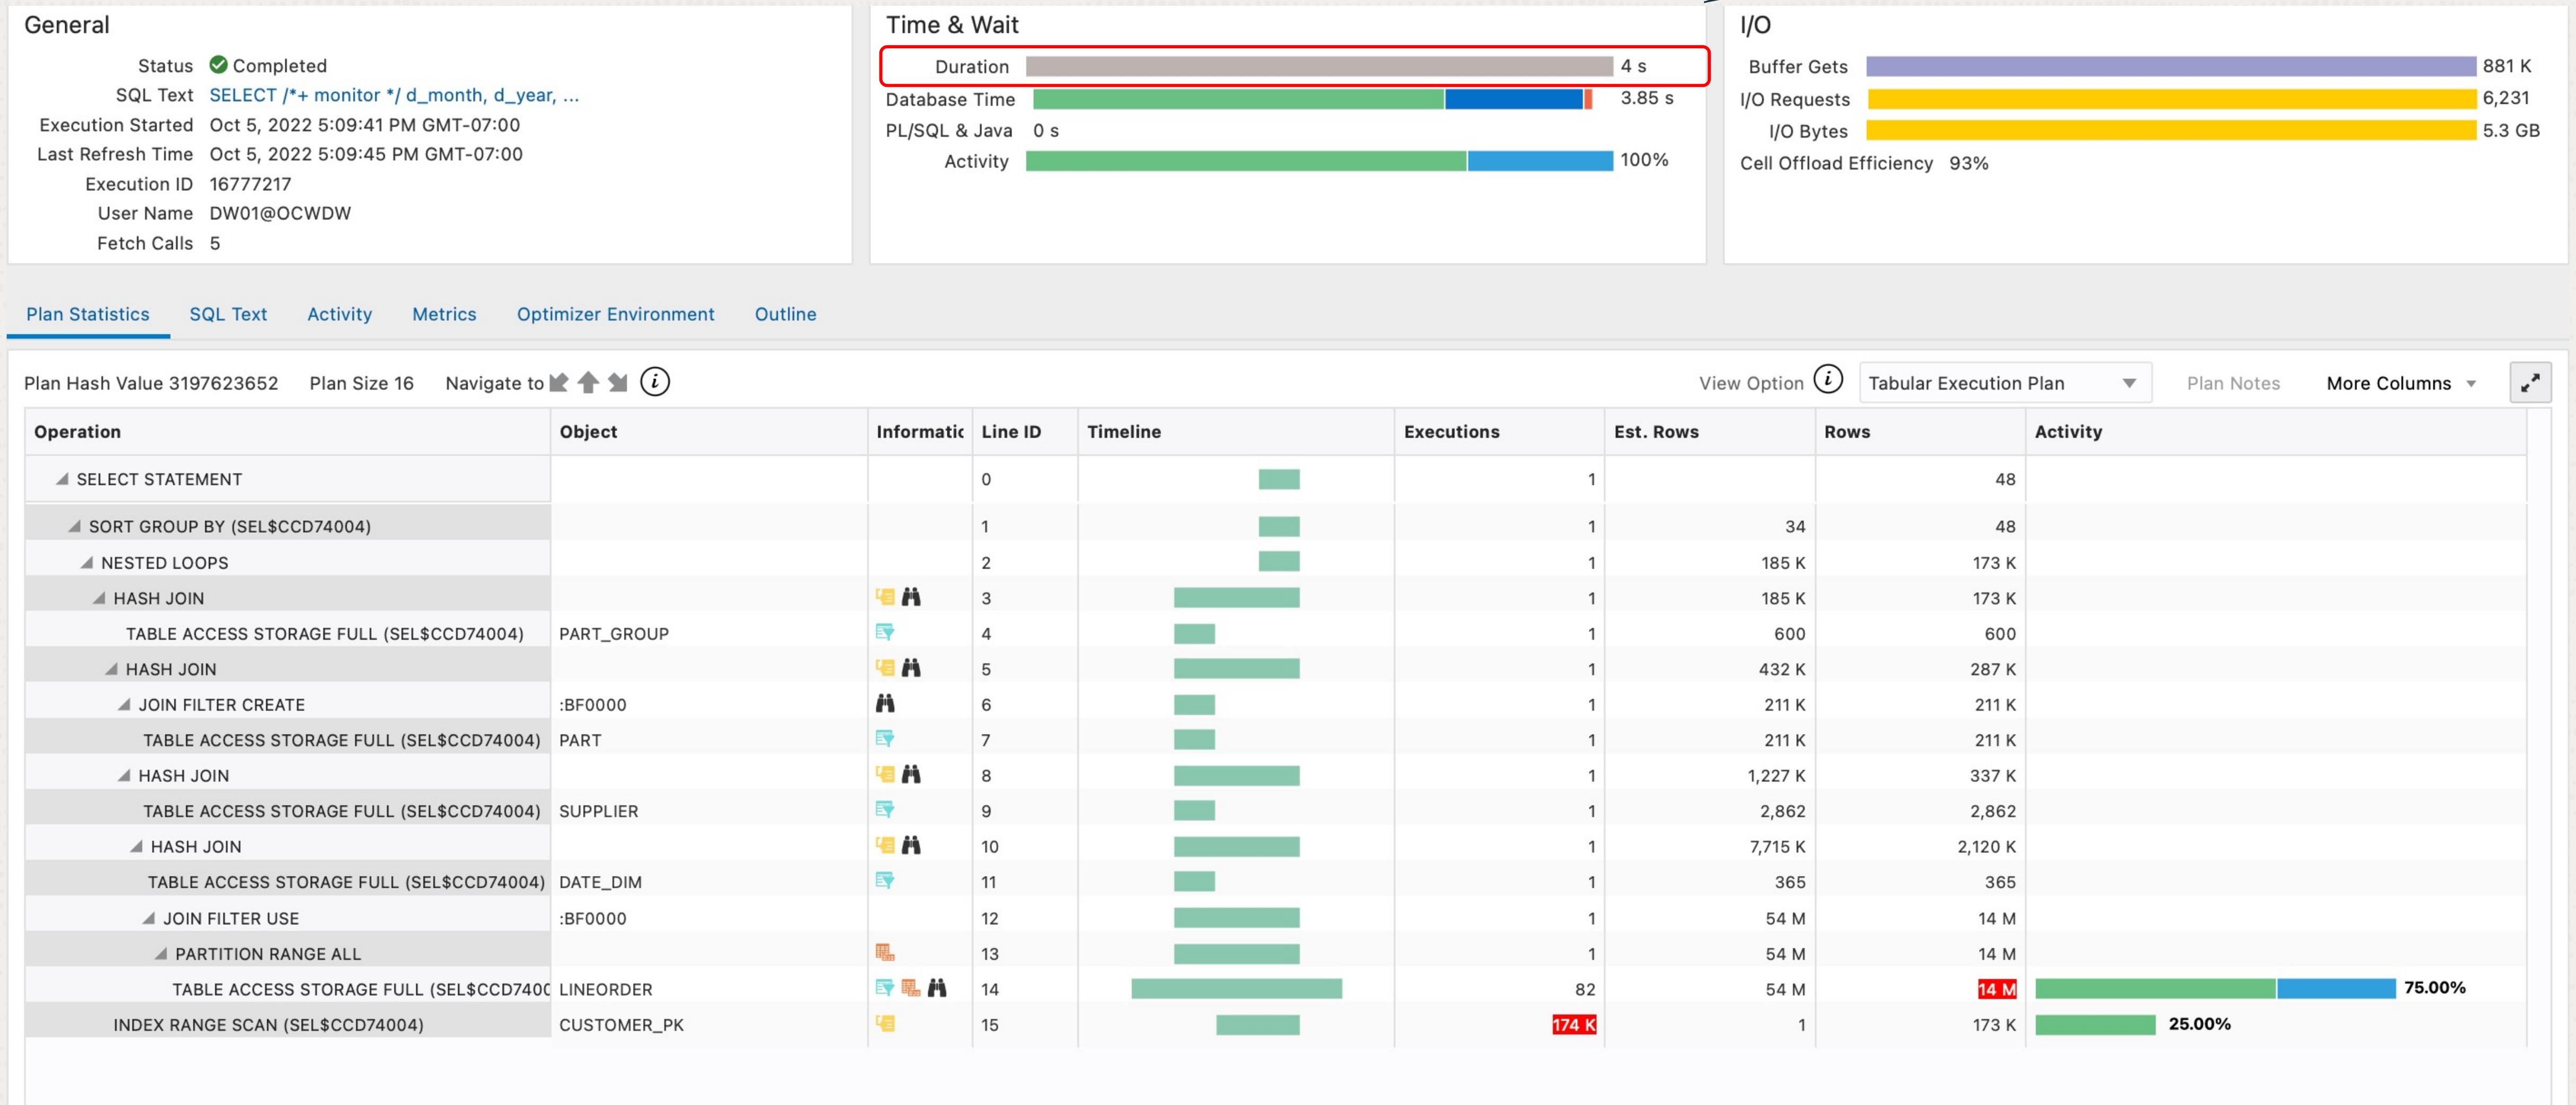
You'll need to tell the optimizer it can trust constraints in RELY state by setting :

query\_rewrite\_integrity=TRUSTED;

```
ALTER TABLE part_group ADD CONSTRAINT part_group_pk PRIMARY KEY (pg_mfgr, pg_category, pg_brand1) RELY;  
ALTER TABLE part ADD CONSTRAINT part_group_fk FOREIGN KEY(p_mfgr, p_category, p_brand1) REFERENCING part_group RELY;  
ALTER TABLE part ADD CONSTRAINT part_pk PRIMARY KEY(p_partkey) RELY;  
ALTER TABLE lineorder ADD CONSTRAINT lineorder_part_fk FOREIGN KEY (lo_partkey) REFERENCING part RELY DISABLE NOVALIDATE;
```

Step3: running time: 4 seconds

# Step 3: Constraints on PART, PART\_GROUP and LINE ORDER

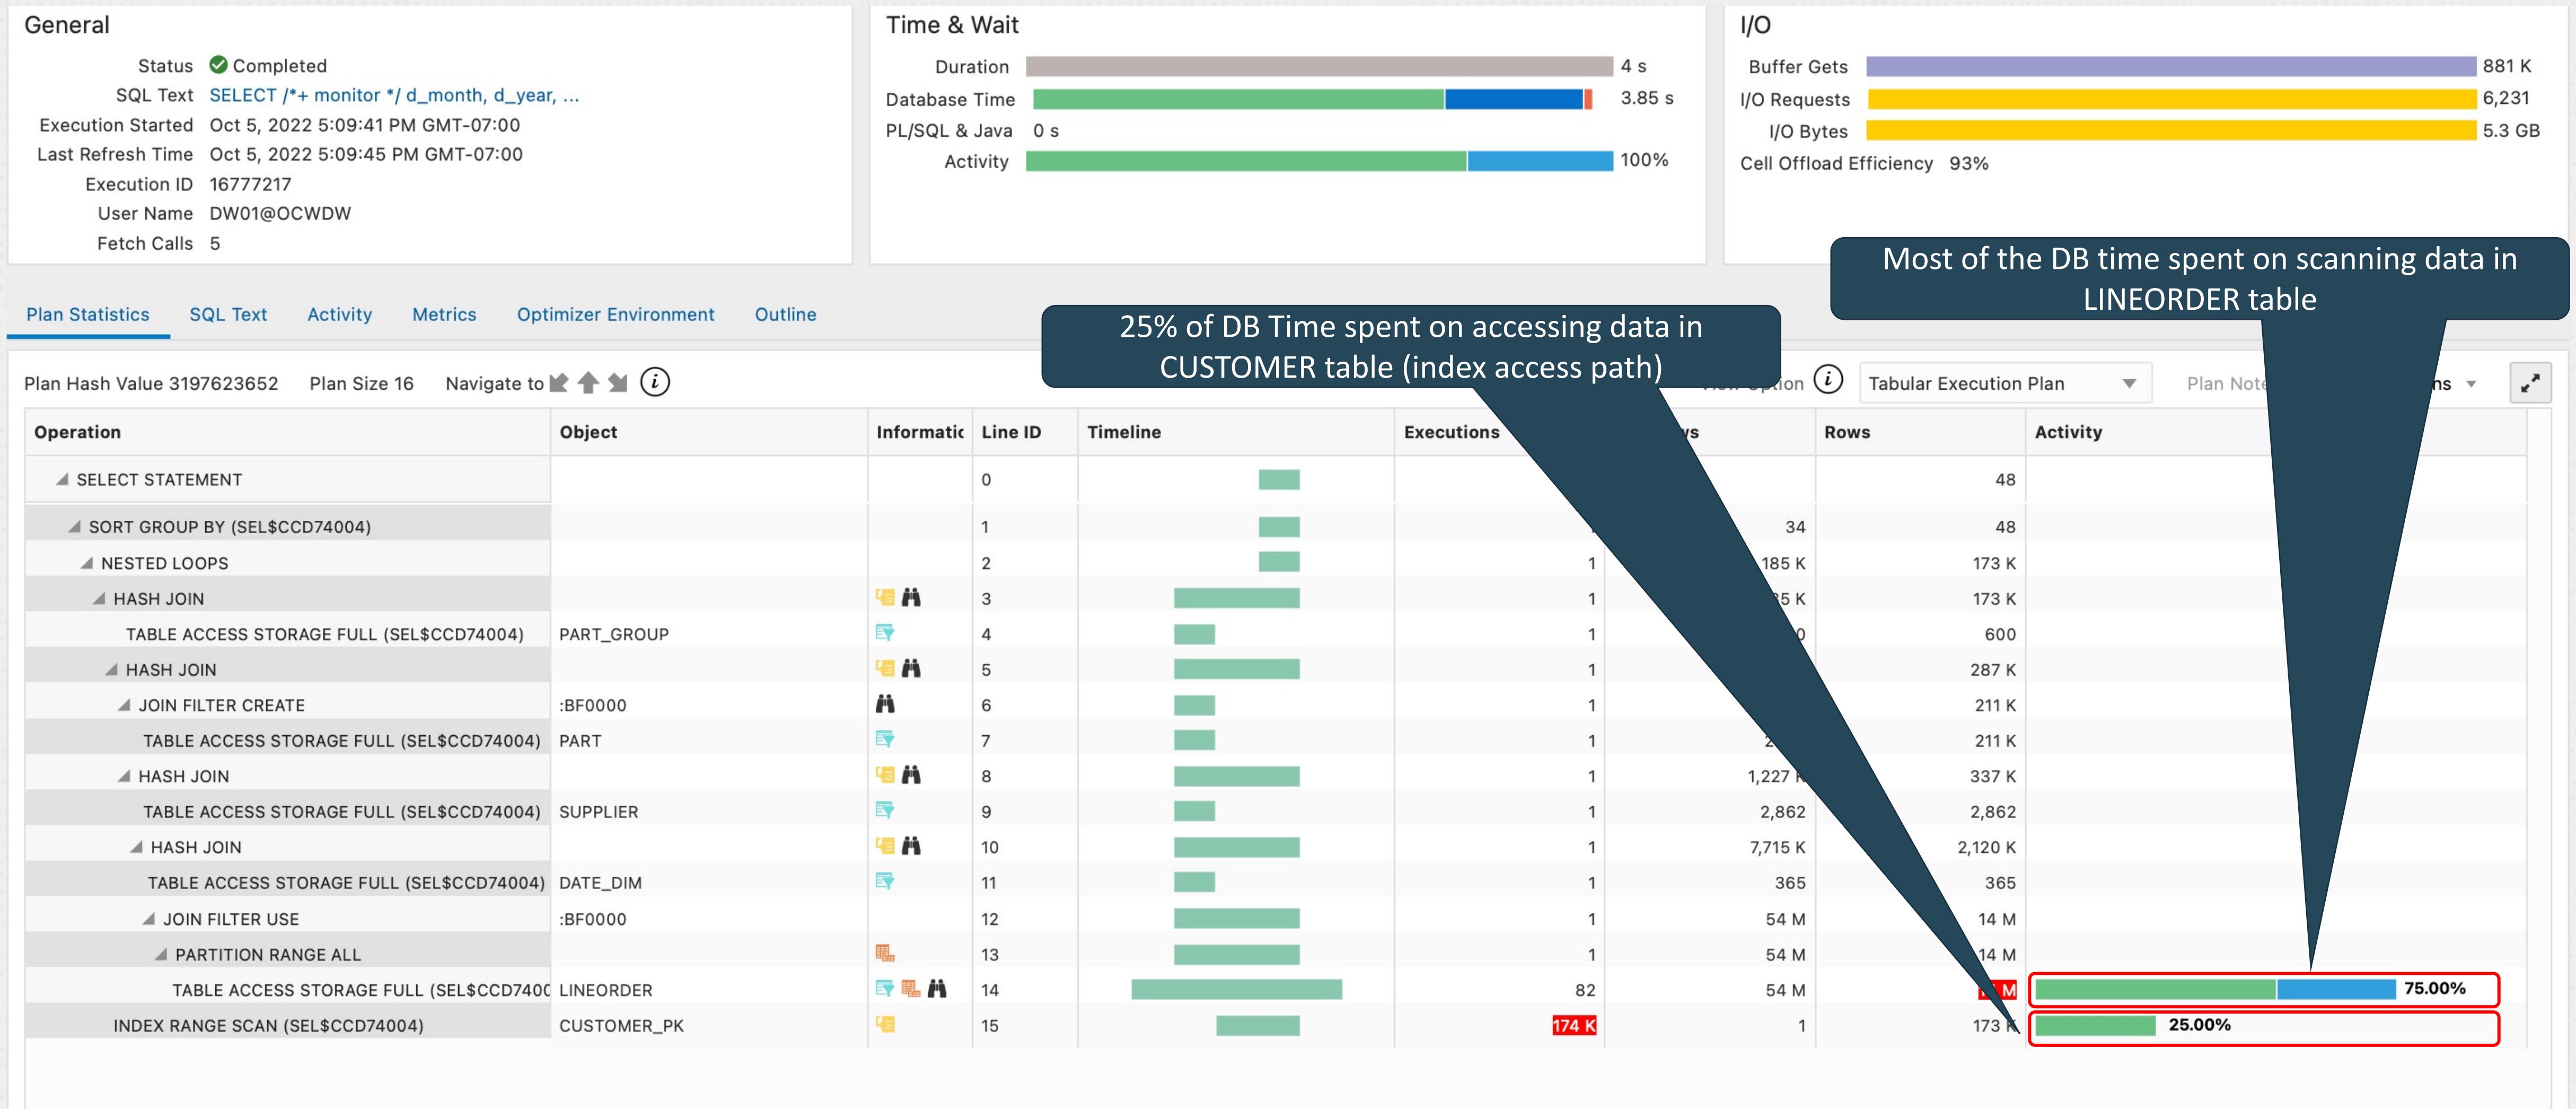


# Worked example: Progress

	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds
Step 2: Column group on SUPPLIER	Statistics	185 seconds
Step 3: Constraints on PART, PART_GROUP and LINEORDER	Constraints	4 seconds



# Step 3: Constraints on PART, PART\_GROUP and LINE ORDER



# Step 3: Constraints on PART, PART\_GROUP and LINE ORDER

**General**

- Status: Completed
- SQL Text: `SELECT /*+ monitor */ d_month, d_year, ...`
- Execution Started: Oct 5, 2022 5:09:41 PM GMT-07:00
- Last Refresh Time: Oct 5, 2022 5:09:45 PM GMT-07:00
- Execution ID: 16777217
- User Name: DW01@OCWDW
- Fetch Calls: 5

**Time & Wait**

Metric	Value
Duration	4 s
Database Time	3.85 s
PL/SQL & Java	0 s
Activity	100%

**I/O**

Metric	Value
Buffer Gets	881 K
I/O Requests	6,231
I/O Bytes	5.3 GB
Cell Offload Efficiency	93%

**Plan Statistics** **SQL Text** **Activity** **Metrics** **Optimizer Environment** **Outline**

**SQL Text:**

```

SELECT
/*+ monitor */
d_month, d_year, s_nation,
SUM(lo_quantity) lo_quantity,
SUM(lo_revenue) lo_revenue,
SUM(lo_supplycost) lo_supplycost
FROM lineorder l
    JOIN customer      ON lo_custkey = c_custkey
    JOIN date_dim      ON lo_orderdate = d_datekey
    JOIN part          ON lo_partkey = p_partkey
    JOIN part_group pg ON p_mfgr=pg_mfgr AND p_category=pg_category AND p_brand1=pg_brand1
    JOIN supplier      ON lo_suppkey = s_suppkey
WHERE d_year = 2007
AND s_region = 'ASIA'
AND s_nation IN ('CHINA', 'INDIA', 'JAPAN', 'VIETNAM')
AND pg_type IN (1,2,3)
AND p_size<12
GROUP BY d_month, d_year, s_nation
ORDER BY d_month, d_year, s_nation

```

**Save to Script**

CUSTOMER table joined to LINEORDER. But, there are no other columns from CUSTOMER table used in the query.

Is the join to CUSTOMER table needed in the query?

Can the optimizer eliminate the join automatically?

Yes if there is FK from LINEORDER to CUSTOMER.

- Join elimination transformation

# Step 4: Eliminate the join to CUSTOMER

## Help the optimizer safely eliminate the join to CUSTOMER

- Create a primary key on CUSTOMER
- Create a foreign key on LINEORDER referencing CUSTOMER

You'll need to tell the optimizer it can trust constraints in RELY state by setting :

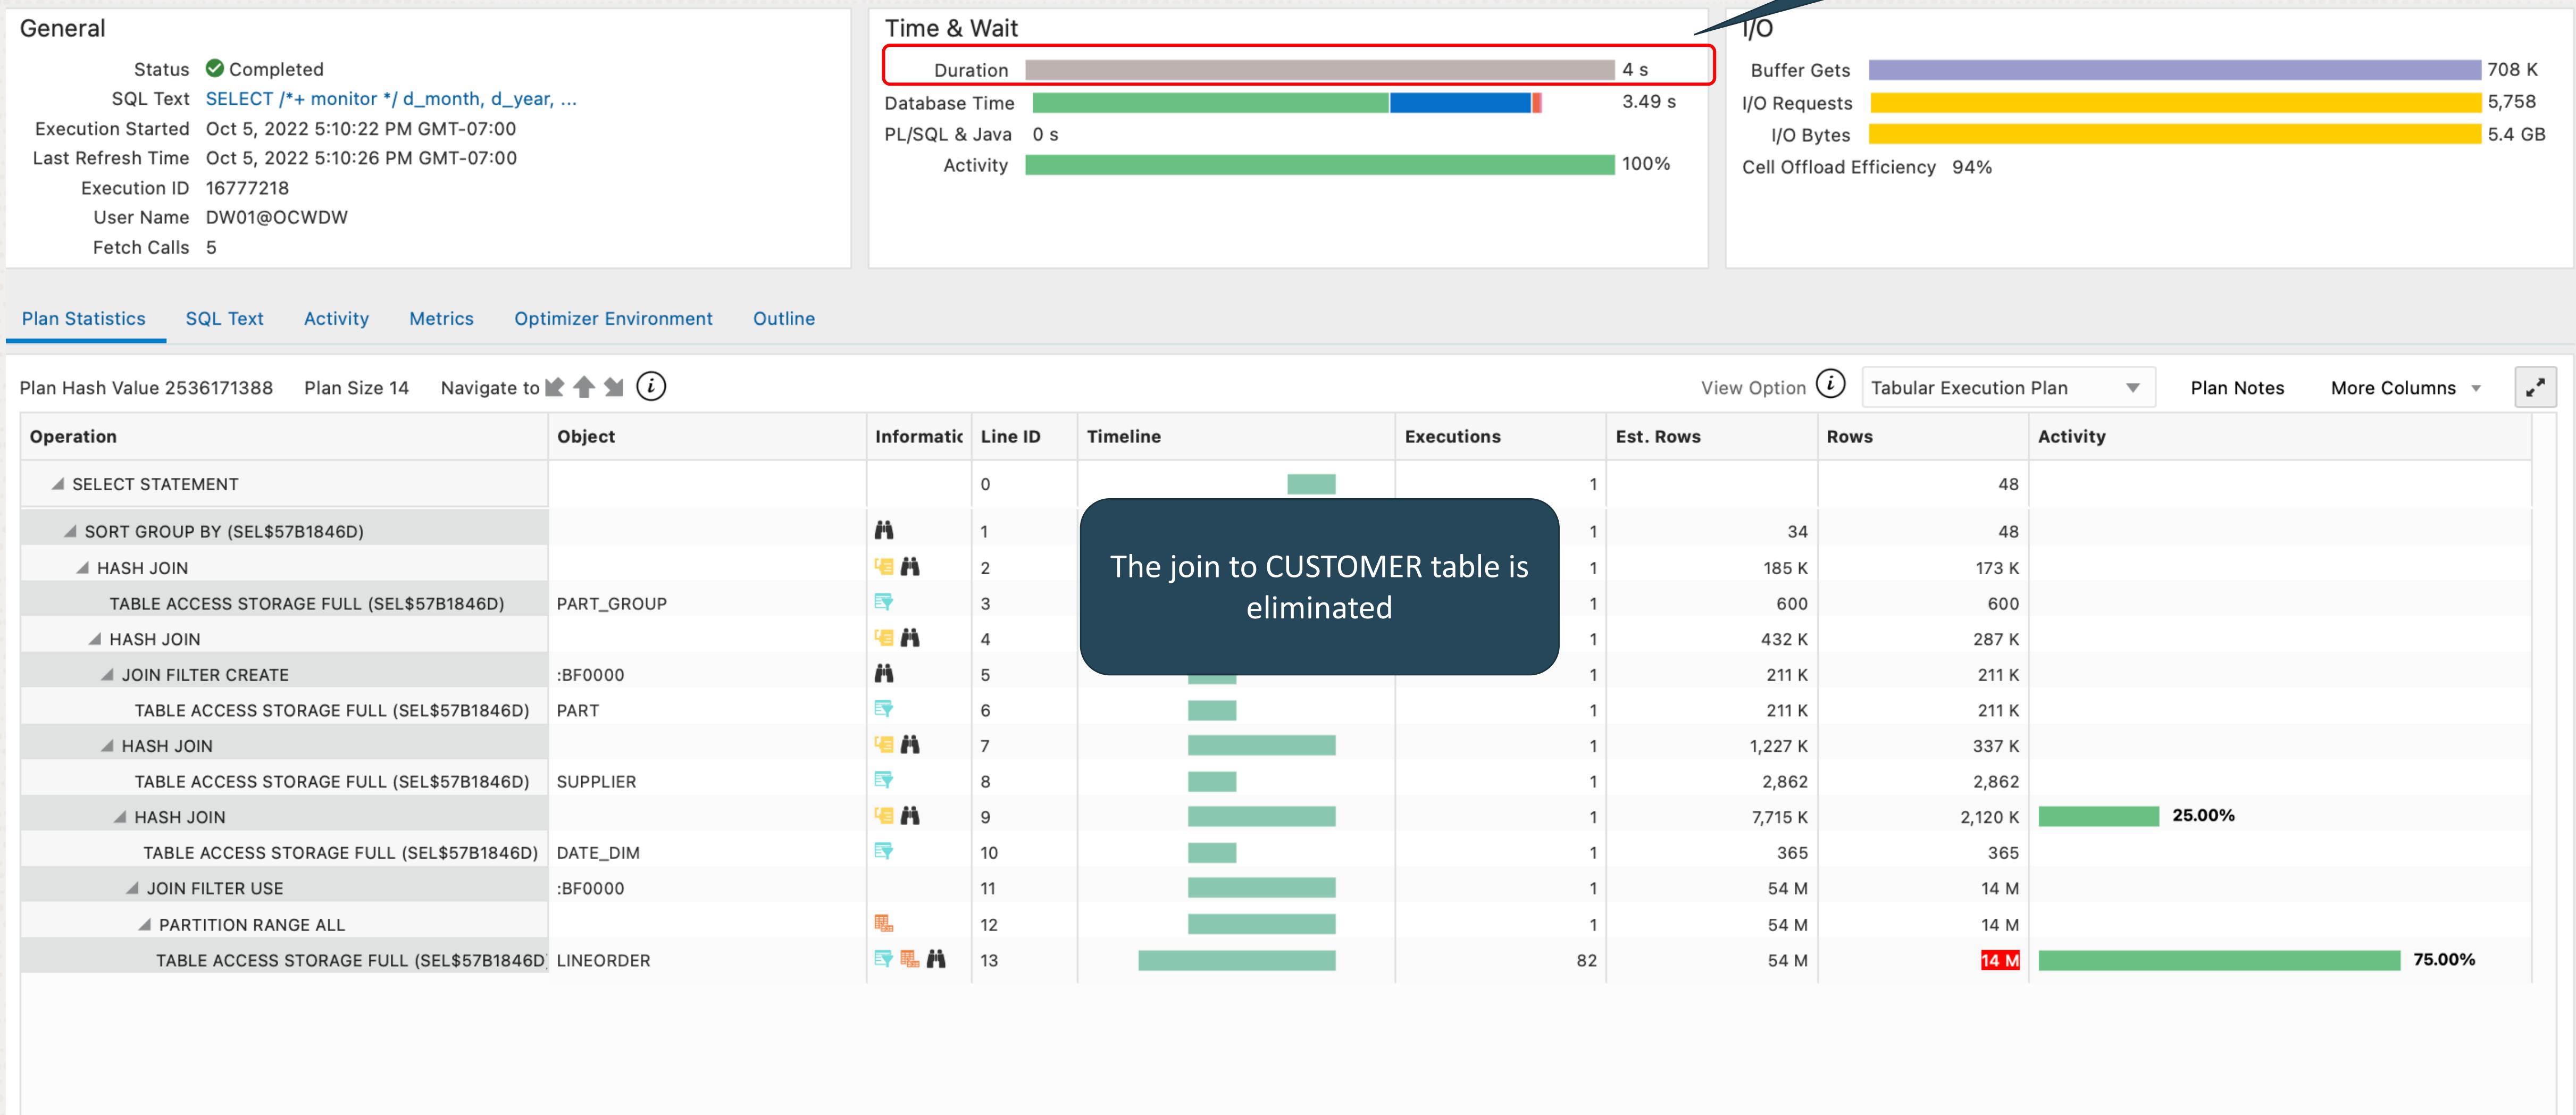
query\_rewrite\_integrity=TRUSTED;

```
ALTER TABLE customer ADD CONSTRAINT customer_pk PRIMARY KEY (c_custkey) RELY;
```

```
ALTER TABLE lineorder ADD CONSTRAINT lineorder_cust_fk FOREIGN KEY (lo_custkey) REFERENCING customer RELY DISABLE NOVALIDATE;
```

Step4: running time: 4 seconds

# Step 4: PK on CUSTOMER – FK on LINEORDER

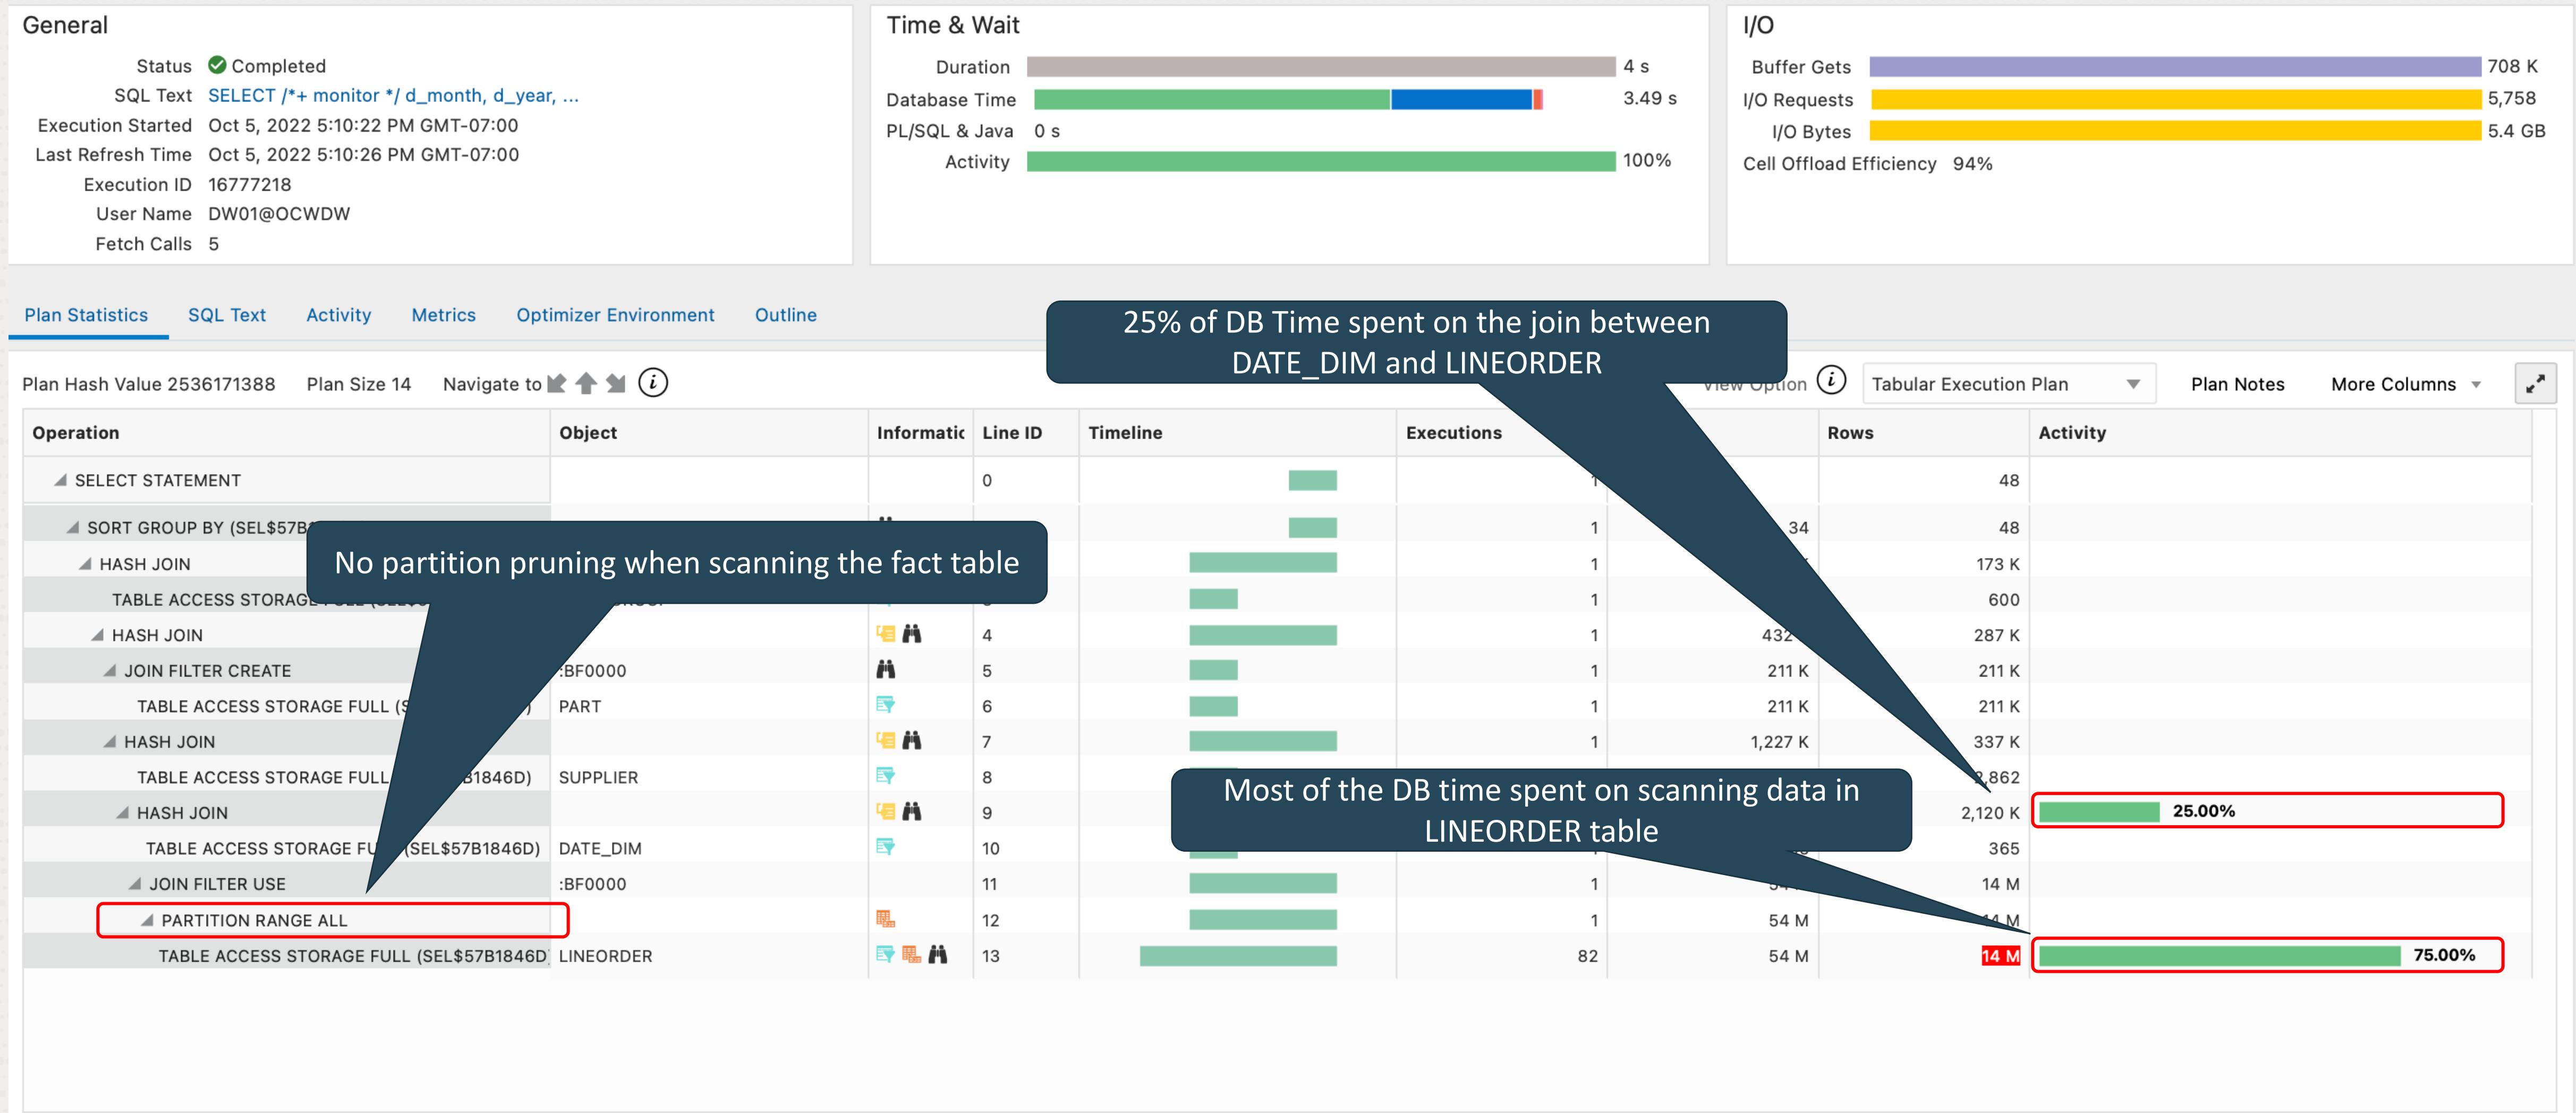


# Worked example: Progress

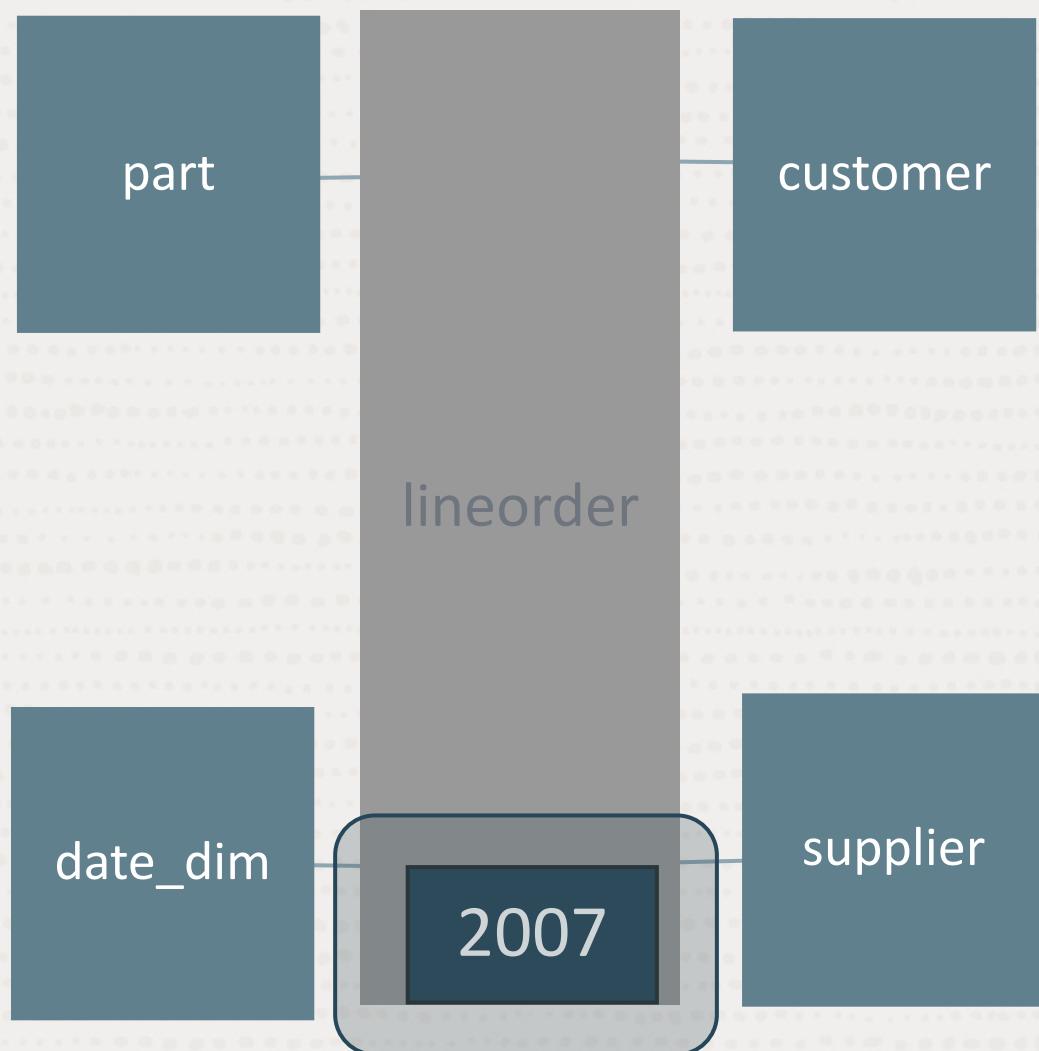
	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds
Step 2: Column group on SUPPLIER	Statistics	185 seconds
Step 3: Constraints on PART, PART_GROUP and LINEORDER	Constraints	4 seconds
Step 4: PK on CUSTOMER and FK on LINEORDER referencing CUSTOMER	Constraints	4 seconds



# Step 4: PK on CUSTOMER – FK on LINEORDER



# What You Must Do: Partitioning



- Typically RANGE or INTERVAL
  - Typically fact tables in star schemas are partitioned on a column of DATE data type
- Reduces the number of rows extracted from the fact table (i.e., early filtering)
- Improves manageability

# Step 5: Correct the partitioning strategy

The query is interested only in data from 2007 ( $d\_year=2007$ )

The partition pruning reduces the amount of data being scanned from the fact table

- LINEORDER is partitioned on LO\_COMMITDATE

```
SELECT * FROM USER_PART_KEY_COLUMNS WHERE name='LINEORDER';  
  
NAME          OBJECT_TYPE COLUMN_NAME      COLUMN_POSITION COLLATED_COLUMN_ID  
---  
LINEORDER     TABLE       LO_COMMITDATE    1
```

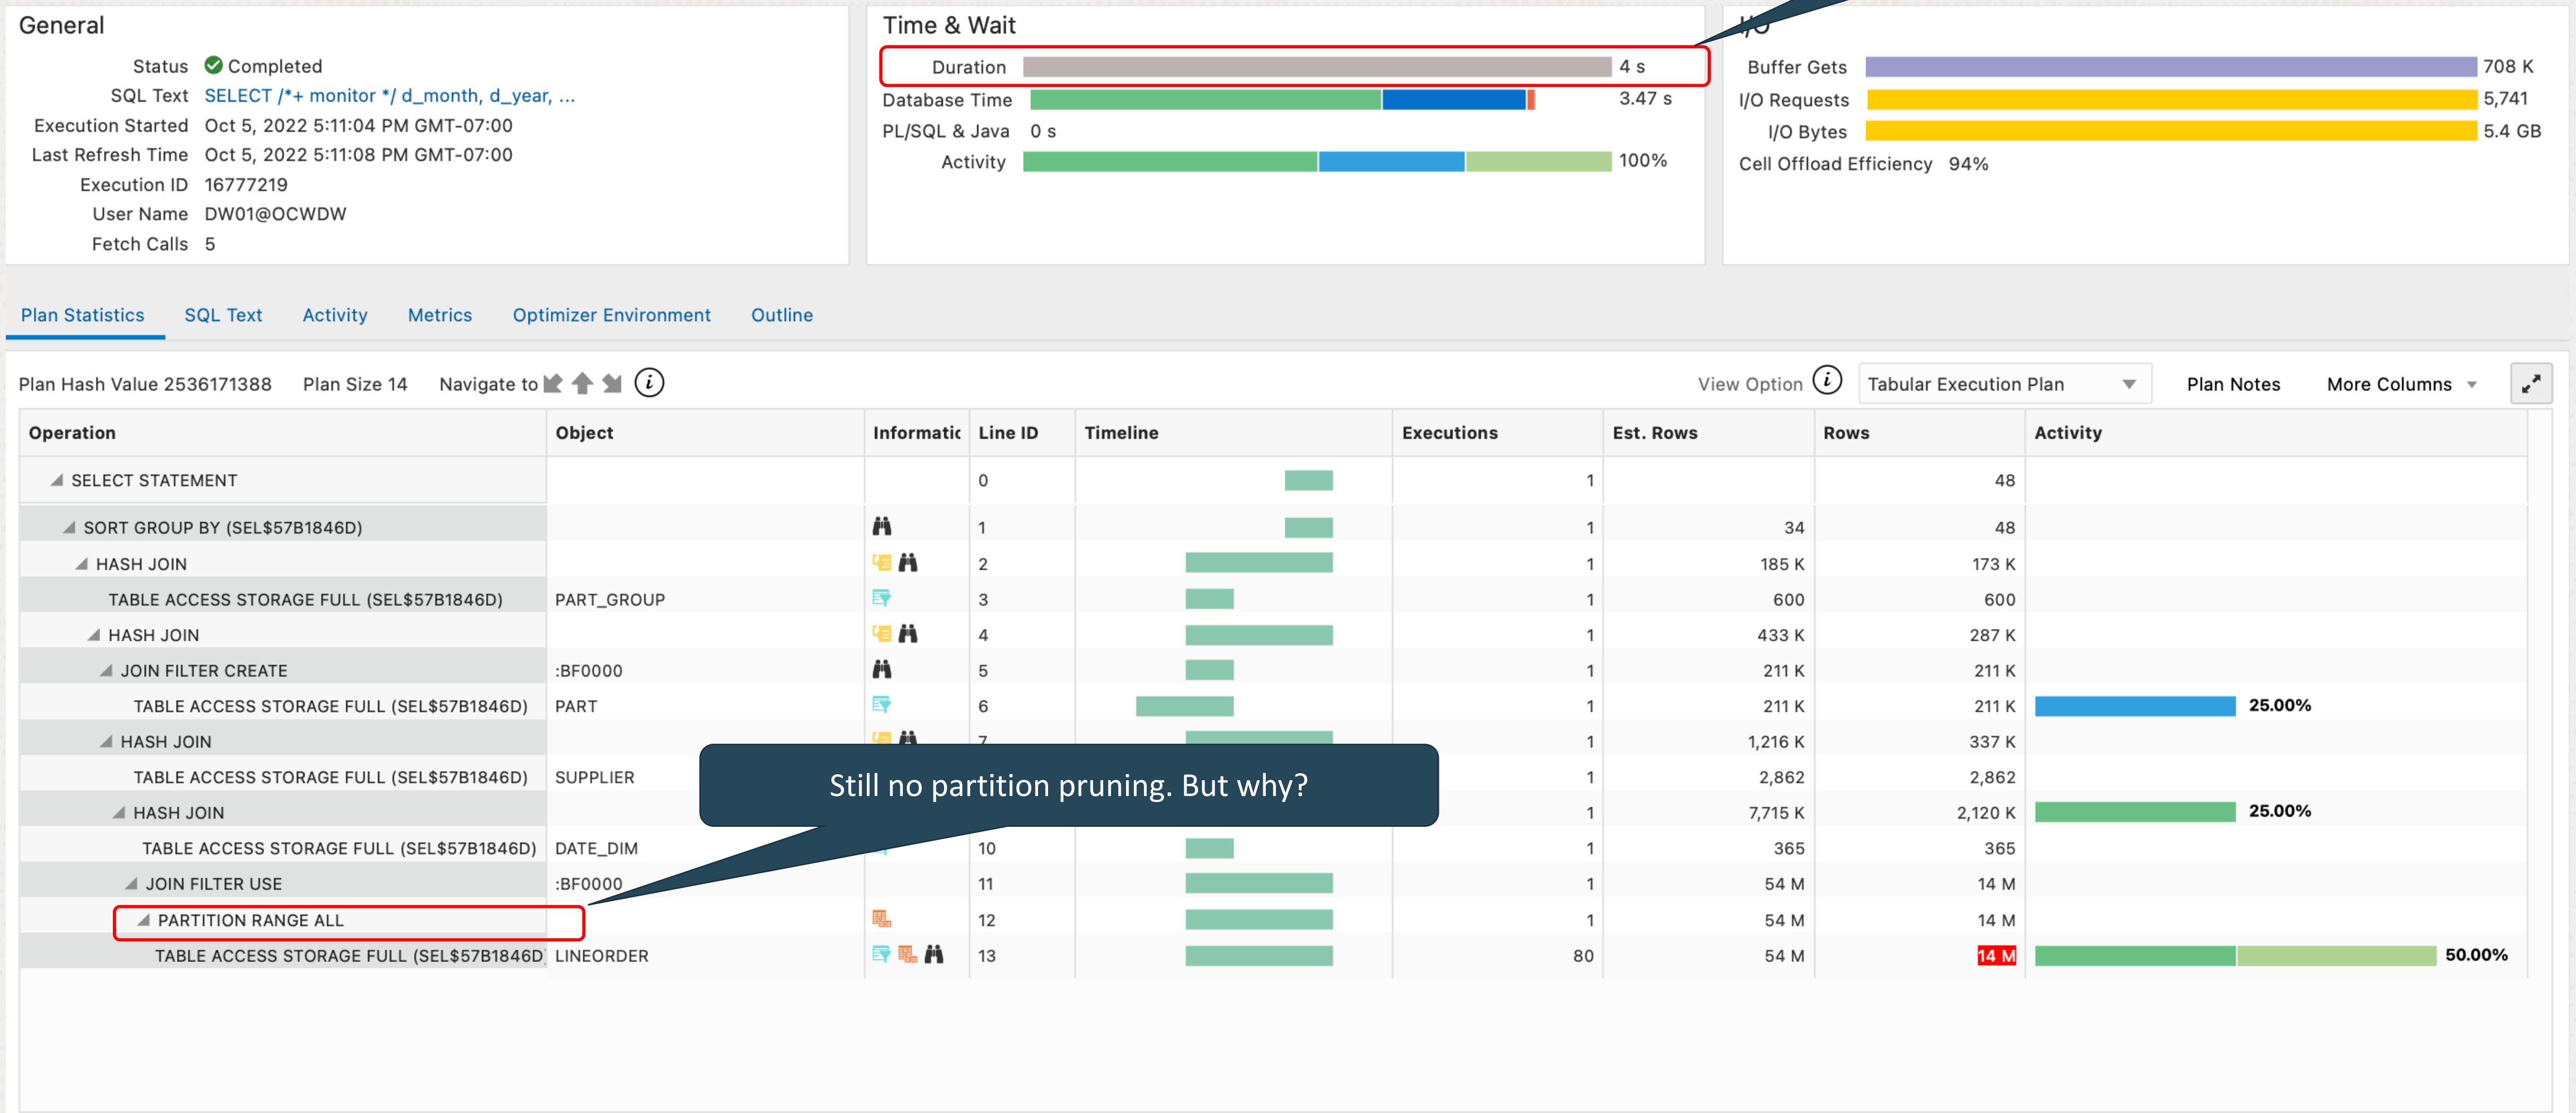
Partition pruning doesn't take place because LINEORDER is not partitioned on LO\_ORDERDATE

- LINEORDER is joined to DATE\_DIM on LO\_ORDERDATE

```
SELECT * FROM USER_PART_KEY_COLUMNS WHERE name='LINEORDER';  
  
NAME          OBJECT_TYPE COLUMN_NAME      COLUMN_POSITION COLLATED_COLUMN_ID  
---  
LINEORDER     TABLE       LO_COMMITDATE    1
```

Step5: running time: 4 seconds

# Step 5: LINEORDER partitioned on LO\_ORDERID

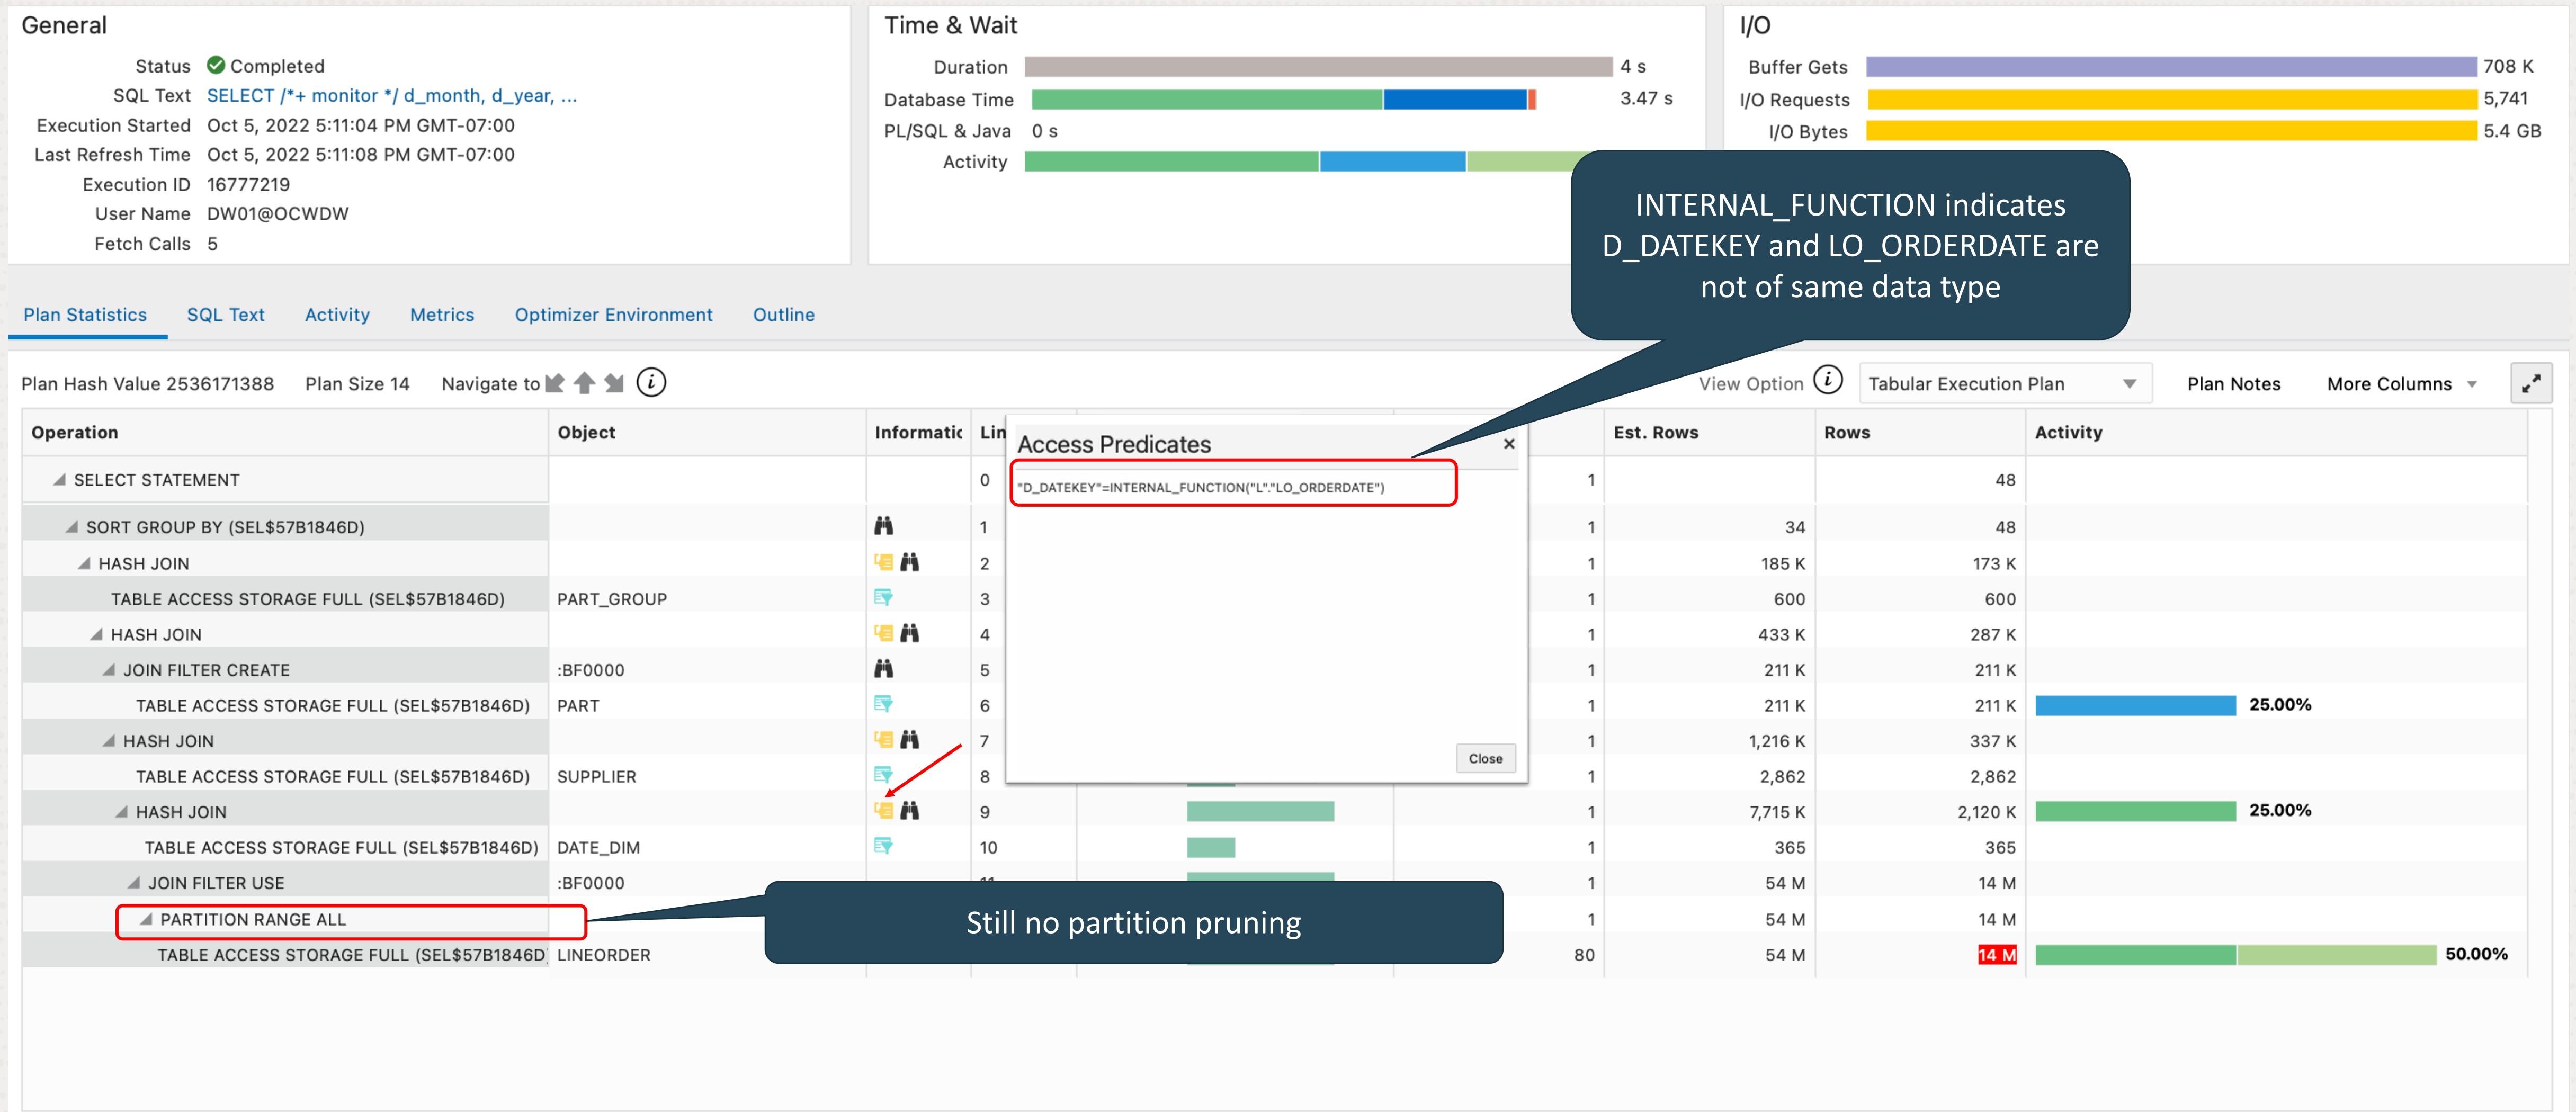


# Worked example: Progress

	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds
Step 2: Column group on SUPPLIER	Statistics	185 seconds
Step 3: Constraints on PART, PART_GROUP and LINEORDER	Constraints	4 seconds
Step 4: PK on CUSTOMER and FK on LINEORDER referencing CUSTOMER	Constraints	4 seconds
Step 5: LINEORDER partitioned on LO_ORDERDATE	Partitioning	4 seconds



# Step 5: LINEORDER partitioned on LO\_ORDERDATE



# What You Must Do: Data Types

- Data types need to be the same on Primary Key and Foreign Key columns
- Data type precision needs to be the same on Primary Key and Foreign Key columns
- Avoid runtime data type conversion

```
SELECT column_name, data_type FROM USER_TAB_COLUMNS WHERE table_name='LINEORDER' AND column_name='LO_ORDERDATE';
```

COLUMN_NAME	DATA_TYPE
-----	-----
LO_ORDERDATE	DATE

```
SELECT column_name, data_type FROM USER_TAB_COLUMNS WHERE table_name='DATE_DIM' AND column_name='D_DATEKEY';
```

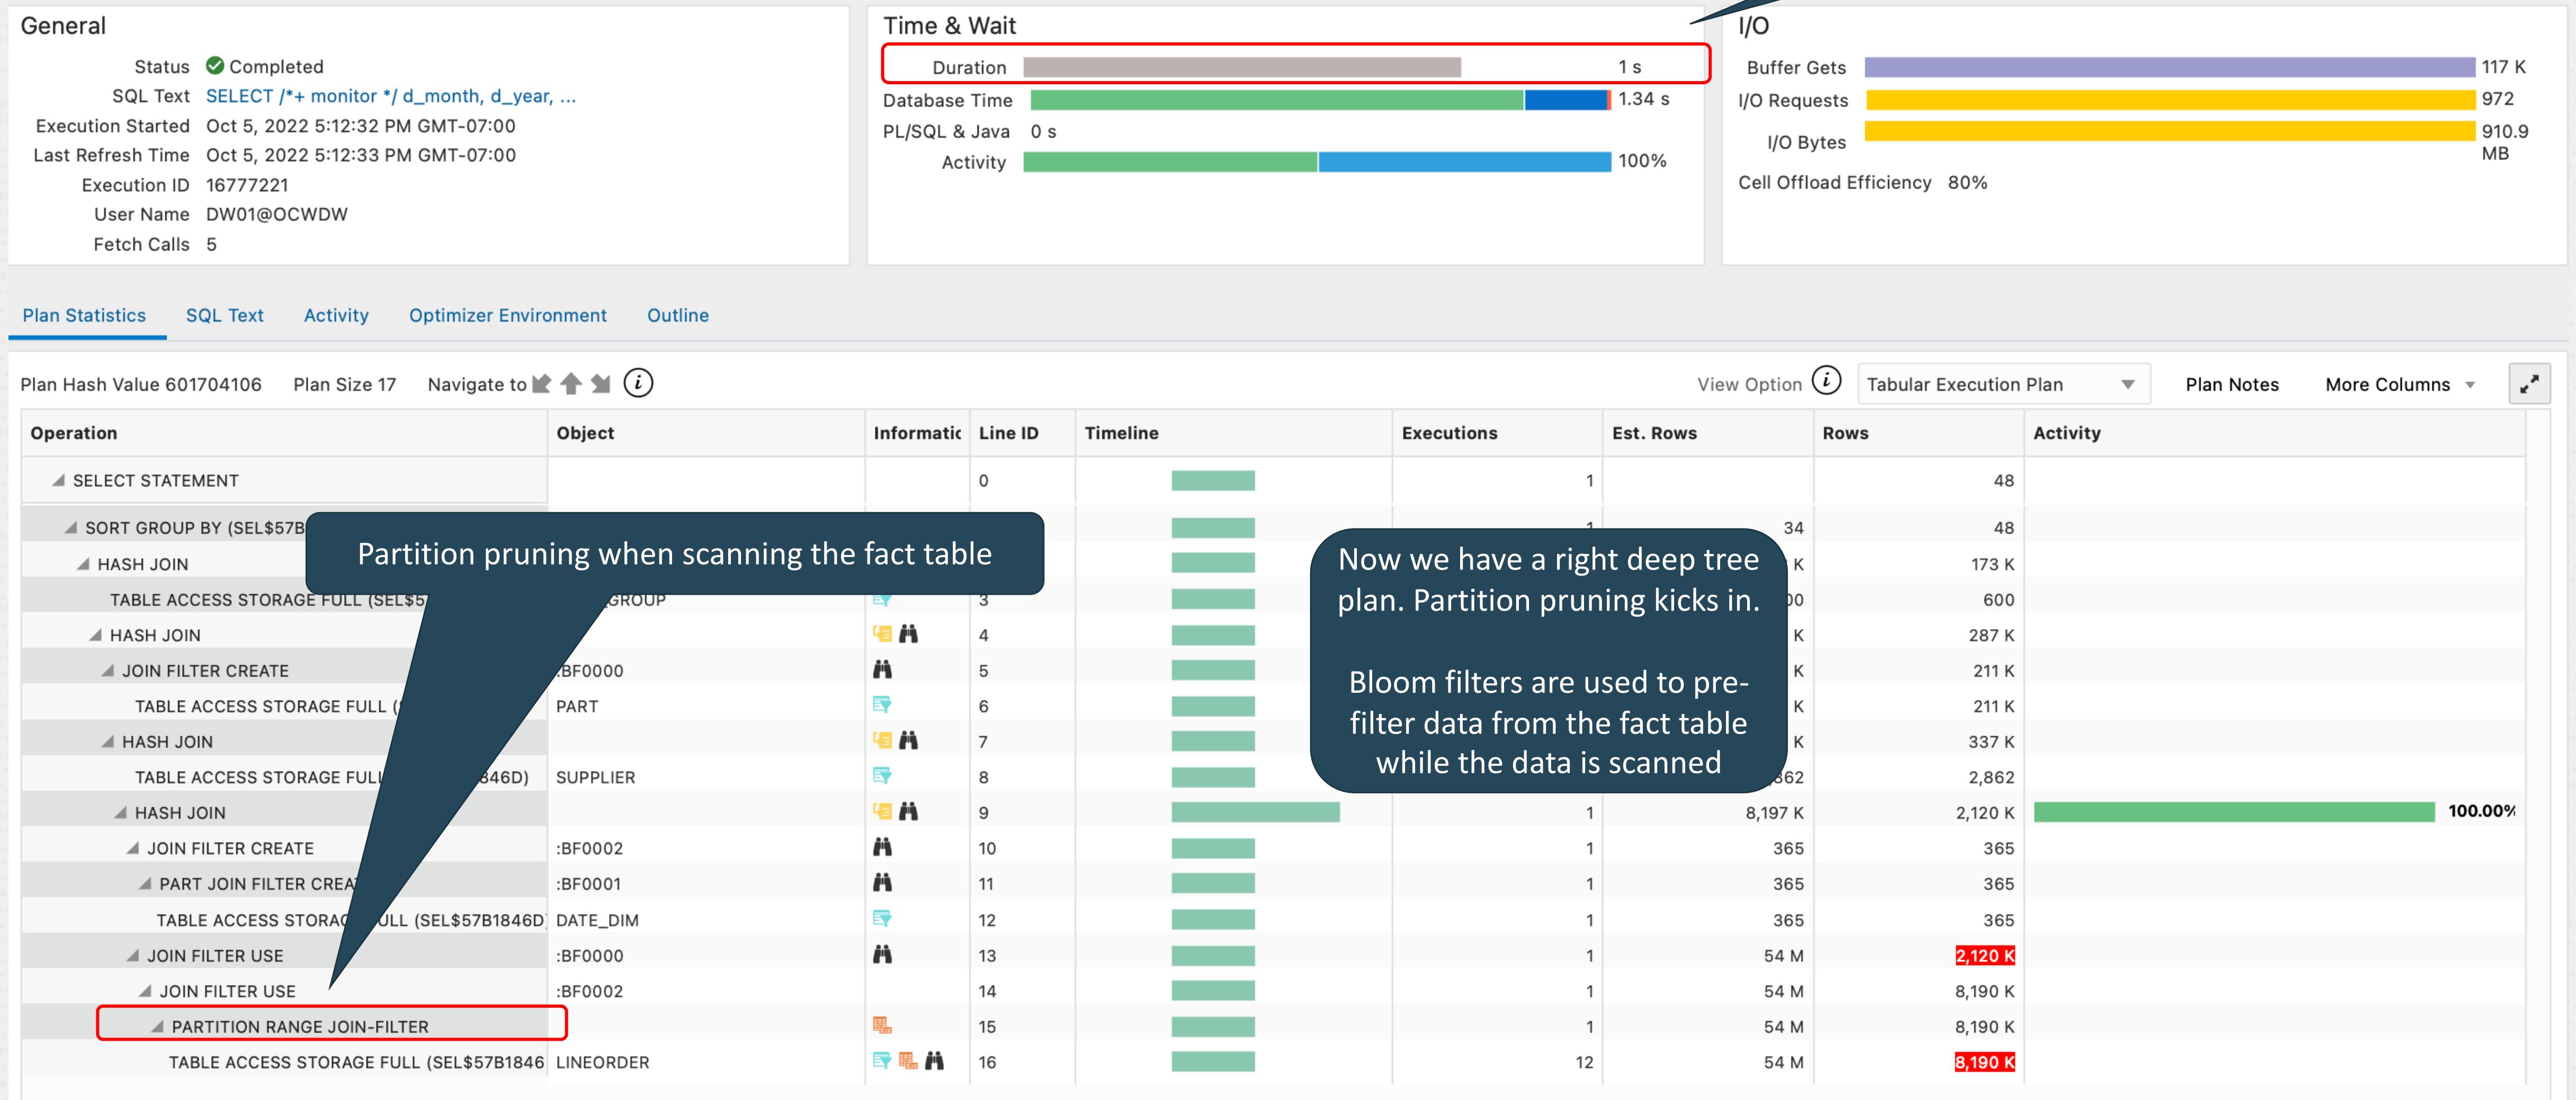
COLUMN_NAME	DATA_TYPE
-----	-----
D_DATEKEY	TIMESTAMP (6)

Data type mismatch on the join  
columns between LINEORDER and  
DATE\_DIM



Step6: running time: 1 second

# Step 6: Correct data type on DATE.D\_DATEKEY



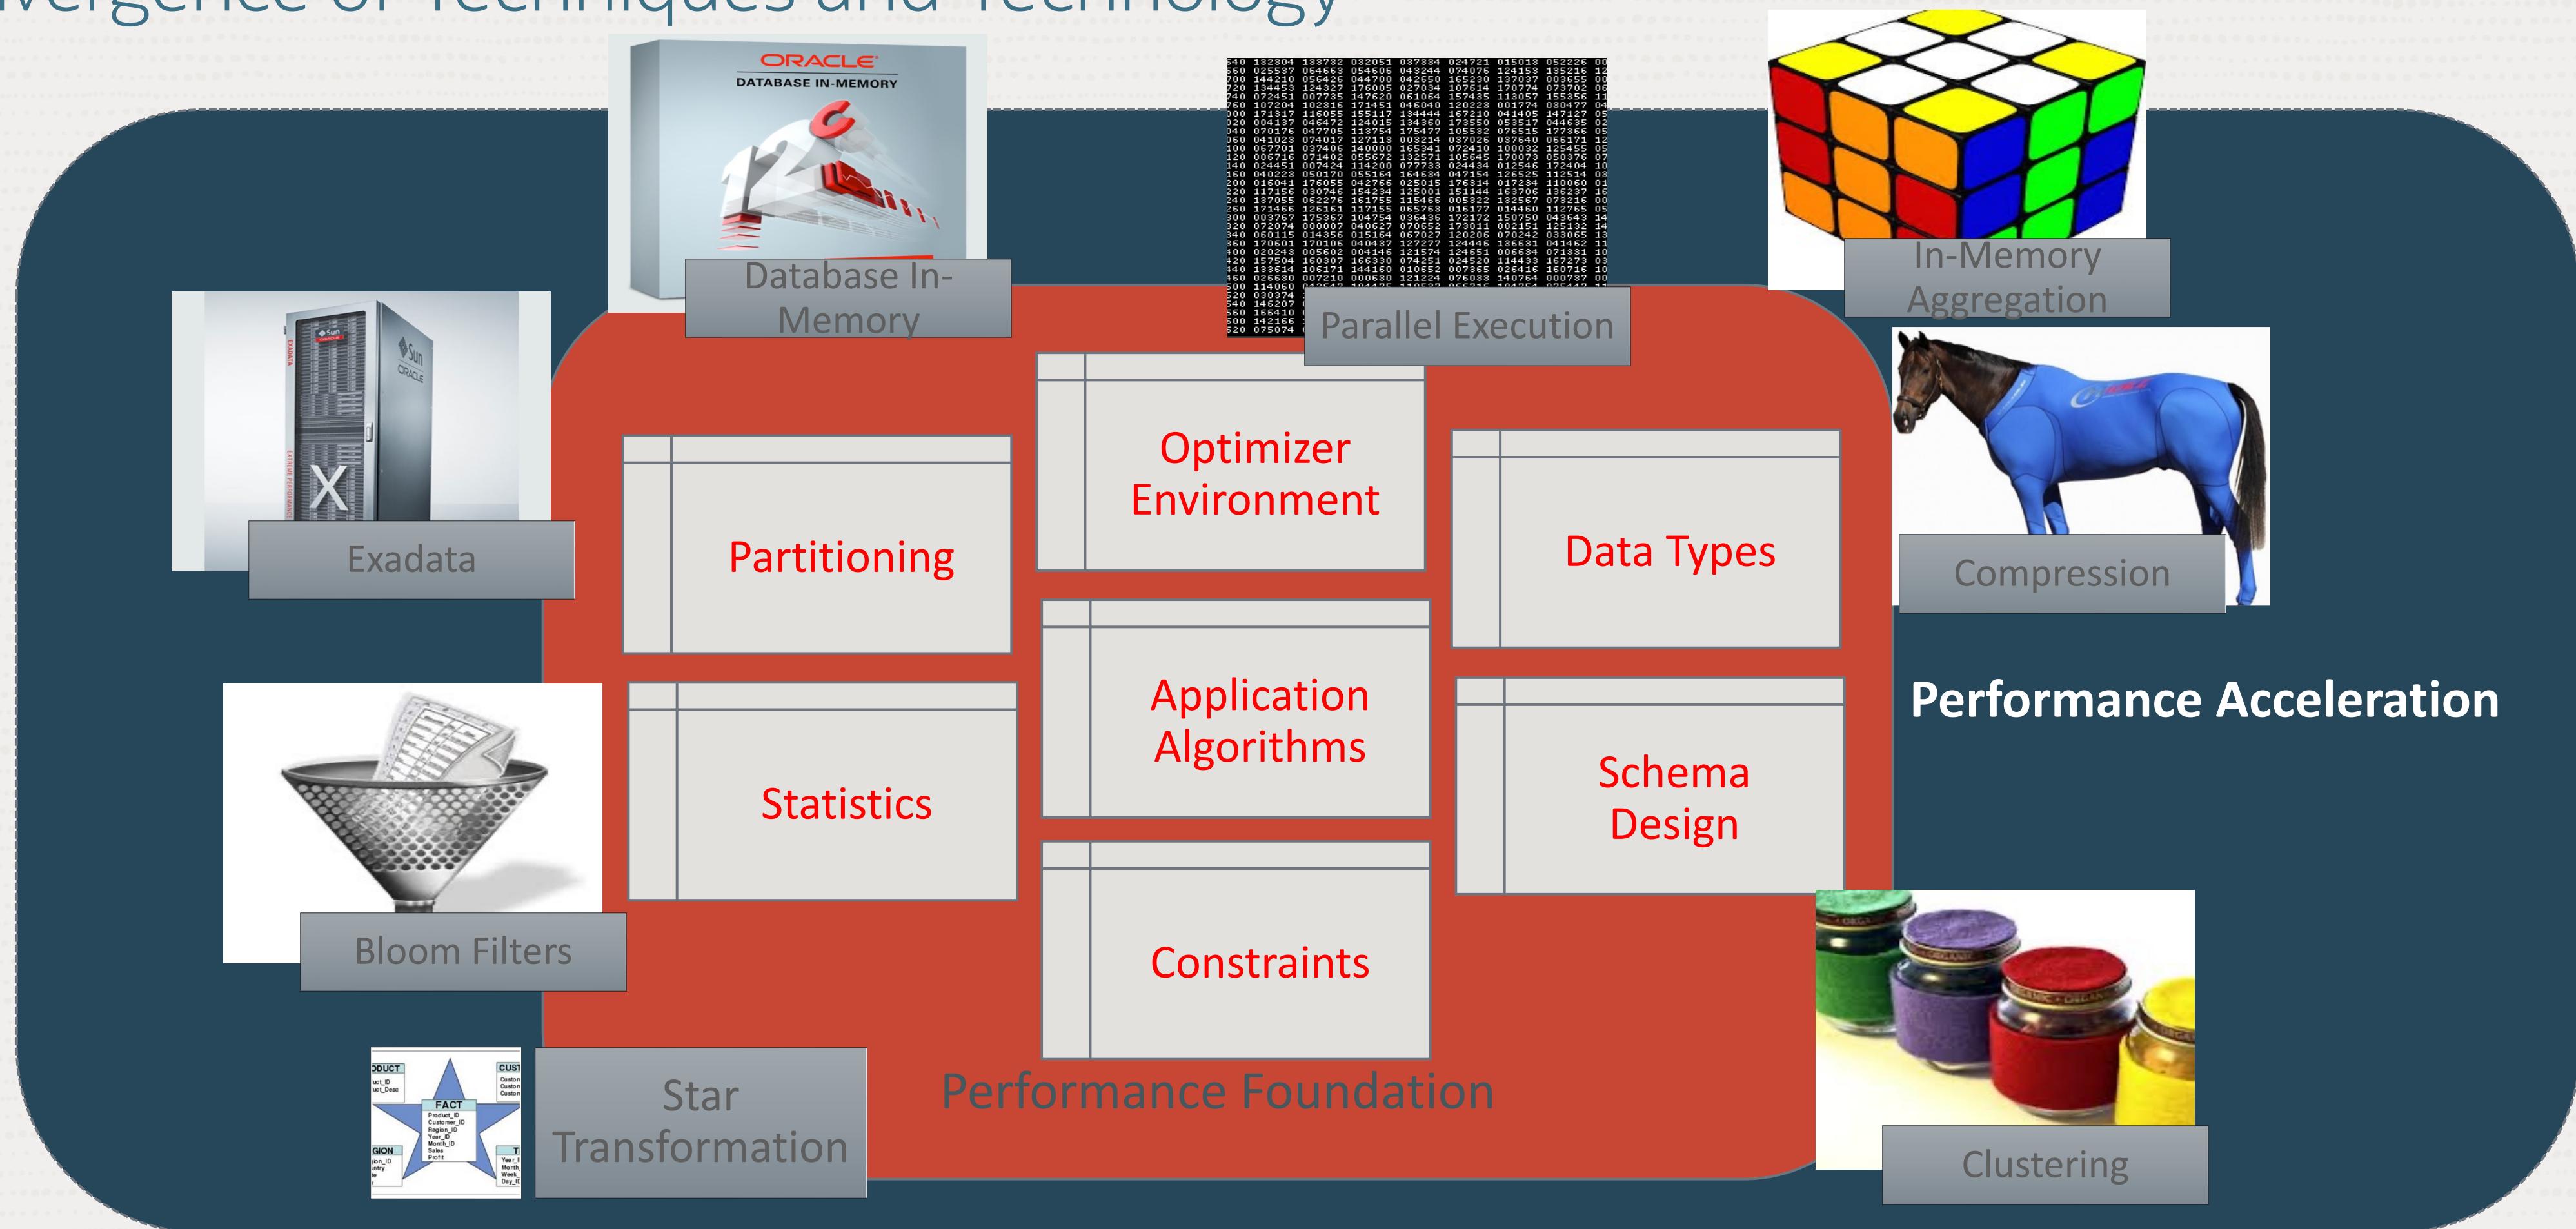
# Worked example: Progress

	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds
Step 2: Column group on SUPPLIER	Statistics	185 seconds
Step 3: Constraints on PART, PART_GROUP and LINEORDER	Constraints	4 seconds
Step 4: PK on CUSTOMER and FK on LINEORDER referencing CUSTOMER	Constraints	4 seconds
Step 5: LINEORDER partitioned on LO_ORDERDATE	Partitioning	4 seconds
Step 6: Correct data types	Data types	1 second



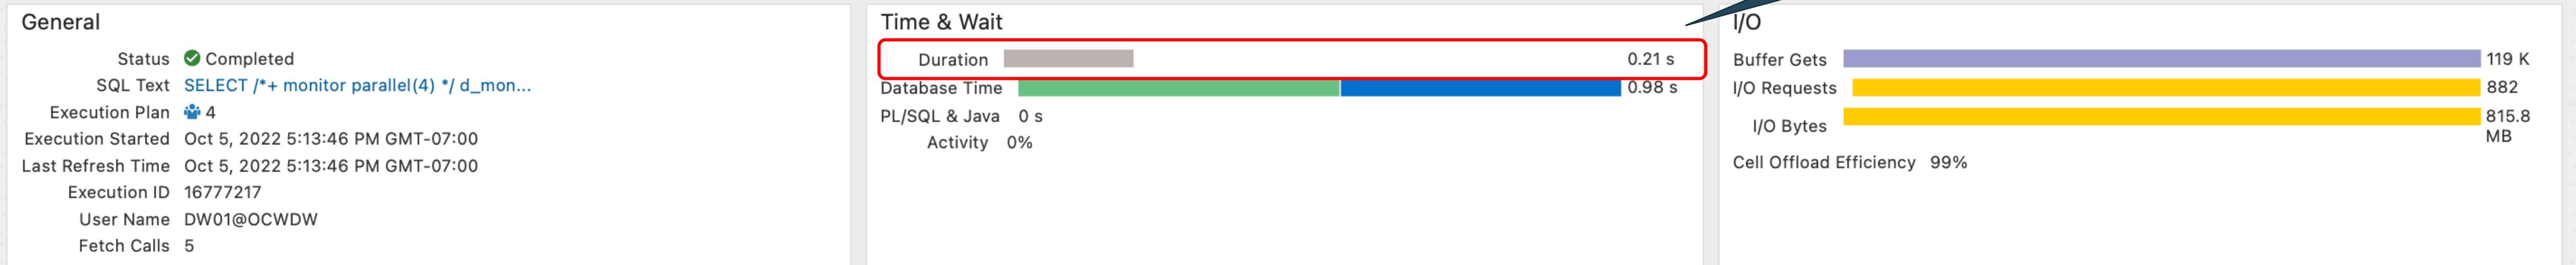
# The Dimensional Model

## Convergence of Techniques and Technology

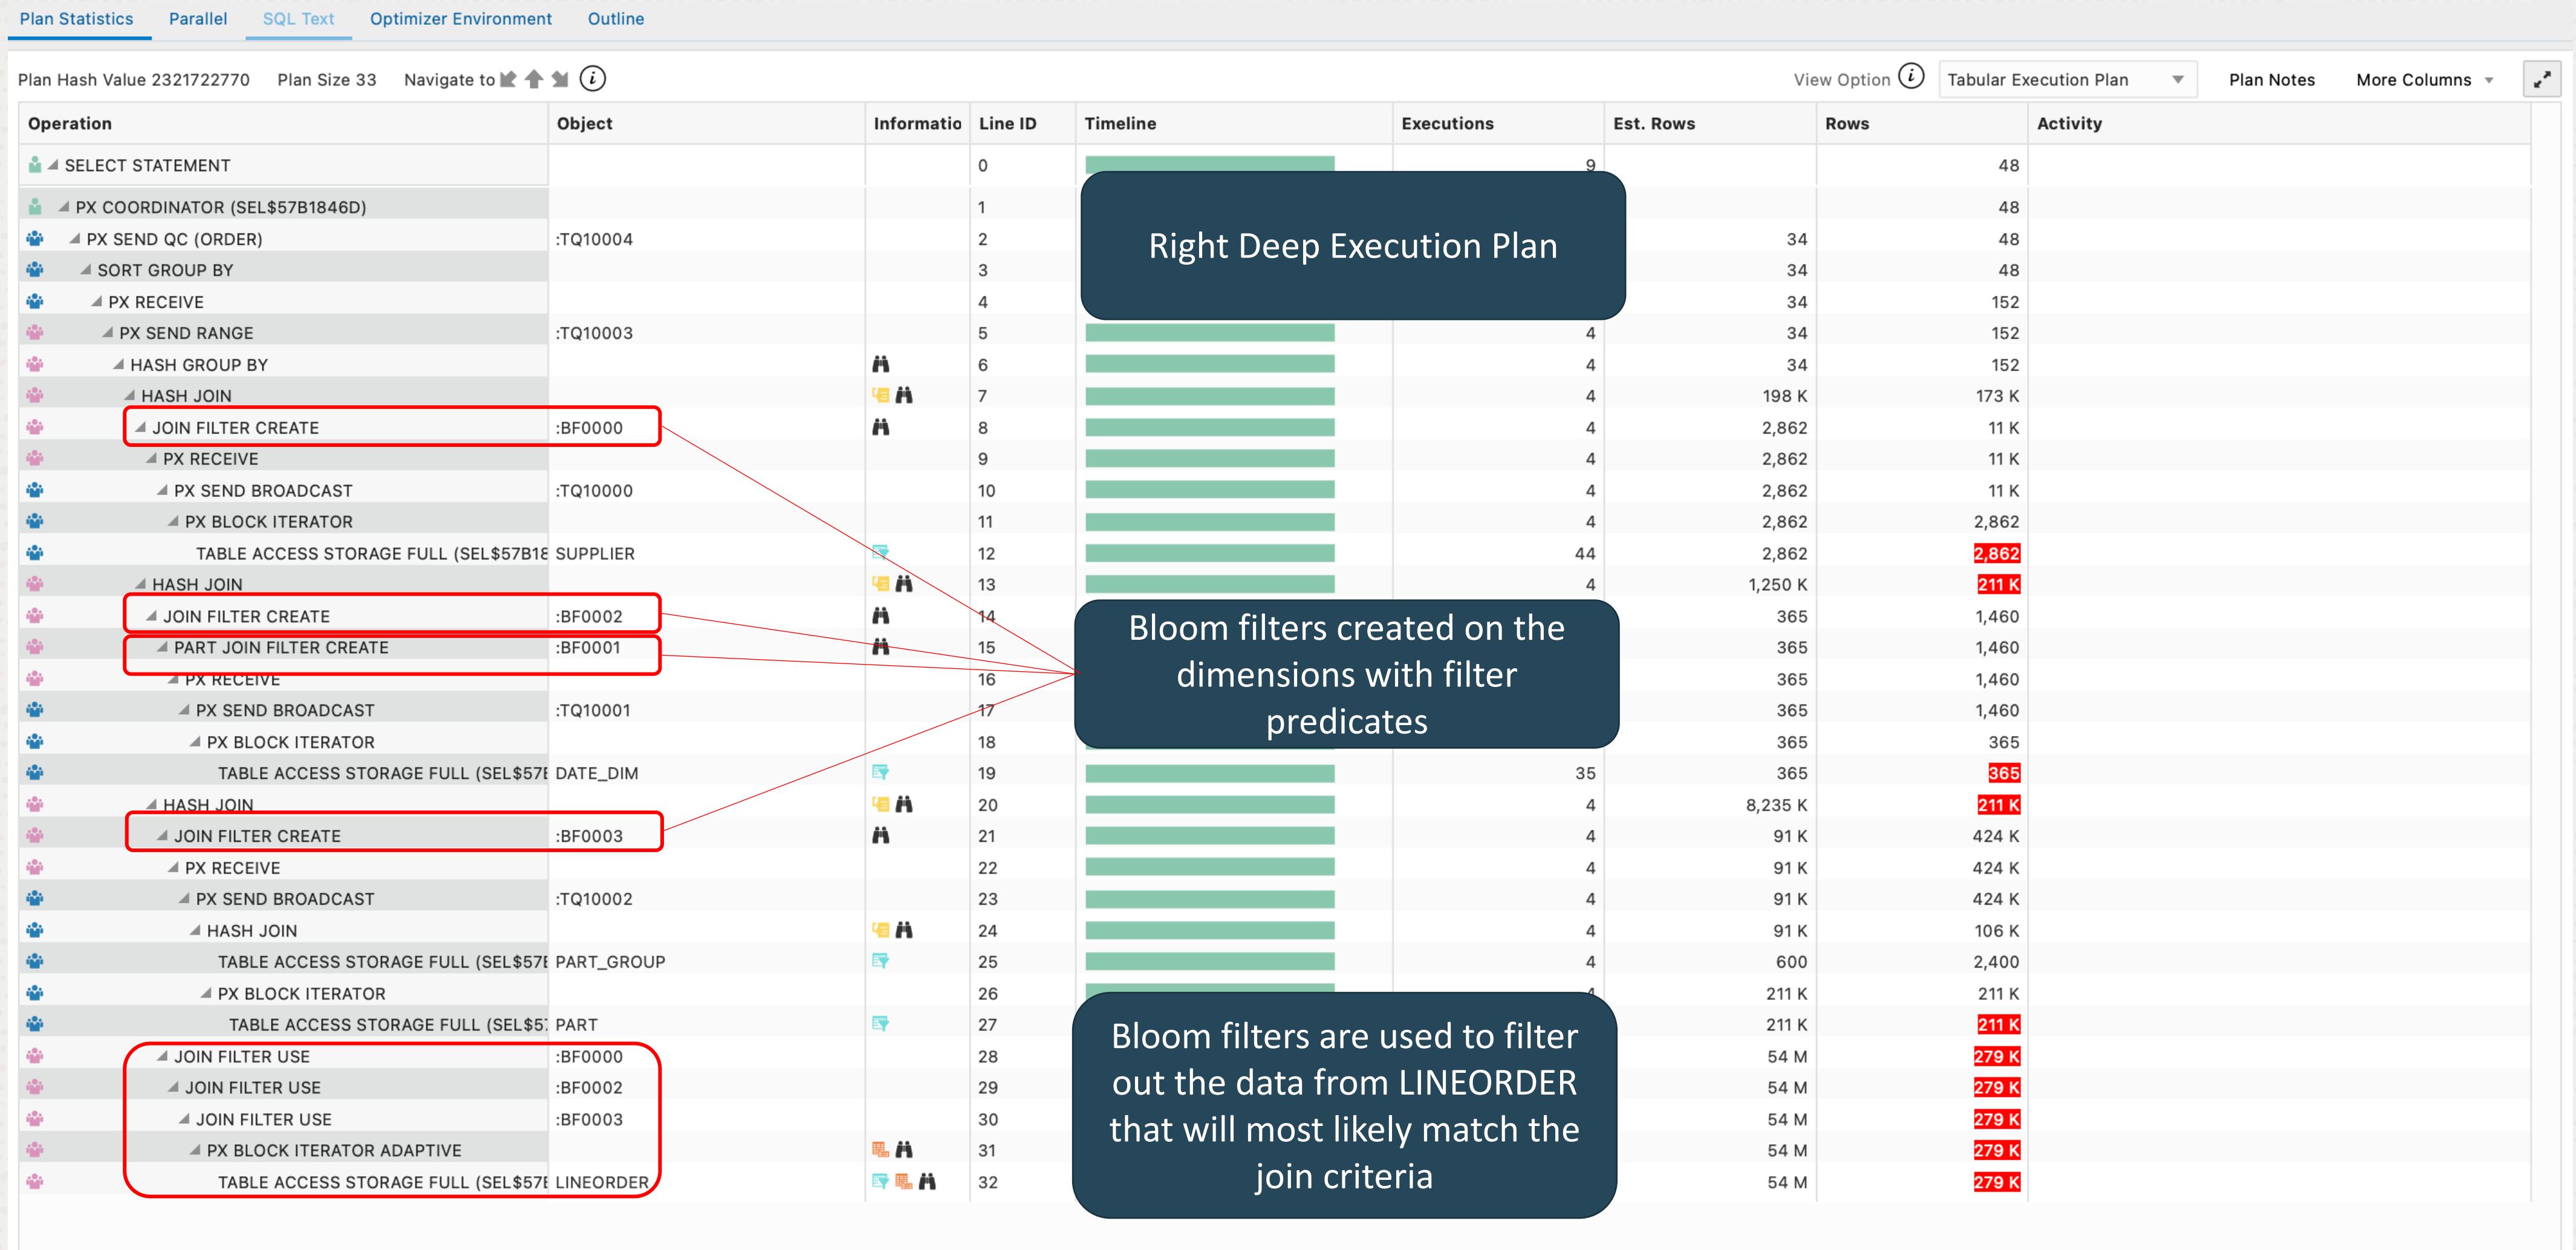


# Step 7: Parallel execution

Step7: running time: 0.21 seconds



# Step 7: Parallel execution



# Worked example: Progress

	Fix Category	Running Time
Initial Performance		256 seconds
Step 1: Statistics on PART	Statistics	211 seconds
Step 2: Column group on SUPPLIER	Statistics	185 seconds
Step 3: Constraints on PART, PART_GROUP and LINEORDER	Constraints	4 seconds
Step 4: PK on CUSTOMER and FK on LINEORDER referencing CUSTOMER	Constraints	4 seconds
Step 5: LINEORDER partitioned on LO_ORDERDATE	Partitioning	4 seconds
Step 6: Correct data types	Data types	1 second
Step 7: Parallel Execution	Other technologies	0.21s



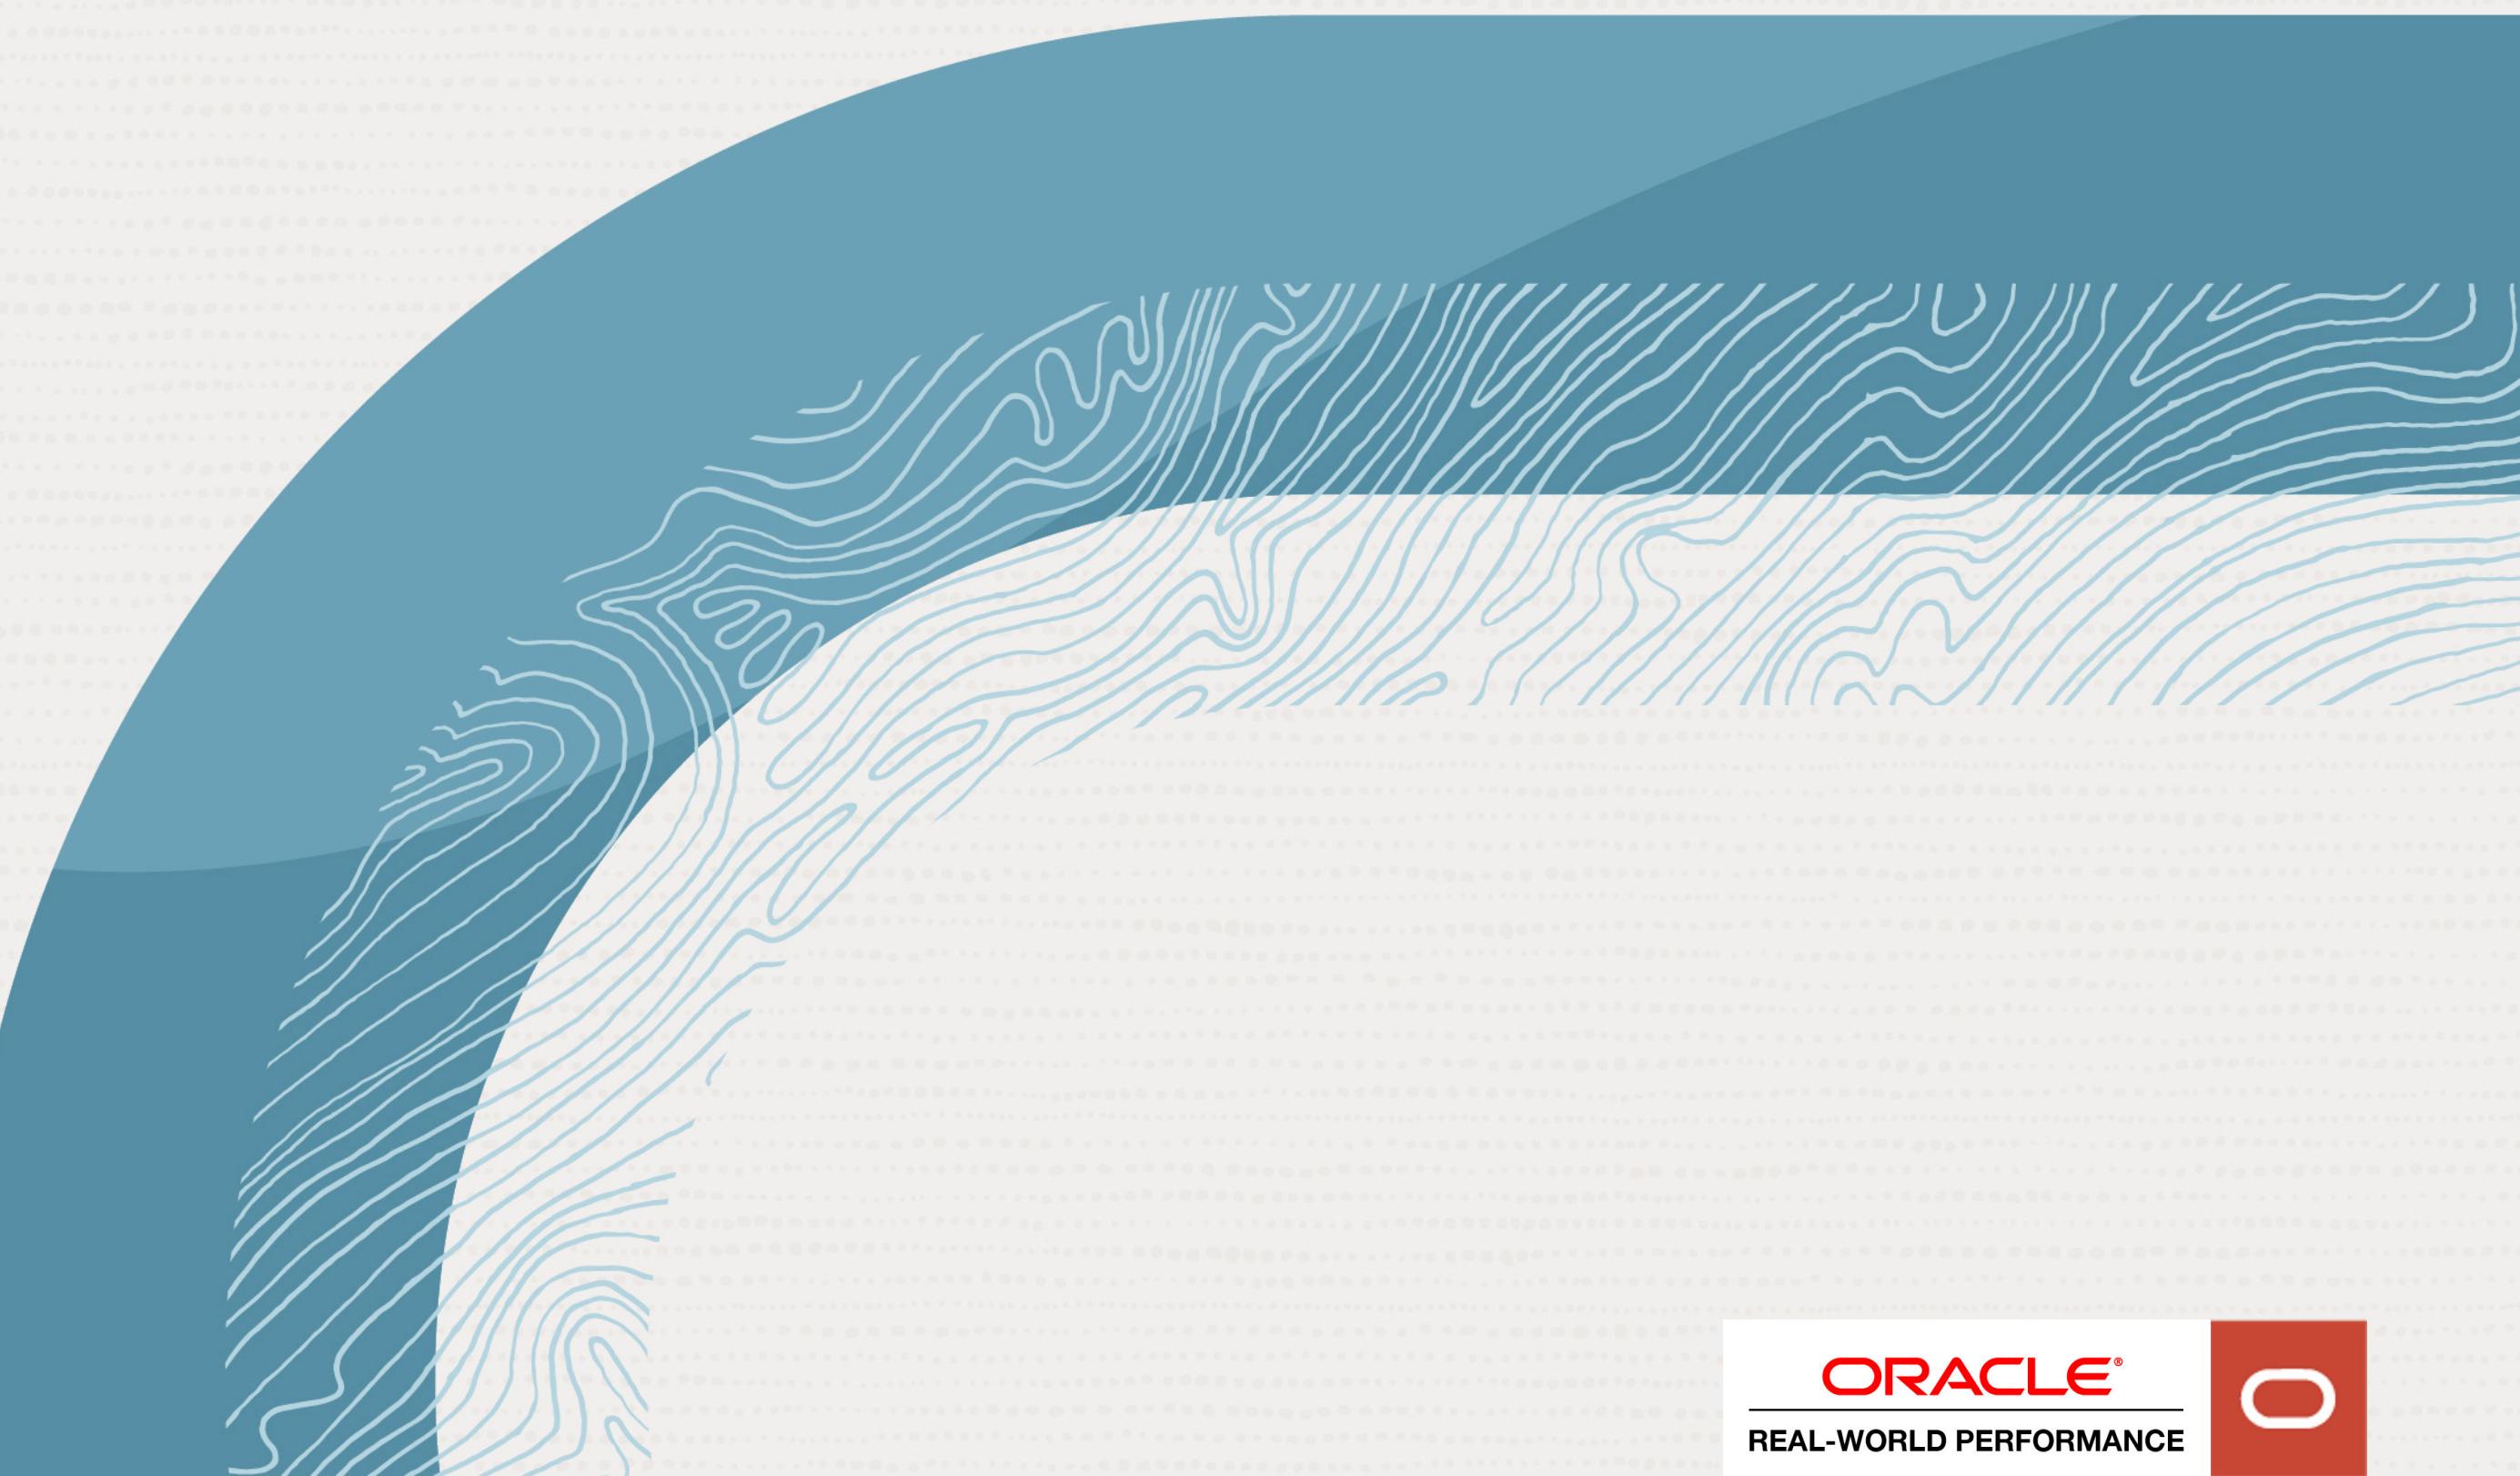
# What Do You Gain by Following the Prescription?

- Better cardinality estimates
- Better execution plans
- More access paths available
- Ability for the optimizer to perform many transformations and optimizations (join elimination, materialized view rewrites, In-Memory Aggregation transformation, and many more)
- Partition pruning
- Exploit other technologies for optimal performance
  - Parallel Execution
  - Materialized Views
  - Compression
  - Database In-Memory



# Star Query Execution on Autonomous Data Warehouse

Analysis of the default performance of our worked example



# Star Query on Autonomous Data Warehouse

Analyze out of the box performance of the worked example when running on Autonomous Data Warehouse (ADW)

The session that executes the query will use two different services

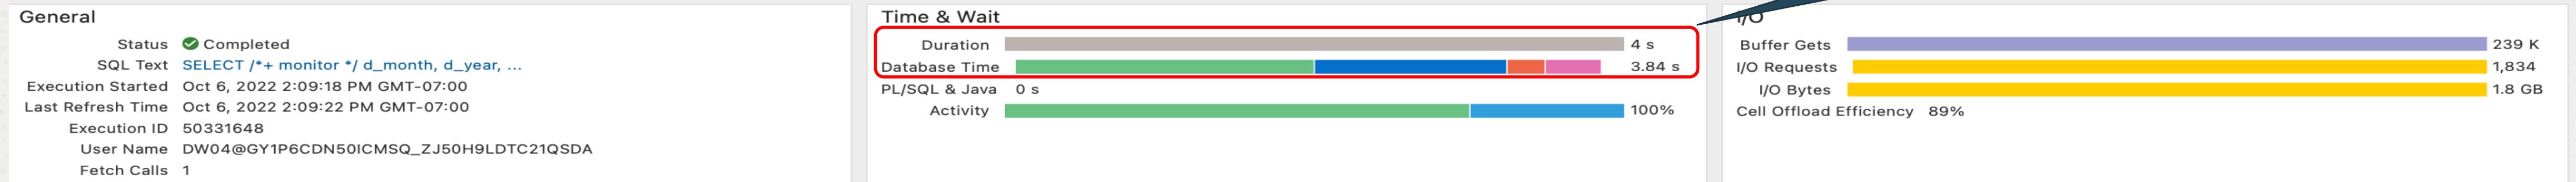
- Low Service (serial execution)
- Medium Service (parallel execution with DoP 4)



# ADW: Out of the Box Performance (Serial)

- The query is executed by a session connected to the Low service
  - SQL statements executed using the Low service run serially

Initial running time: 4 seconds



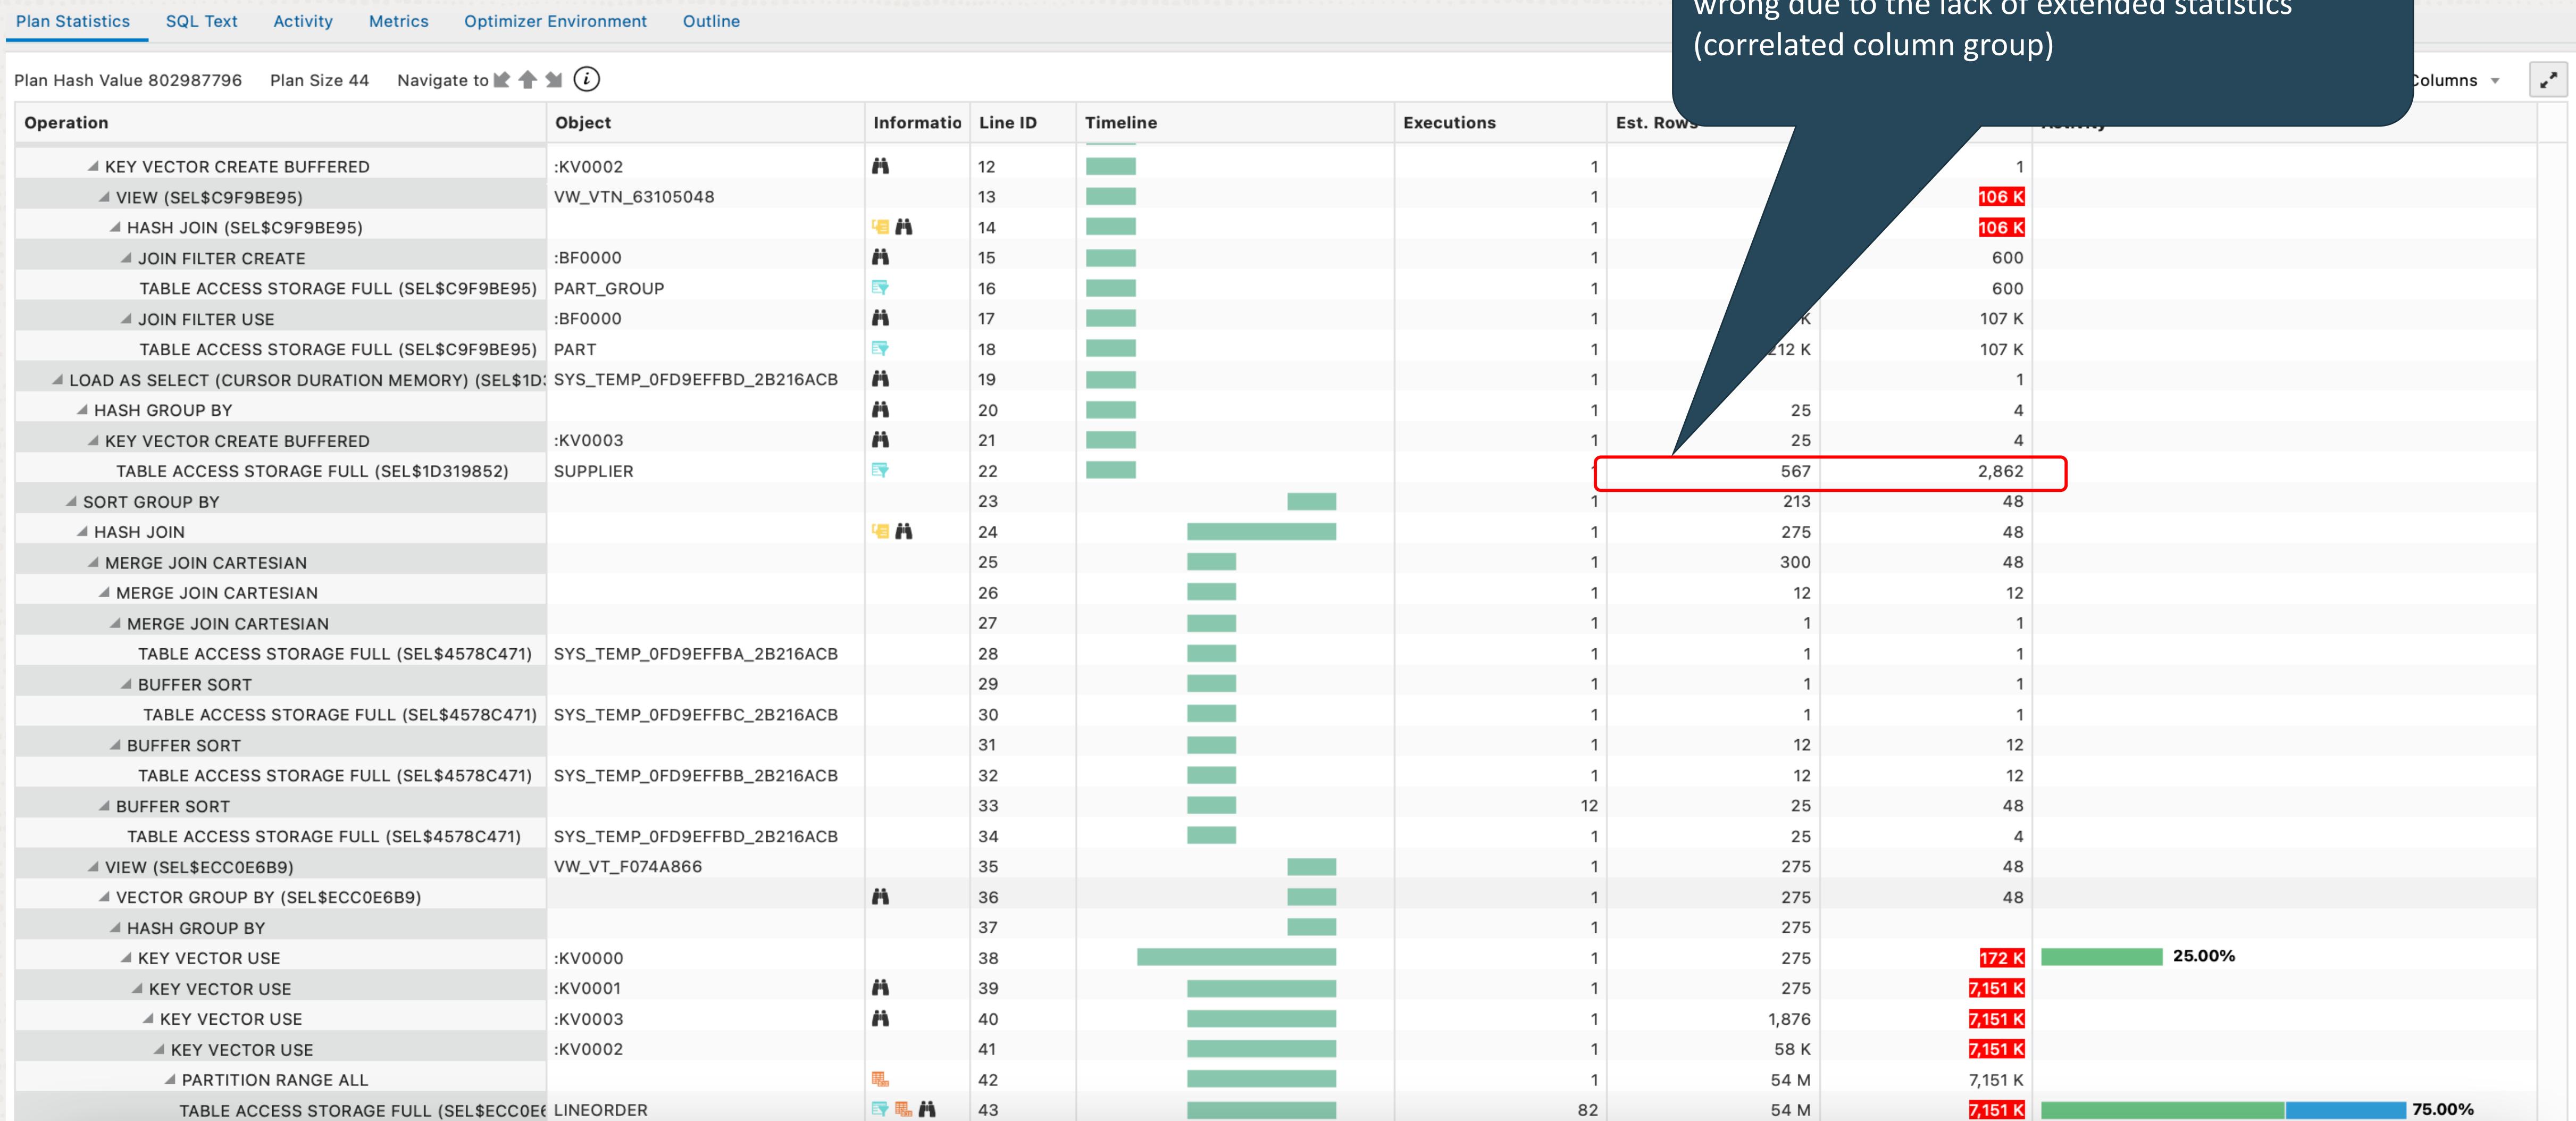
# ADW: Out of the Box Performance (Serial)

Cardinality estimate for PART is spot on:  
 - High frequency stats reduce the chance for stale stats

The execution plan is different and uses In-Memory Aggregation (Vector Transformation)  
 In ADW IMA is enabled and can take advantage of the data stored in columnar store that resides on flash disks

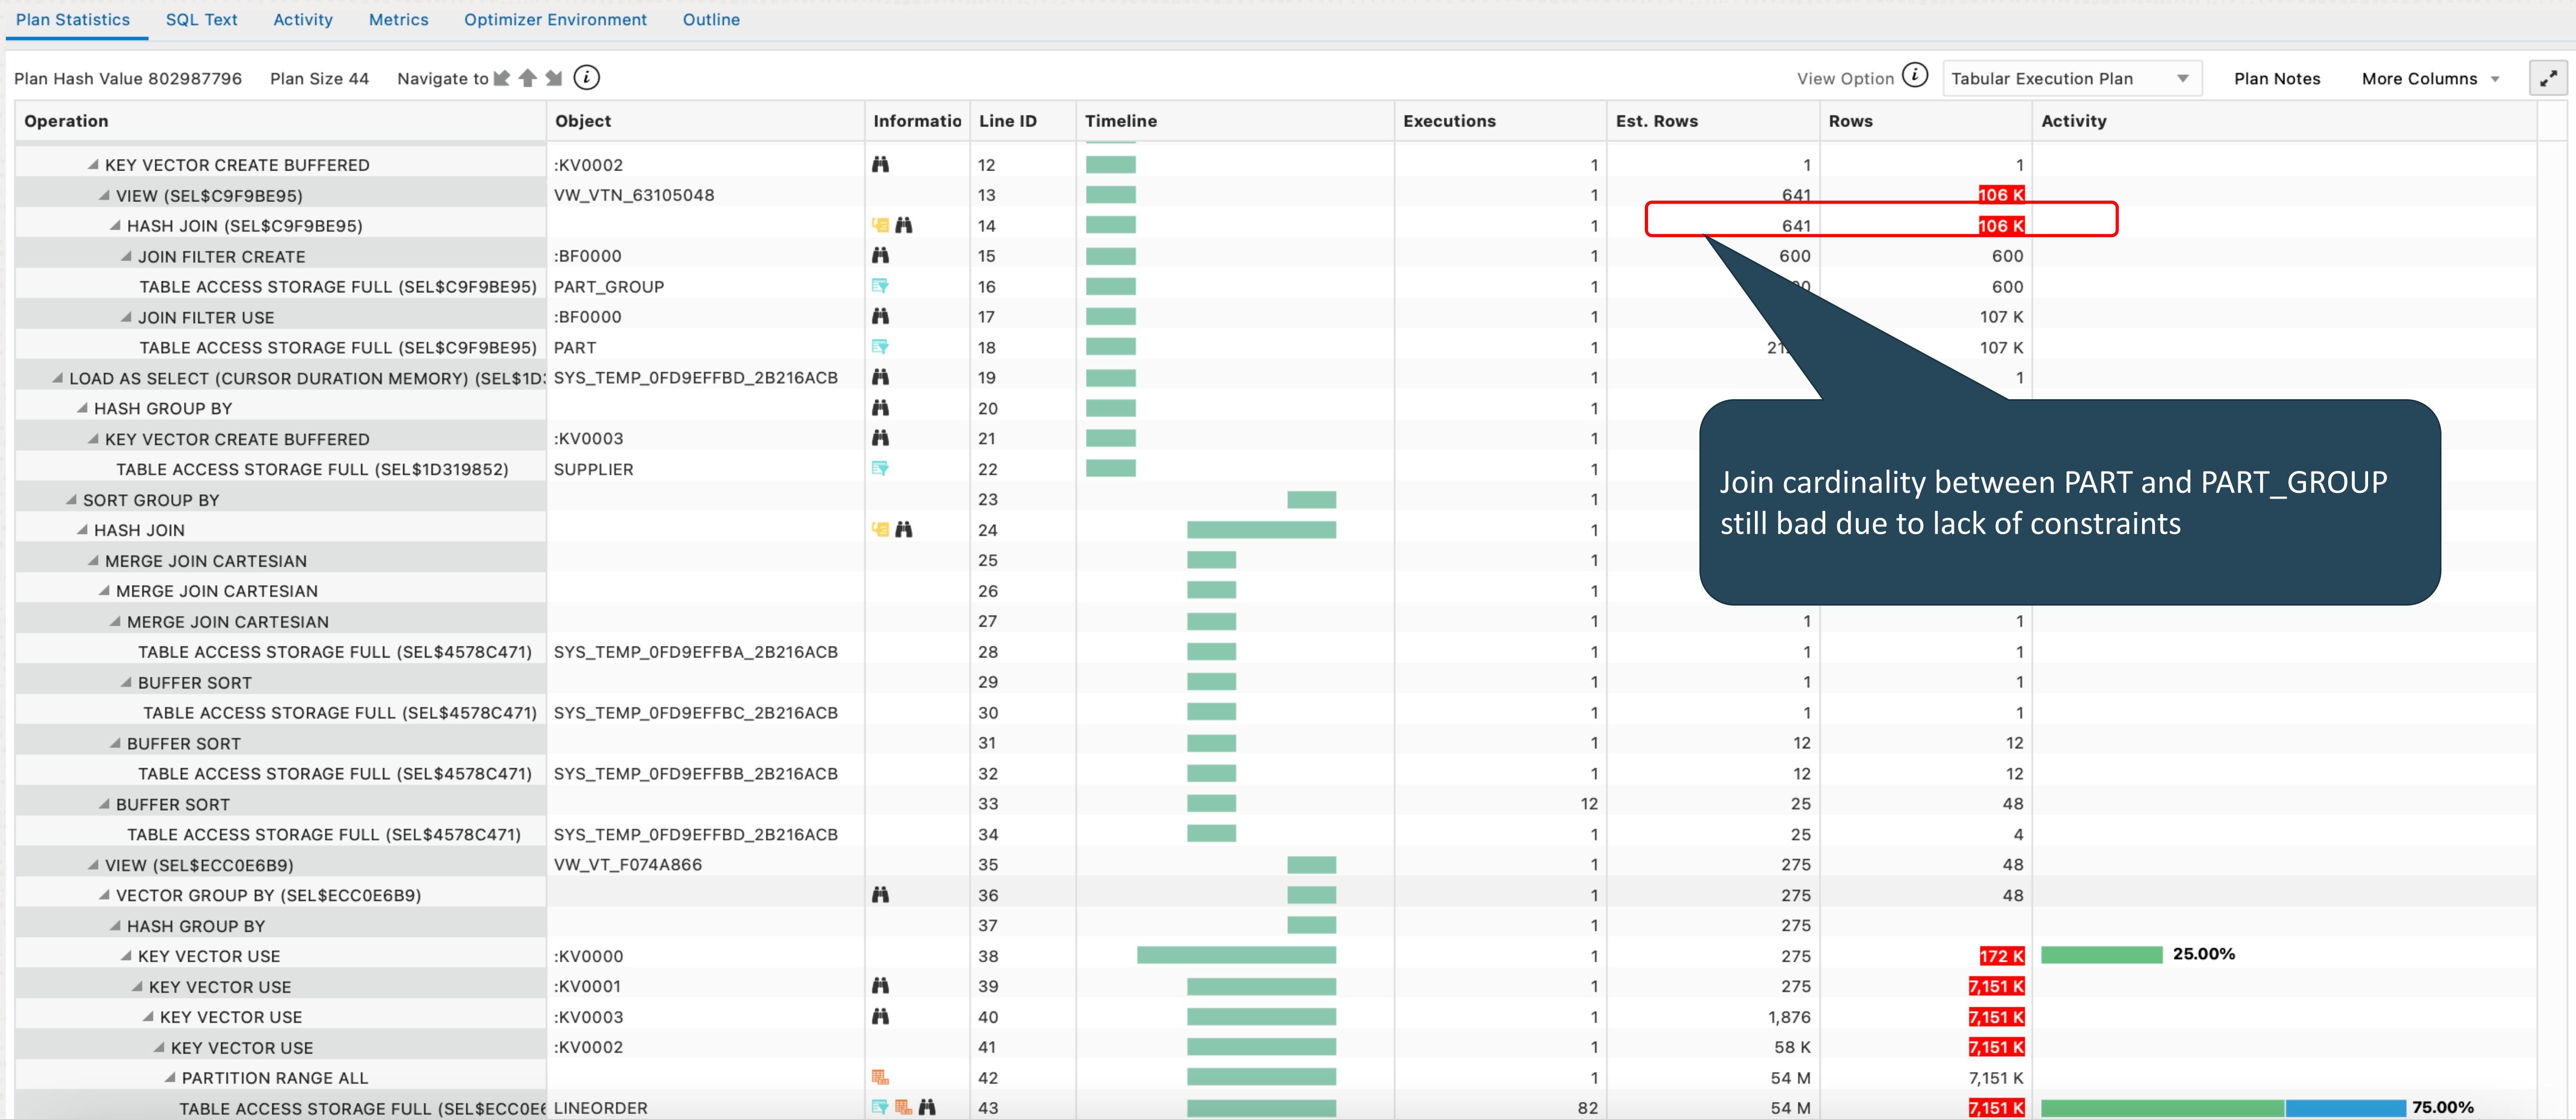
Operation	Object	Information	Line Number	Activity
KEY VECTOR CREATE BUFFERED	:KV0002	12		
VIEW (SEL\$C9F9BE95)	VW_VTN_63105048	13		
HASH JOIN (SEL\$C9F9BE95)		14		
JOIN FILTER CREATE	:BF0000	15		
TABLE ACCESS STORAGE FULL (SEL\$C9F9BE95)	PART_GROUP	16		
JOIN FILTER USE	:BF0000	17		
TABLE ACCESS STORAGE FULL (SEL\$C9F9BE95)	PART	18		
LOAD AS SELECT (CURSOR DURATION MEMORY) (SEL\$1D319852)	SYS_TEMP_0FD9EFFBD_2B216ACB	19		
HASH GROUP BY		20		
KEY VECTOR CREATE BUFFERED	:KV0003	21		
TABLE ACCESS STORAGE FULL (SEL\$1D319852)	SUPPLIER	22		
SORT GROUP BY		23		
HASH JOIN		24		
MERGE JOIN CARTESIAN		25		
MERGE JOIN CARTESIAN		26		
MERGE JOIN CARTESIAN		27		
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBA_2B216ACB	28		
BUFFER SORT		29		
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBC_2B216ACB	30		
BUFFER SORT		31		
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBB_2B216ACB	32		
BUFFER SORT		33		
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBD_2B216ACB	34		
VIEW (SEL\$ECC0E6B9)	VW_VT_F074A866	35		
VECTOR GROUP BY (SEL\$ECC0E6B9)		36		
HASH GROUP BY		37		
KEY VECTOR USE	:KV0000	38		
KEY VECTOR USE	:KV0001	39		
KEY VECTOR USE	:KV0003	40		
KEY VECTOR USE	:KV0002	41		
PARTITION RANGE ALL		42		
TABLE ACCESS STORAGE FULL (SEL\$ECC0E6E)	LINEORDER	43		

# ADW: Out of the Box Performance (Serial)



But, the cardinality estimate for SUPPLIER is still wrong due to the lack of extended statistics (correlated column group)

# ADW: Out of the Box Performance (Serial)



# ADW: Out of the Box Performance (Serial)

Plan Statistics   SQL Text   Activity   Metrics   Optimizer Environment   Outline

Plan Hash Value 802987796   Plan Size 44   Navigate to ↕ ↑ ↓ (i)   View Option (i)   Tabular Execution Plan ▾   Plan Notes   More Columns ▾

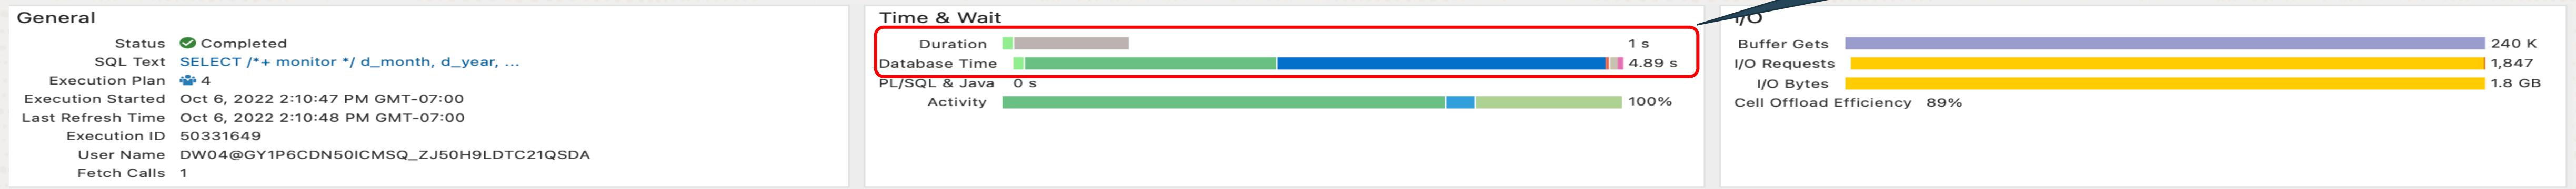
Operation	Object	Information	Line ID	Timeline	Executions	Est. Rows	Rows	Activity
▲ KEY VECTOR CREATE BUFFERED	:KV0002	双眼	12		1	1	1	1
▲ VIEW (SEL\$C9F9BE95)	VW_VTN_63105048		13		1	641	641	106 K
▲ HASH JOIN (SEL\$C9F9BE95)		双眼	14		1	641	641	106 K
▲ JOIN FILTER CREATE	:BF0000	双眼	15		1	600	600	600
TABLE ACCESS STORAGE FULL (SEL\$C9F9BE95)	PART_GROUP	文件夹	16		1	600	600	600
▲ JOIN FILTER USE	:BF0000	双眼	17		1	212 K	212 K	107 K
TABLE ACCESS STORAGE FULL (SEL\$C9F9BE95)	PART	文件夹	18		1	212 K	212 K	107 K
▲ LOAD AS SELECT (CURSOR DURATION MEMORY) (SEL\$1D319852)	SYS_TEMP_0FD9EFFBD_2B216ACB	双眼	19		1	1	1	
▲ HASH GROUP BY		双眼	20		1	25	25	4
▲ KEY VECTOR CREATE BUFFERED	:KV0003	双眼	21		1	25	25	4
TABLE ACCESS STORAGE FULL (SEL\$1D319852)	SUPPLIER	文件夹	22		1	567	567	2,862
▲ SORT GROUP BY			23		1	213	213	48
▲ HASH JOIN		双眼	24		1	275	275	48
▲ MERGE JOIN CARTESIAN			25		1	300	300	48
▲ MERGE JOIN CARTESIAN			26		1	12	12	12
▲ MERGE JOIN CARTESIAN			27		1	1	1	1
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBA_2B216ACB		28		1	1	1	1
▲ BUFFER SORT			29		1	1	1	1
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBC_2B216ACB		30		1	1	1	1
▲ BUFFER SORT			31		1	12	12	12
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBB_2B216ACB		32		1	12	12	12
▲ BUFFER SORT			33		12	25	25	48
TABLE ACCESS STORAGE FULL (SEL\$4578C471)	SYS_TEMP_0FD9EFFBD_2B216ACB		34		1	25	25	4
▲ VIEW (SEL\$ECC0E6B9)	VW_VT_F074A866		35		1	275	275	48
▲ VECTOR GROUP BY (SEL\$ECC0E6B9)			36		1	275	275	48
▲ HASH GROUP BY			37		1	275	275	
▲ KEY VECTOR USE	:KV0000	双眼	38		1	275	172 K	172 K
▲ KEY VECTOR USE	:KV0001	双眼	39		1	275	275	7,151 K
▲ KEY VECTOR USE	:KV0003	双眼	40		1	1,876	1,876	7,151 K
▲ KEY VECTOR USE	:KV0002	双眼	41		1	58 K	58 K	7,151 K
▲ PARTITION RANGE ALL		方块	42		1	54 M	54 M	7,151 K
TABLE ACCESS STORAGE FULL (SEL\$ECC0E6E)	LINEORDER	文件夹	43		82	54 M	54 M	7,151 K

No partition pruning due to bad partitioning key

# ADW: Out of the Box Performance (Parallel)

- The query is executed by a session connected to the Medium service
  - SQL statements executed using the Medium service run with DoP 4

PX running time: 1 seconds  
(out-of-the-box)



# ADW: Out of the Box Performance (Parallel)

Plan Statistics   Parallel   SQL Text   Activity   Optimizer Environment   Outline

Plan Hash Value 523409758   Plan Size 75   Navigate to ↺ ↑ ↻ (i)   Columns ▾

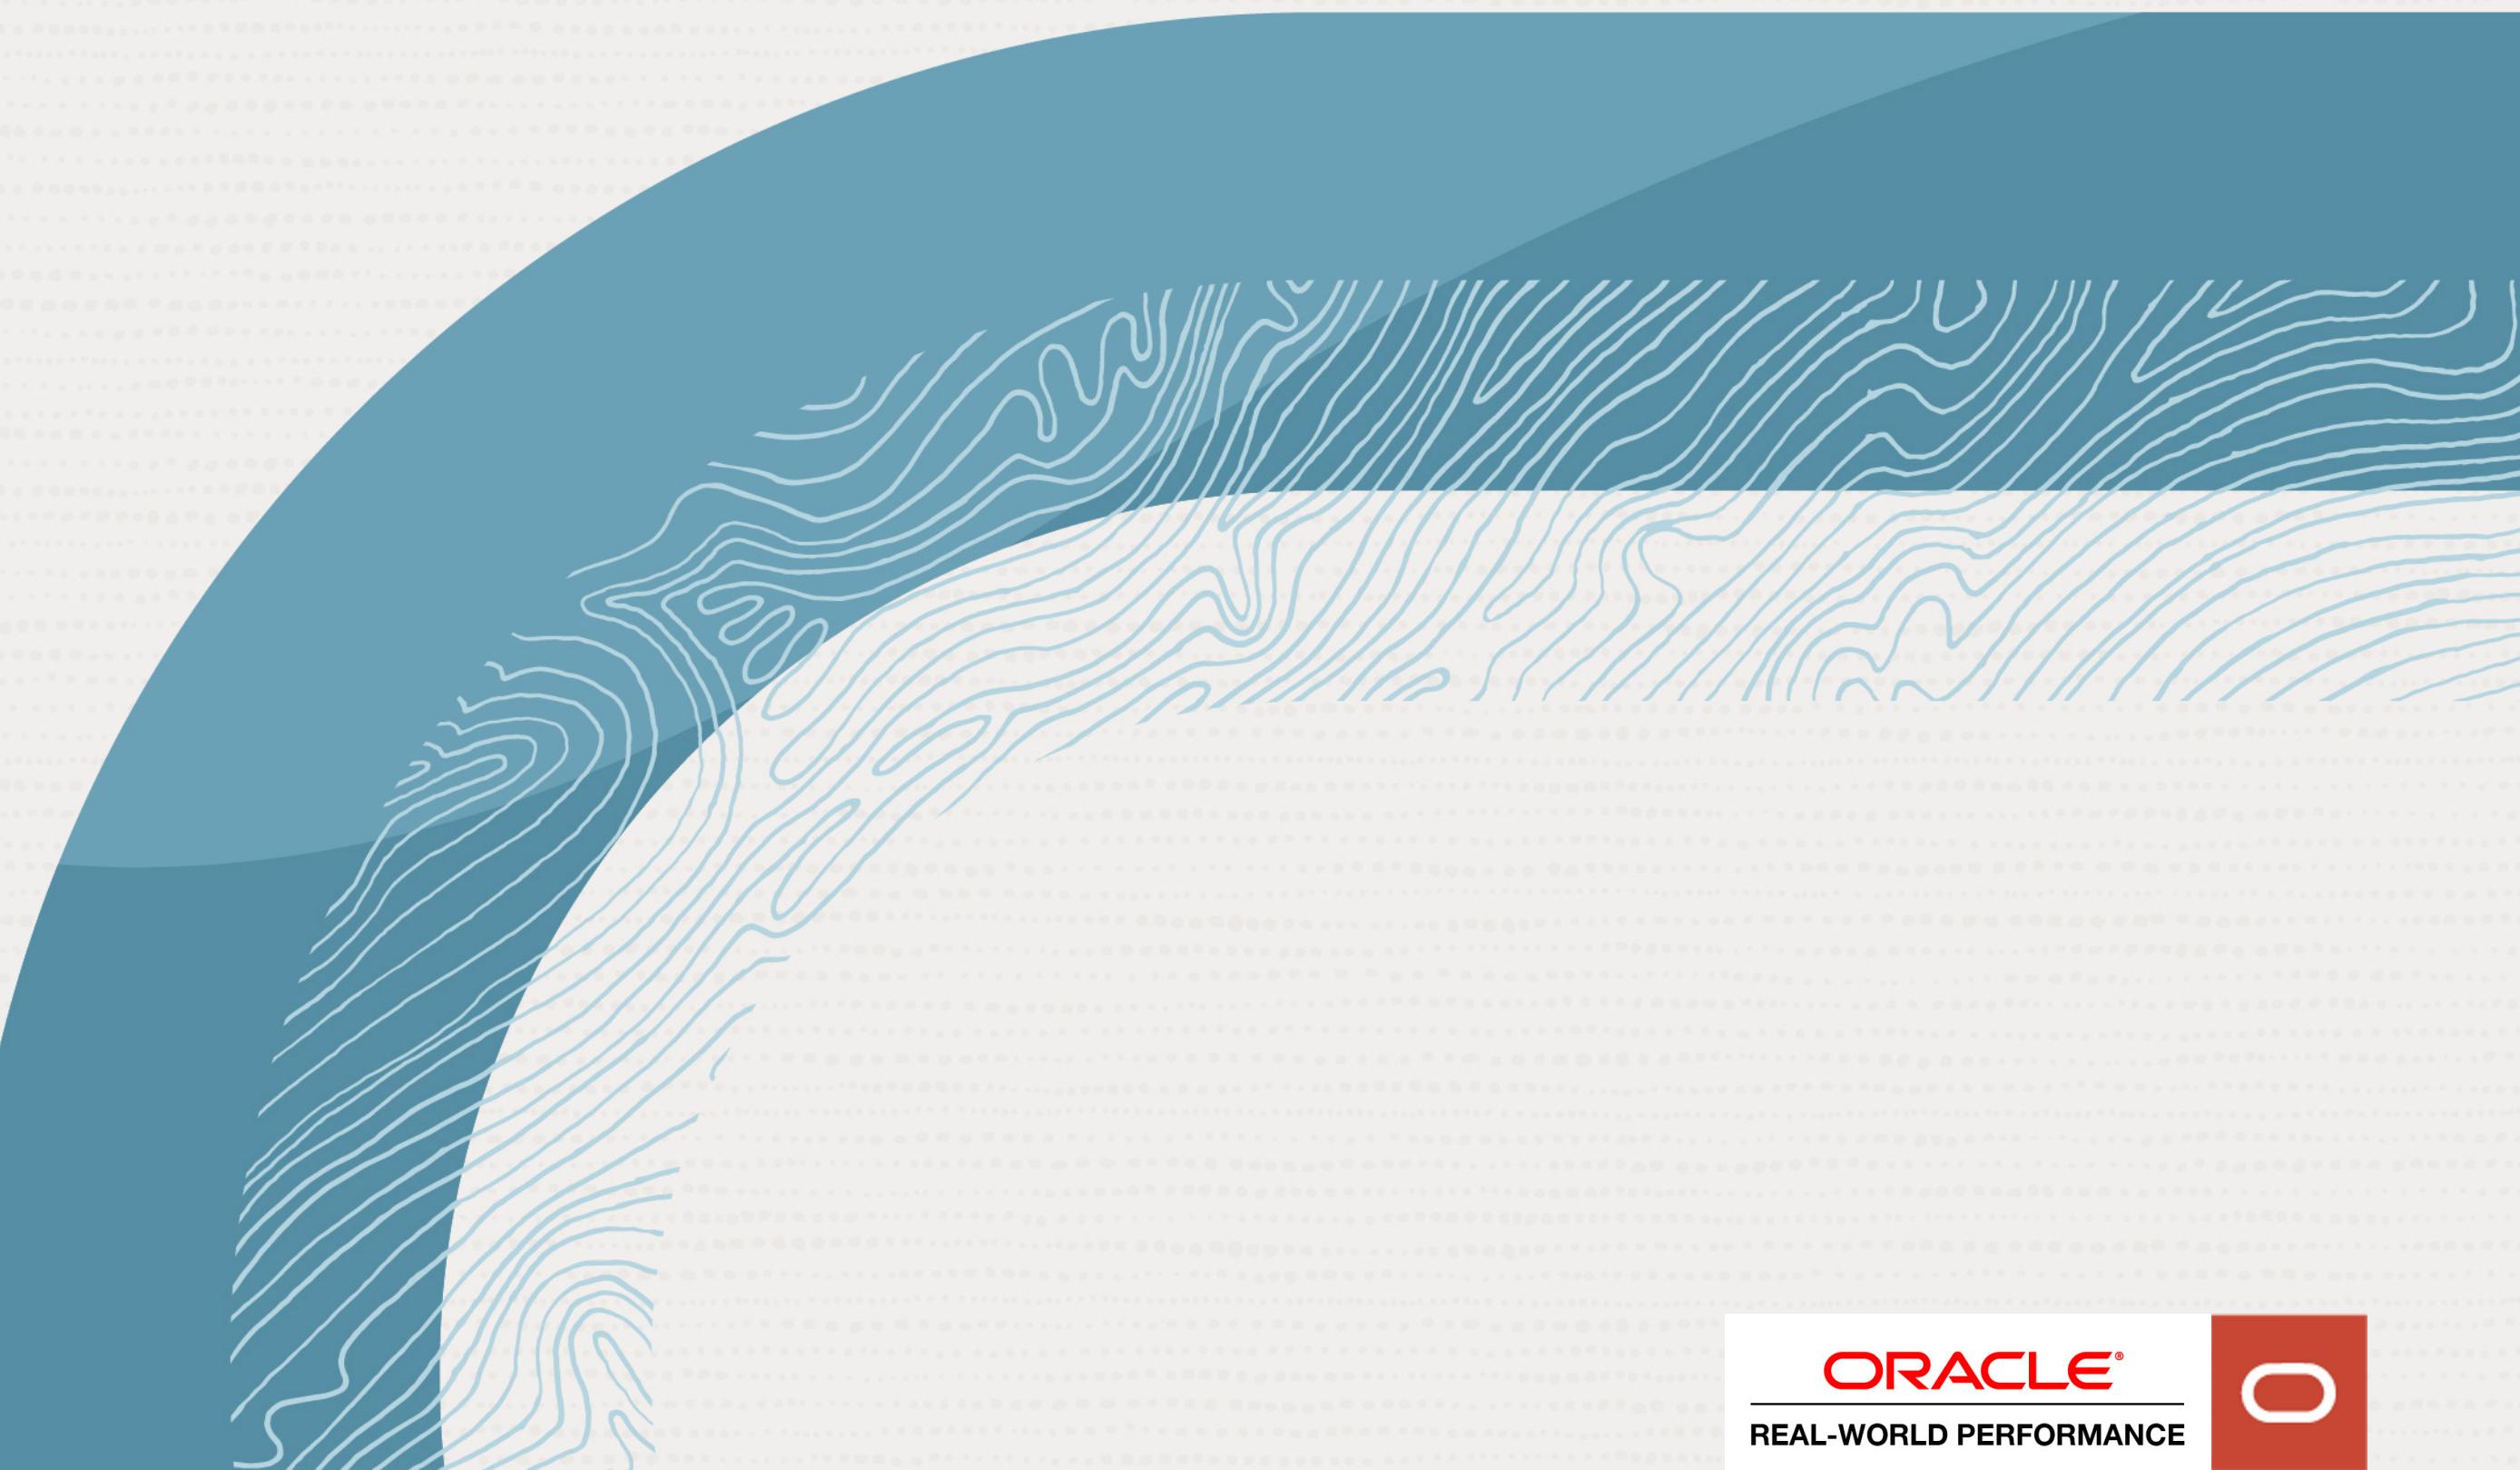
Operation	Object	Informatic	Line	Cost	Rows	Cardinality
▶ PX COORDINATOR			22			
▶ PX SEND QC (RANDOM)	:TQ30001		23			
▶ HASH GROUP BY		⠄⠄	24			
▶ PX RECEIVE			25			
▶ PX SEND HASH	:TQ30000		26			
▶ KEY VECTOR CREATE BUFFERED	:KV0002		27			
▶ VIEW (SEL\$C9F9BE95)	VW_VTN_63105048		28			
▶ HASH JOIN (SEL\$C9F9BE95)		⠄⠄	29	4	6	106 K
▶ JOIN FILTER CREATE	:BF0000	⠄⠄	30	4	2,400	2,400
TABLE ACCESS STORAGE FULL (SEL\$C9F9BE95)	PART_GROUP	⠄	31	4	2,400	2,400
▶ JOIN FILTER USE	:BF0000	⠄⠄	32	4	177 K	177 K
▶ PX BLOCK ITERATOR			33	4	177 K	177 K
TABLE ACCESS STORAGE FULL (SEL\$C9F9BE95)	PART	⠄	34	17	177 K	177 K
▶ LOAD AS SELECT (SEL\$1D319852)	SYS_TEMP_0FD9EFFCB_2B216ACB		35	1	2	
▶ PX COORDINATOR			36	9	4	
▶ PX SEND QC (RANDOM)	:TQ40001		37	4	25	4
▶ HASH GROUP BY		⠄⠄	38	4	25	4
▶ PX RECEIVE			39	4	25	8
▶ PX SEND HASH	:TQ40000		40	4	25	8
▶ KEY VECTOR CREATE BUFFERED	:KV0003		41	4	25	8
▶ PX BLOCK ITERATOR			42	4	2,862	2,862
TABLE ACCESS STORAGE FULL (SEL\$1D319852)	SUPPLIER	⠄	43	2	2,862	2,862
▶ PX COORDINATOR			44	9	48	
▶ PX SEND QC (ORDER)	:TQ50004		45	4	213	48
▶ SORT ORDER BY			46	4	213	48
▶ PX RECEIVE			47	4	275	48
▶ PX SEND RANGE	:TQ50003	⠄⠄	48	4	275	48
▶ HASH JOIN BUFFERED			49	4	275	48
▶ PX RECEIVE			50	4	1	4
▶ PX SEND BROADCAST	:TQ50000		51	4	1	4
▶ PX BLOCK ITERATOR			52	4	1	1
TABLE ACCESS STORAGE FULL (SEL\$45780)	SYS_TEMP_0FD9EFFC8_2B216ACB		53	1	1	1

The execution plan is different and uses In-Memory Aggregation (Vector Transformation)  
In ADW IMA is enabled and can take advantage of the data stored in columnar store that resides on flash disks

Cardinality estimate for SUPPLIER is now better because dynamic statistics kicked in automatically for every parallel query.

... and that's pretty much it

# Summary



# Summary

Get the fundamentals right and everything else falls into place

Autonomous Database can help address some problems:

- Augments bad/stale stats
- Dynamic statistics may correct some cardinality estimates (do not solely rely on dynamic statistics)
- Stats quality (every column has a histogram)
- Use of features like IMA, Columnar Cache in Flash is available

These features are enabled by default and kick-in automatically.

Where Autonomous Database cannot help:

- Suboptimal schema design
  - Wrong data types on join keys
  - Lack of constraints
  - Sub-optimal partitioning strategy

